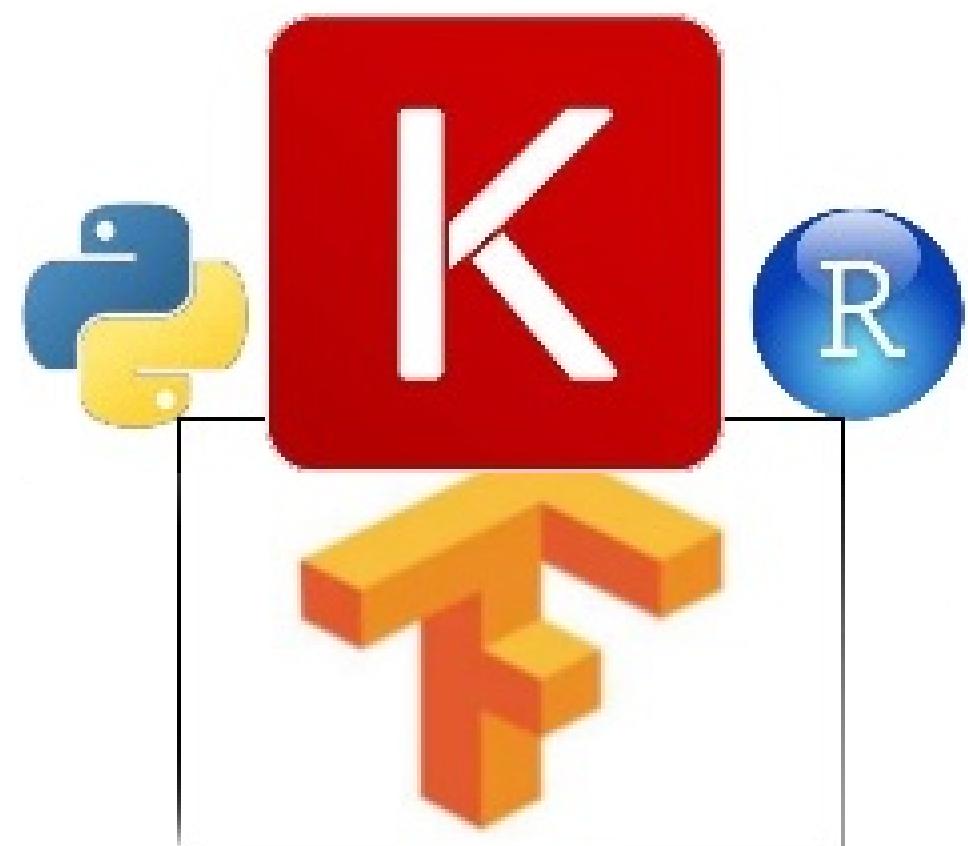


# **Image Classification & Natural Language Processing with Keras and TensorFlow in Python**

## **Workshop**



Dr. Shirin Glander

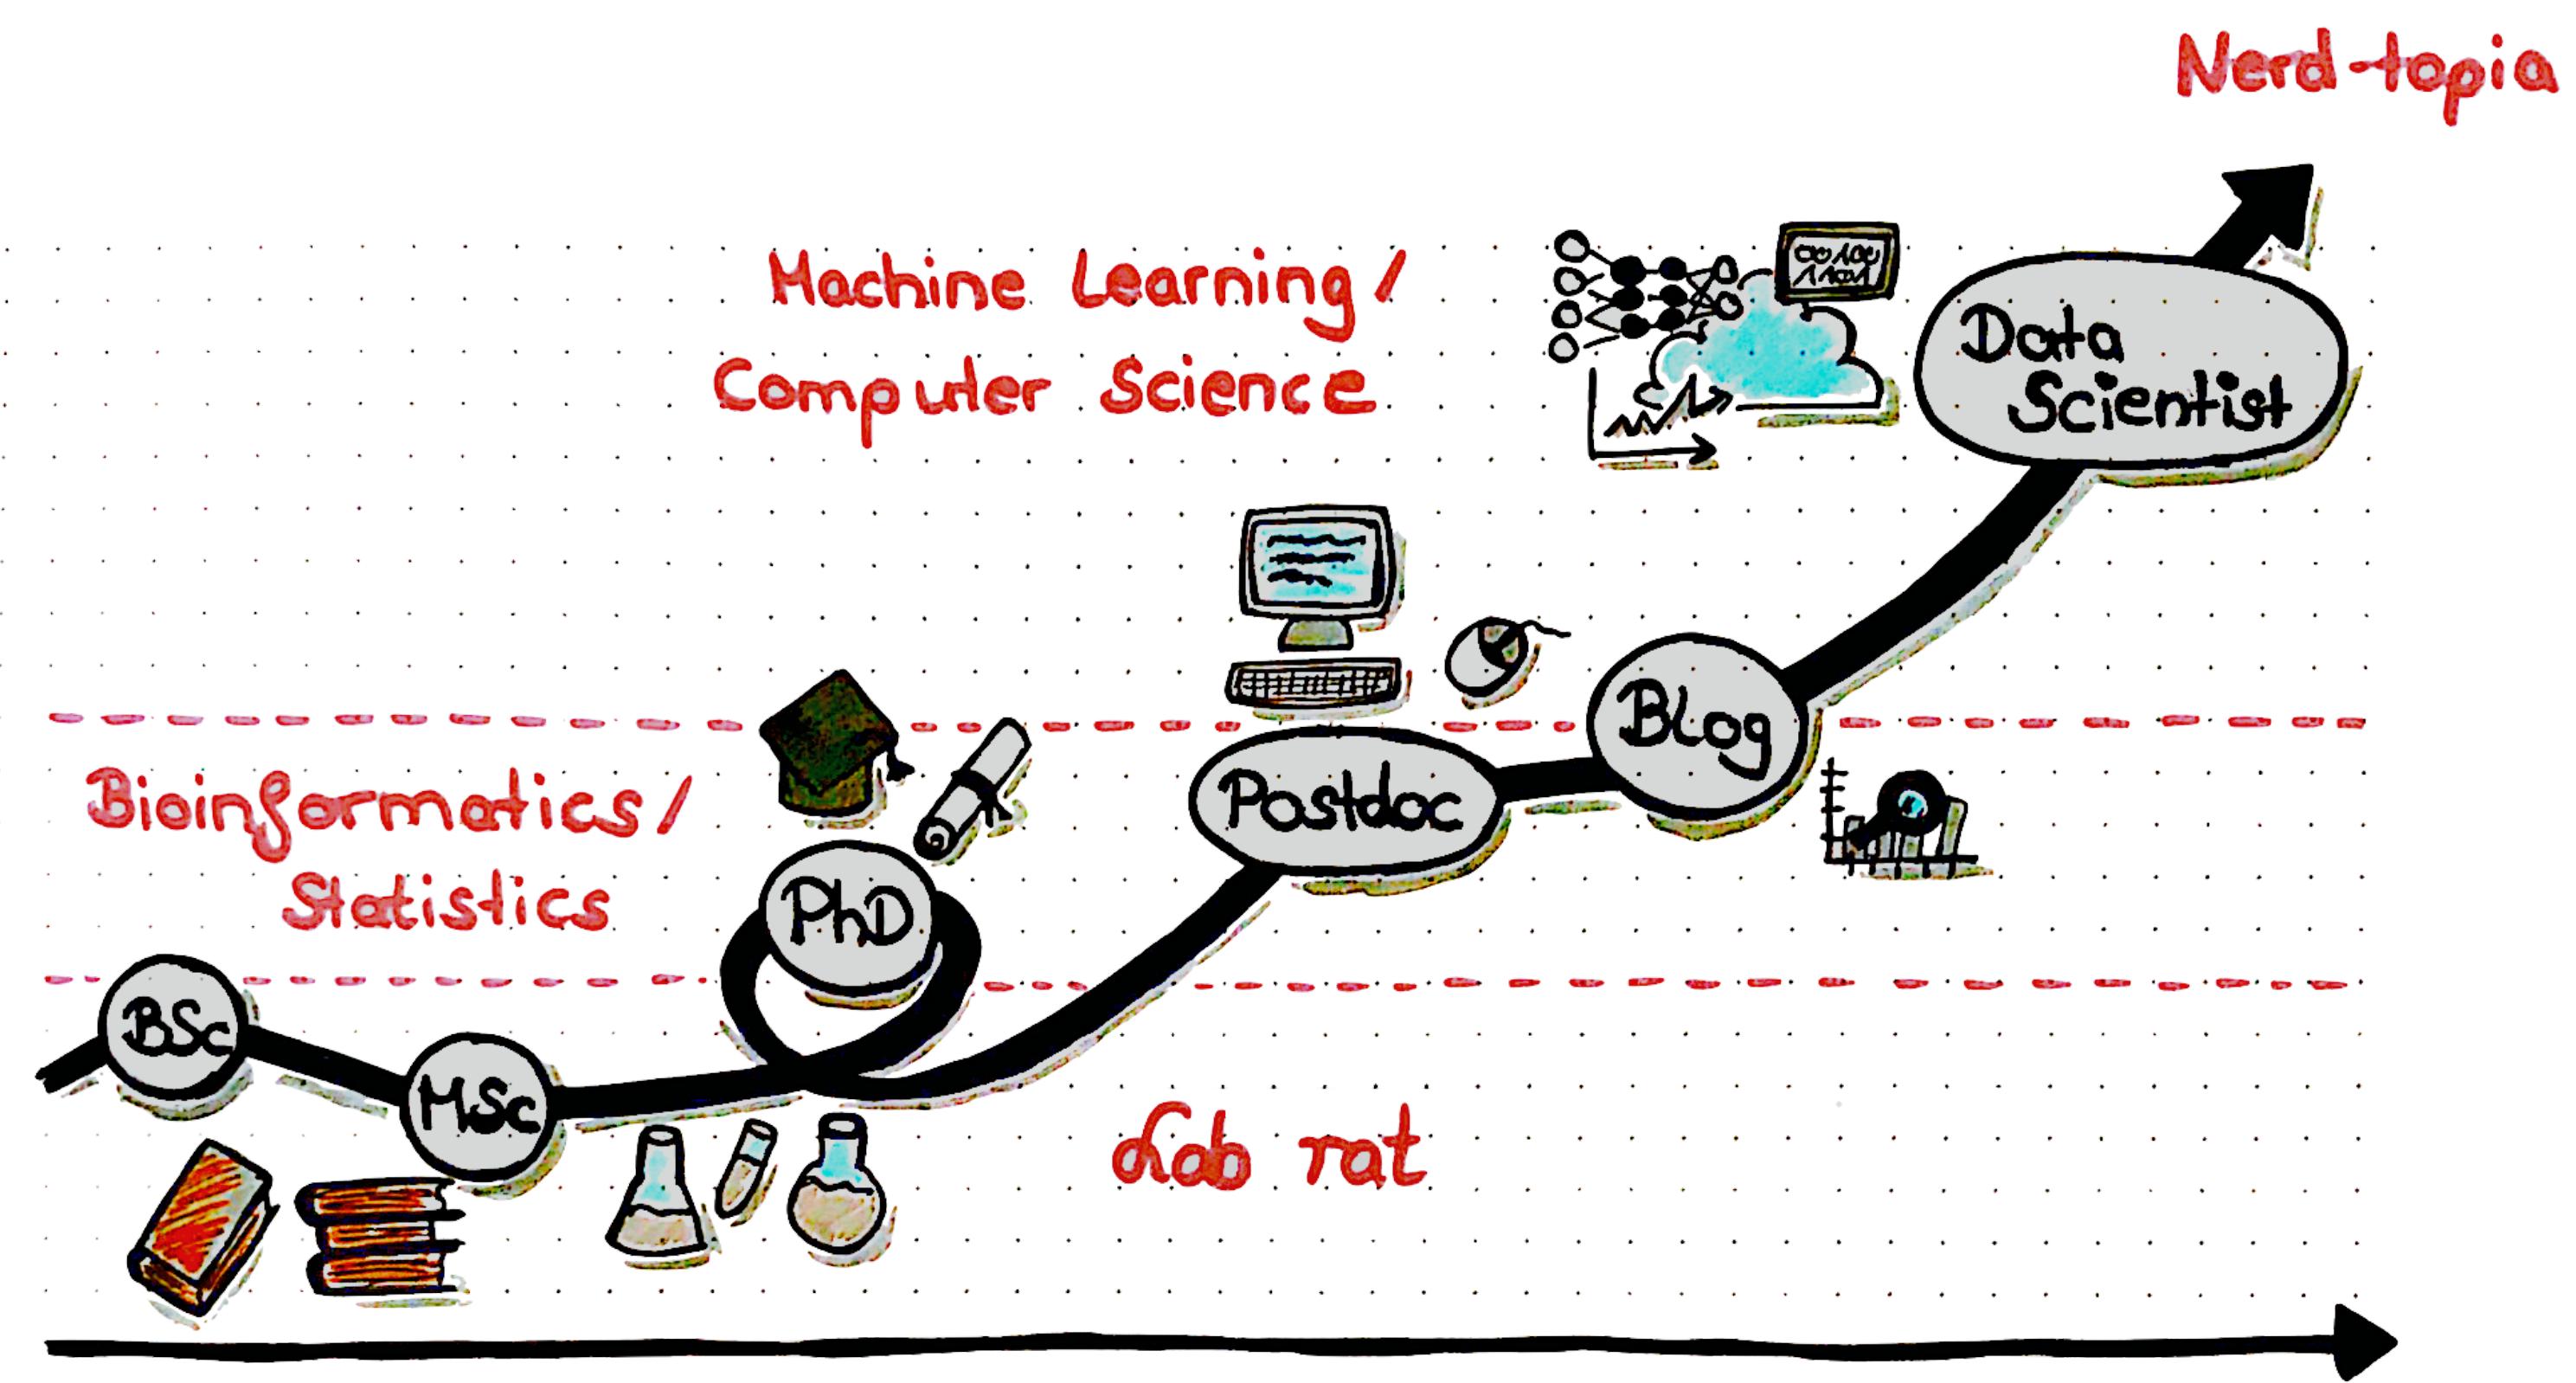
24.10.2018 - Bitmarck

# About this Workshop

- Learn what neural nets are and how they can be used in computer vision & NLP.
- Understand how a computer learns to "see" ...
- ... and how a computer understands natural language.
- Apply pre-trained nets and modify them.
- Build a model from scratch that differentiates between fruits on images.
- Visualise convolutions and layers.
- Train a text classification model with CNNs ...
- ... and with LSTMs.
- (Bonus: Explaining our classifications with LIME)

Material on Github: [https://github.com/ShirinG/image\\_classification\\_keras\\_tf](https://github.com/ShirinG/image_classification_keras_tf)

# About me



# Why Keras + TensorFlow?

## The Keras philosophy

- Rapid prototyping
- Very fast turnover from idea to model

## Keras is now the default TensorFlow API

- Keras is the first high-level library added to core TensorFlow at Google

## Start simple & small

- Always start simple and small to generate a baseline model against which to compare!

# How does a computer learn to "see"?

## Convolutional Neural Nets



- image classification  
class = tree
- object detection  
flower

## Computer Vision

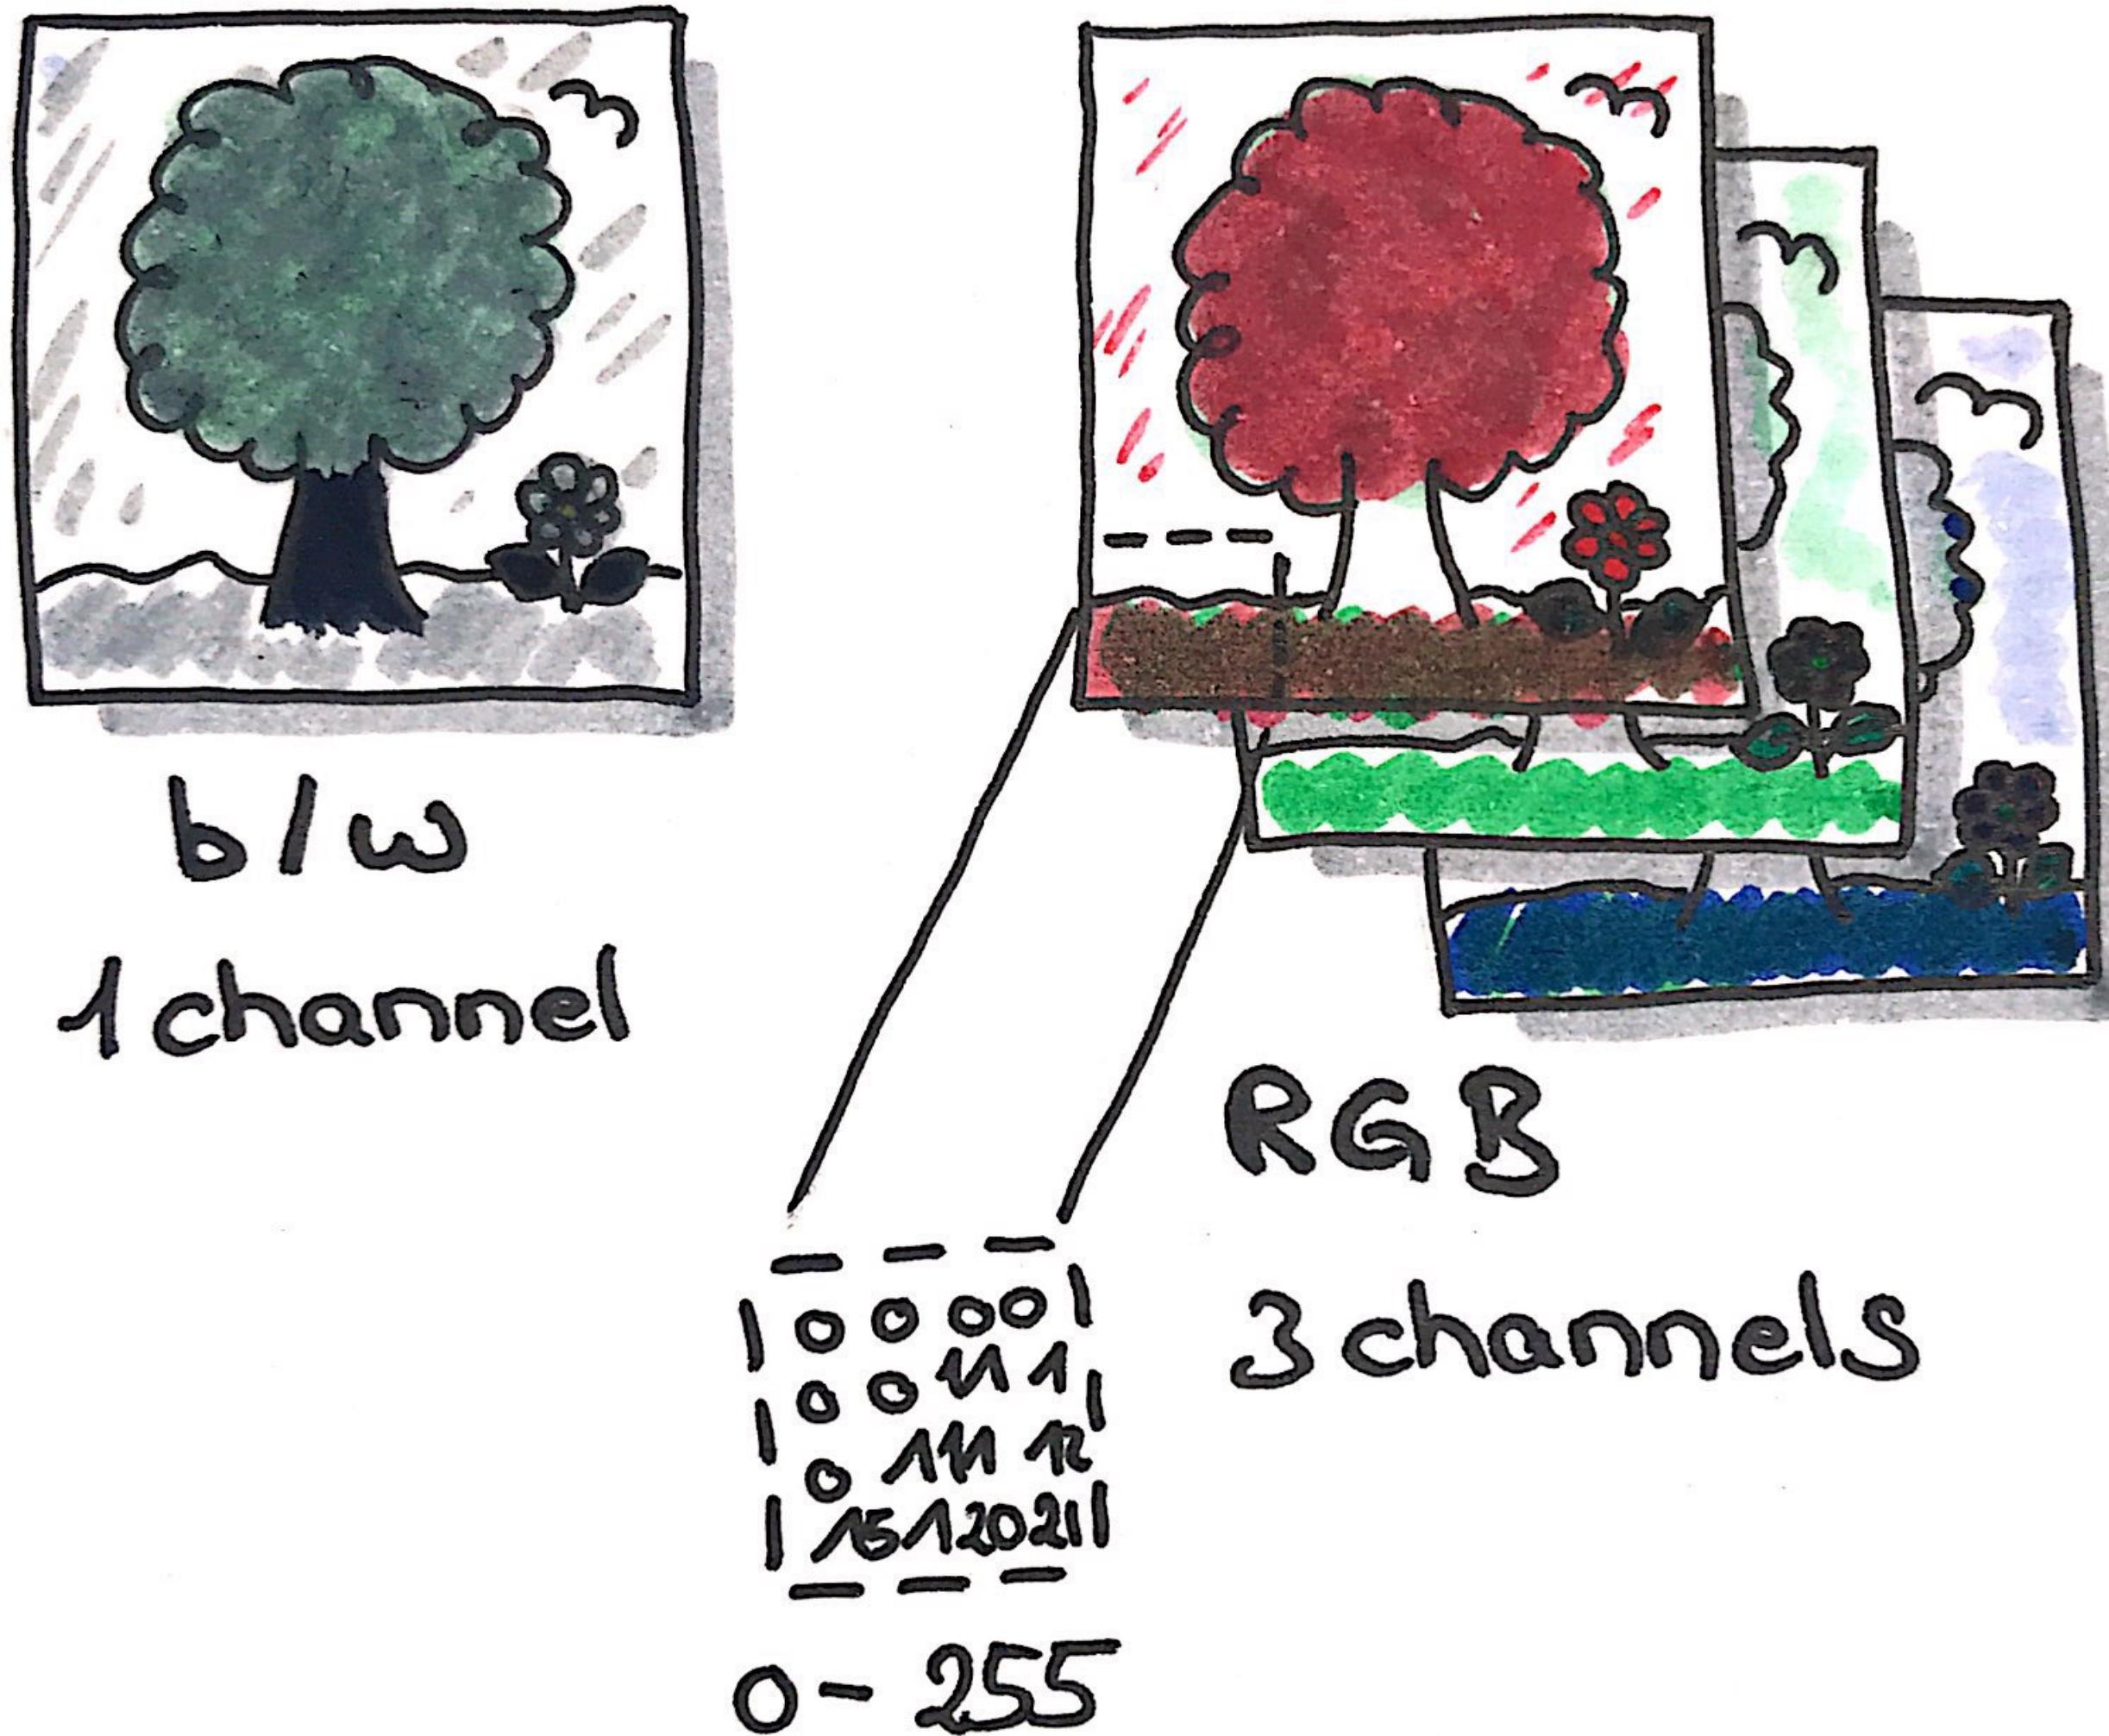
See also [this video from codecentric.AI](#) for more on Computer Vision Basics!

## CNNs can also be used for text classification

Dies ist ein Blindtext. An ihm lässt sich vieles über die Schrift ablesen, in der er gesetzt ist. Auf den ersten Blick wird der Grauwert der Schriftfläche sichtbar. Dann kann man prüfen, wie gut die Schrift zu lesen ist und wie sie auf den Leser wirkt.

Dies ist ein Blindtext. An ihm lässt sich vieles über die Schrift ablesen, in der er gesetzt ist. Auf den ersten Blick wird der Grauwert der Schriftfläche sichtbar. Dann kann man prüfen, wie gut die Schrift zu lesen ist und wie sie auf den Leser wirkt.

# How does a computer learn to "see"?



# **How does a computer understand natural language?**

## **Document Term Matrix**

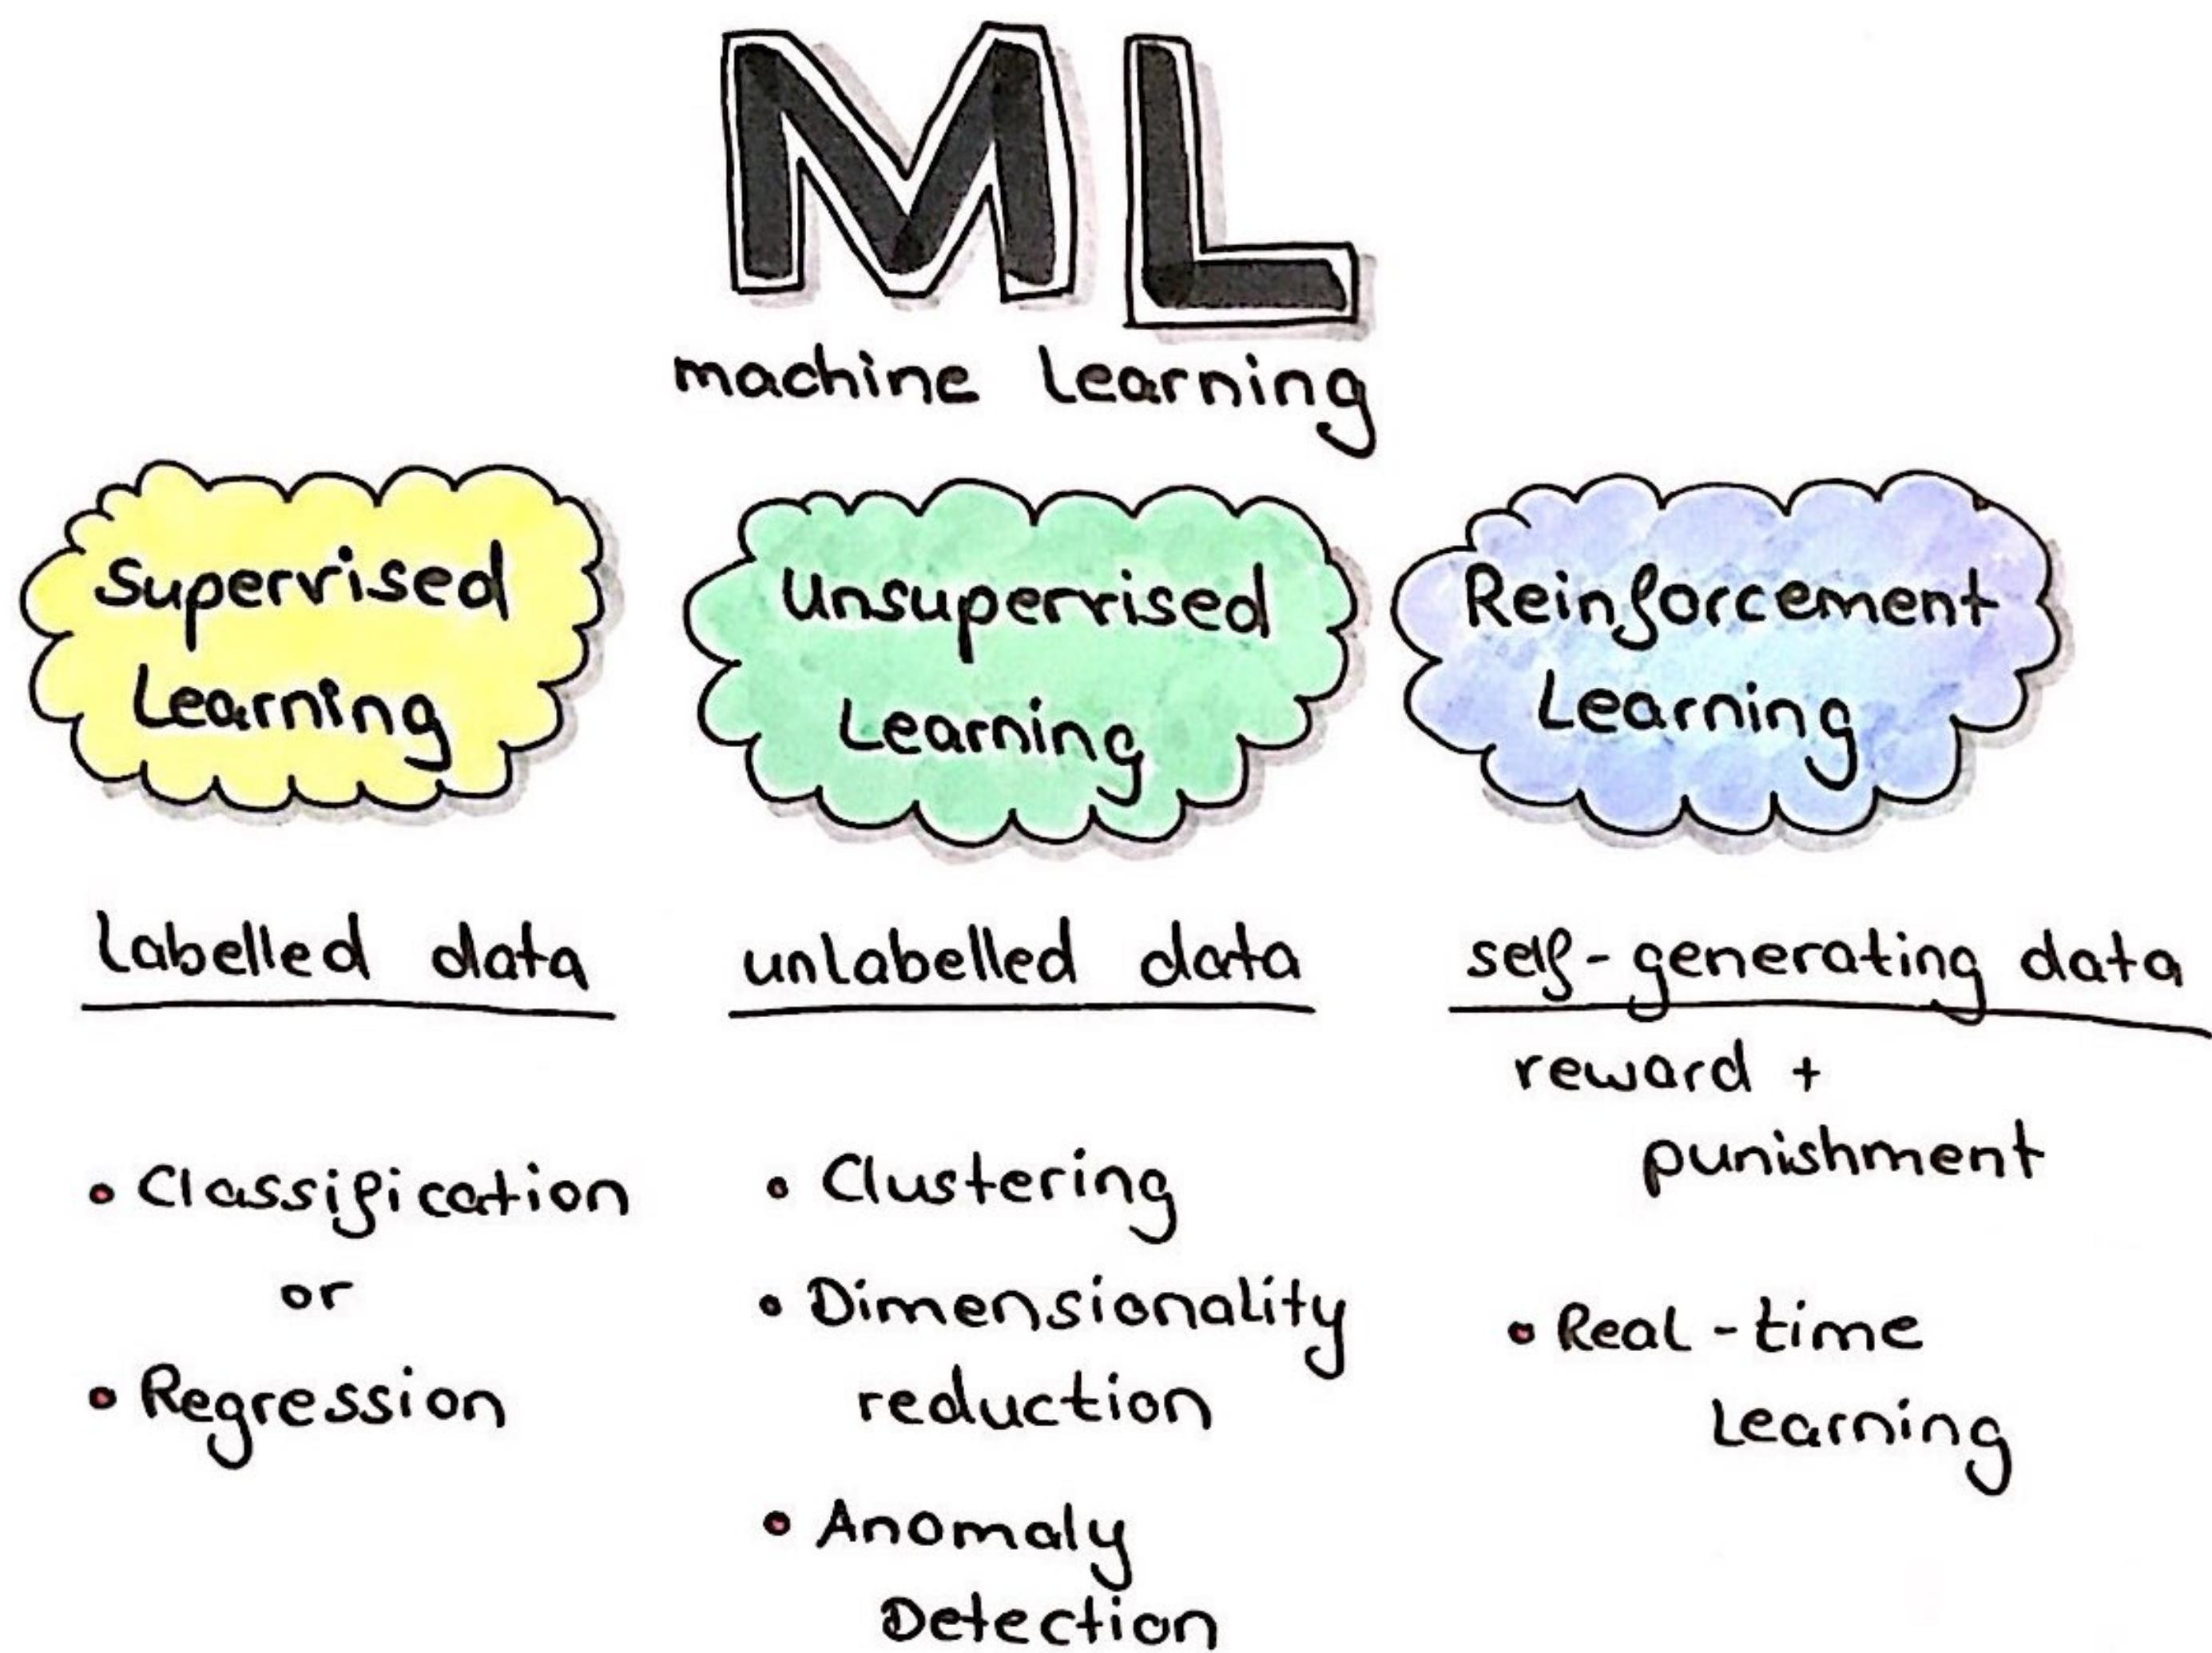
- Generate list of words in the complete corpus
- Count word frequency in each document
- Optional: normalize (e.g. tf/idf)
- Optional: remove stop words, etc.

## **Word embedding matrices**

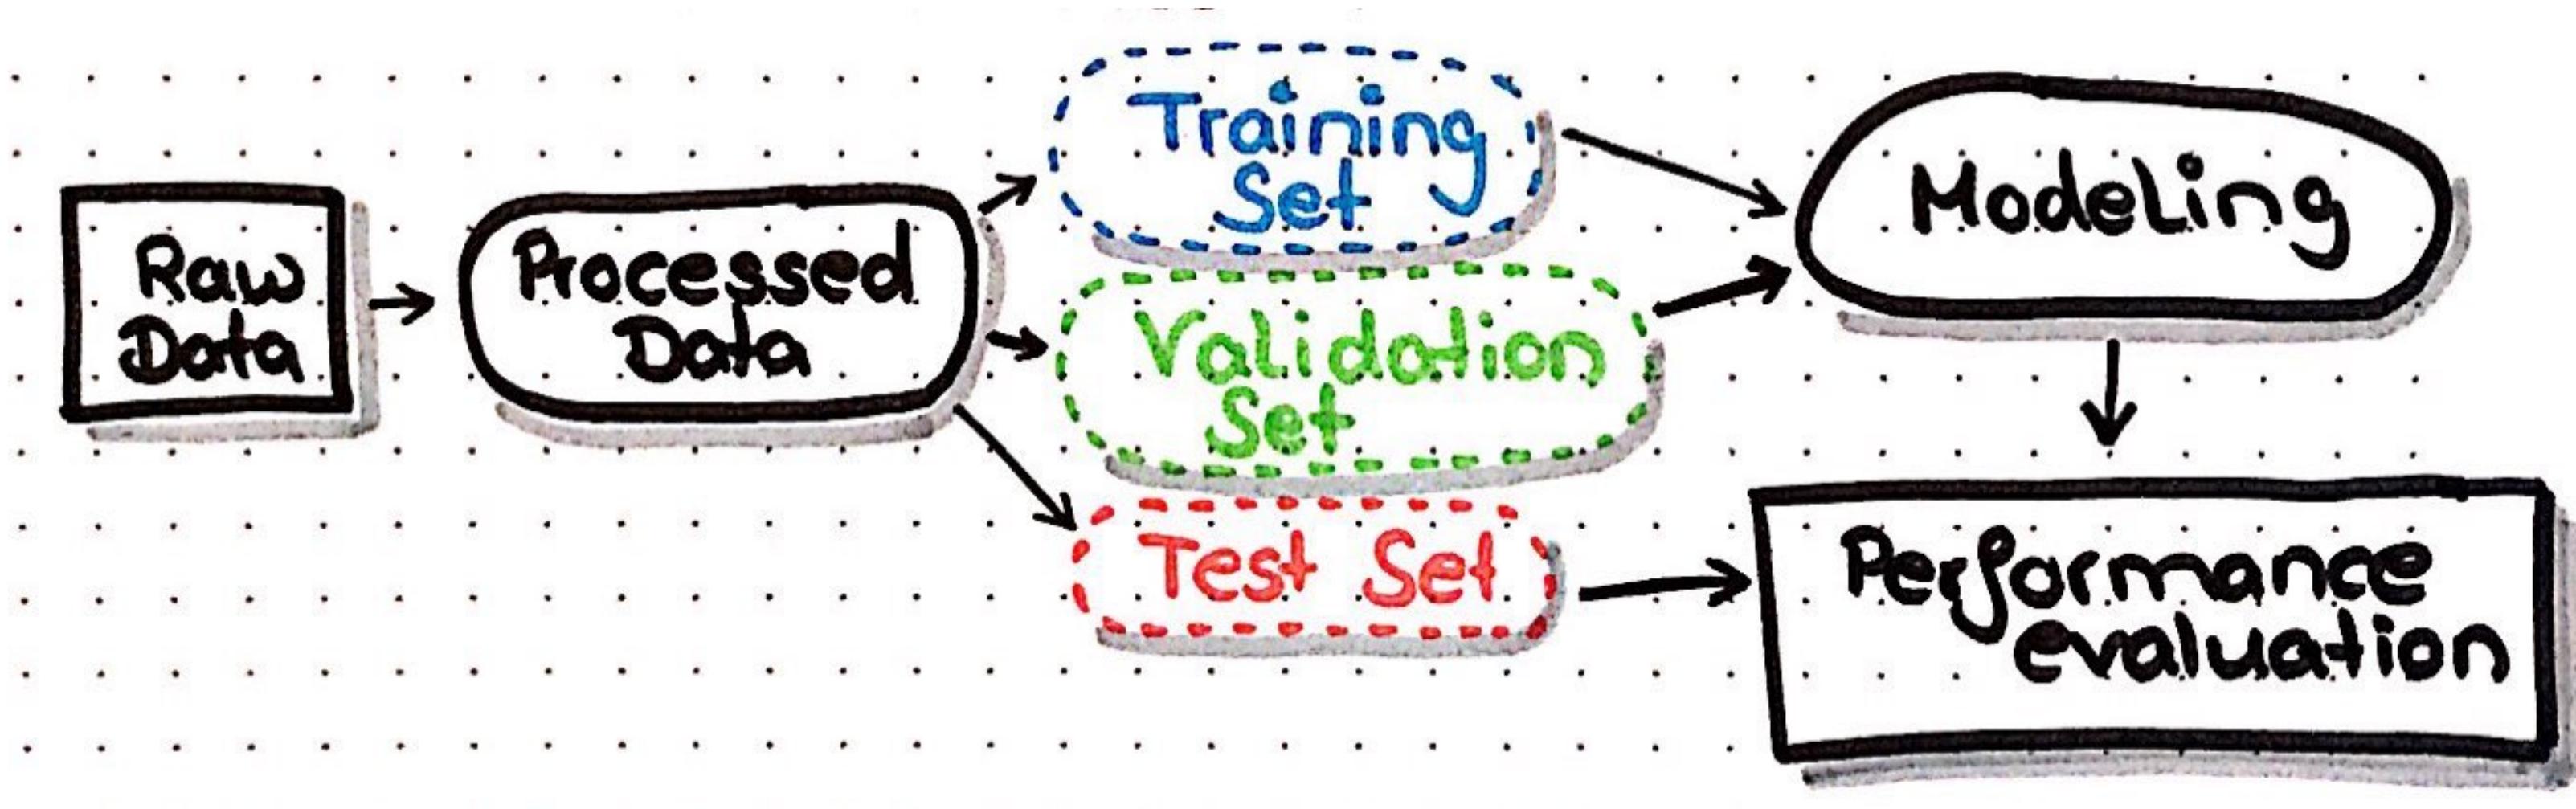
# Word embeddings

- Words/phrases are mapped to vectors
- Can be created with neural nets, PCA/t-SNE/UMAP, co-occurrence matrices
- Vectors represent context
- e.g. word2vec

# What is Machine Learning (ML)



## A typical ML Workflow

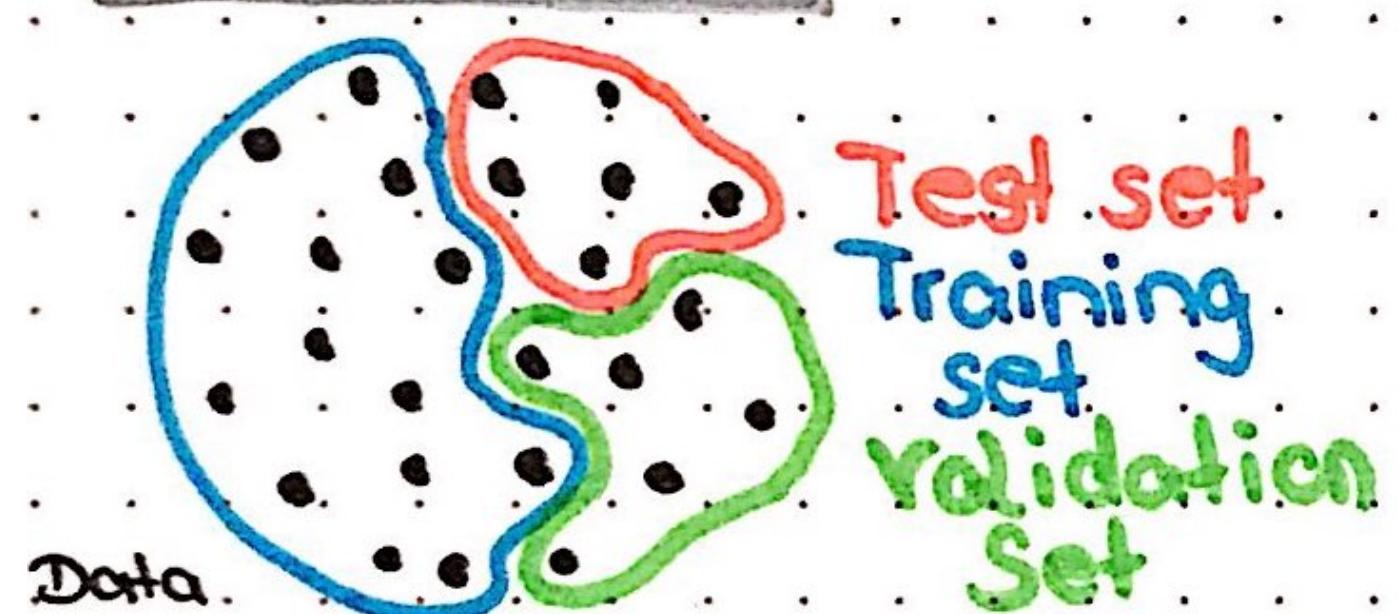


## Why use validation and test data?

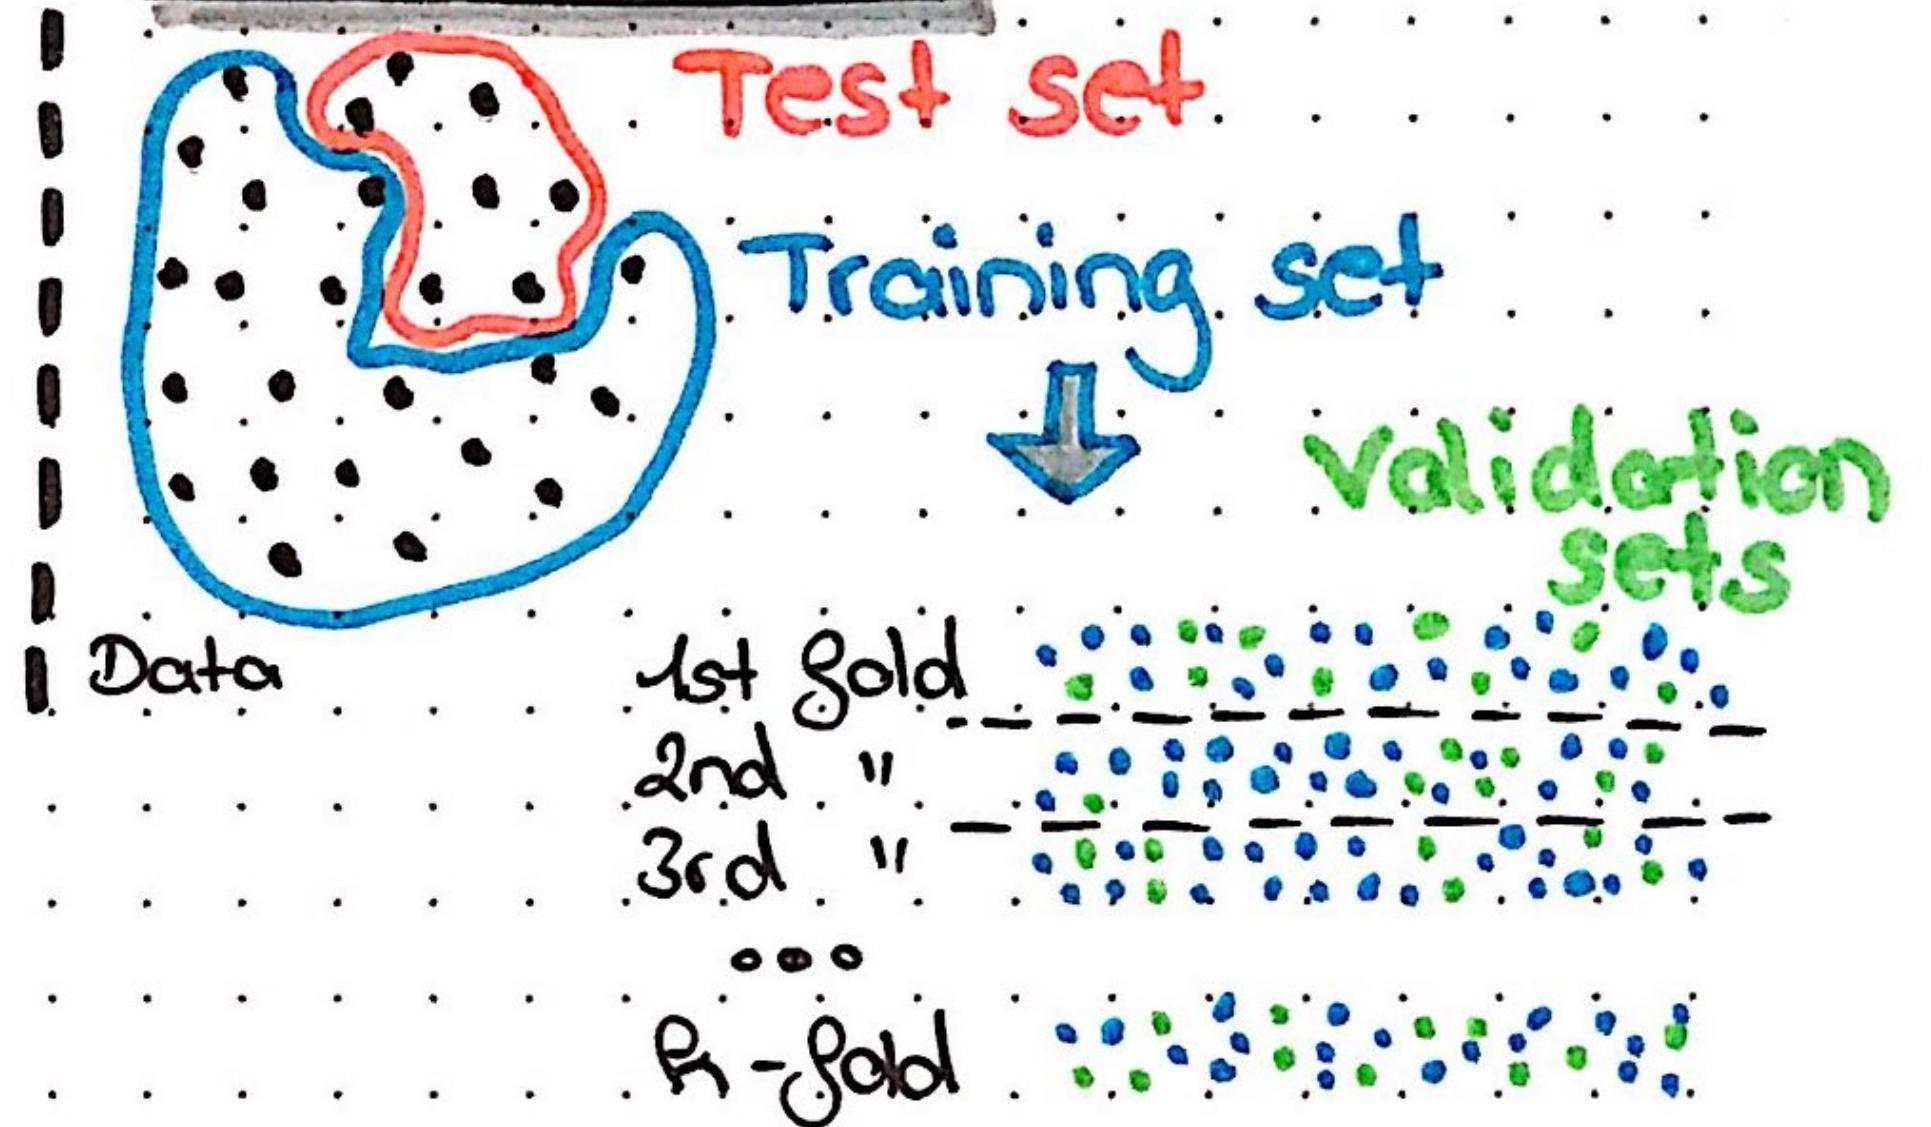


# Validation methods

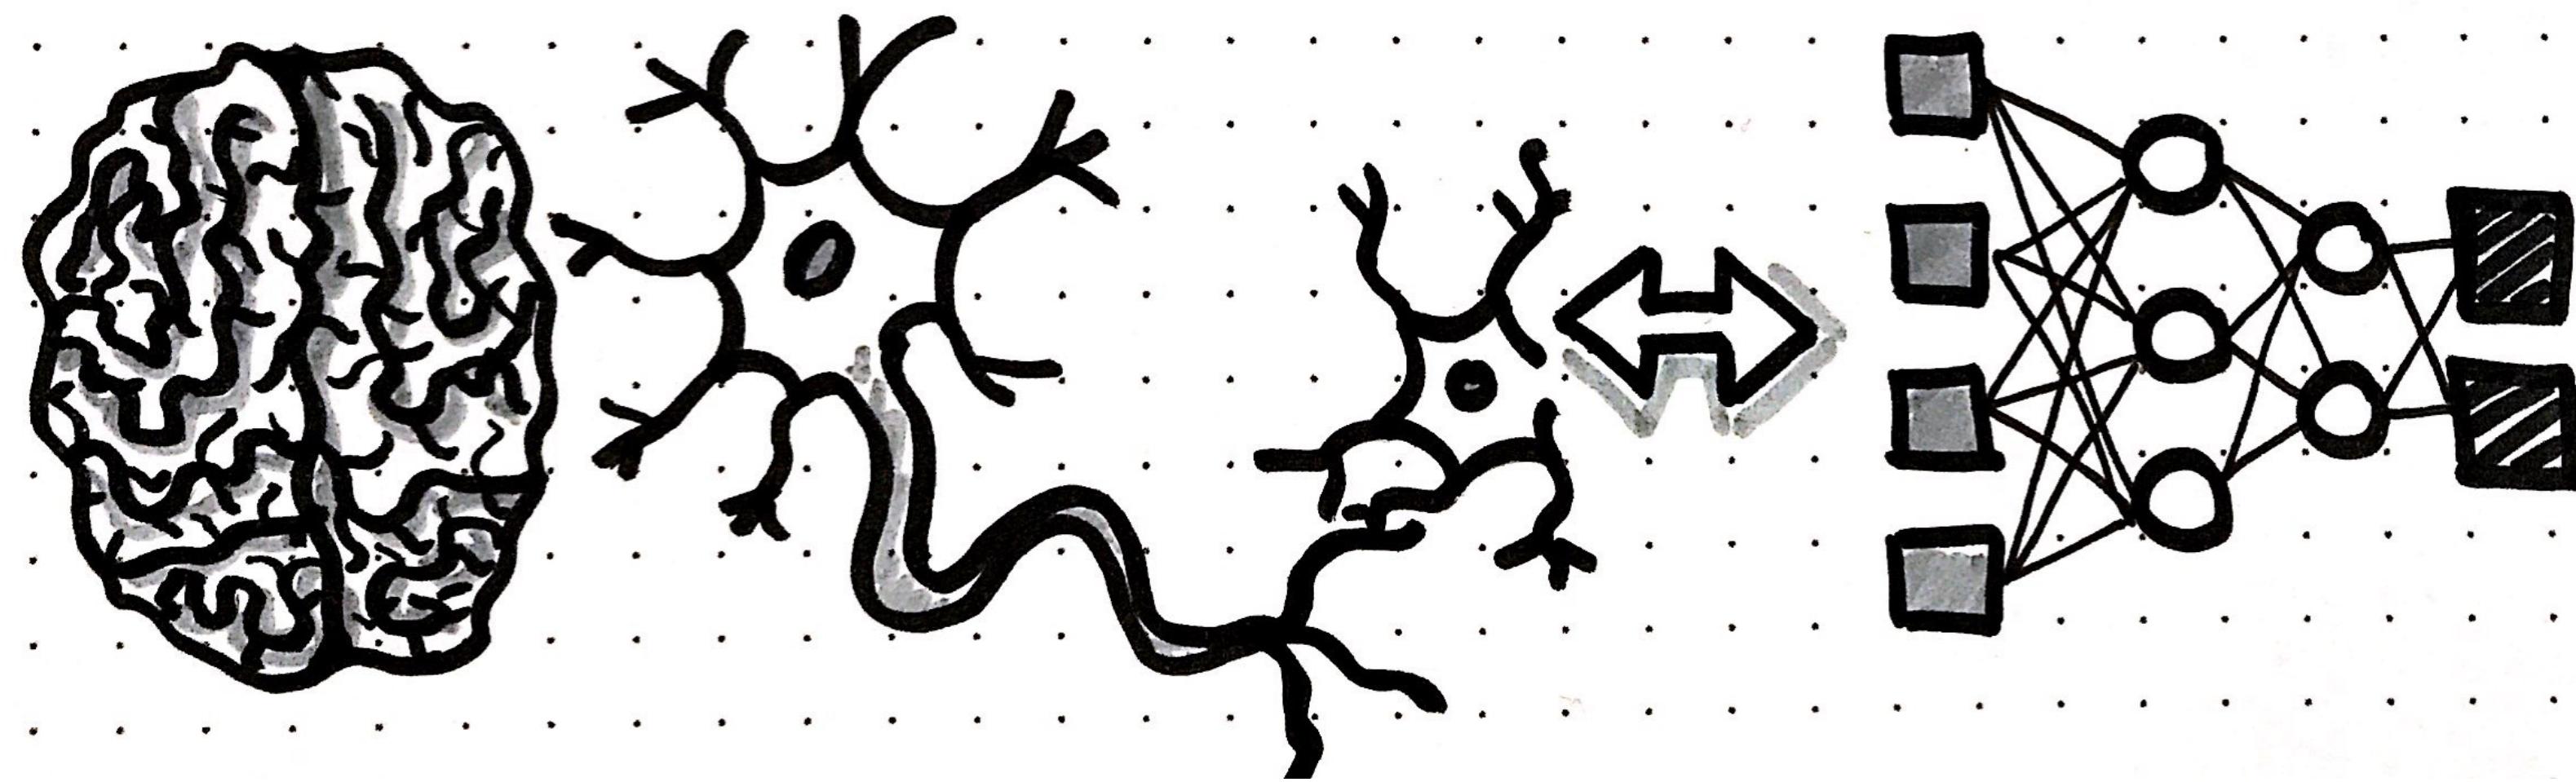
## Hold-out validation



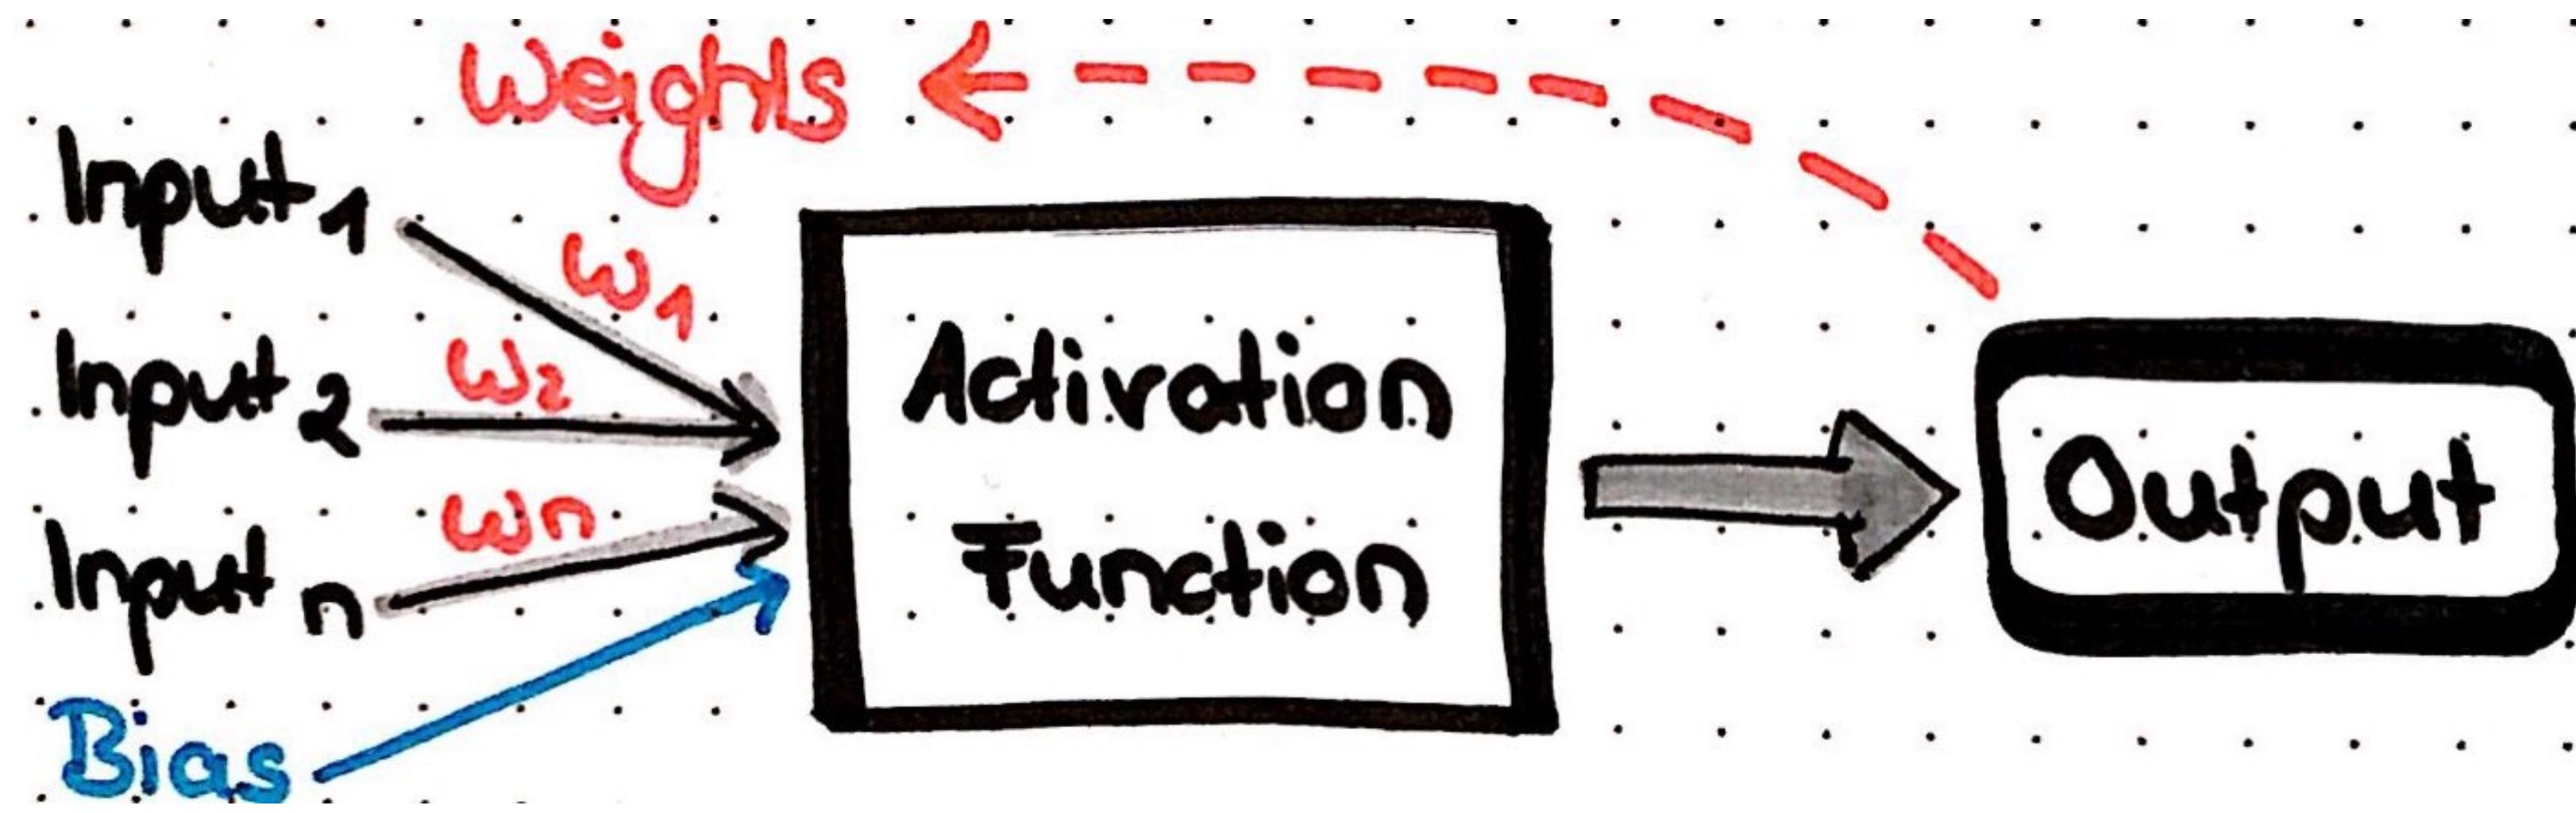
## Cross-validation



# What are neural nets?

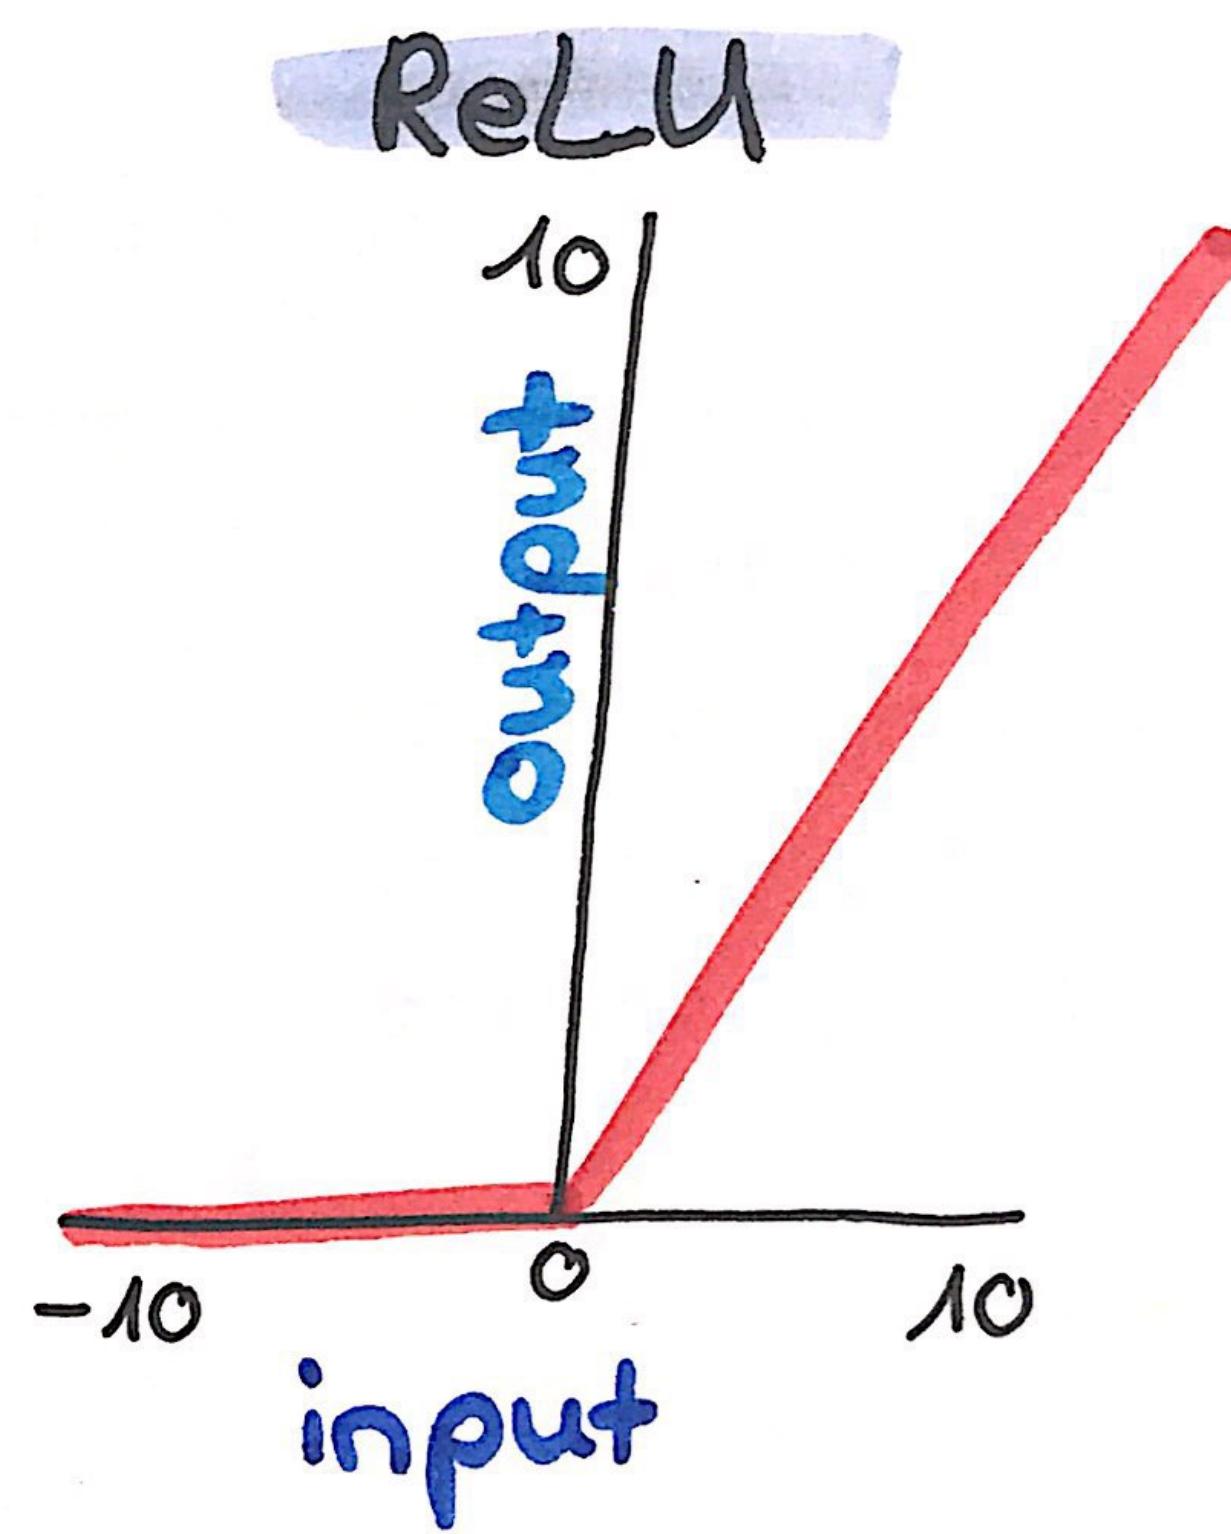
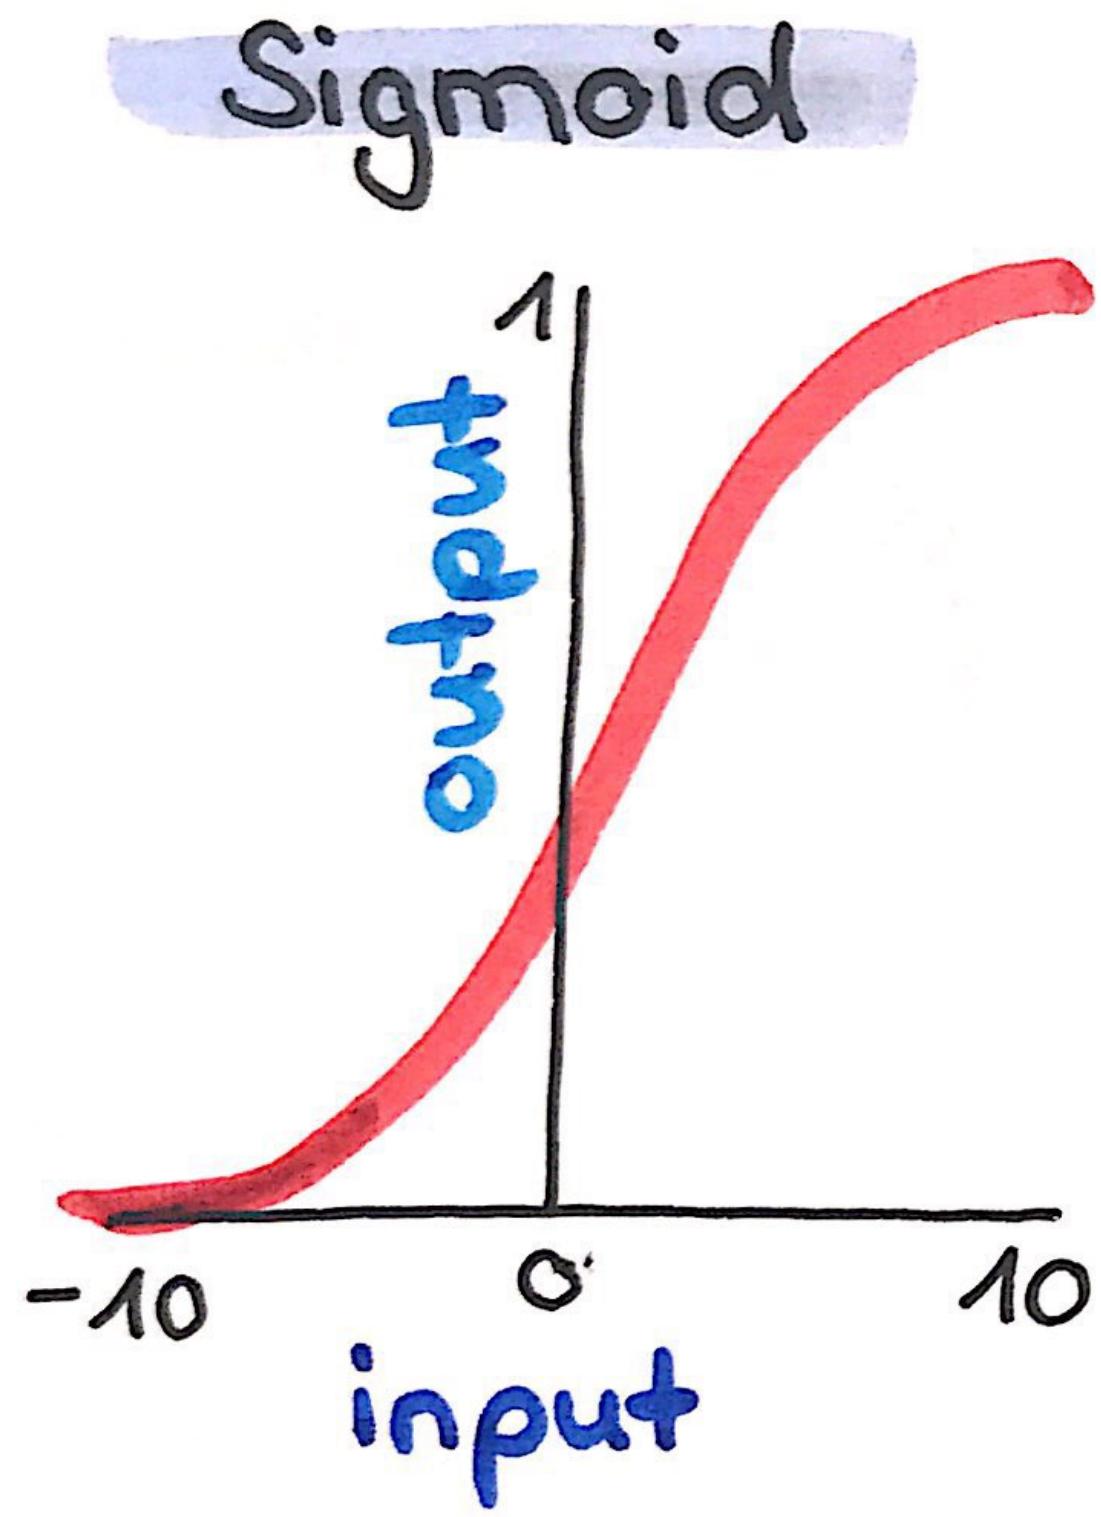


# Perceptrons



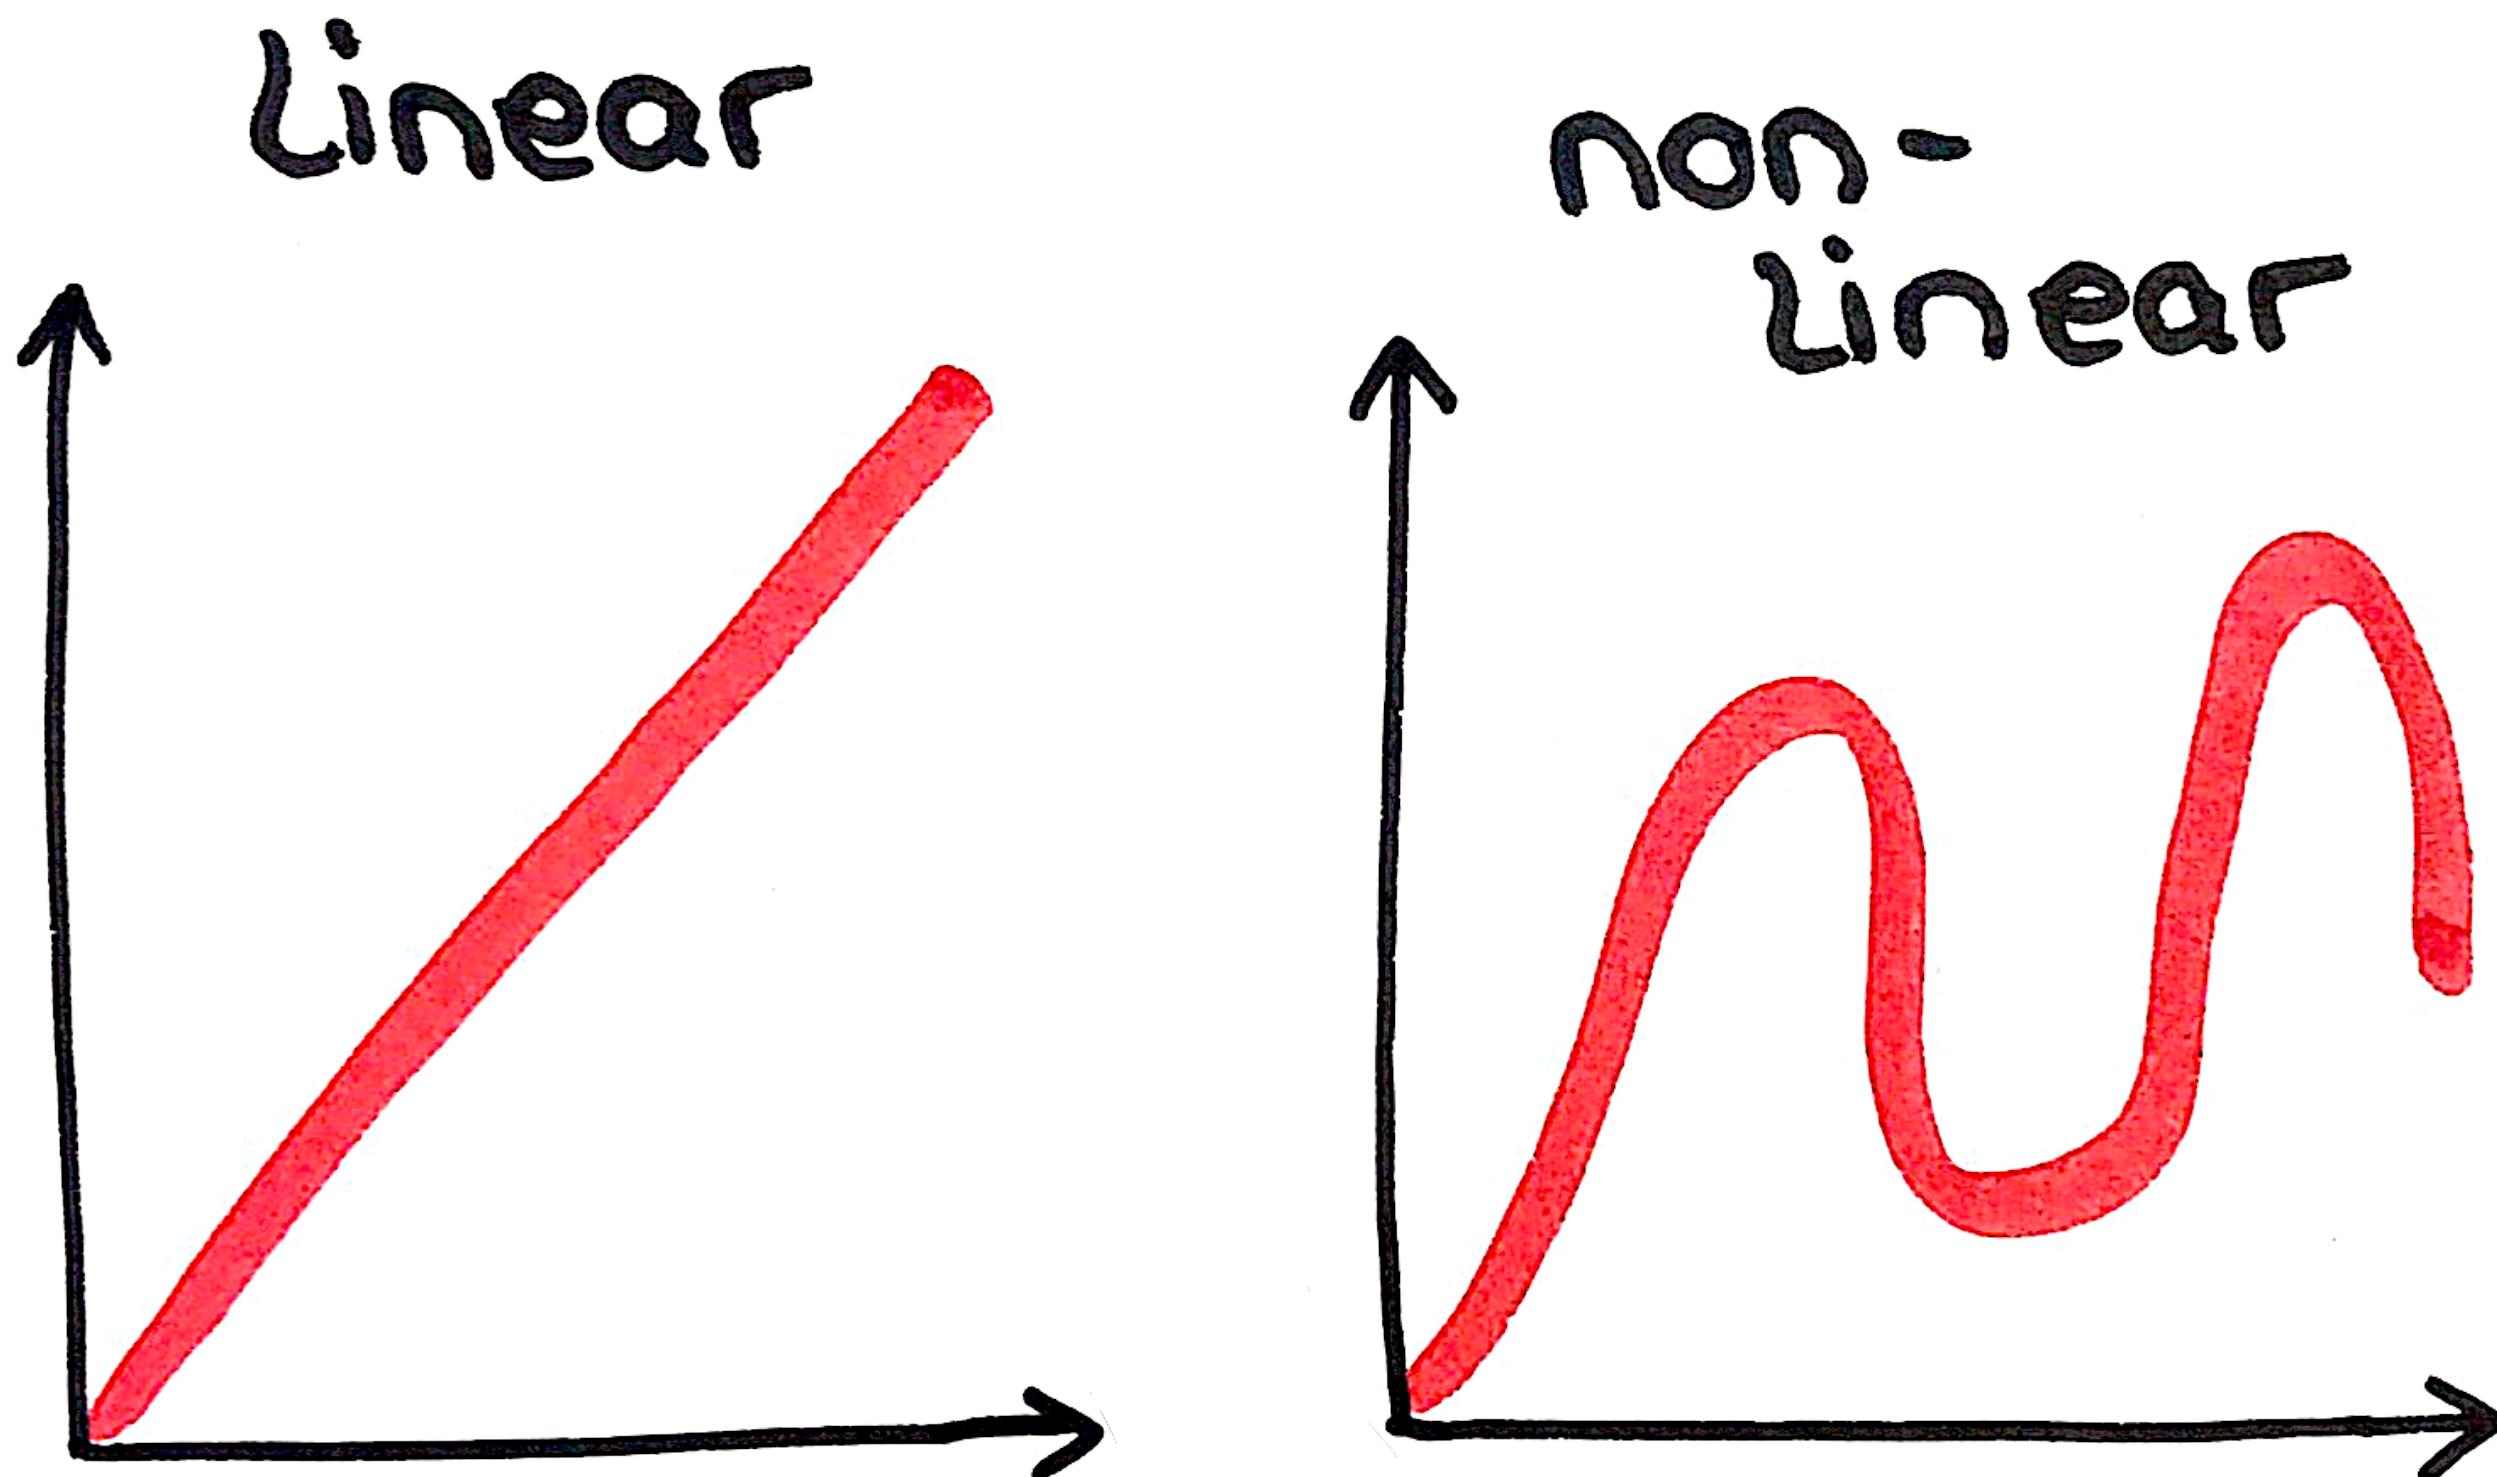
# Activation functions normalise input

- every problem can be described with a mathematical function

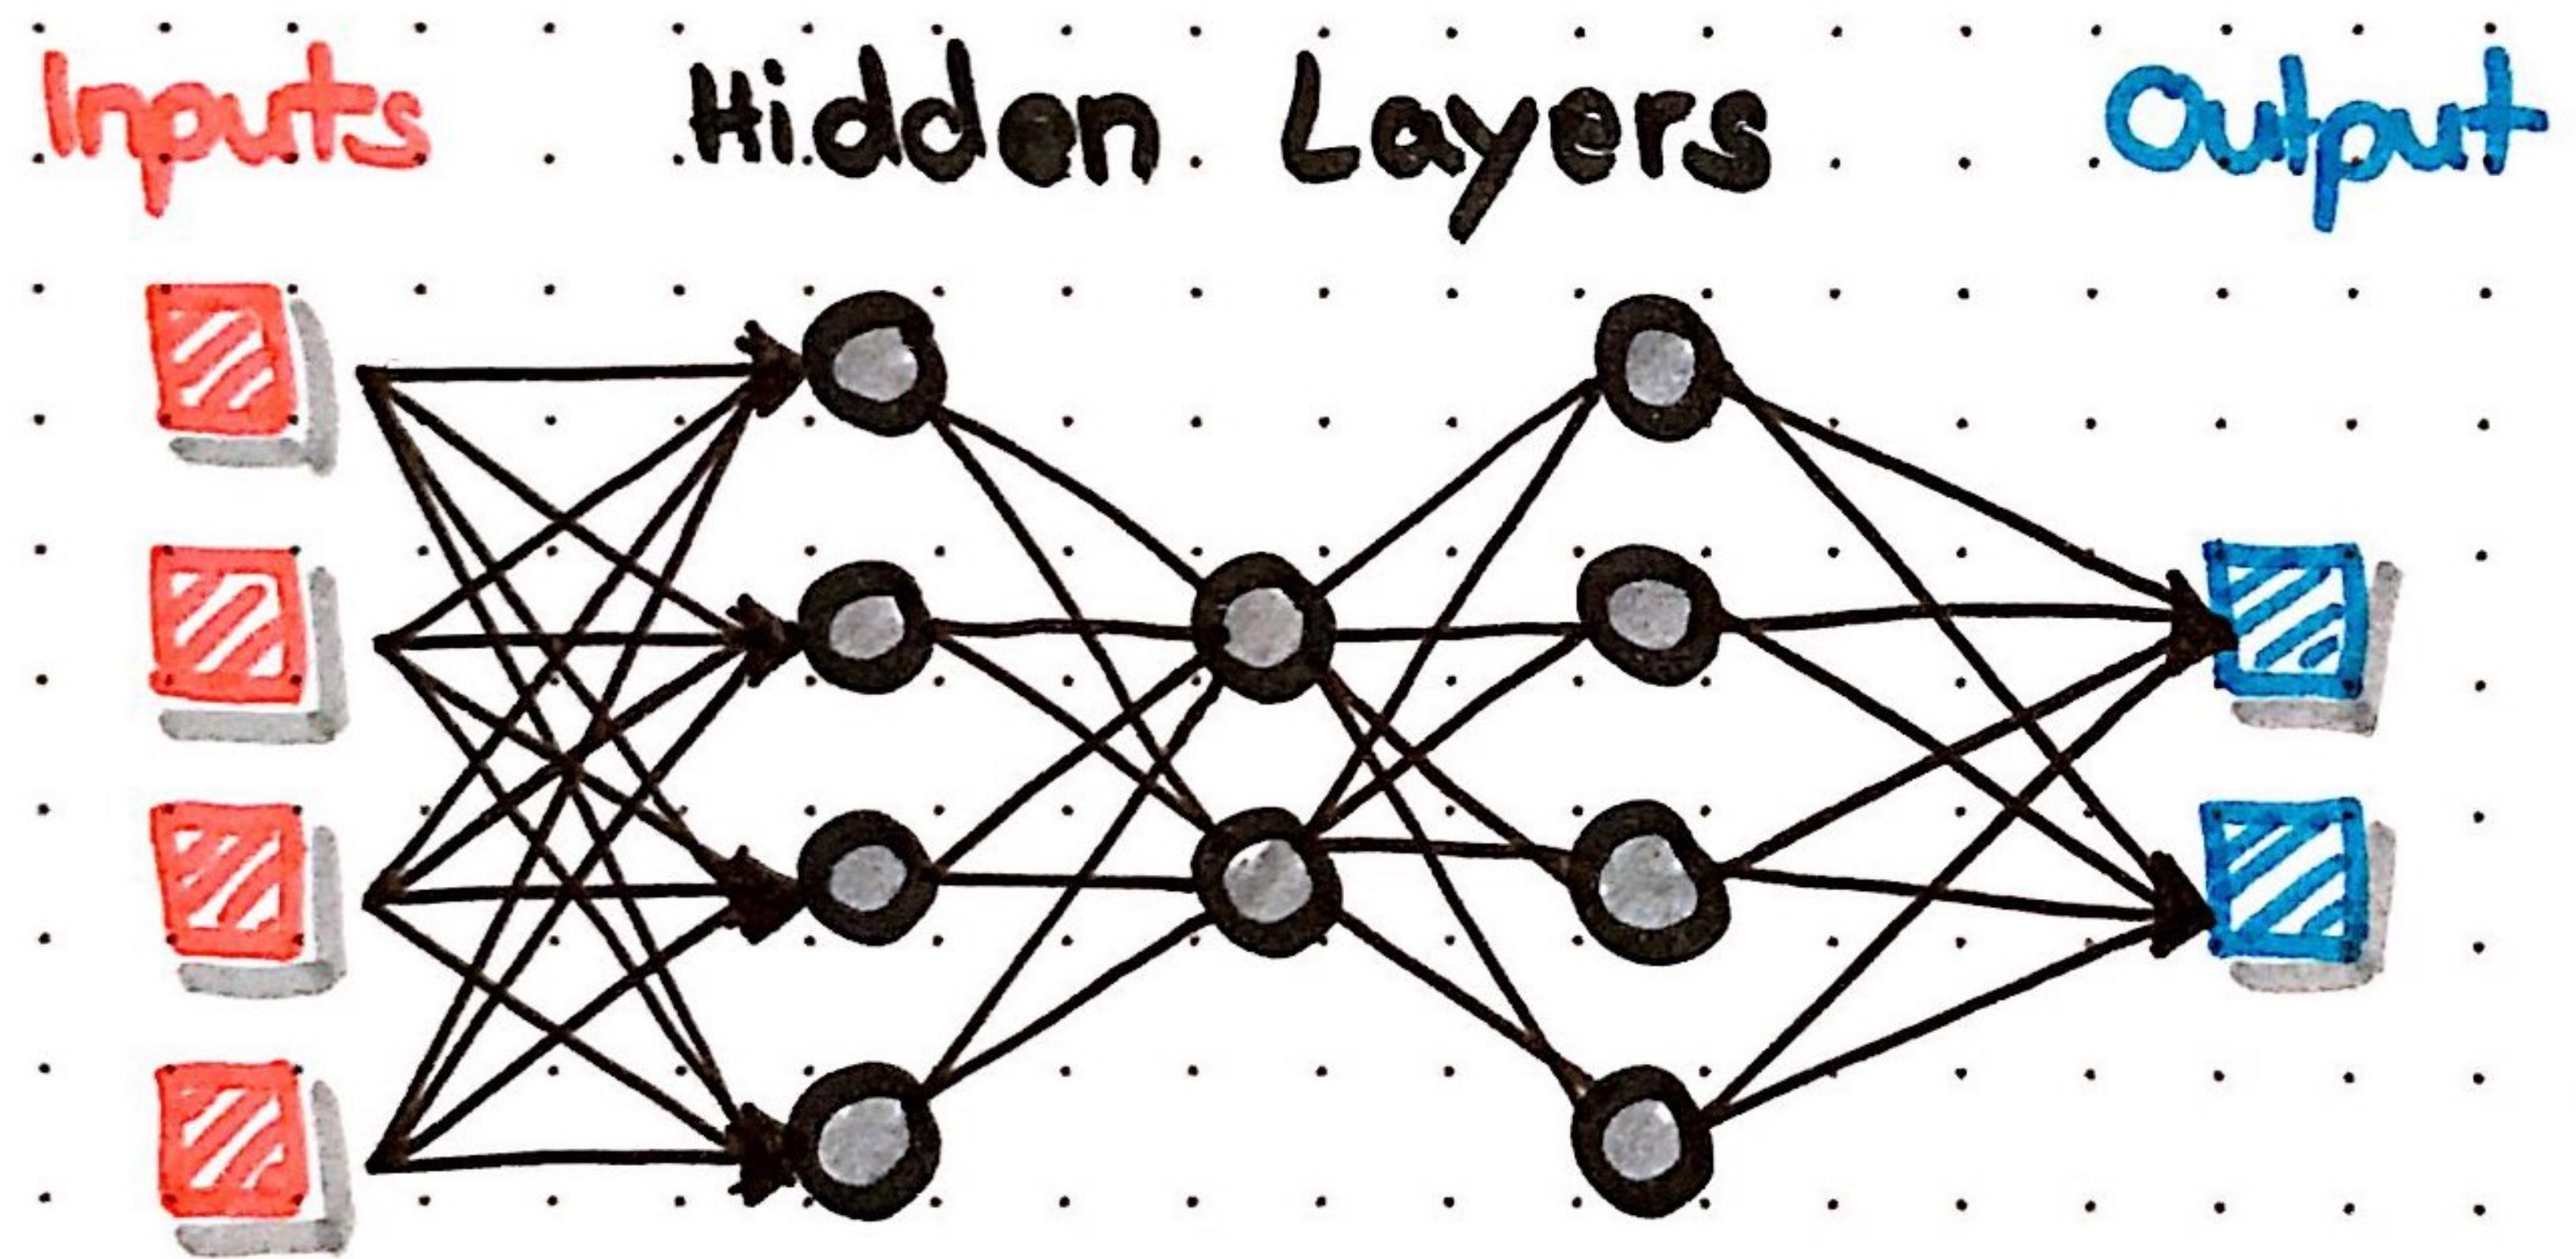


## Activation functions make non-linearity possible

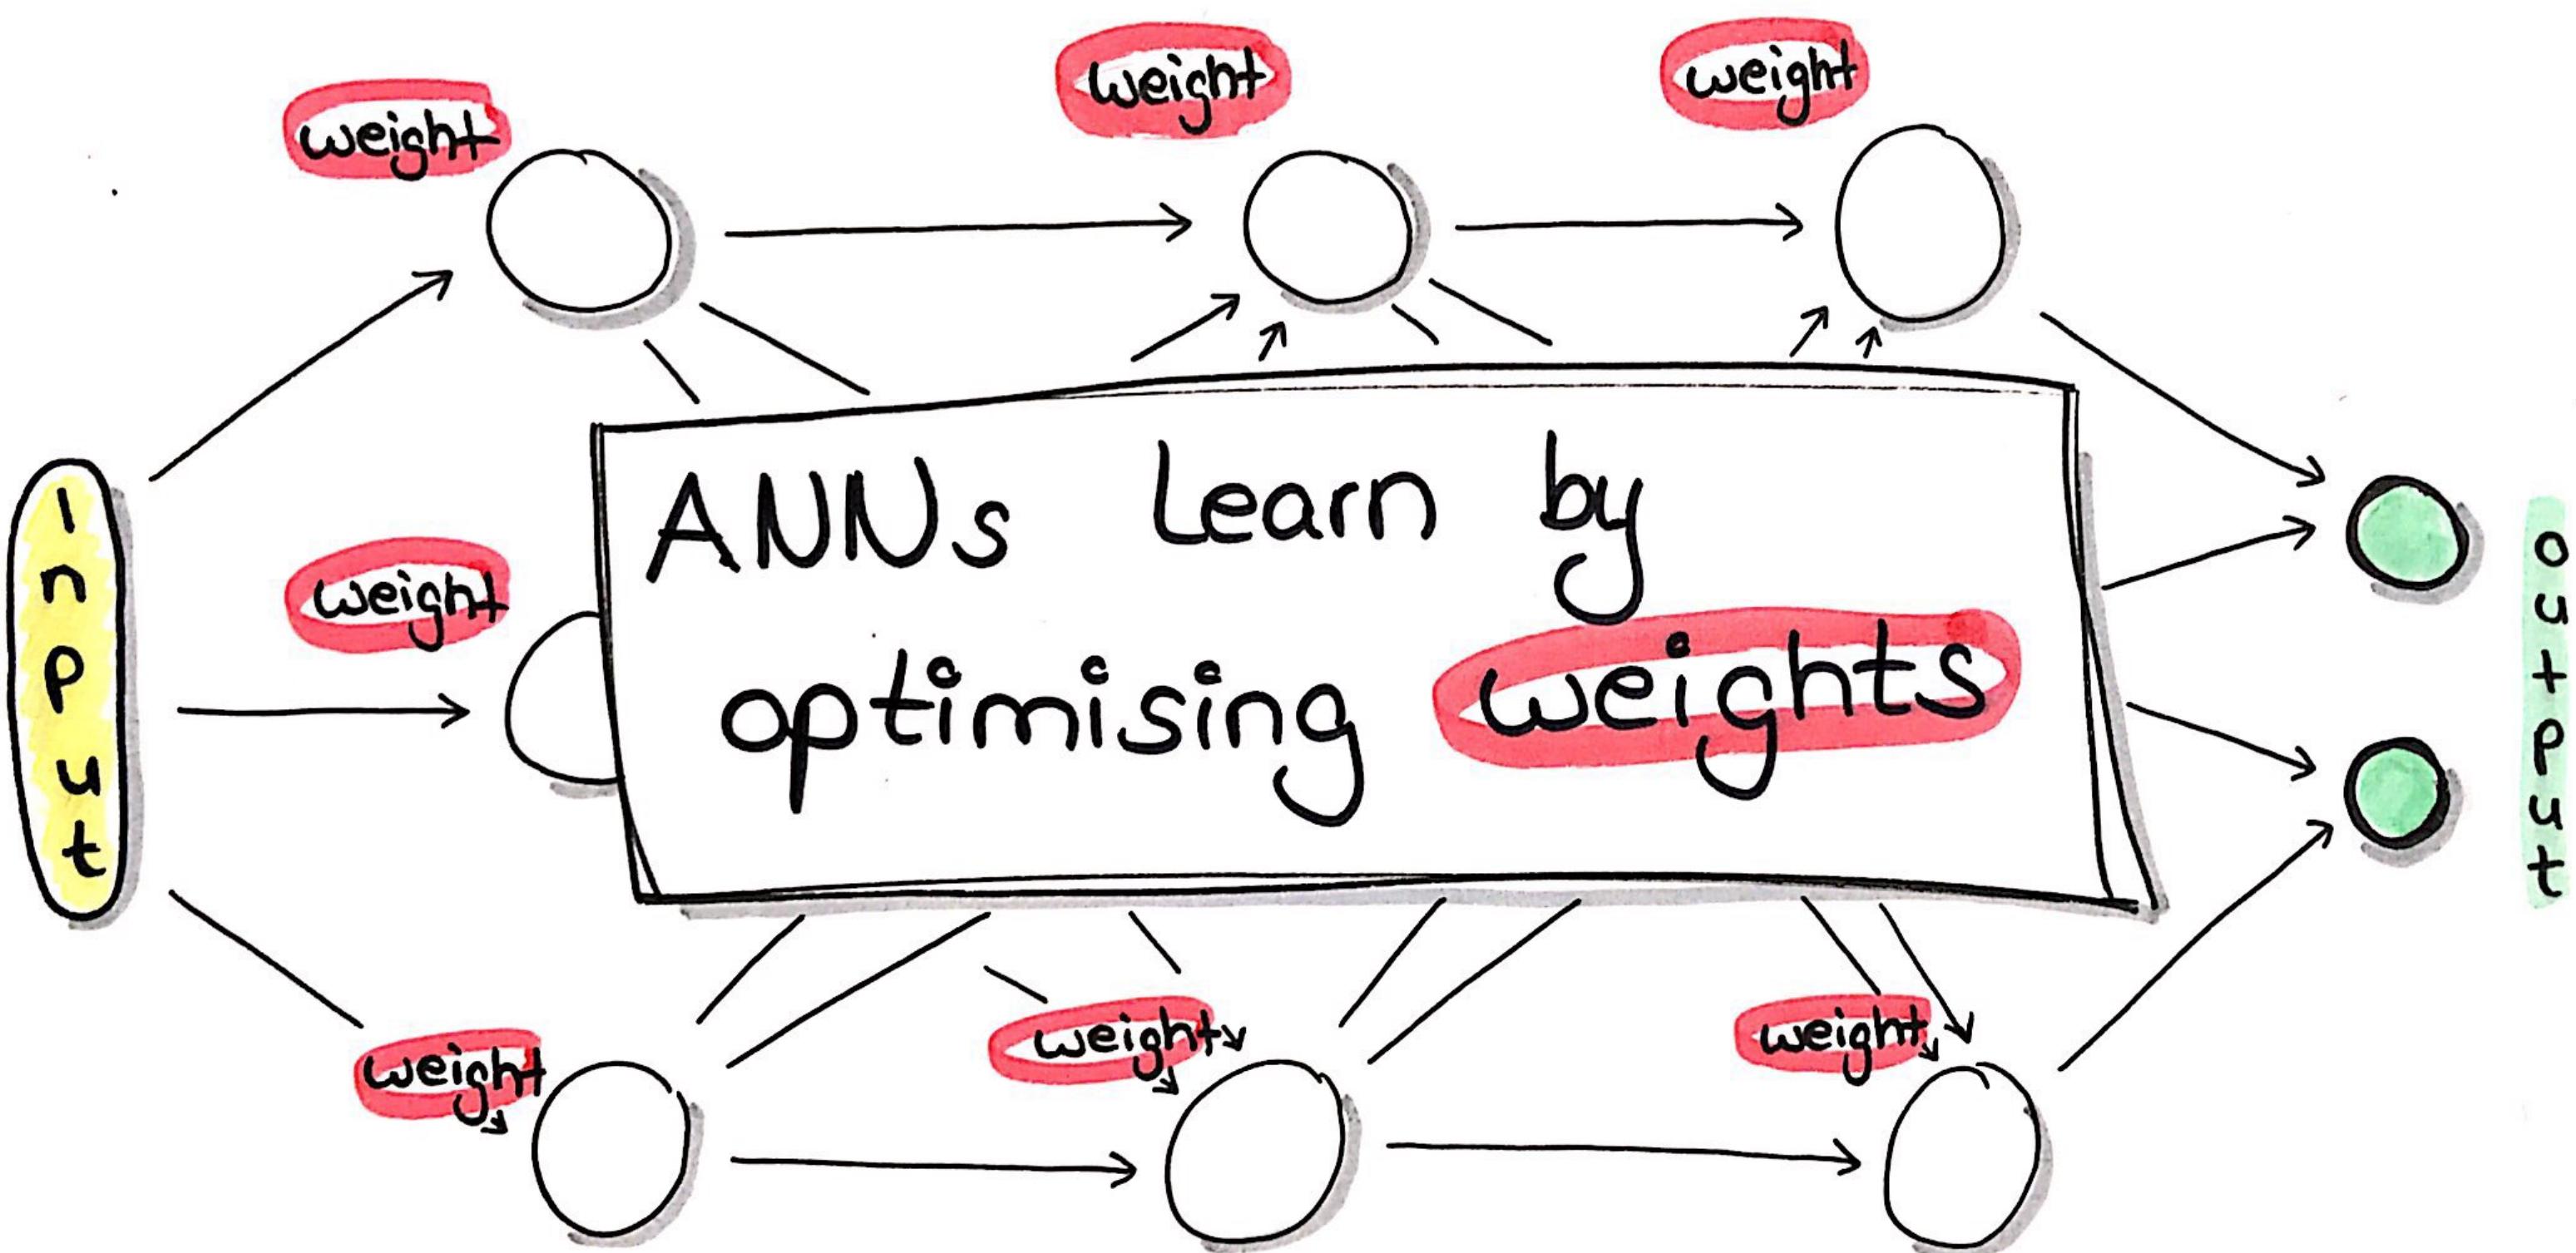
- non-linear activation functions allow us to approximate ANY mathematical formula with neural nets



# Multi-Layer Perceptrons



## How does a neural net learn?



# How does a neural net learn?

## The Softmax function

$$x * \omega + b = y$$

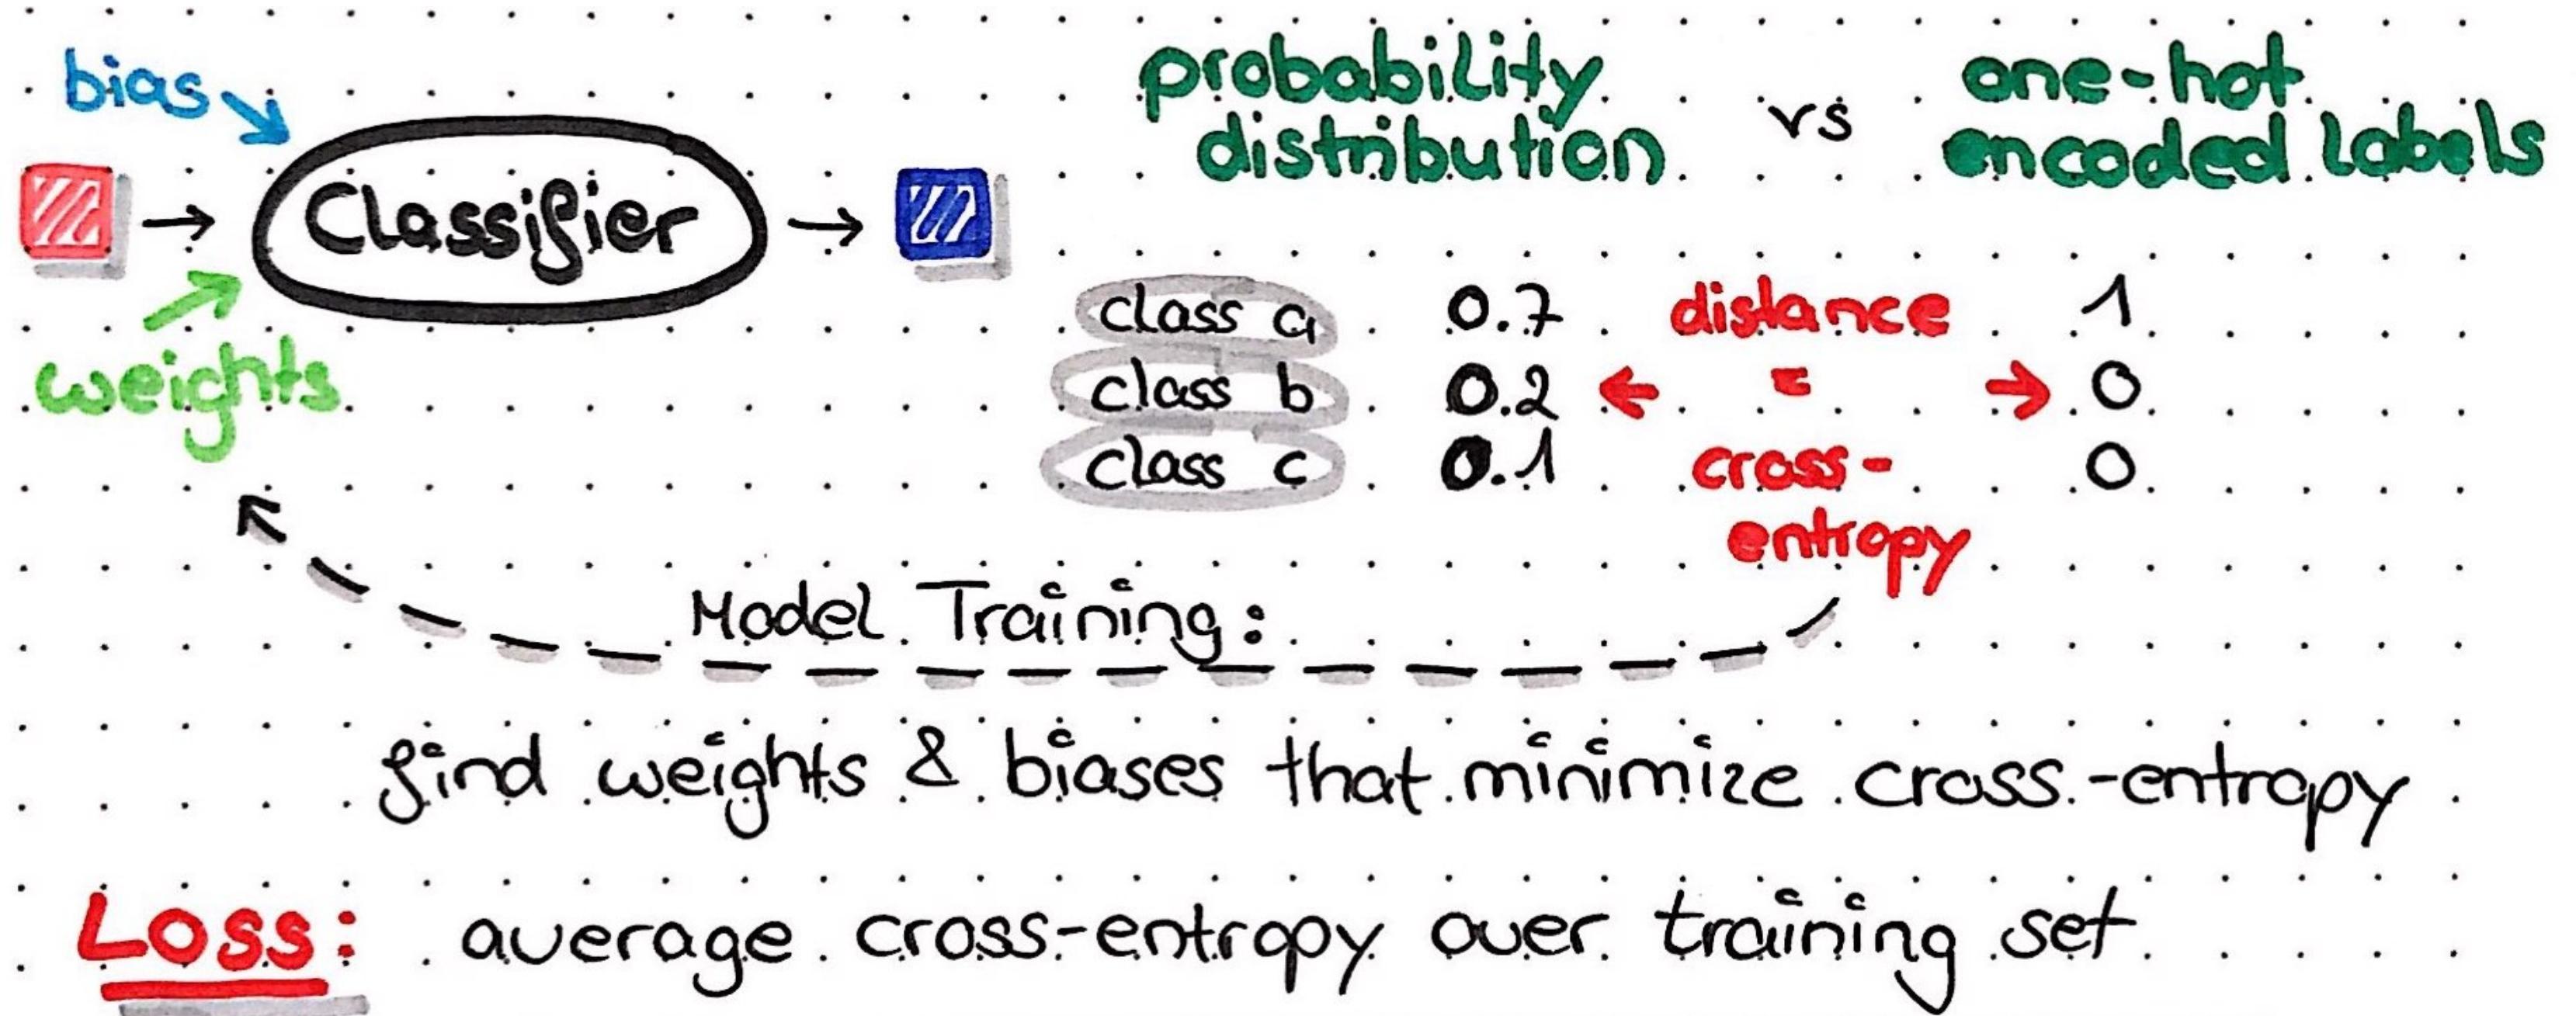
input weight bias      Output

Model Training ~ finding good weights & biases

	score	probability	
class a	2	→ 0.7	correct
class b	1	→ 0.2	
class c	0.1	→ 0.1	
			soft-max

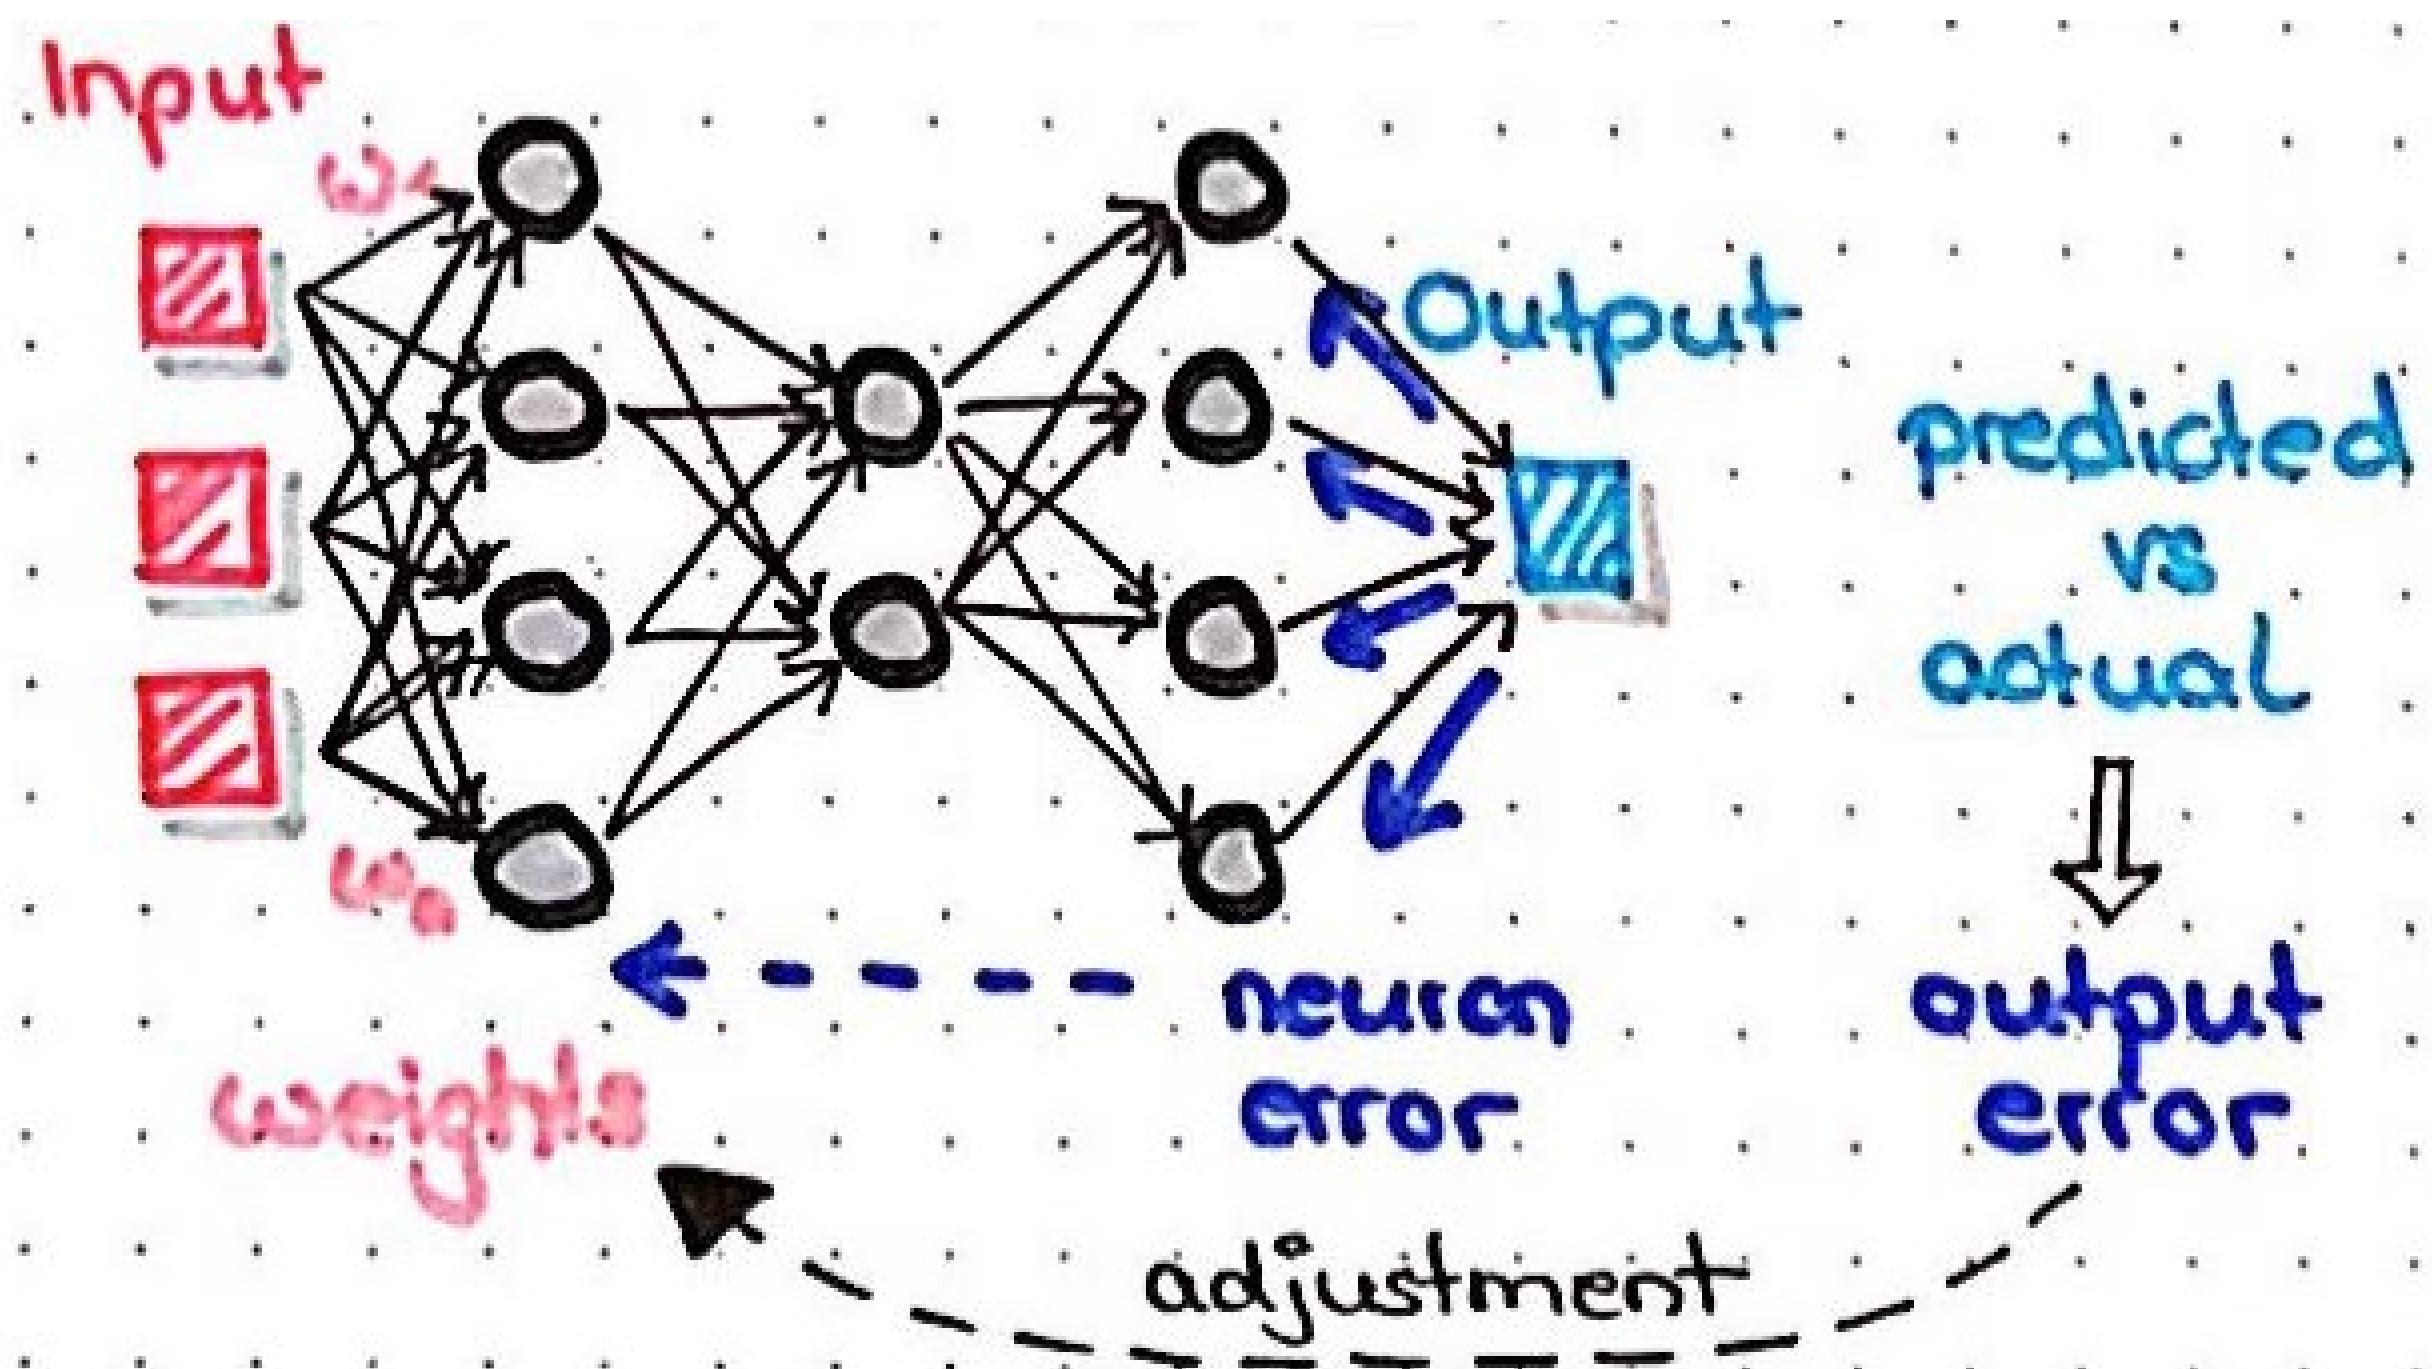
# How does a neural net learn?

## Cross-entropy



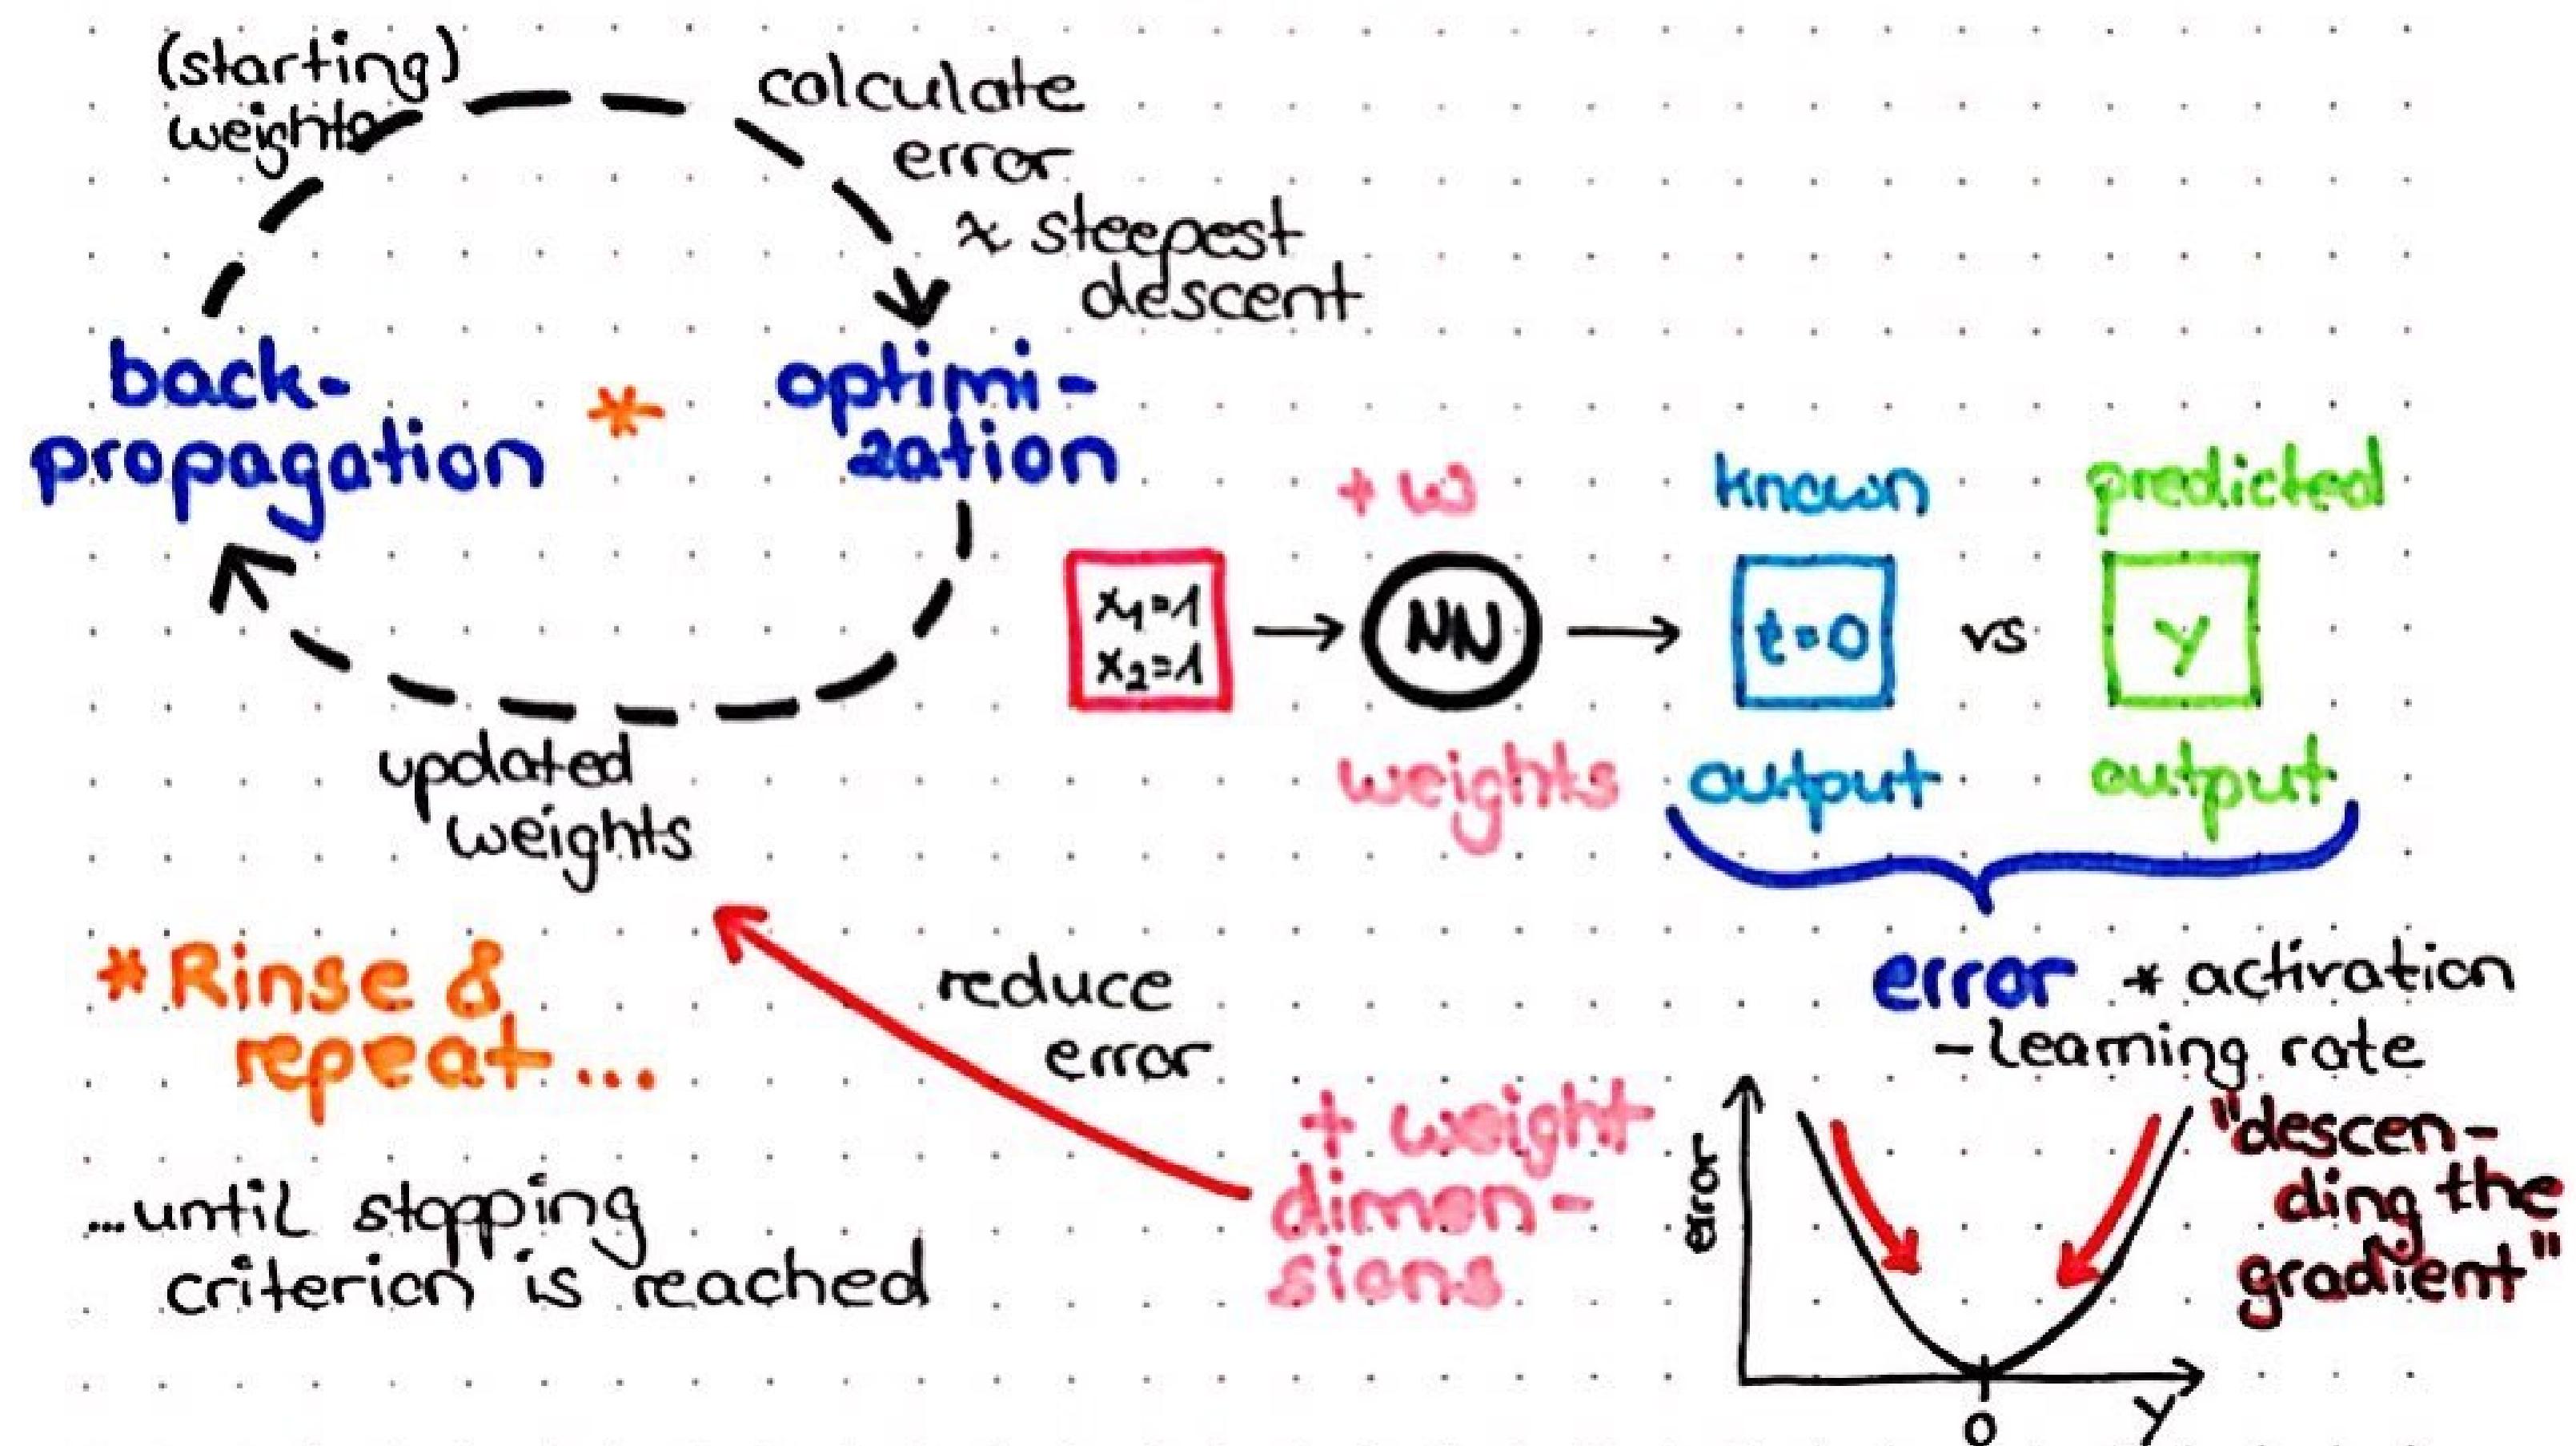
# How does a neural net learn?

## Backpropagation

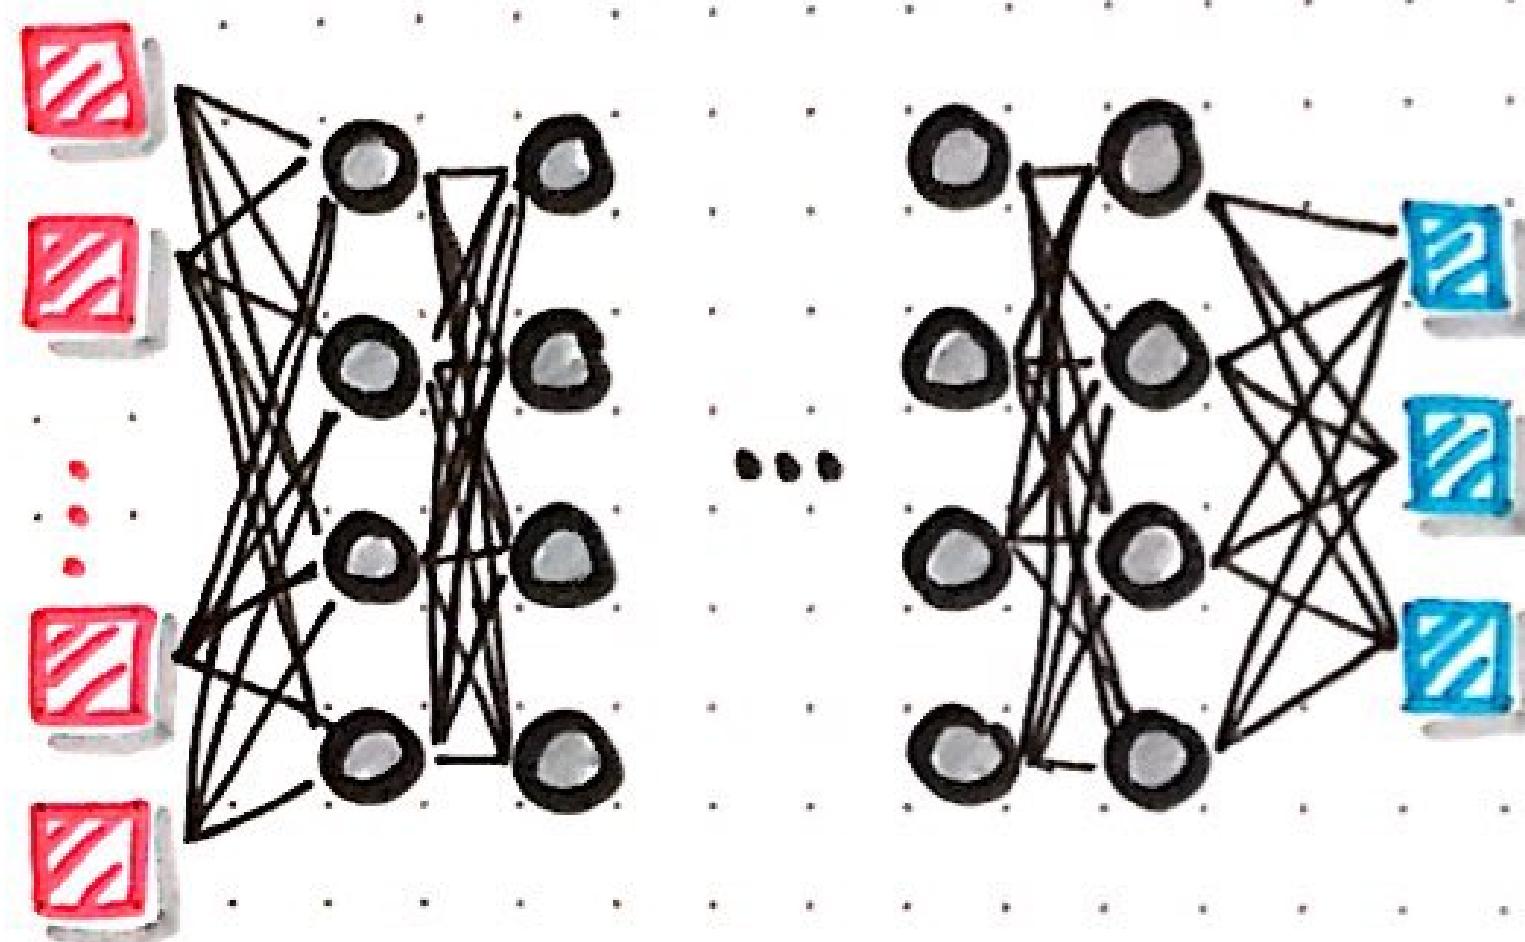


# How does a neural net learn?

## Gradient descent optimization



# Deep Learning



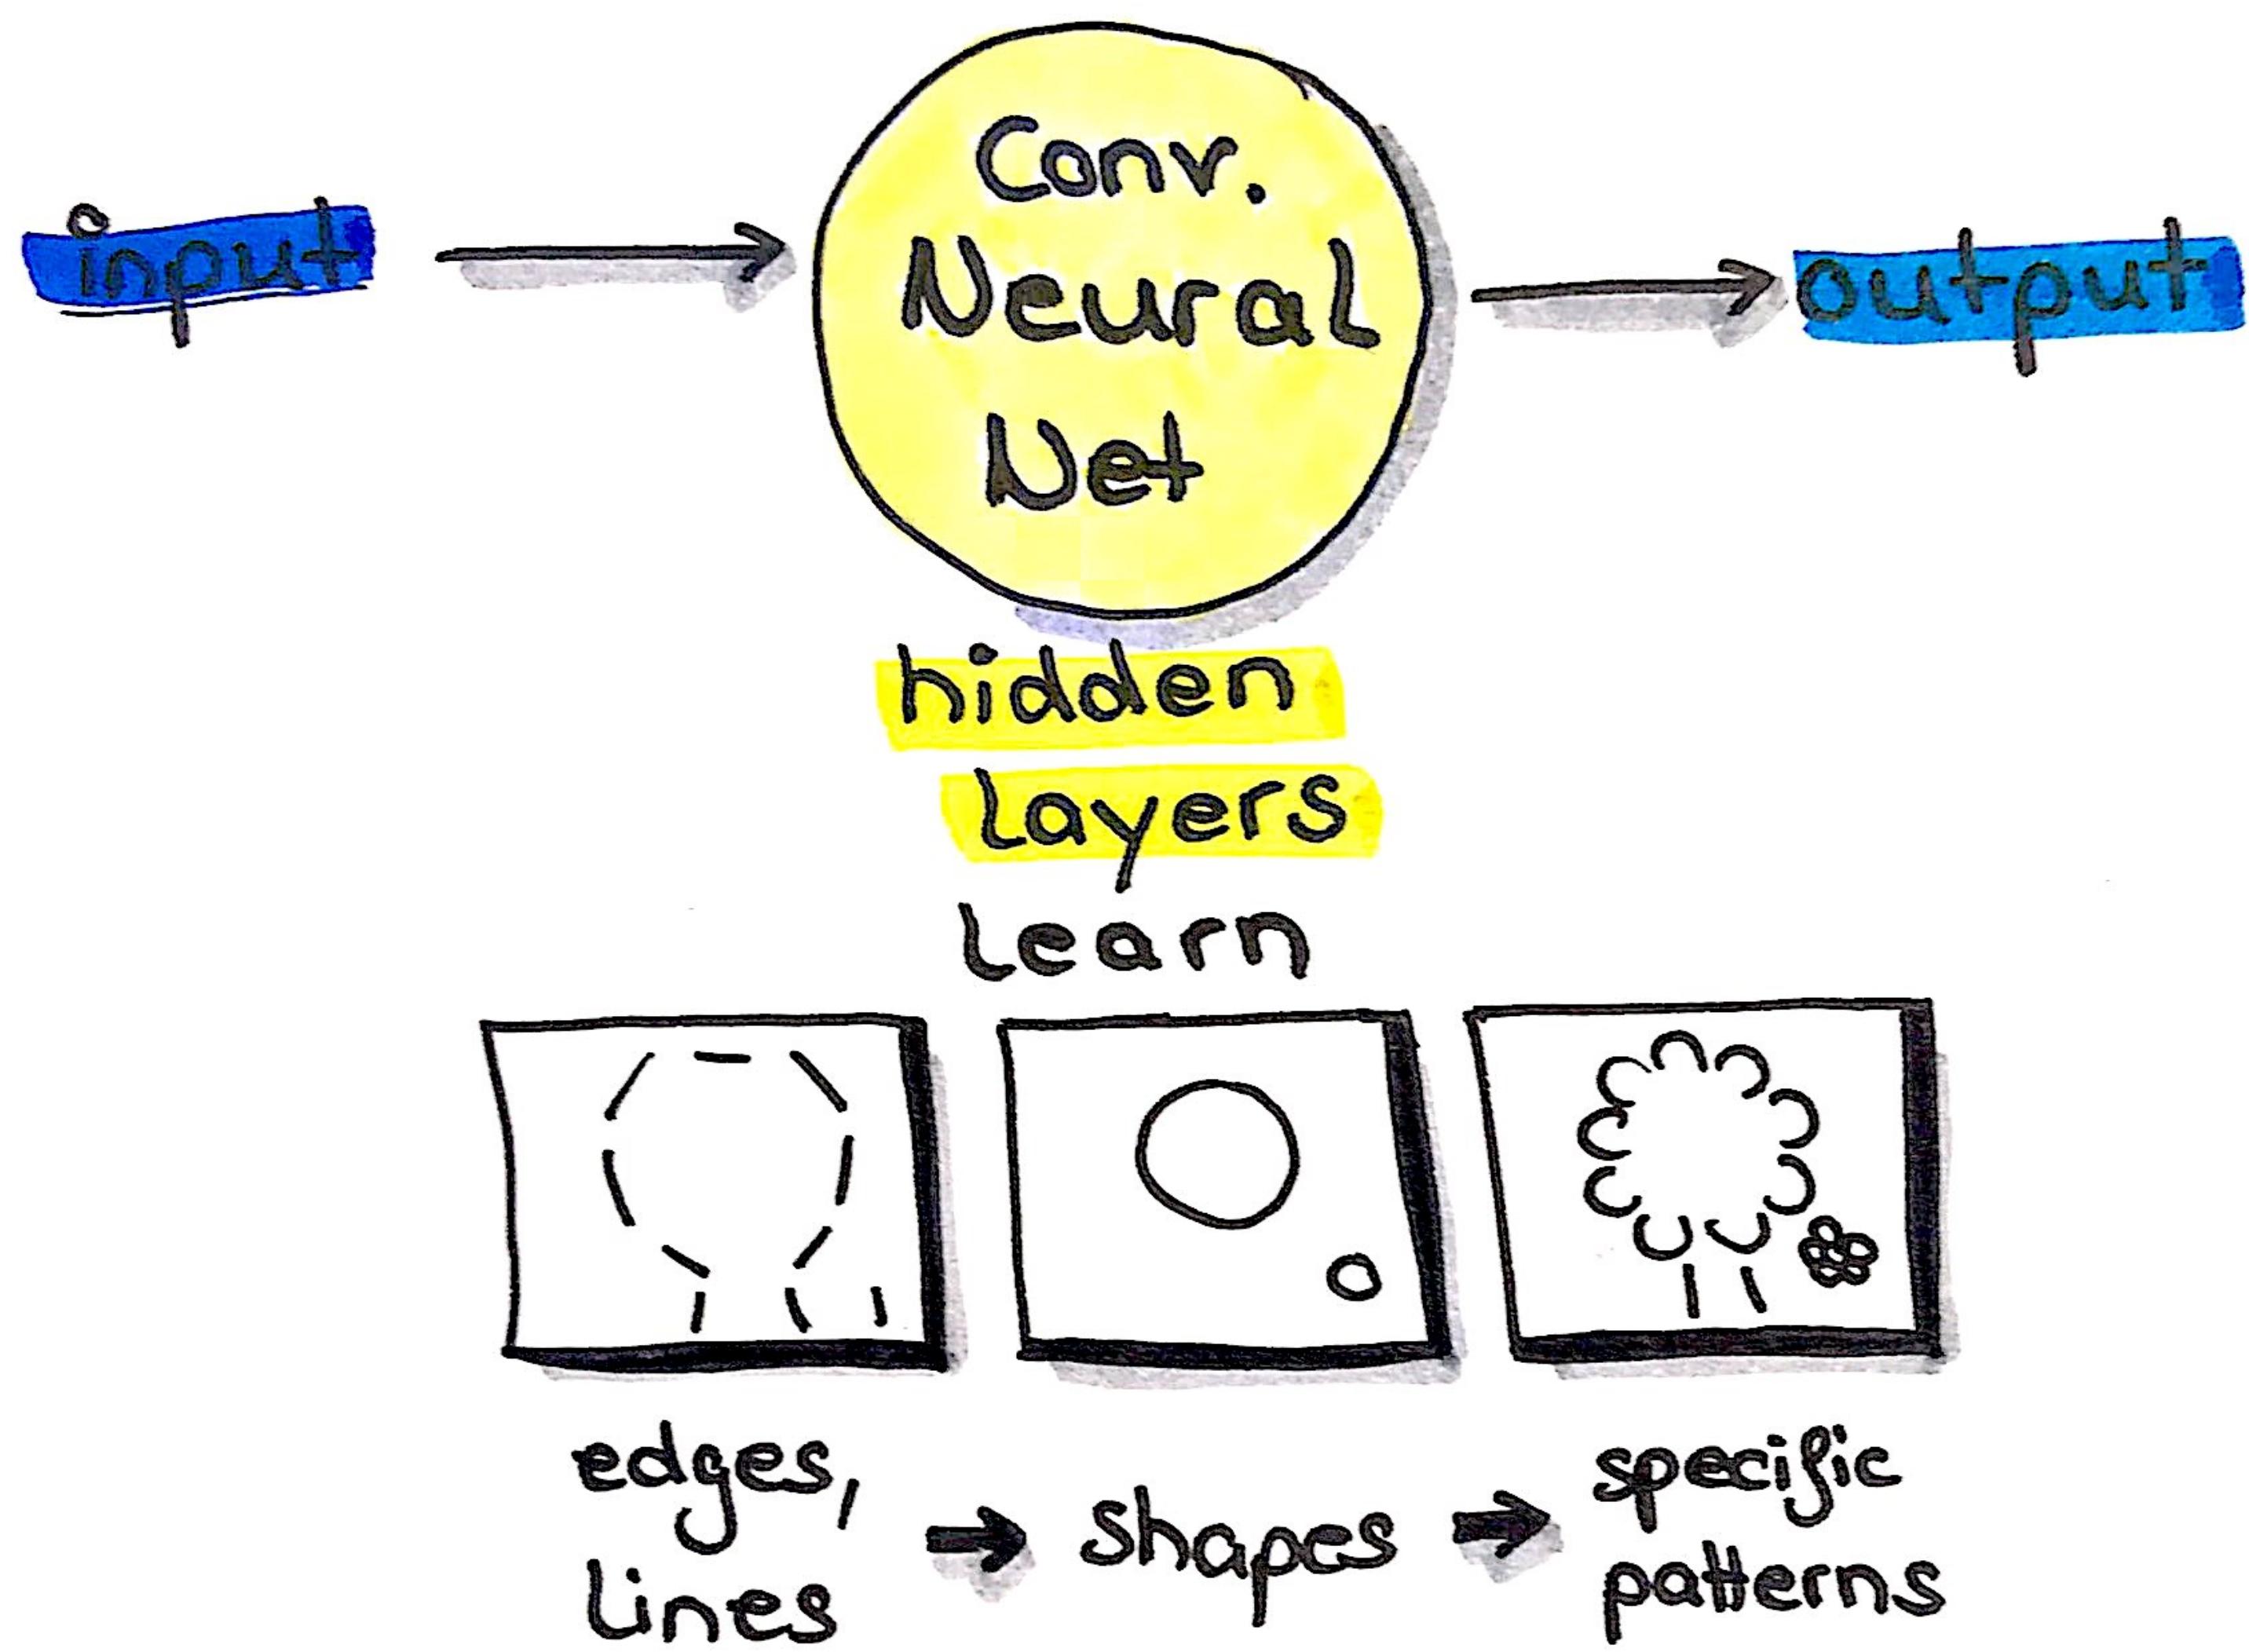
- supervised , semi - supervised or unsupervised
- NLP , speech recognition
- image recognition
- object classification
- recommender systems
- etc.

# Convolutional Neural Nets

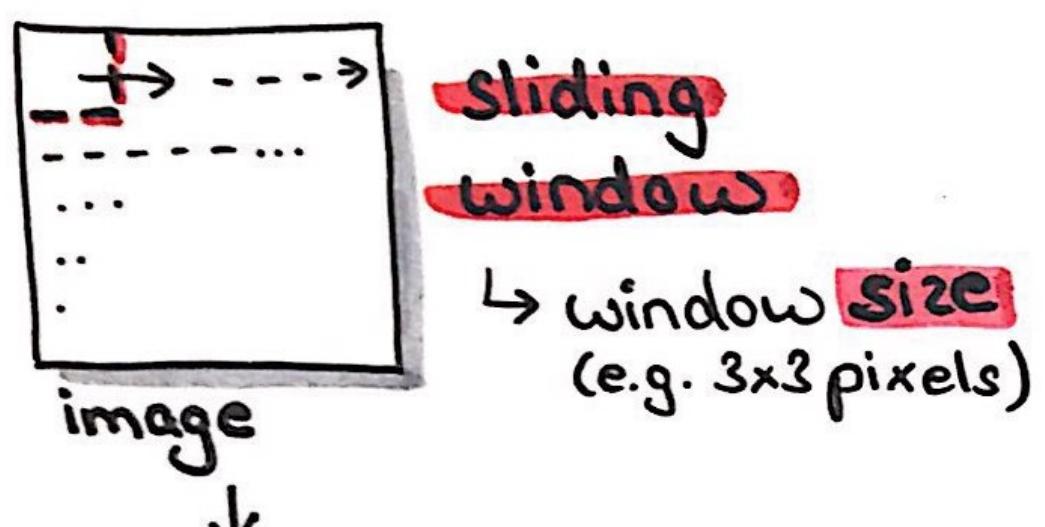
## MLPs vs CNNs

- pixels are considered independent
- computationally faster
- Learned : weights
- pixels are considered as groups of connected information (context)
- analysed as chunks (= windows)
- Learned : filters

# ConvNets



# ConvNets

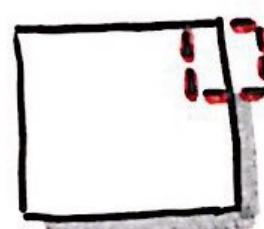


vertical Lines

horizontal Lines

- **Filters**
  - detect shapes & patterns in window chunks
  - multiple filters are combined
  - filters are learned

**padding**



- fake pixels are created at the edge of images to incorporate border pixels

**Convolutional Layers**

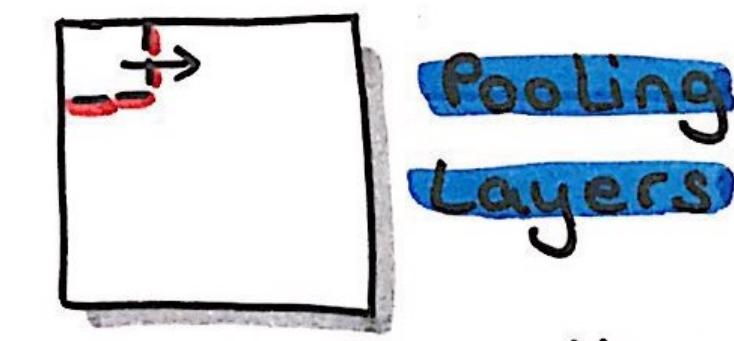
apply filters

↓

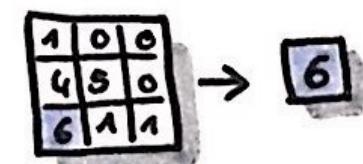
output

=  
**feature maps**

→  
**stacks of feature maps**



e.g. max pooling

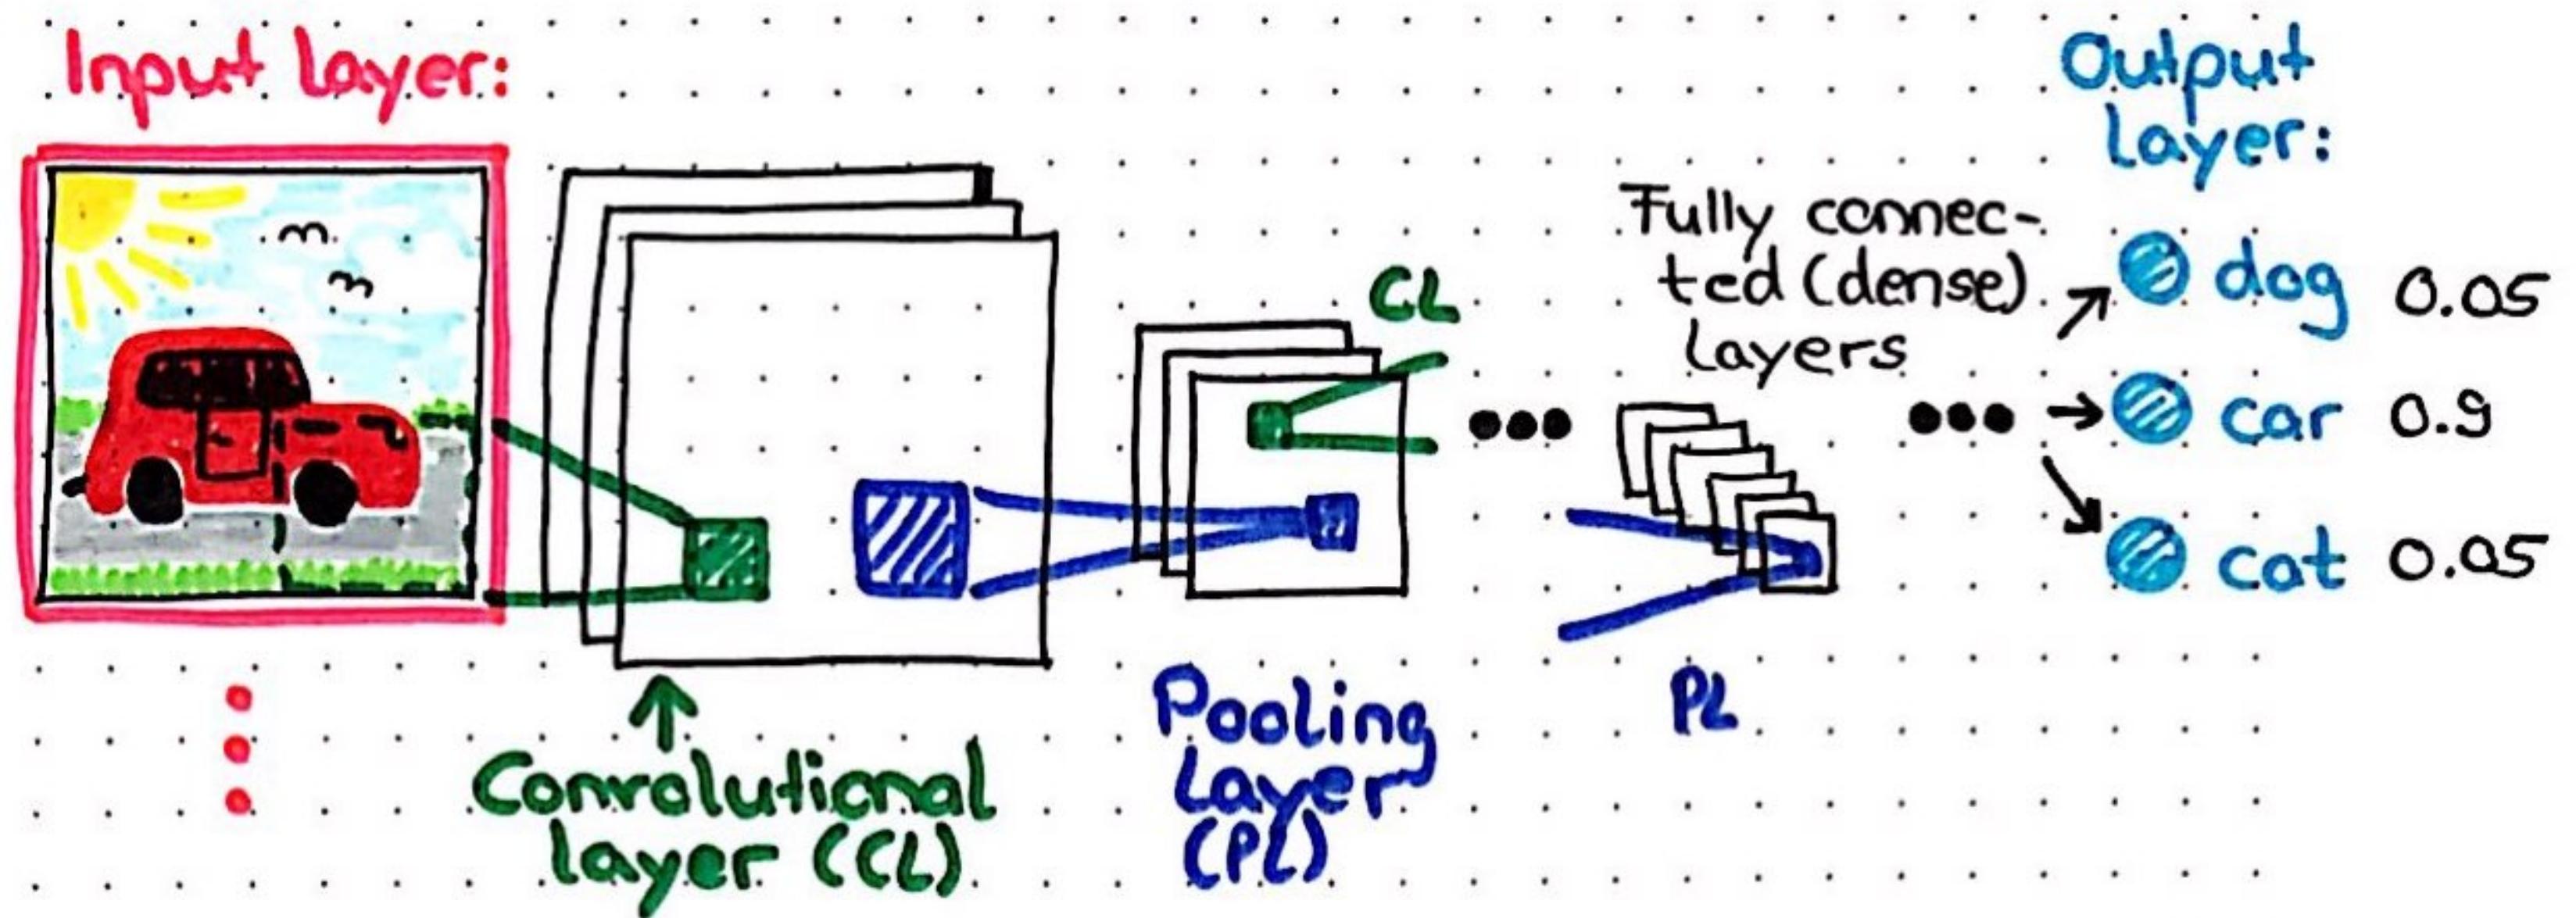


- reduces compute time
- boils information down

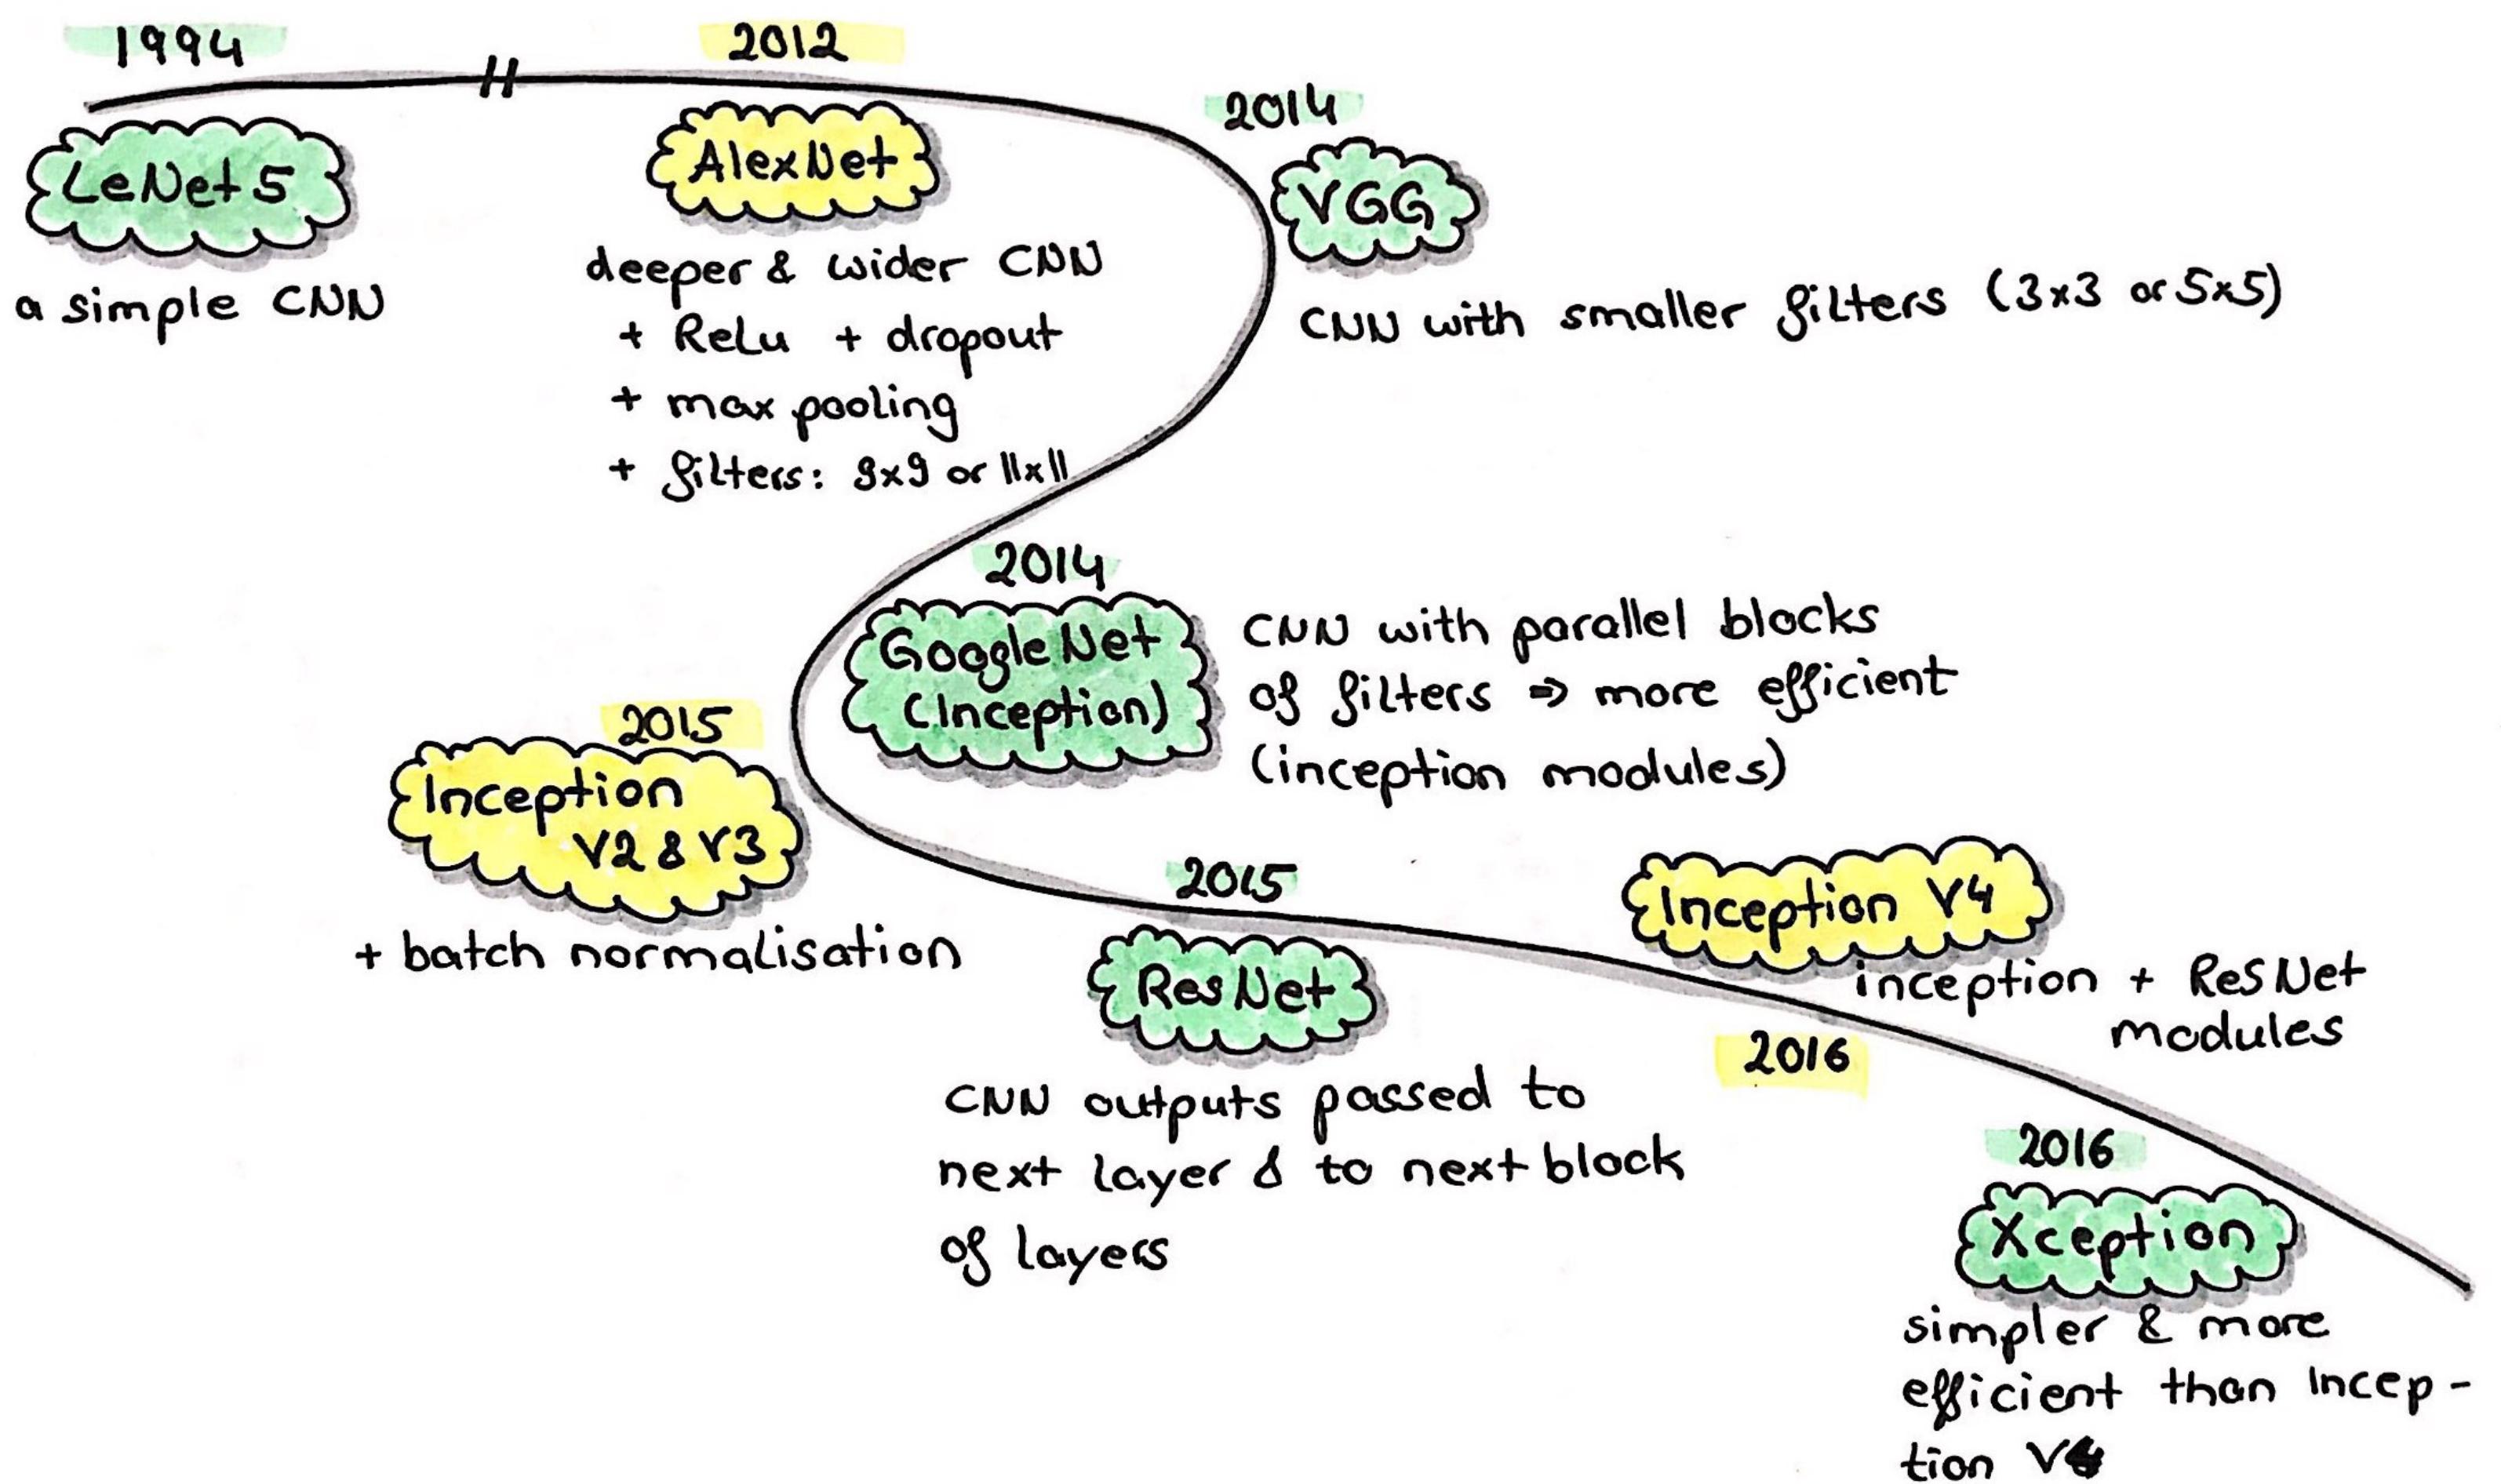
**stride**

- how much overlap to have in sliding window

# ConvNets

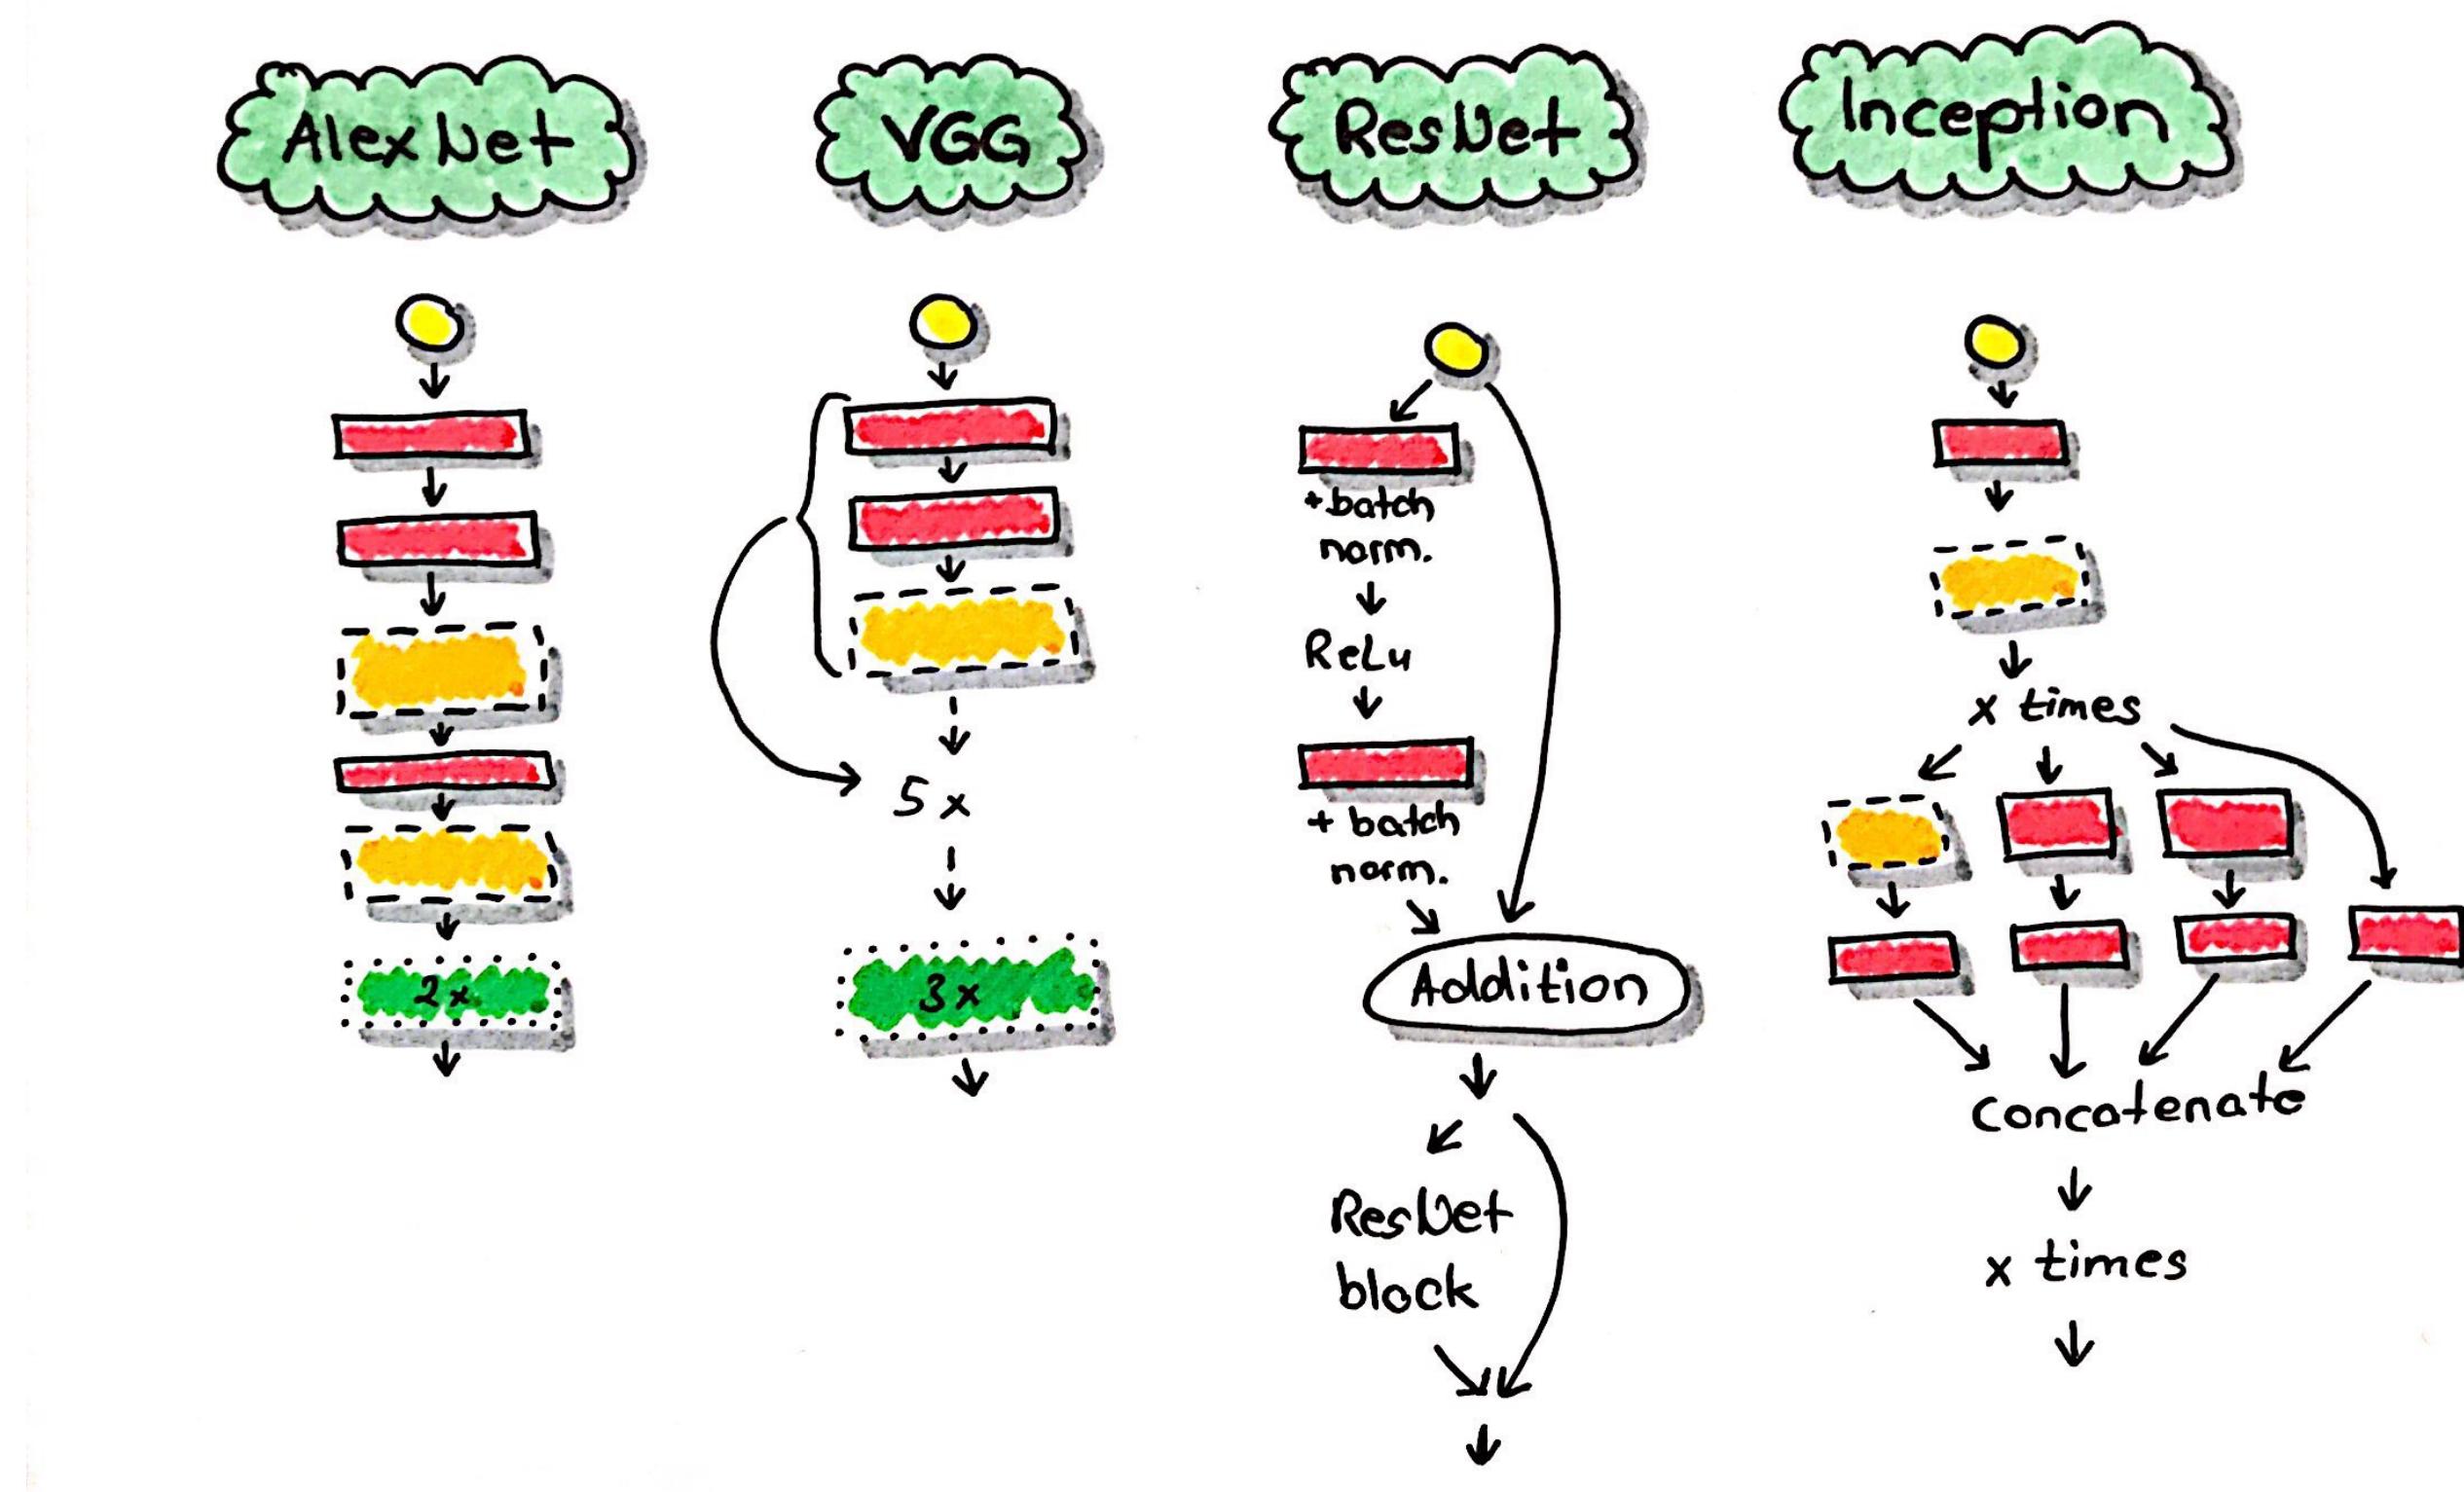


# Evolution of neural nets for image recognition

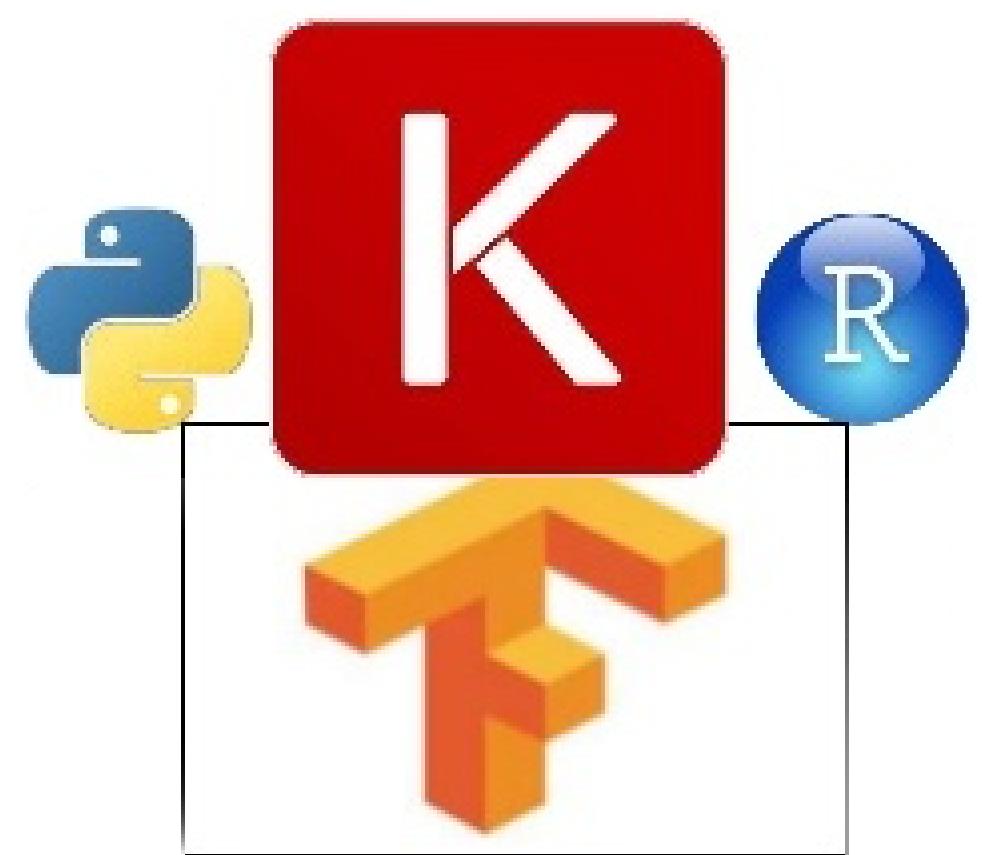


# CNN architectures

- input image
- convolutional layer
- pooling layer
- dense layer



# Introduction to TensorFlow



# What are tensors?



# Tensors = multidimensional arrays

Dimension

1

1	2	3
---	---	---

vector

2

1	2	3
4	5	6
7	8	9

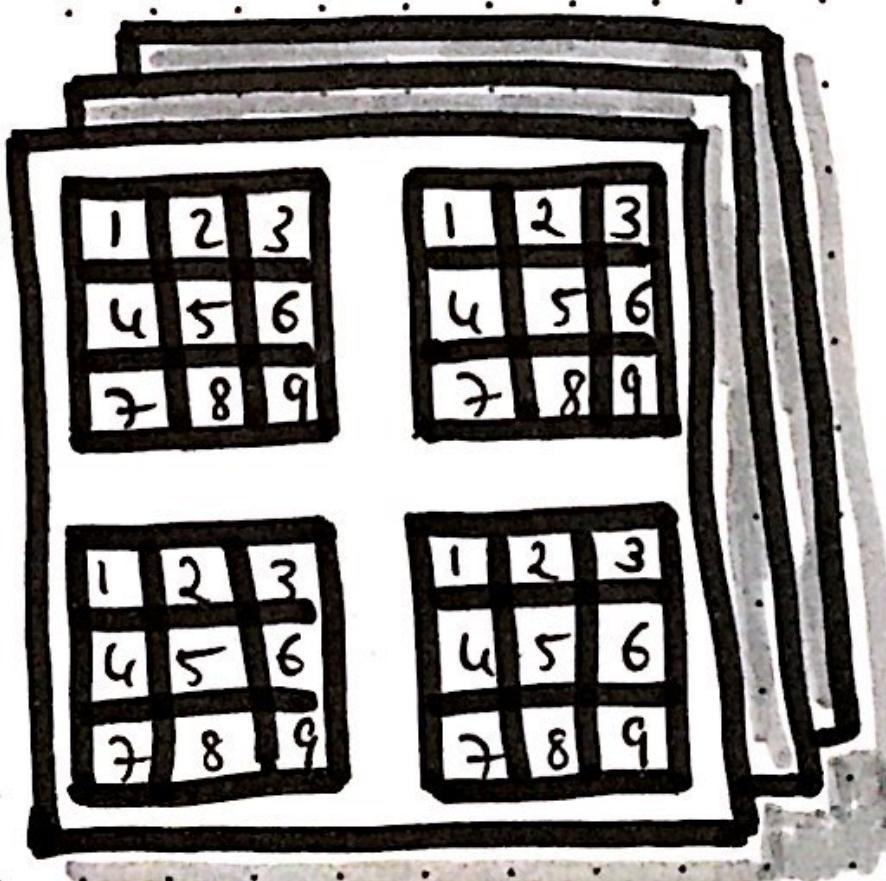
matrix

3

1	2	3
4	5	6
7	8	9

3-dim.  
array

n



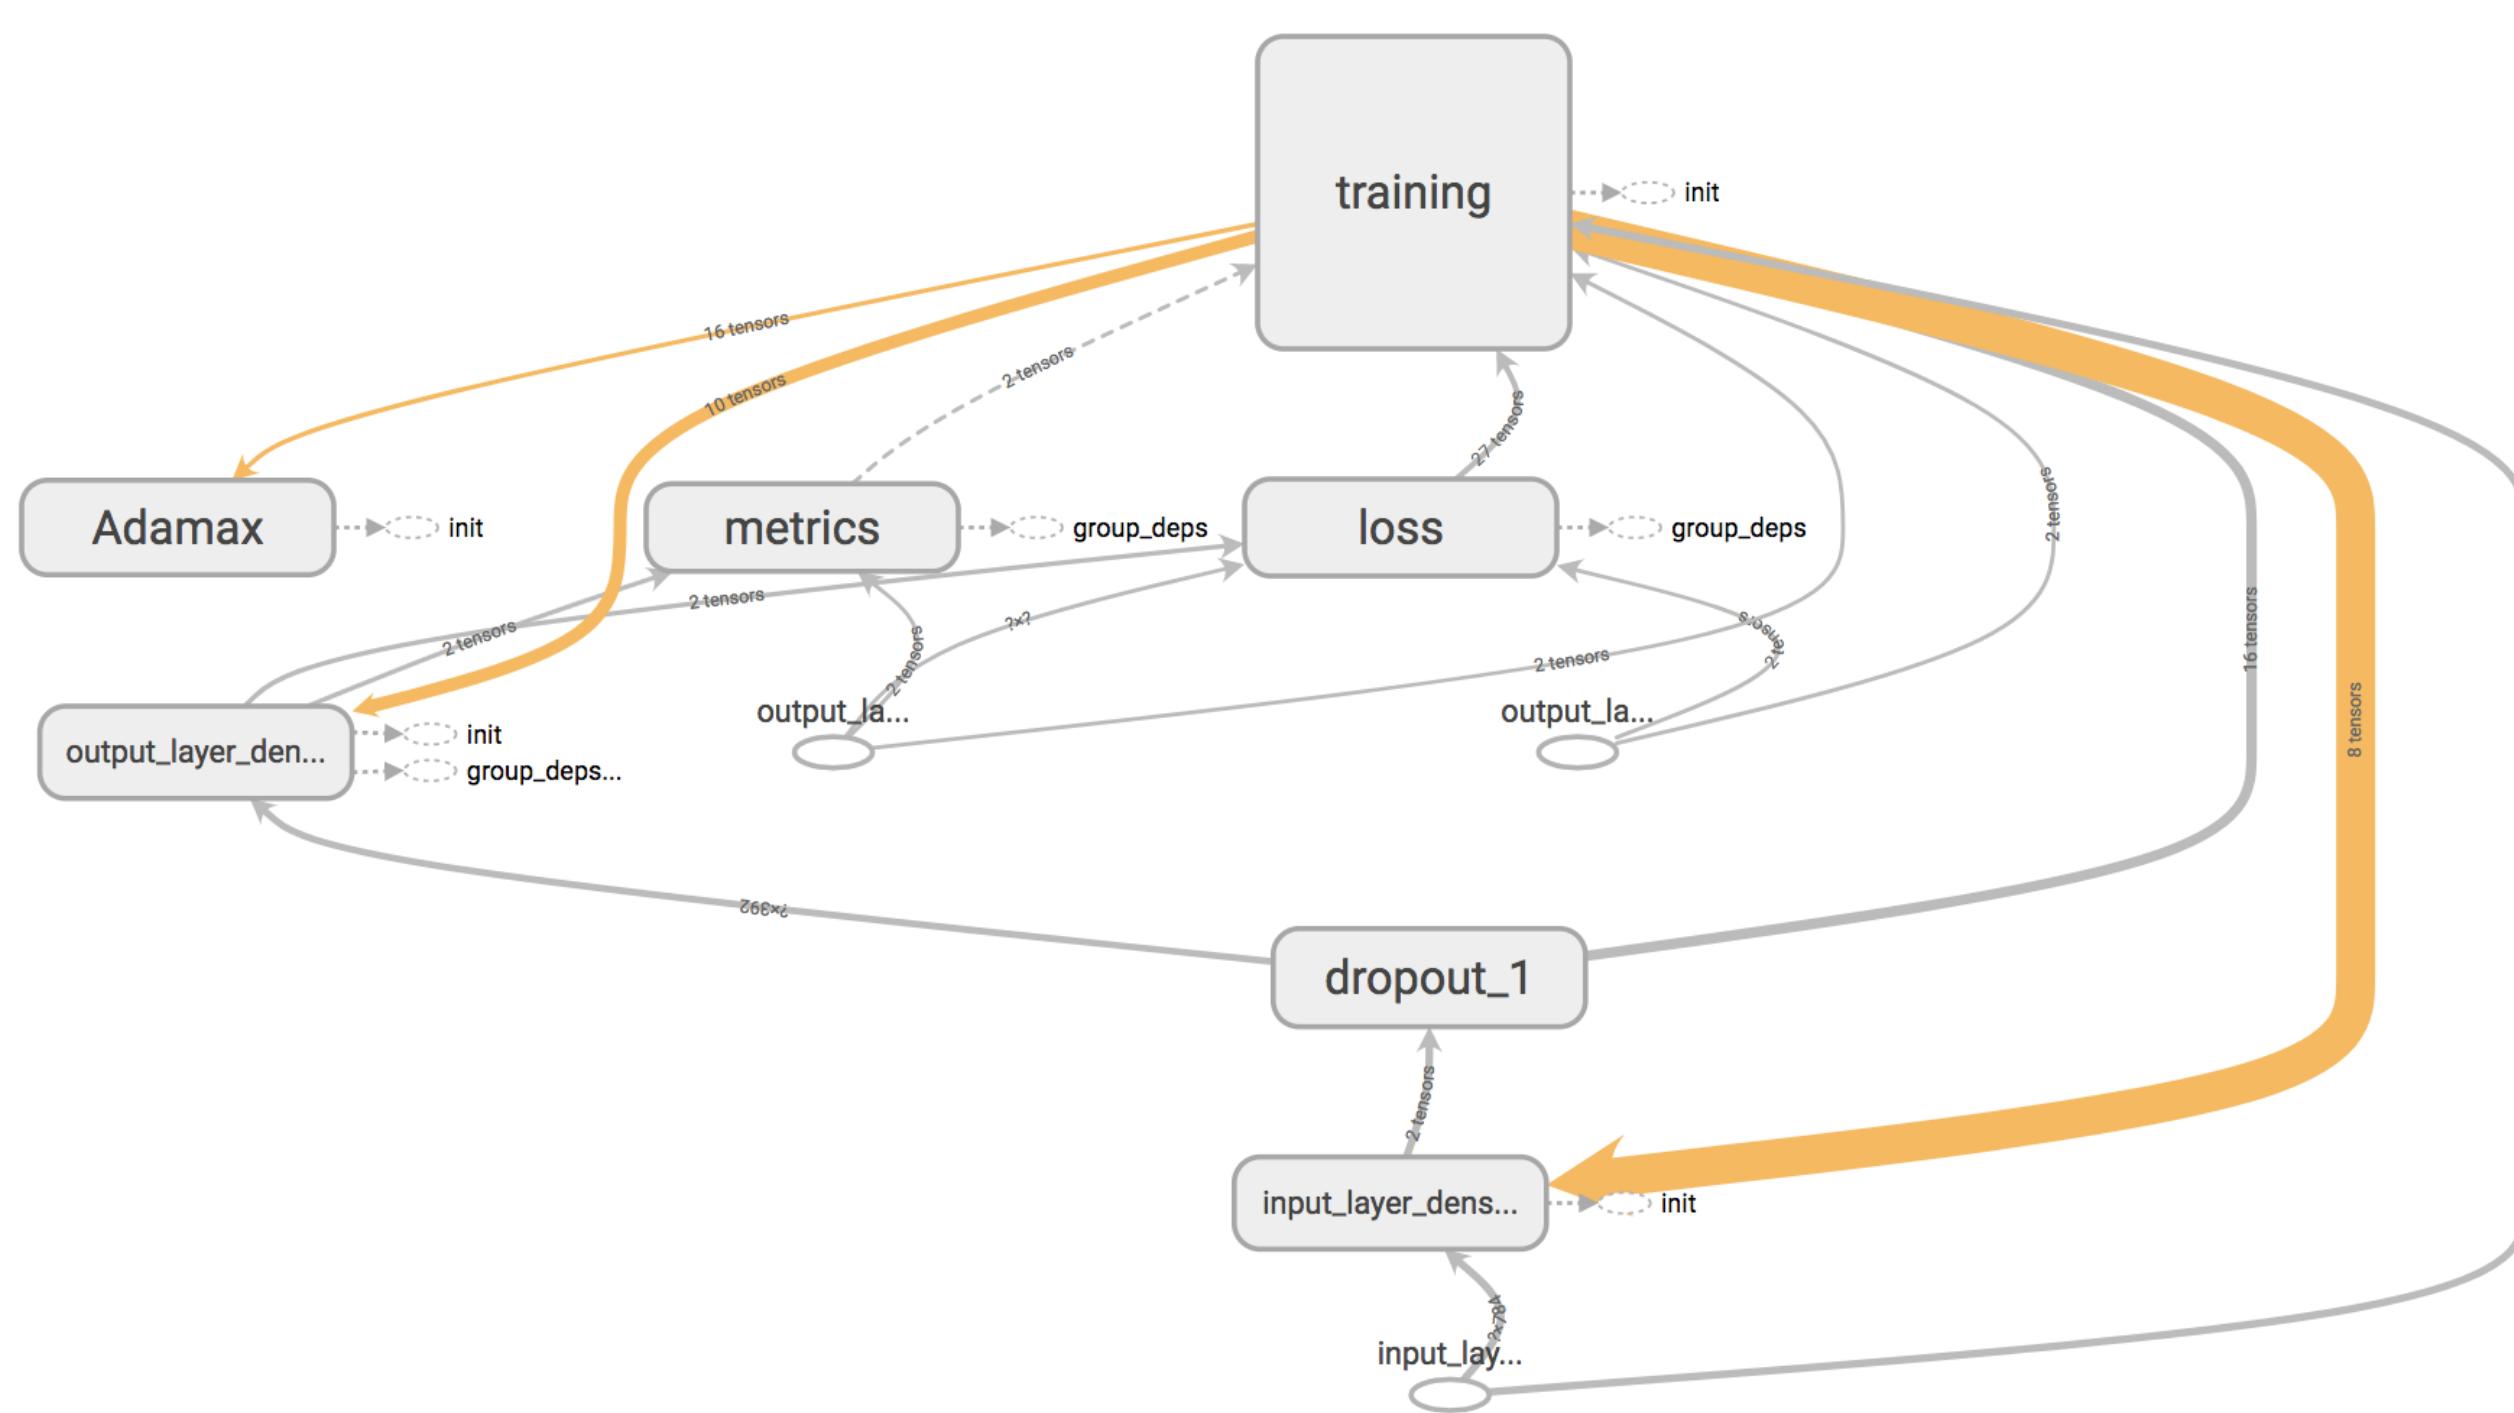
n-dim.  
array

## Tensor "Flow"

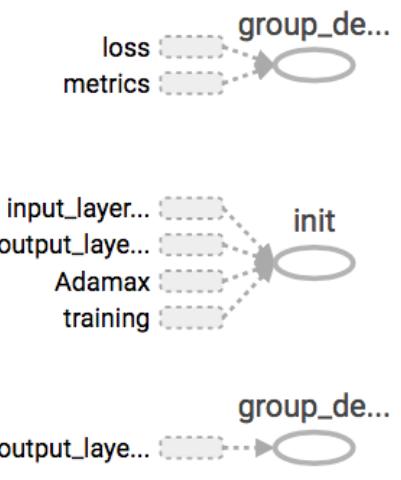


# Graphs in TensorBoard

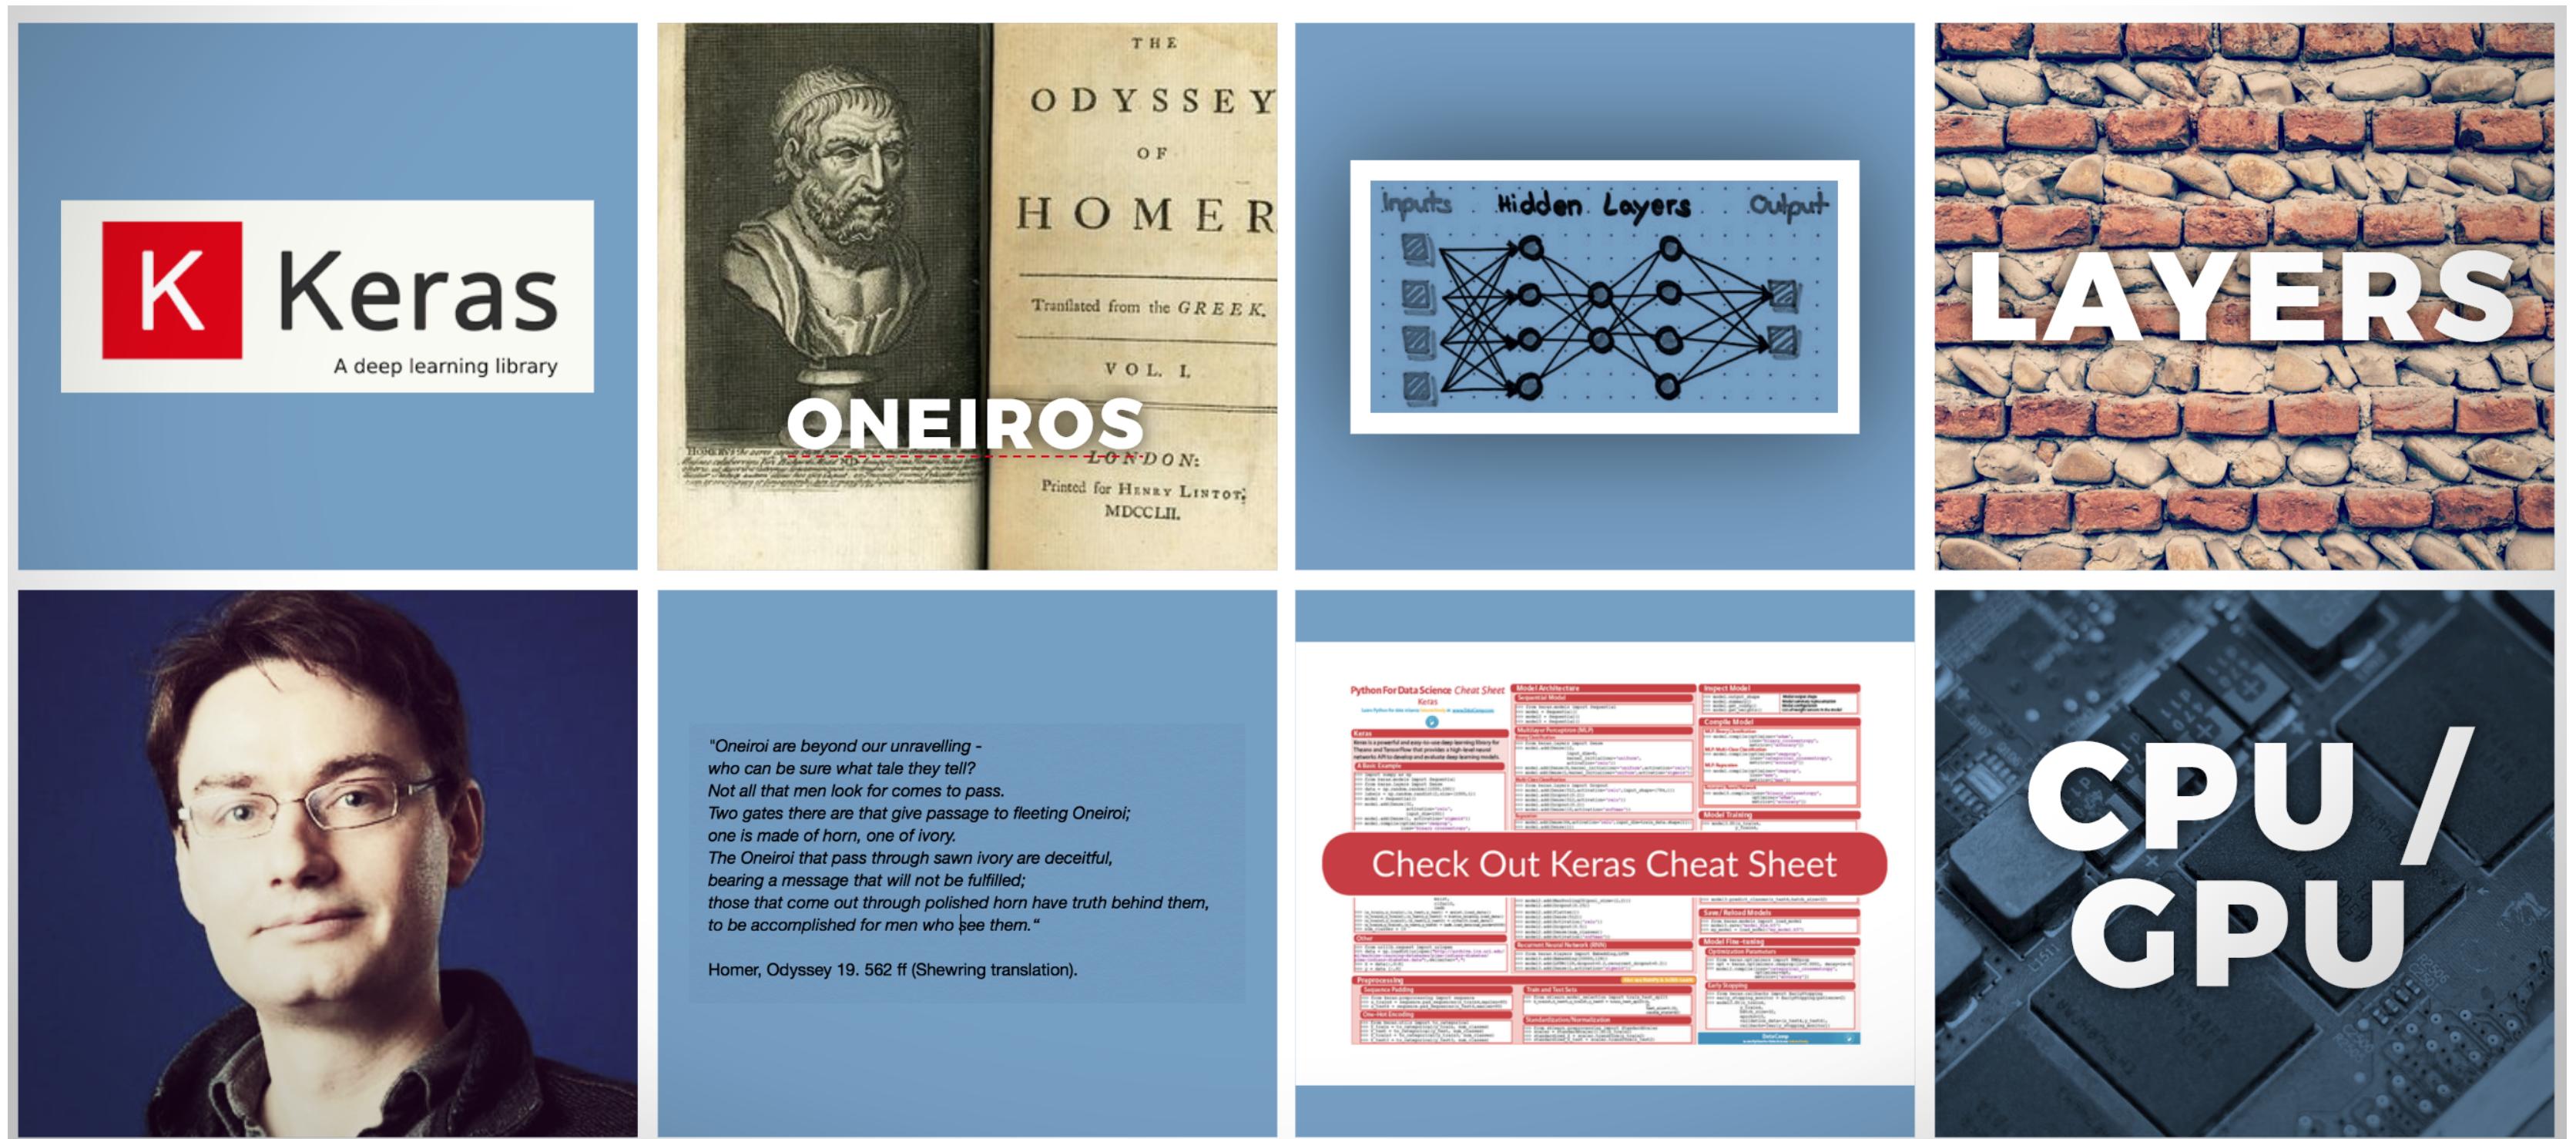
Main Graph



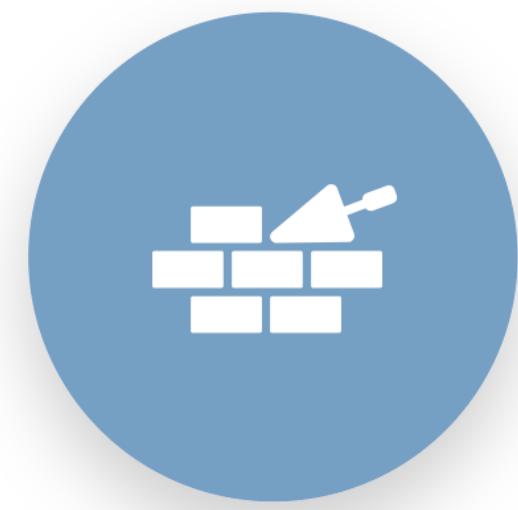
Auxiliary Nodes



# Keras High-Level API for TensorFlow



# Keras APIs



## SEQUENTIAL MODELS

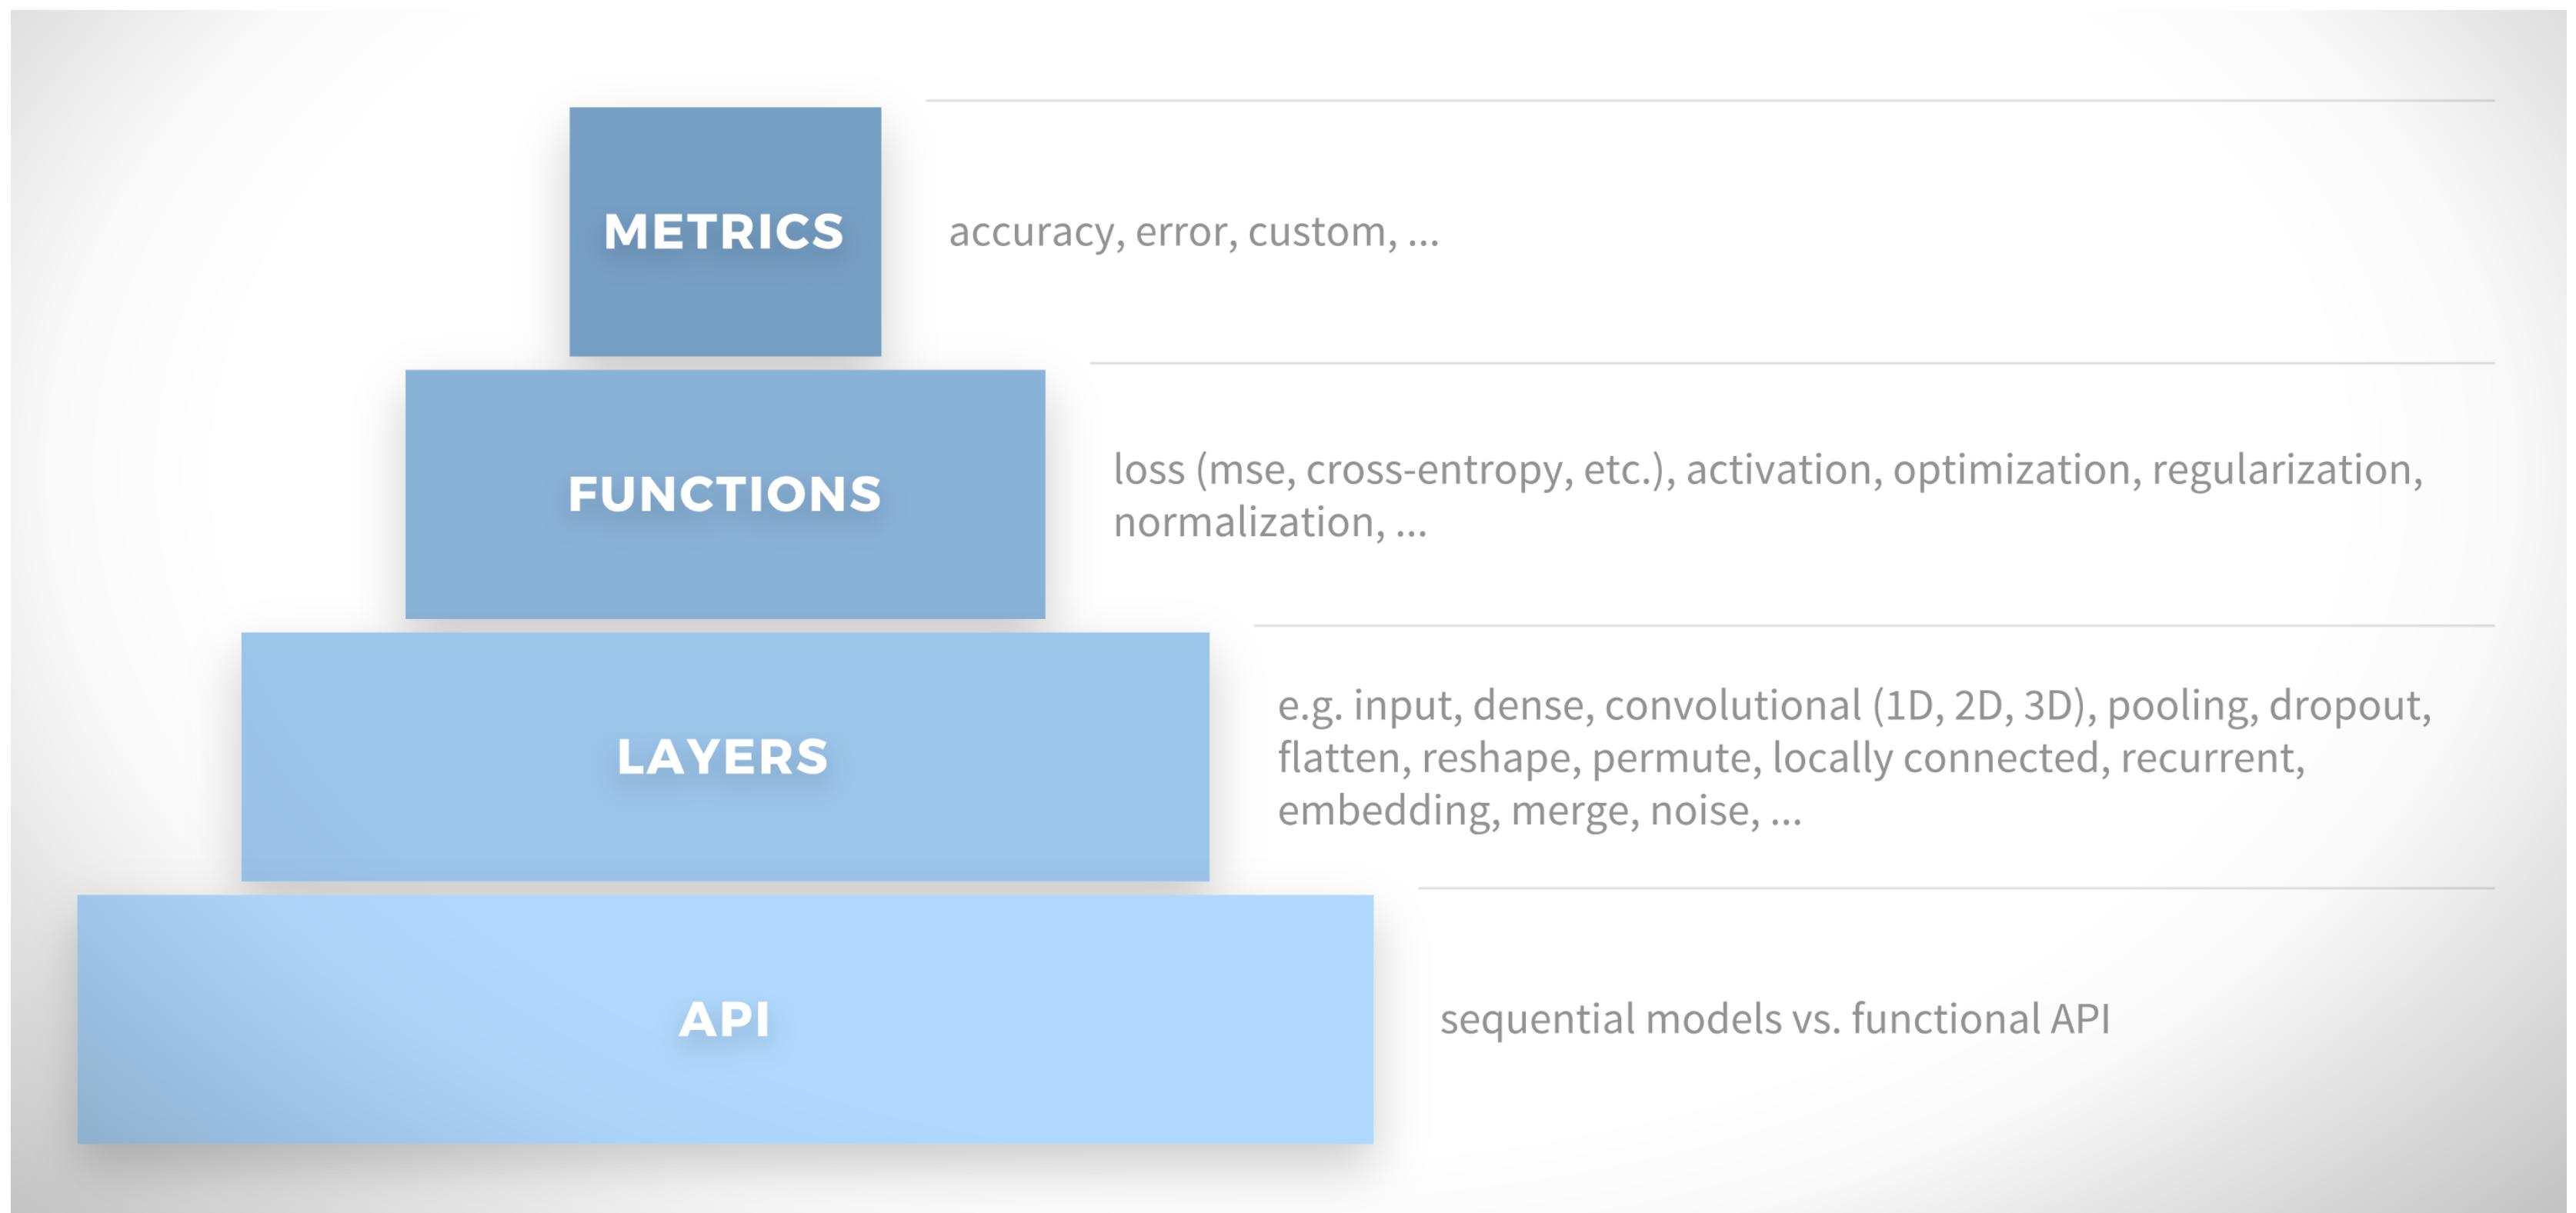
simple  
suitable for most cases  
linear order of layers  
only one direction from input to output



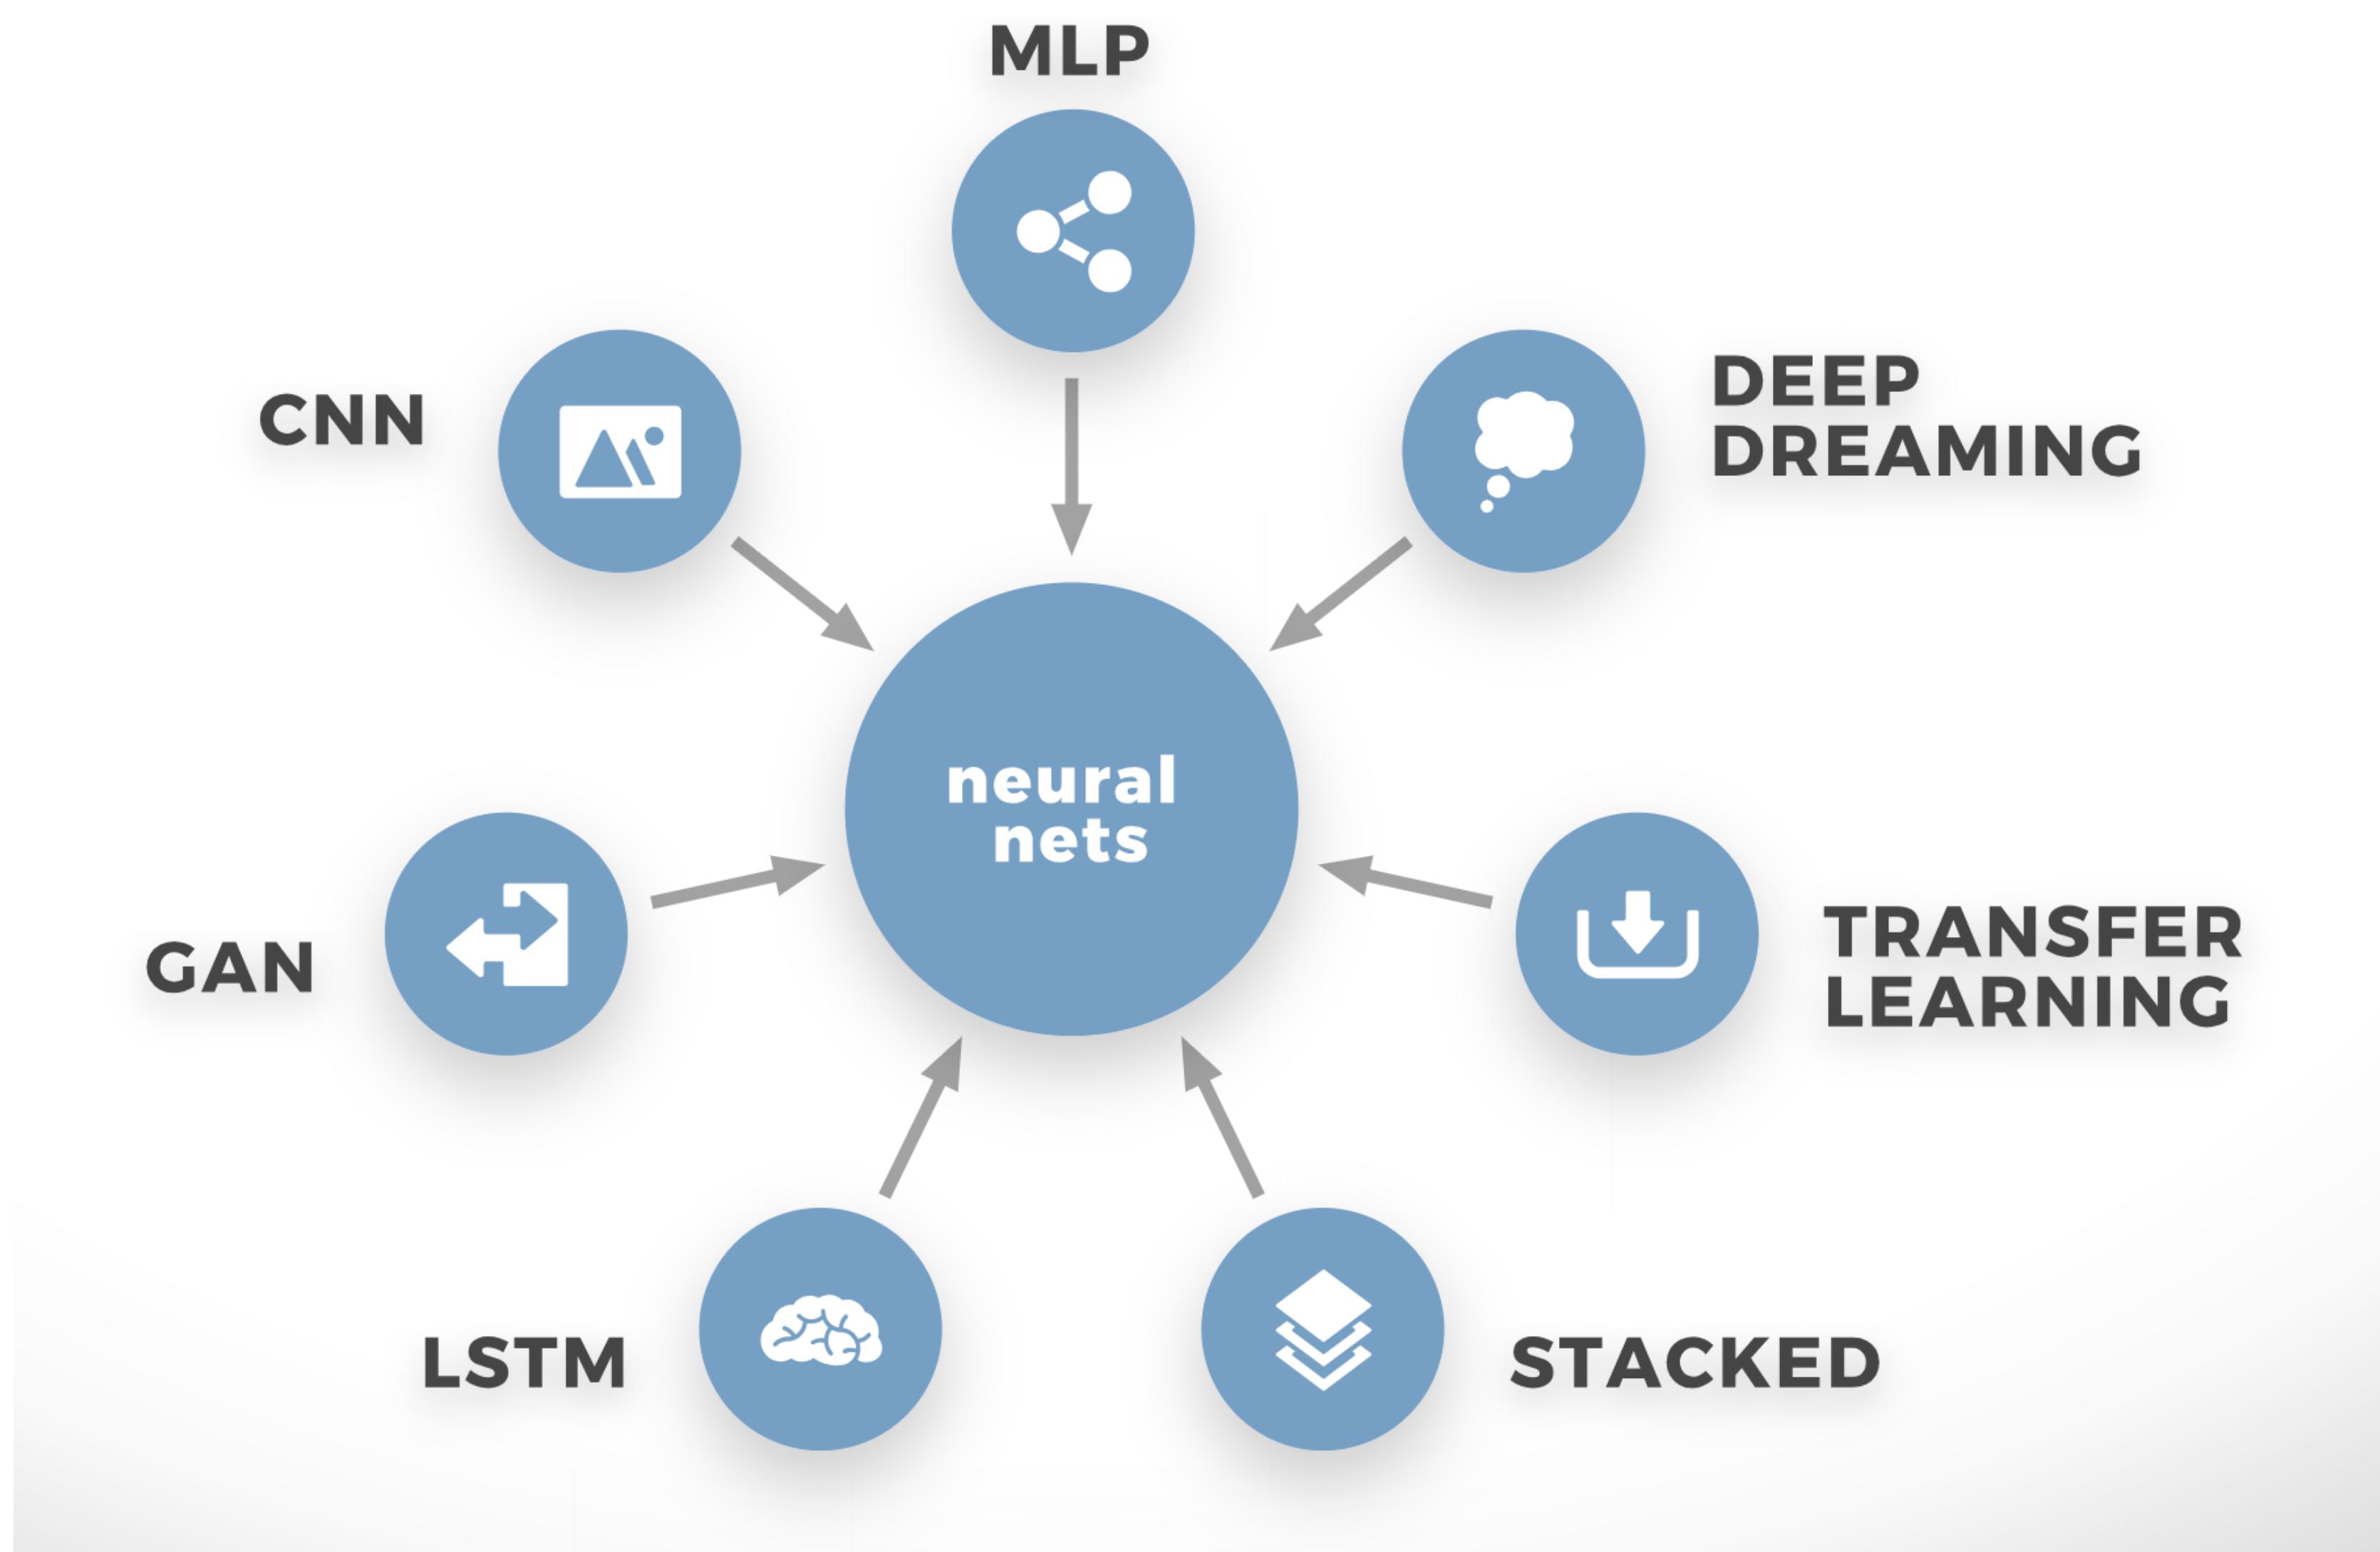
## FUNCTIONAL API

more complex  
suitable for complex models  
can have multiple in- or outputs  
layers can be non-sequential, e.g. LSTM

# Keras layers



# Endless possibilities



# Let's get our hands dirty!

## EDA

[https://github.com/ShirinG/image\\_classification\\_keras\\_tf/blob/master/notebooks/01-eda.ipynb](https://github.com/ShirinG/image_classification_keras_tf/blob/master/notebooks/01-eda.ipynb)

```
In [10]: plot_images(imgs=train_images[:10], rows=2, columns=5)
```



## Apply pretrained nets

- VGG16

[https://github.com/ShirinG/image\\_classification\\_keras\\_tf/blob/master/notebooks/02-pretrained-1.ipynb](https://github.com/ShirinG/image_classification_keras_tf/blob/master/notebooks/02-pretrained-1.ipynb)

- Xception

[https://github.com/ShirinG/image\\_classification\\_keras\\_tf/blob/master/notebooks/02-pretrained-2.ipynb](https://github.com/ShirinG/image_classification_keras_tf/blob/master/notebooks/02-pretrained-2.ipynb)

- Try out: ResNet50

# Modify pretrained nets

[https://github.com/ShirinG/image\\_classification\\_keras\\_tf/blob/master/notebooks/03-modify-pretrained.ipynb](https://github.com/ShirinG/image_classification_keras_tf/blob/master/notebooks/03-modify-pretrained.ipynb)

## Modify VGG16

- transfer learning (freeze all but the penultimate layer and re-train the last Dense layer) and
- fine tuning (un-freeze the lower convolutional layers and retrain more layers)

Validation set: fit\_generator has no option validation\_split

<https://keras.io/applications/#usage-examples-for-image-classification-models>

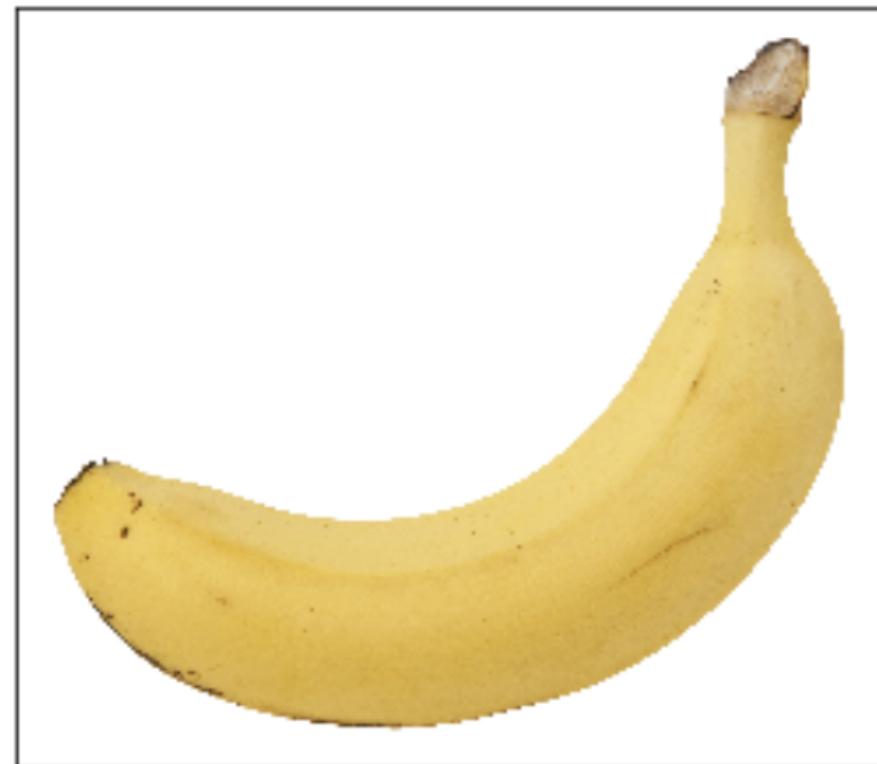
```
In [6]: # important: exclude top layers
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(75, 75, 3))
#base_model.summary()
```

```
In [7]: # Freeze the layers except the last 4 layers
for layer in base_model.layers[:-4]:
    layer.trainable = False
```

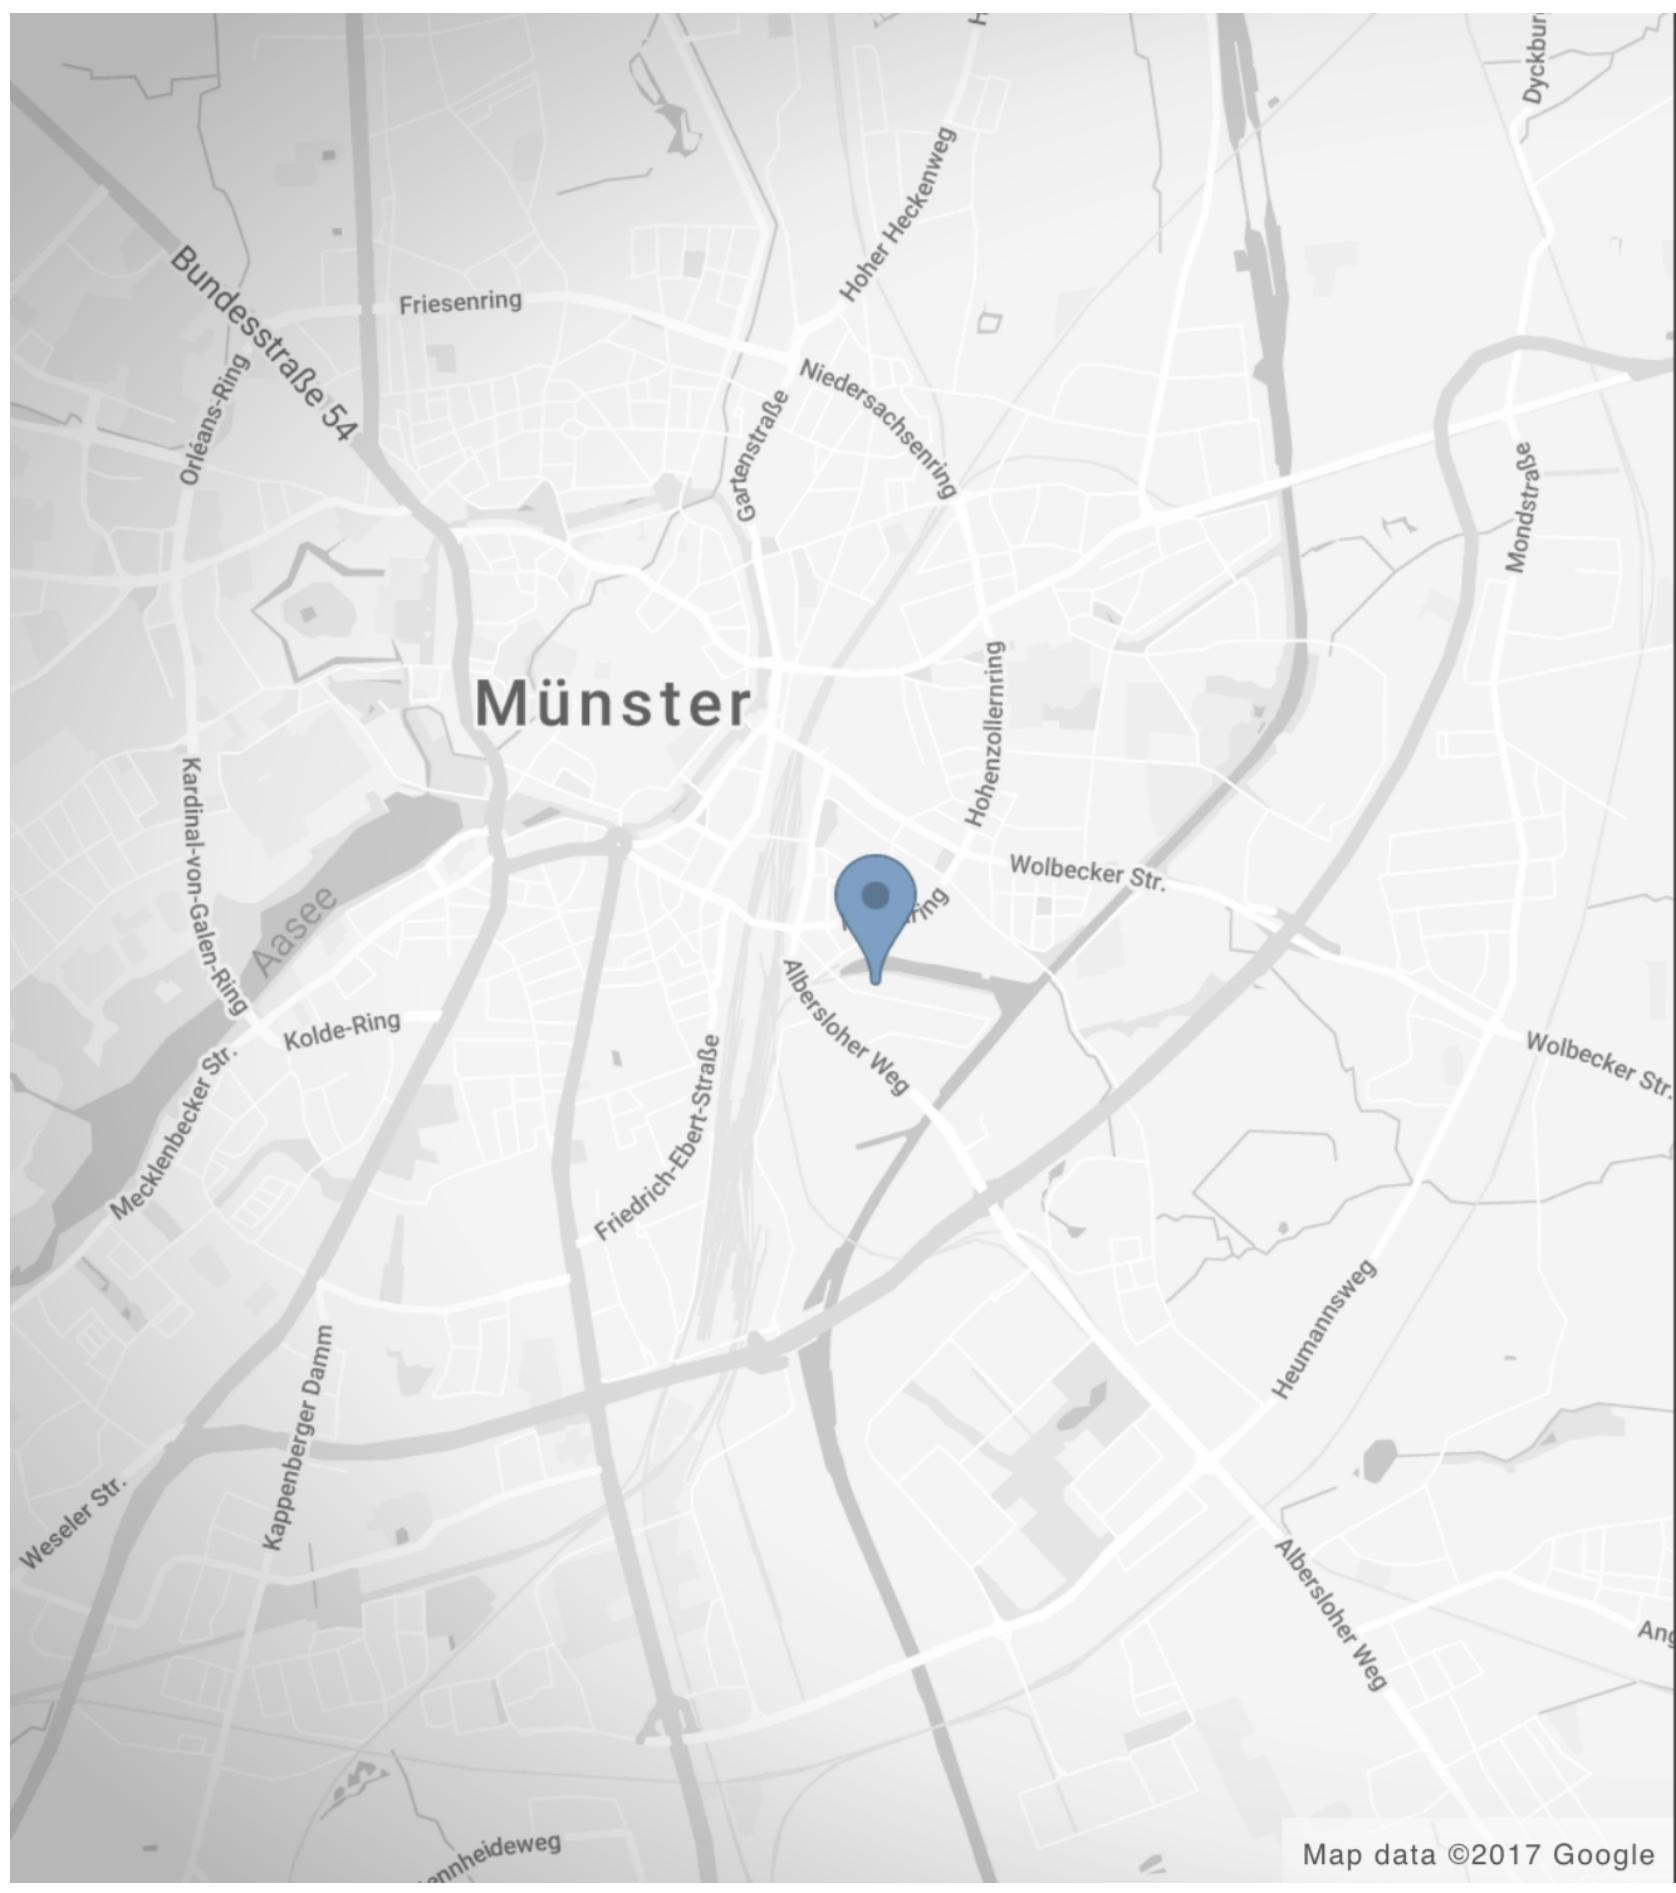
# Build your own CNN

[https://github.com/ShirinG/image\\_classification\\_keras\\_tf/blob/master/notebooks/04-fruits-cnn.ipynb](https://github.com/ShirinG/image_classification_keras_tf/blob/master/notebooks/04-fruits-cnn.ipynb)

```
In [16]: classify_image_model(test_images[0])
```



Predicted class: Banana with probability 99.99985694885254%



**Thank  
you!  
And stay  
connected**

Am Mittelhafen 14, Münster

 ShirinGlander

 shirin.glander@codecentric.de

[www.codecentric.ai](http://www.codecentric.ai)

<https://www.youtube.com/codecentricAI>

[www.shirin-glander.de](http://www.shirin-glander.de)