



BILDKLASSIFIKATION LEICHT GEMACHT – MIT KERAS UND TENSORFLOW

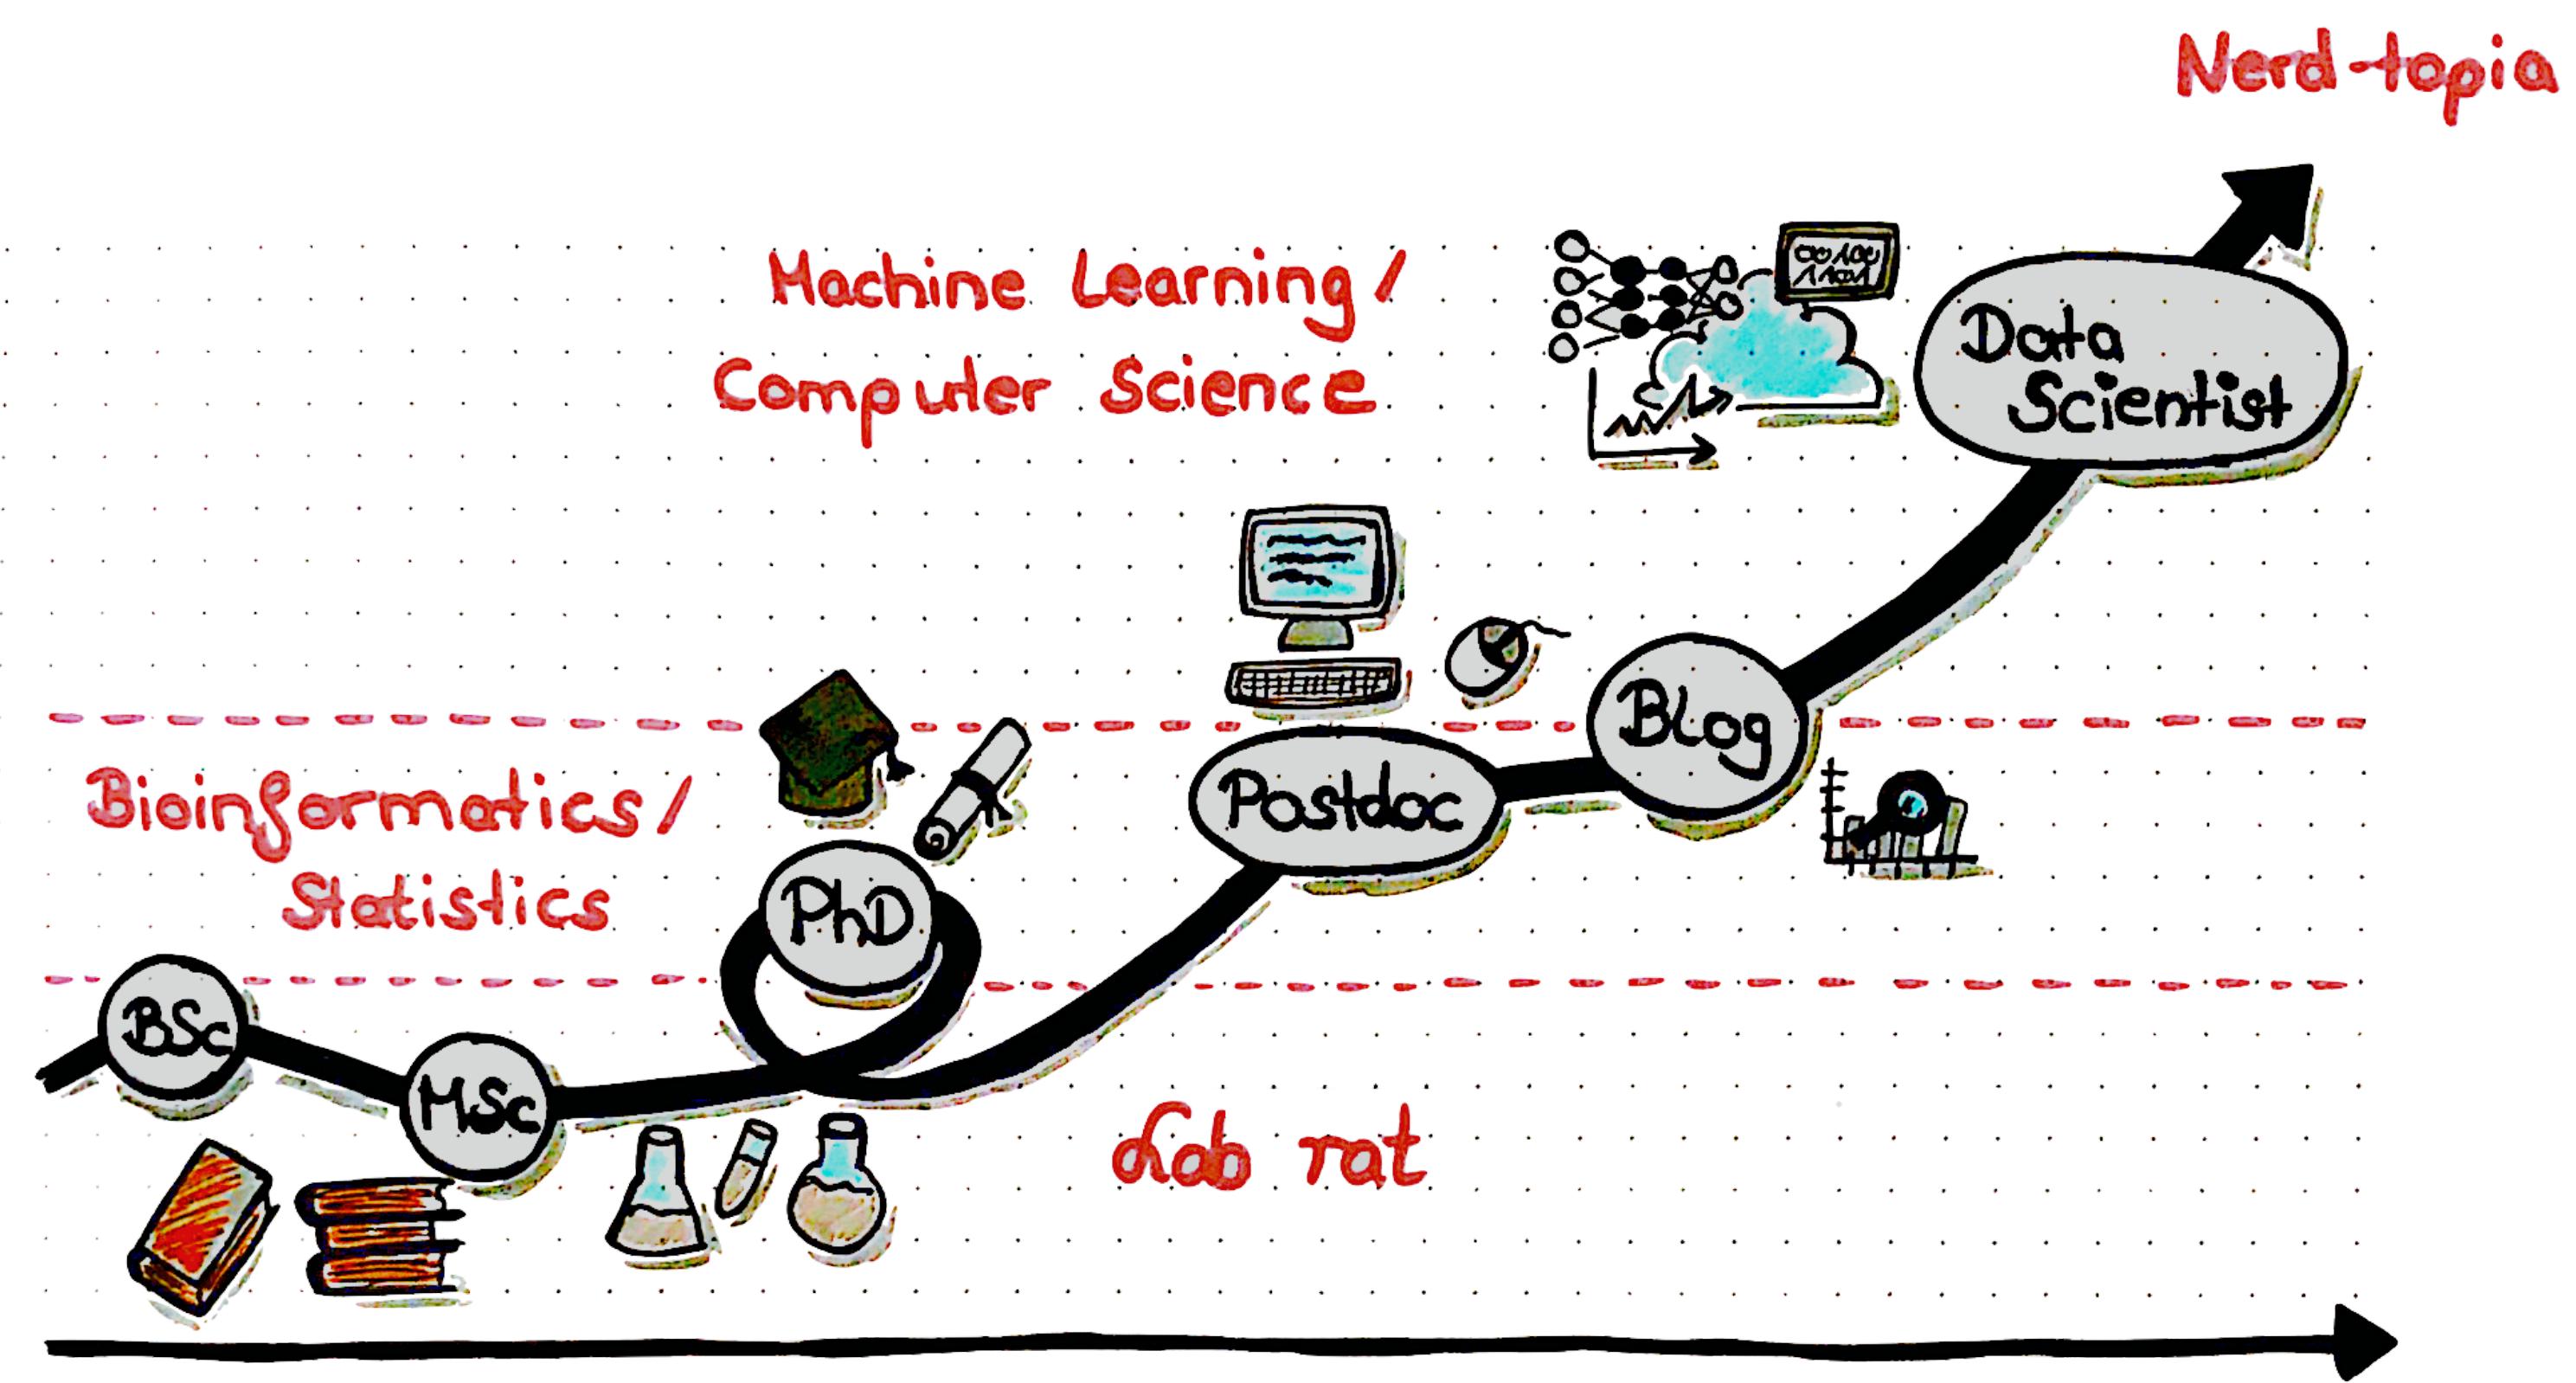
Dr. Shirin Glander
Data Scientist @ codecentric AG

About this Workshop

- Learn what neural nets are and how they can be used in computer vision.
- Understand how a computer learns to "see".
- Apply pre-trained nets and modify them.
- Build a model from scratch that differentiates between fruits on images.
- Visualise convolutions and layers.
- (Bonus: Explaining our classifications with LIME)

Material on Github: https://github.com/ShirinG/image_classification_keras_tf

About me



How does a computer learn to "see"?

Convolutional Neural Nets

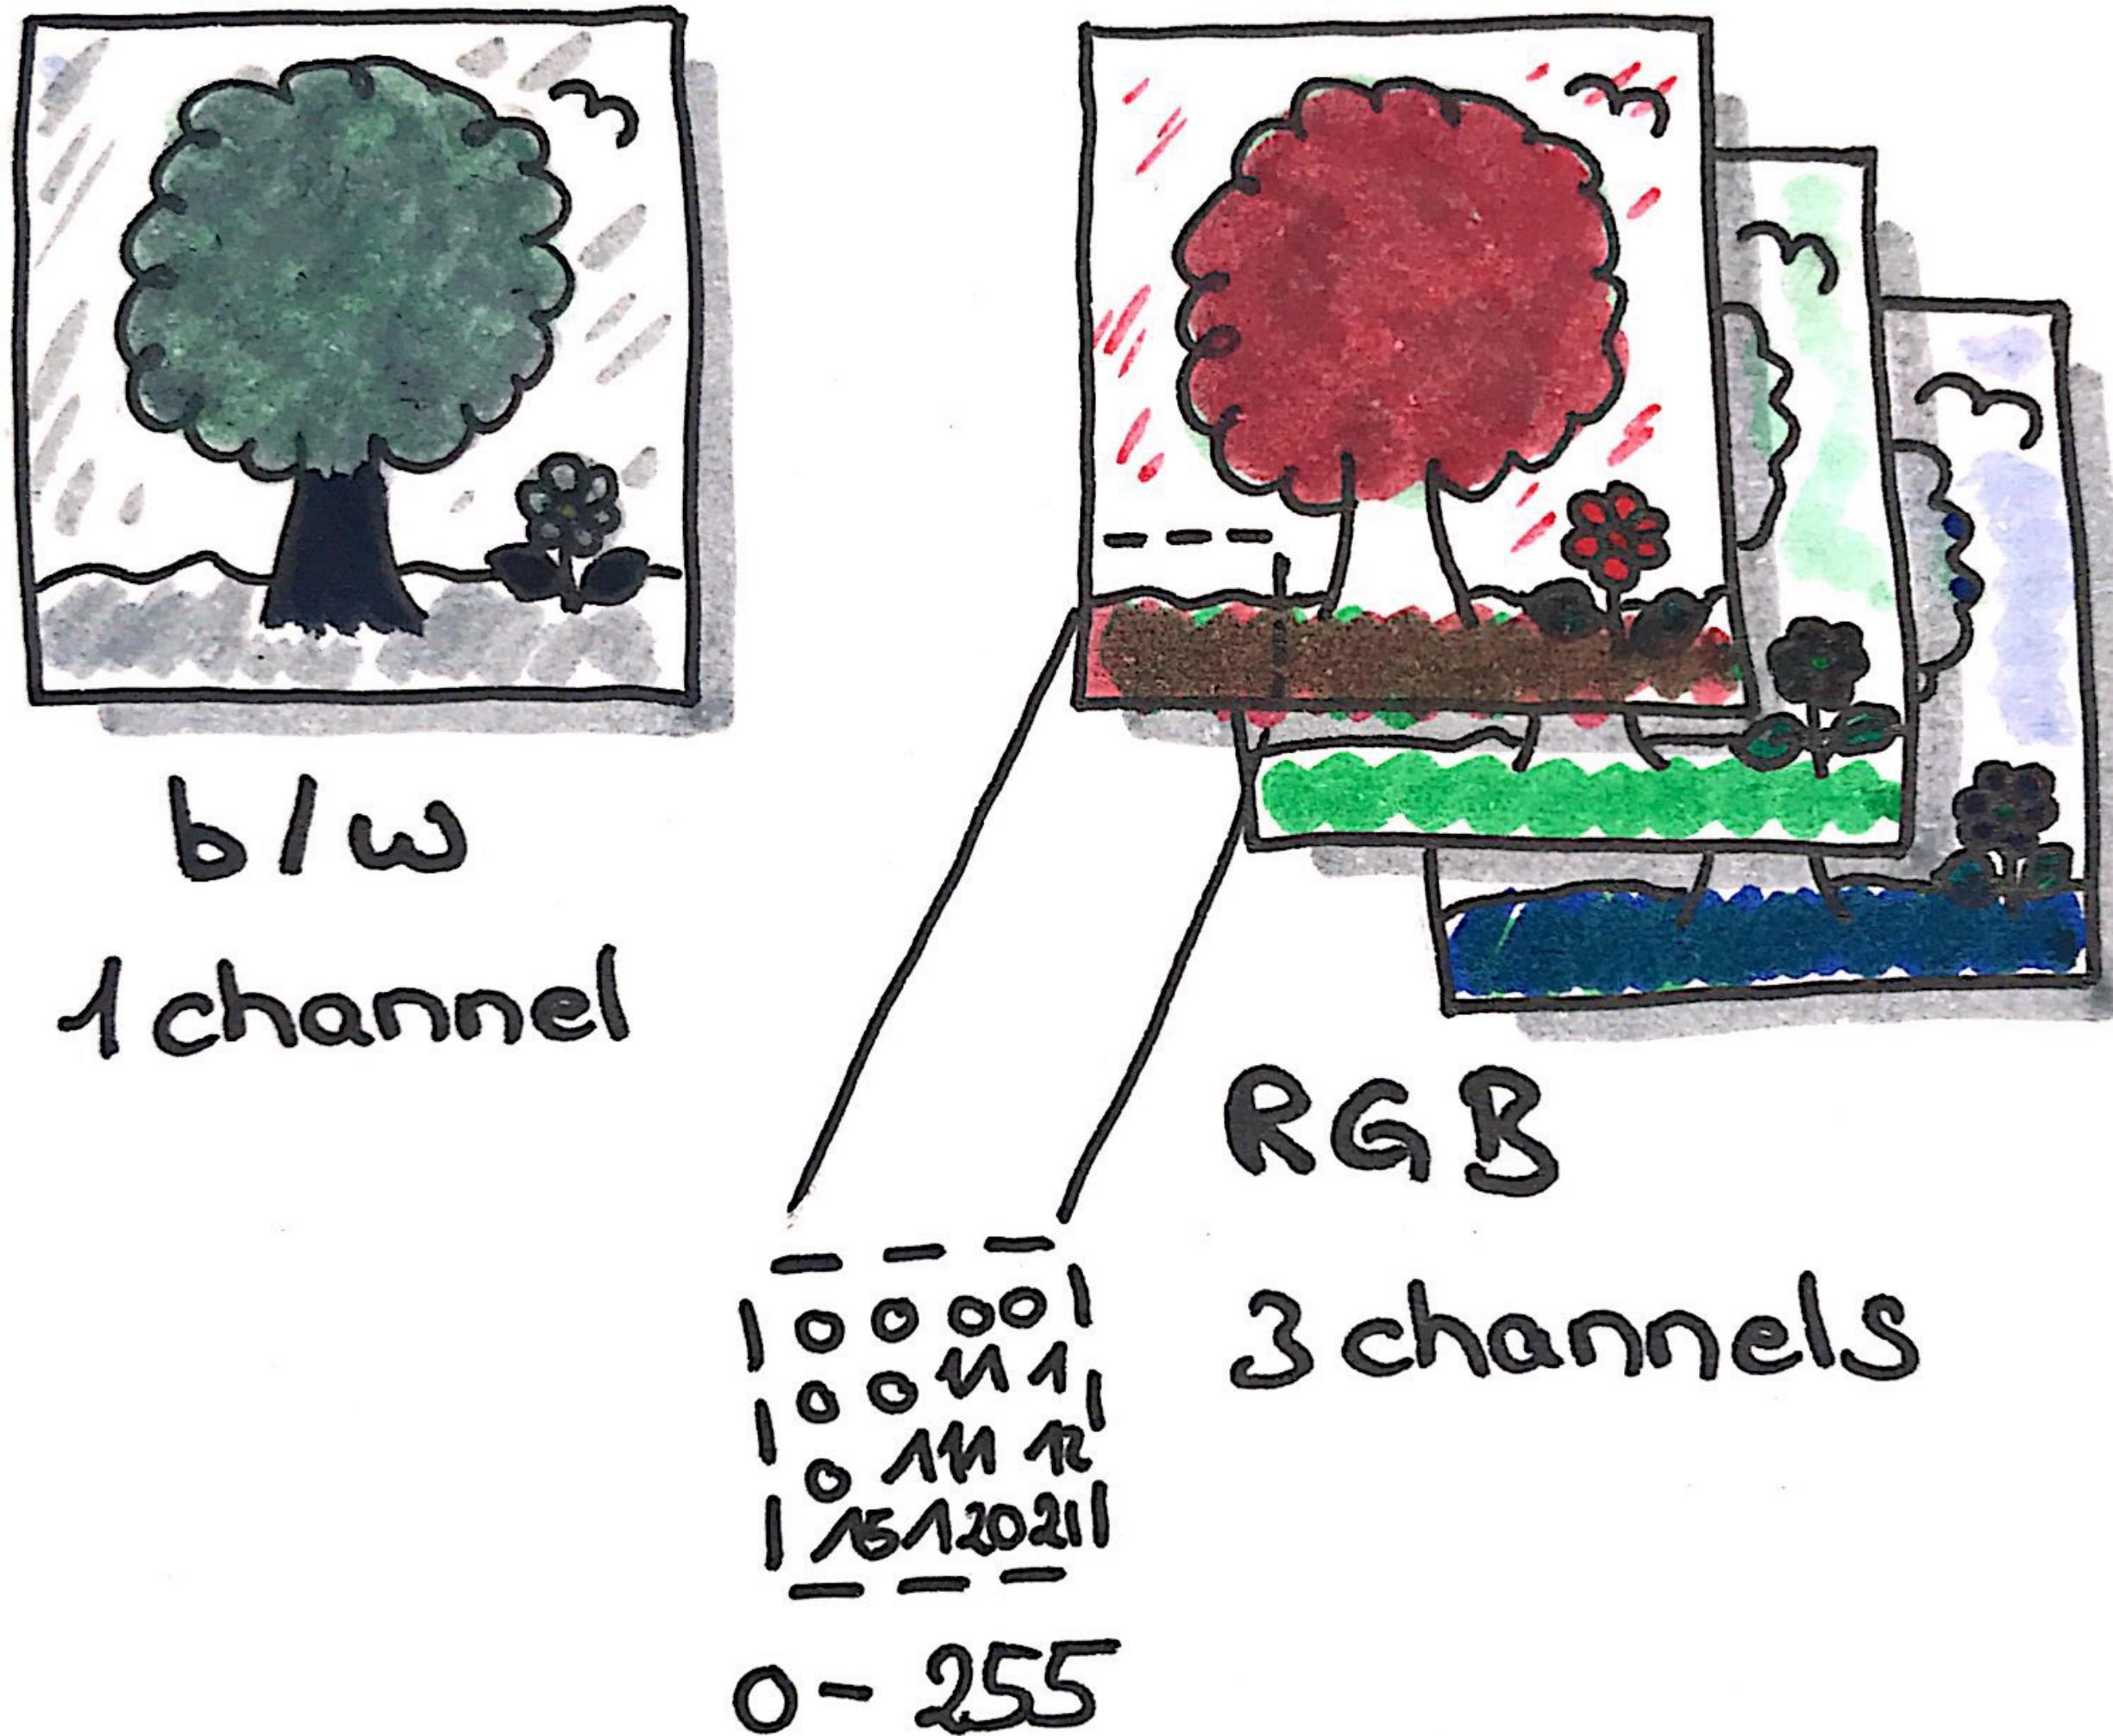


- image classification
class = tree
- object detection
flower

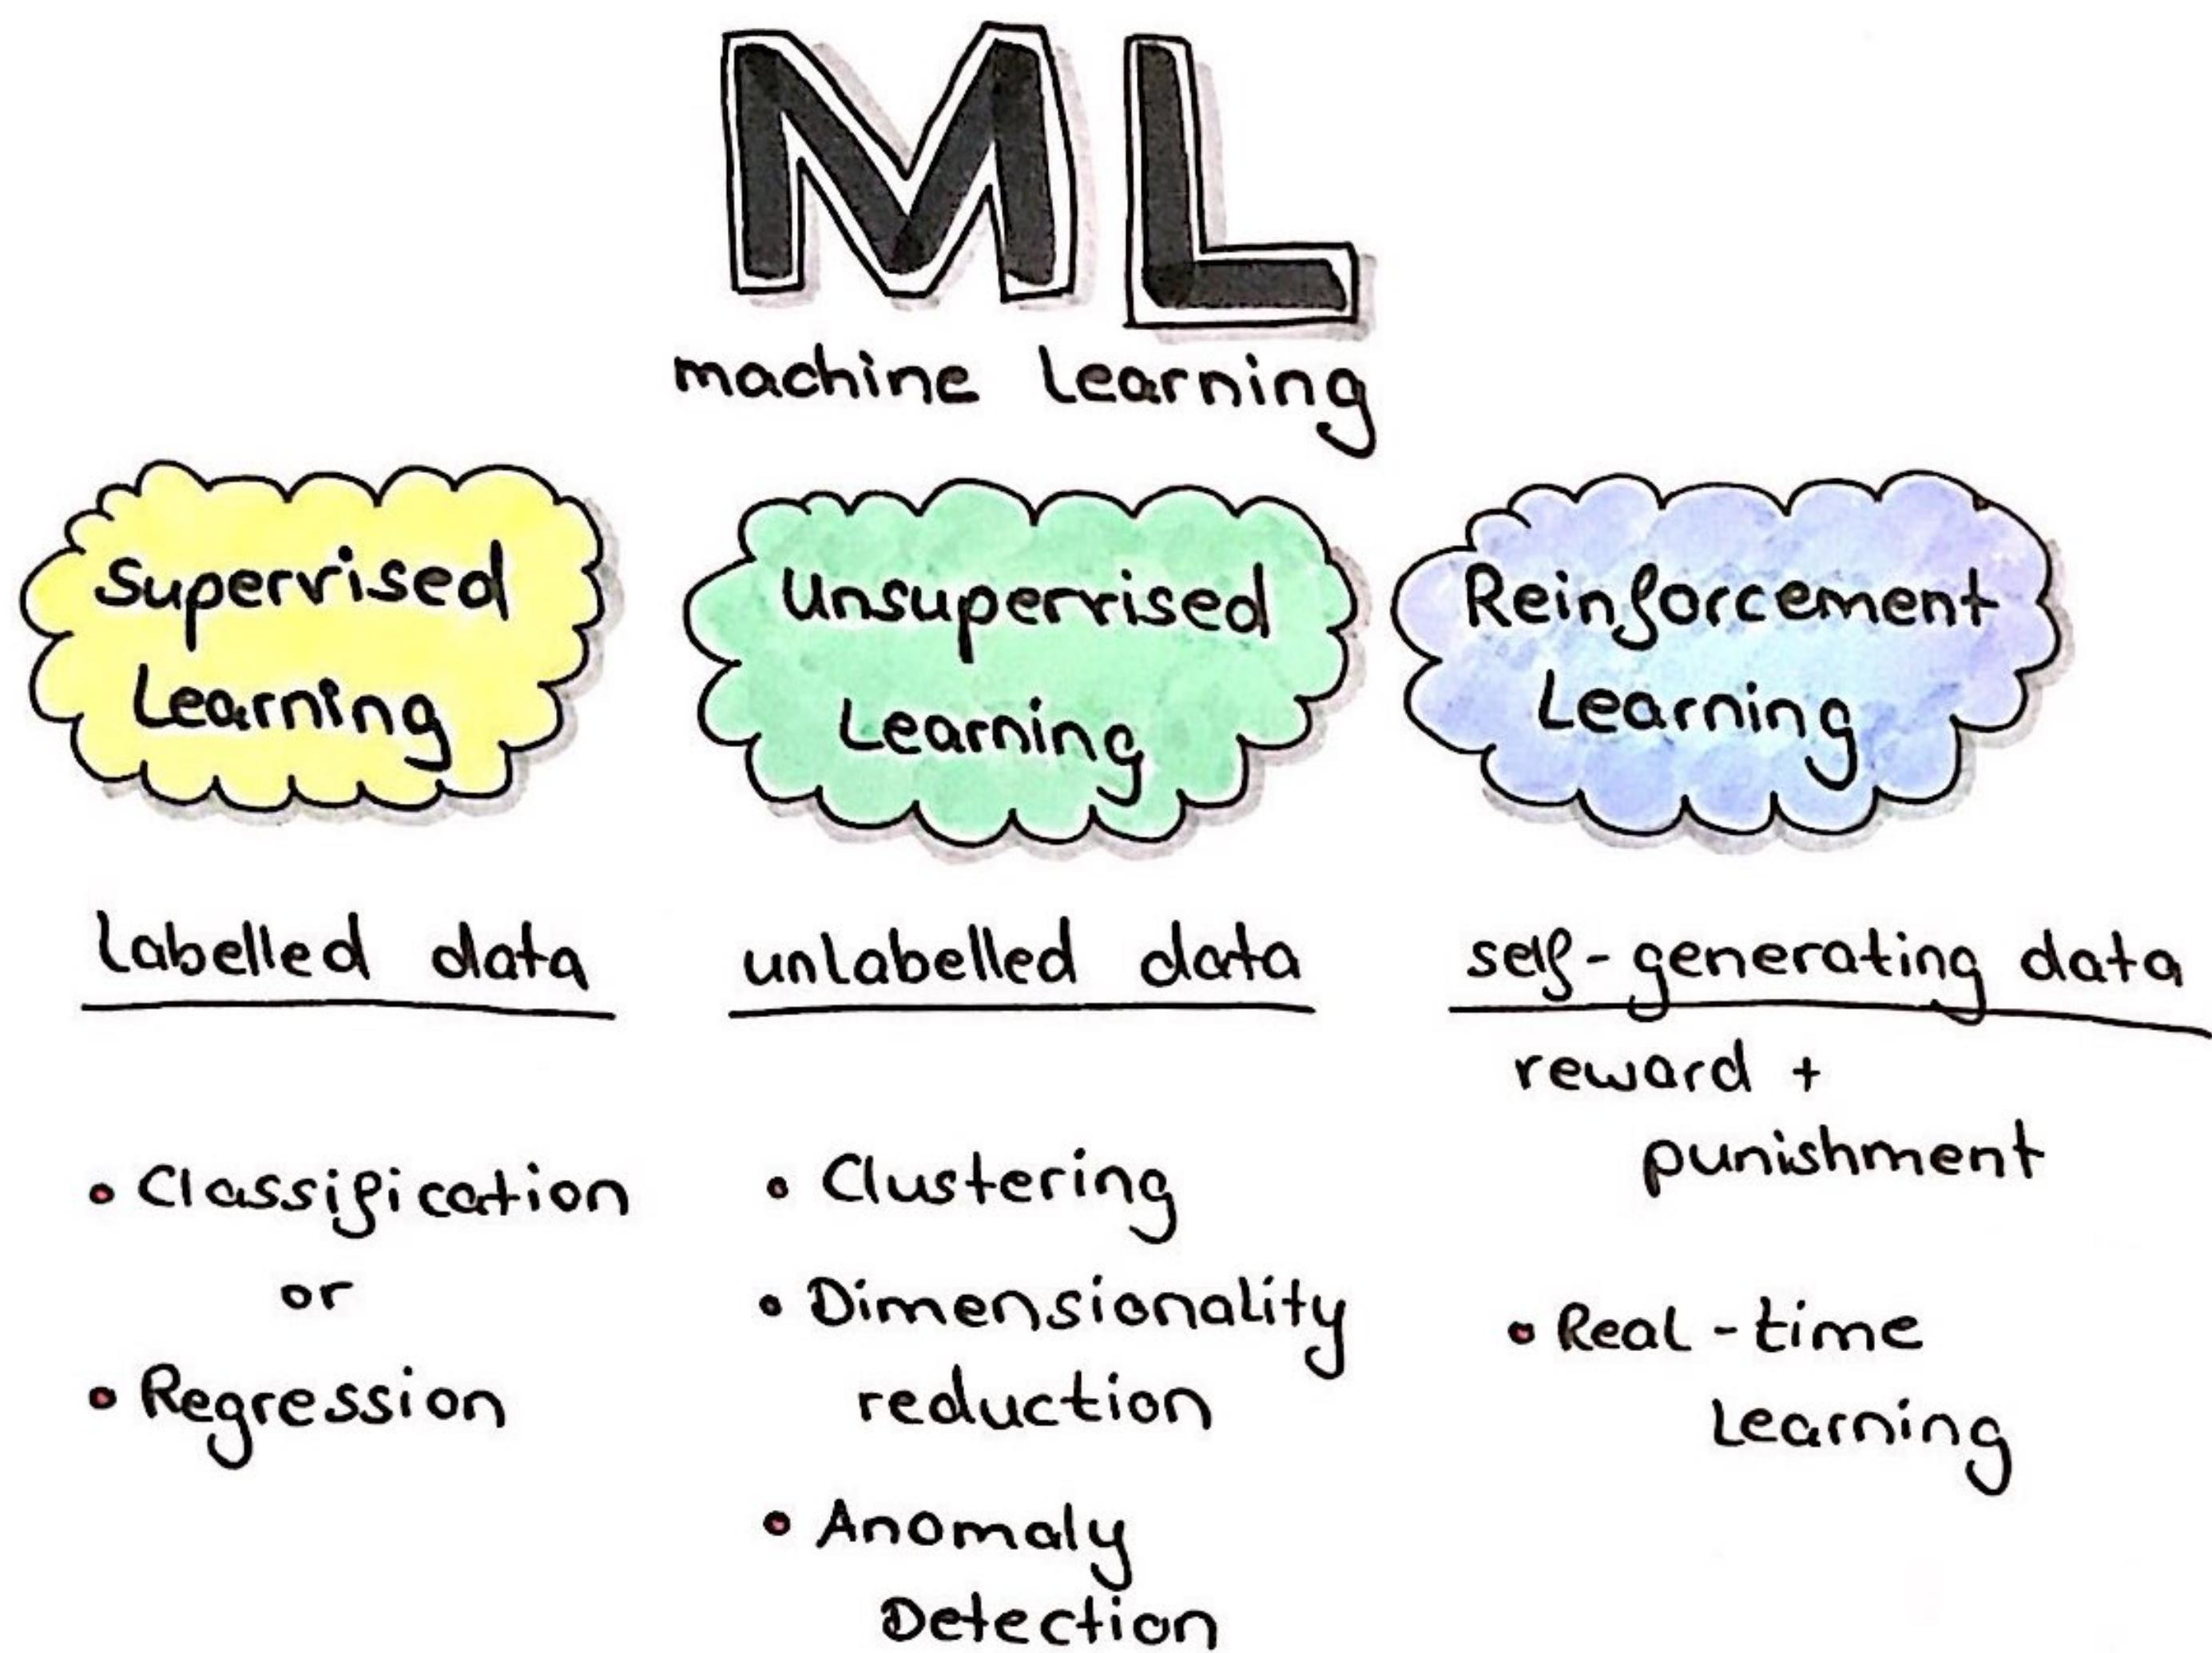
Computer Vision

See also [this video from codecentric.AI](#) for more on Computer Vision Basics!

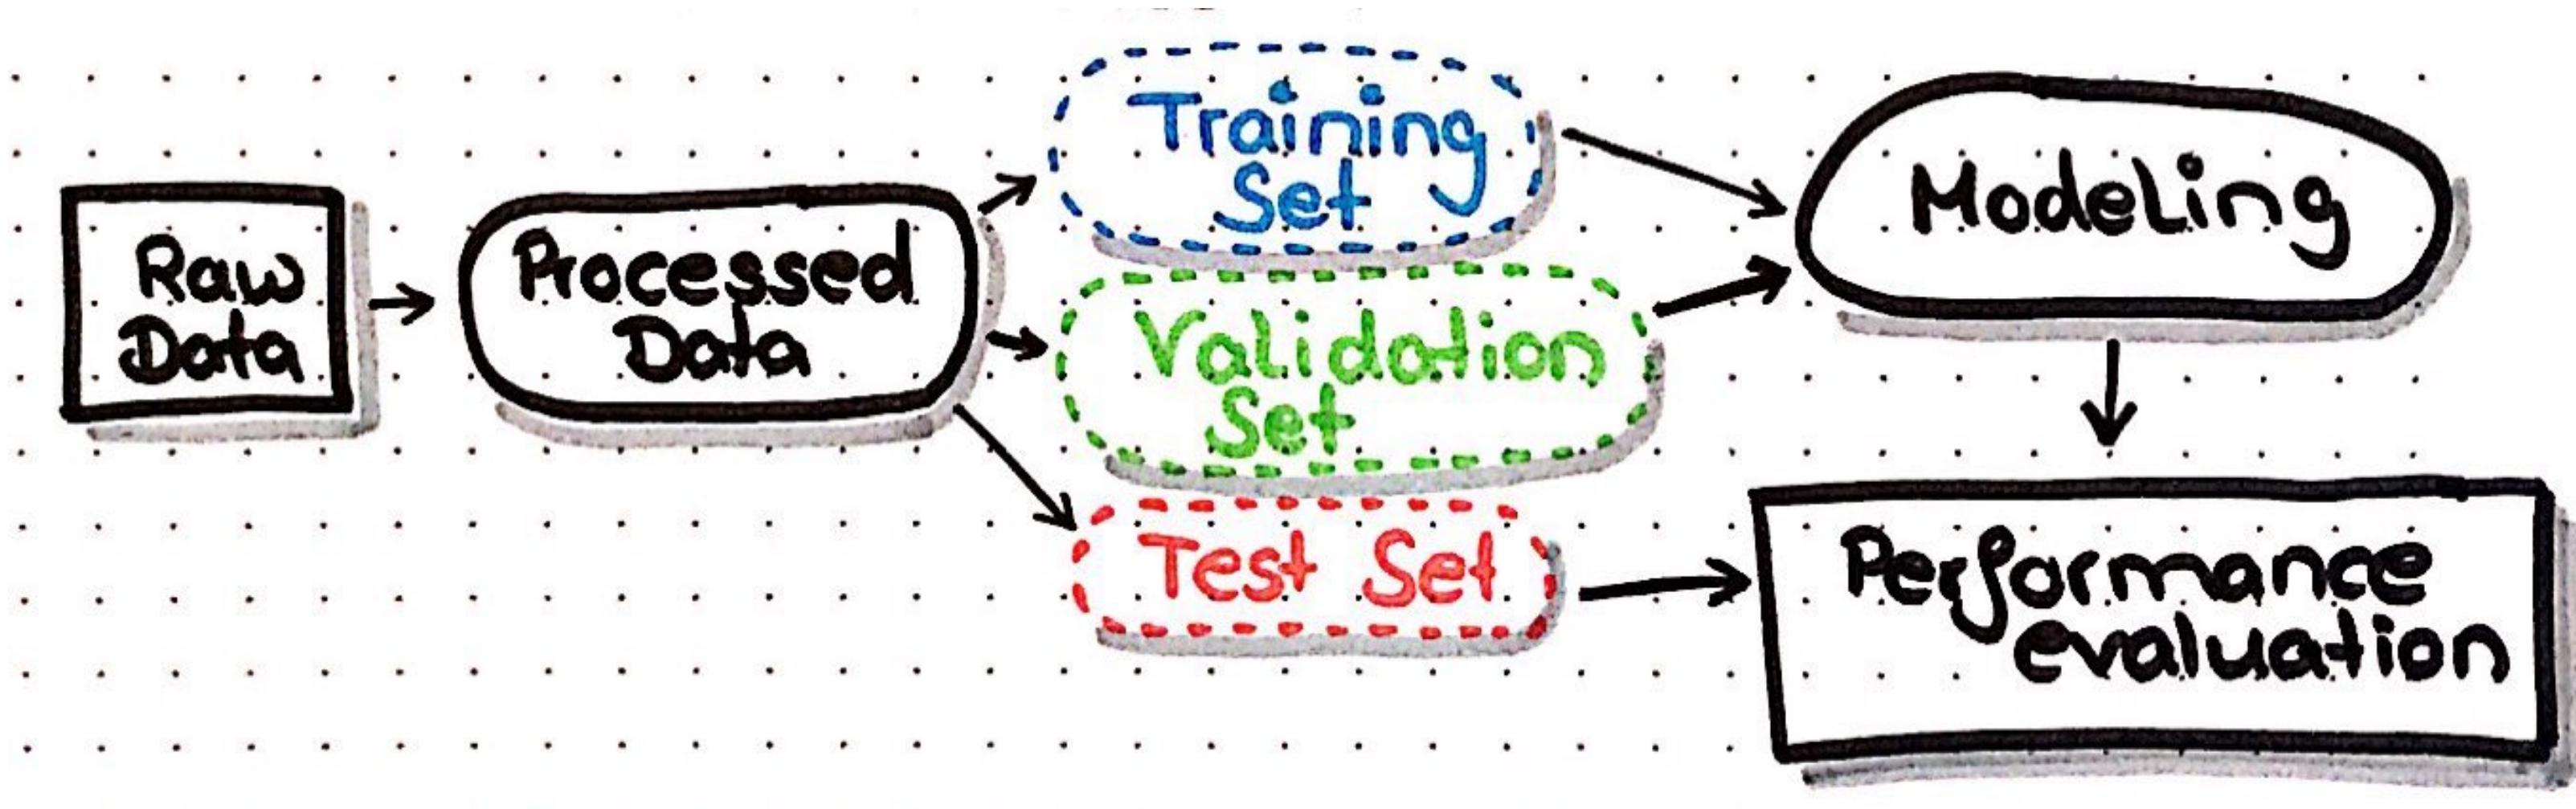
How does a computer learn to "see"?



What is Machine Learning (ML)



A typical ML Workflow

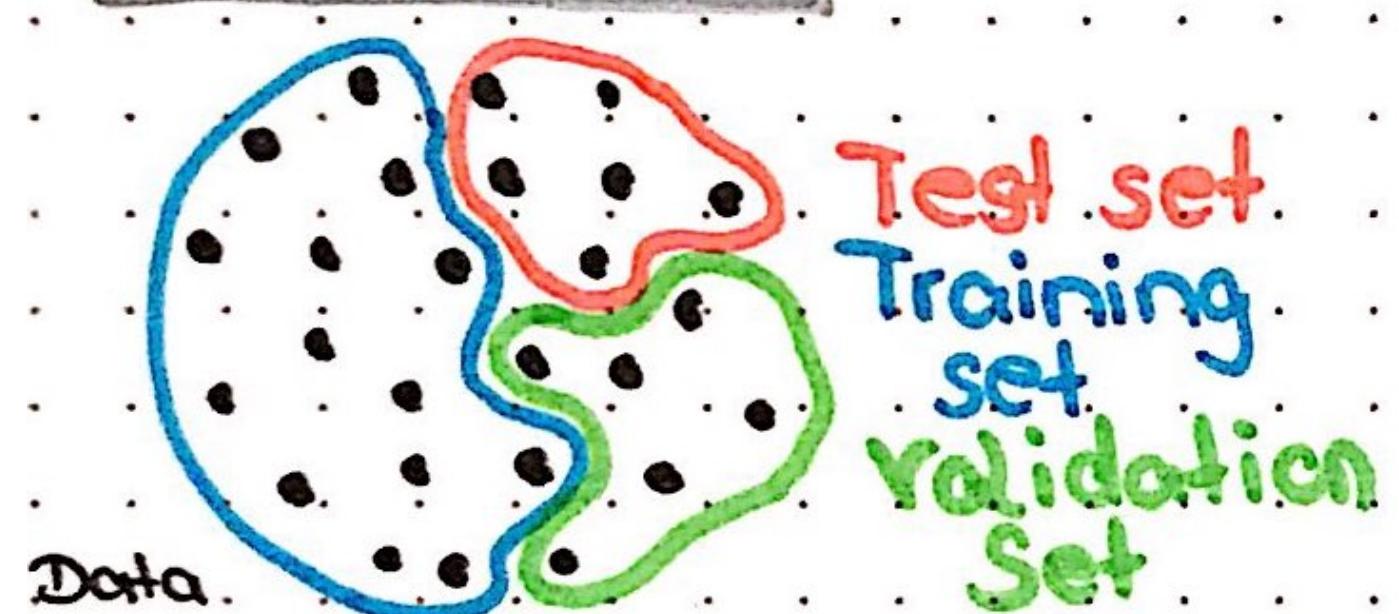


Why use validation and test data?

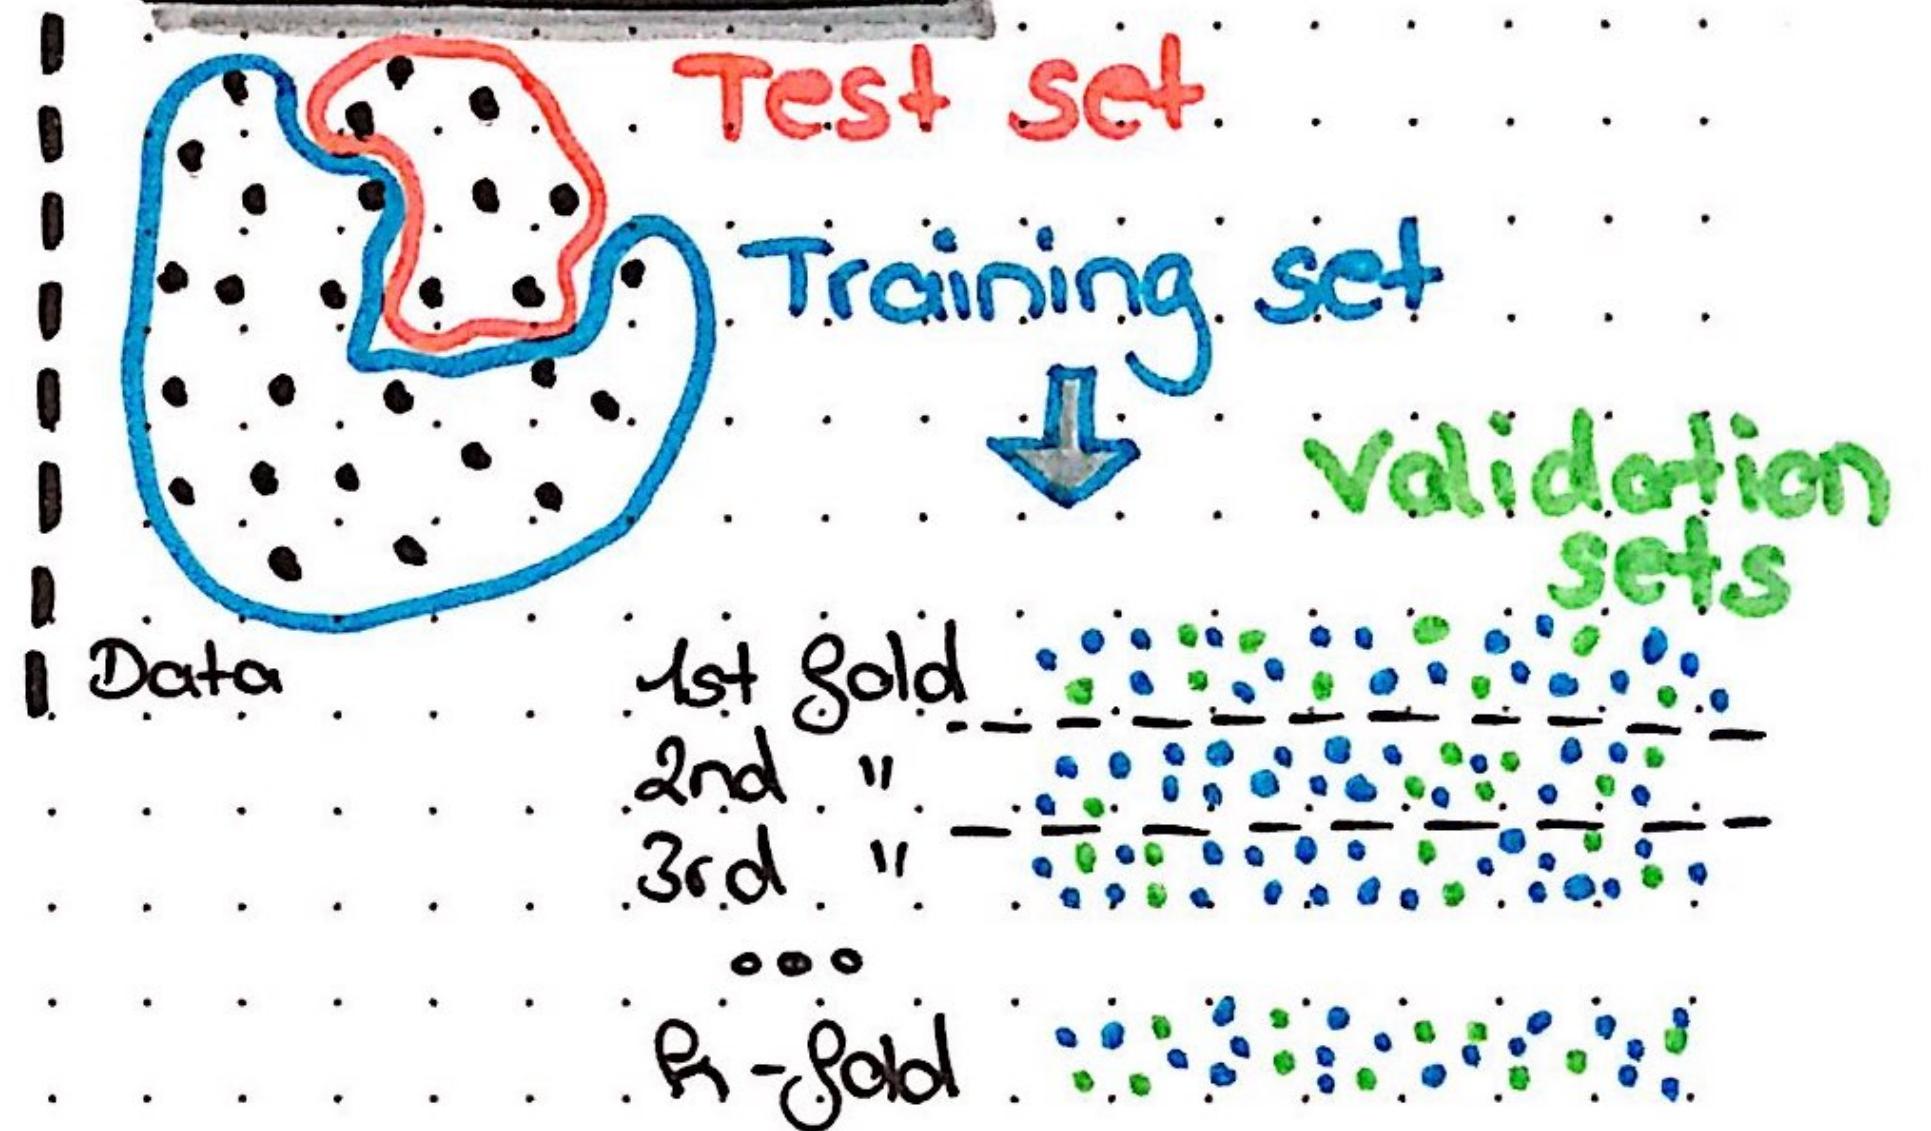


Validation methods

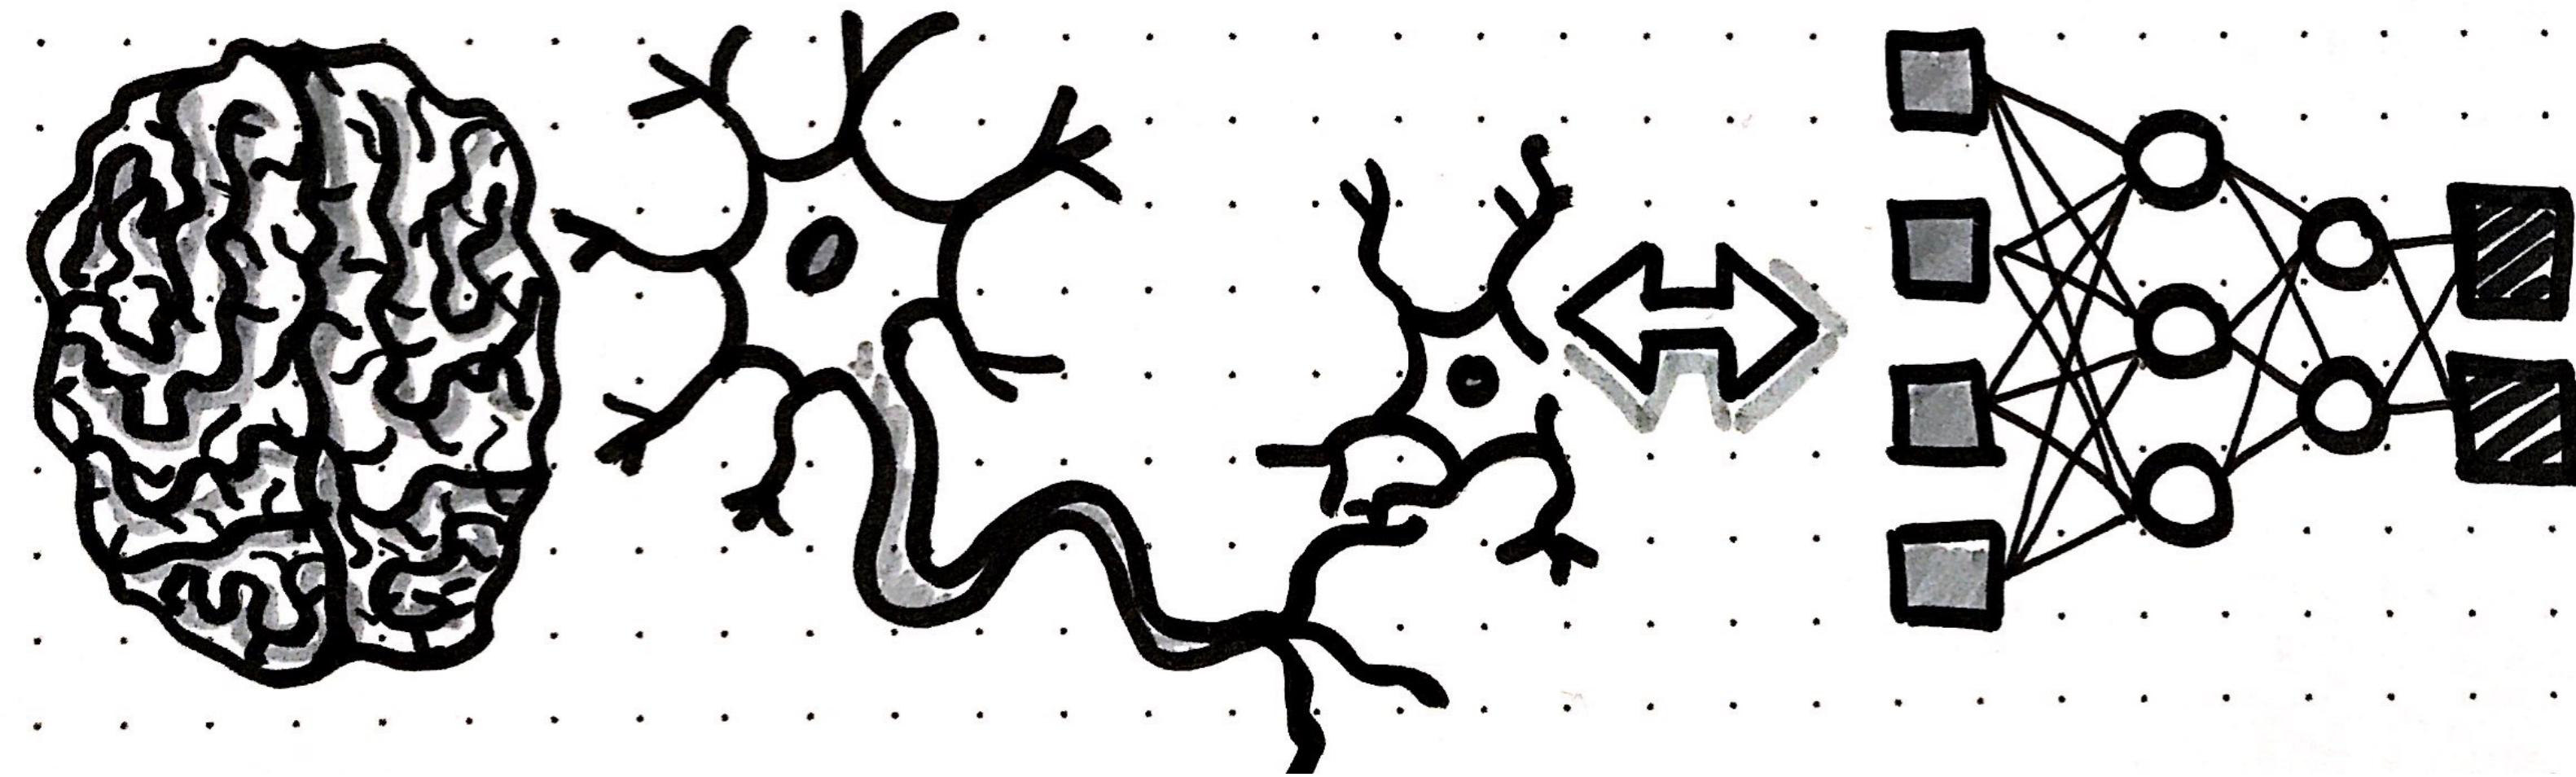
Hold-out validation



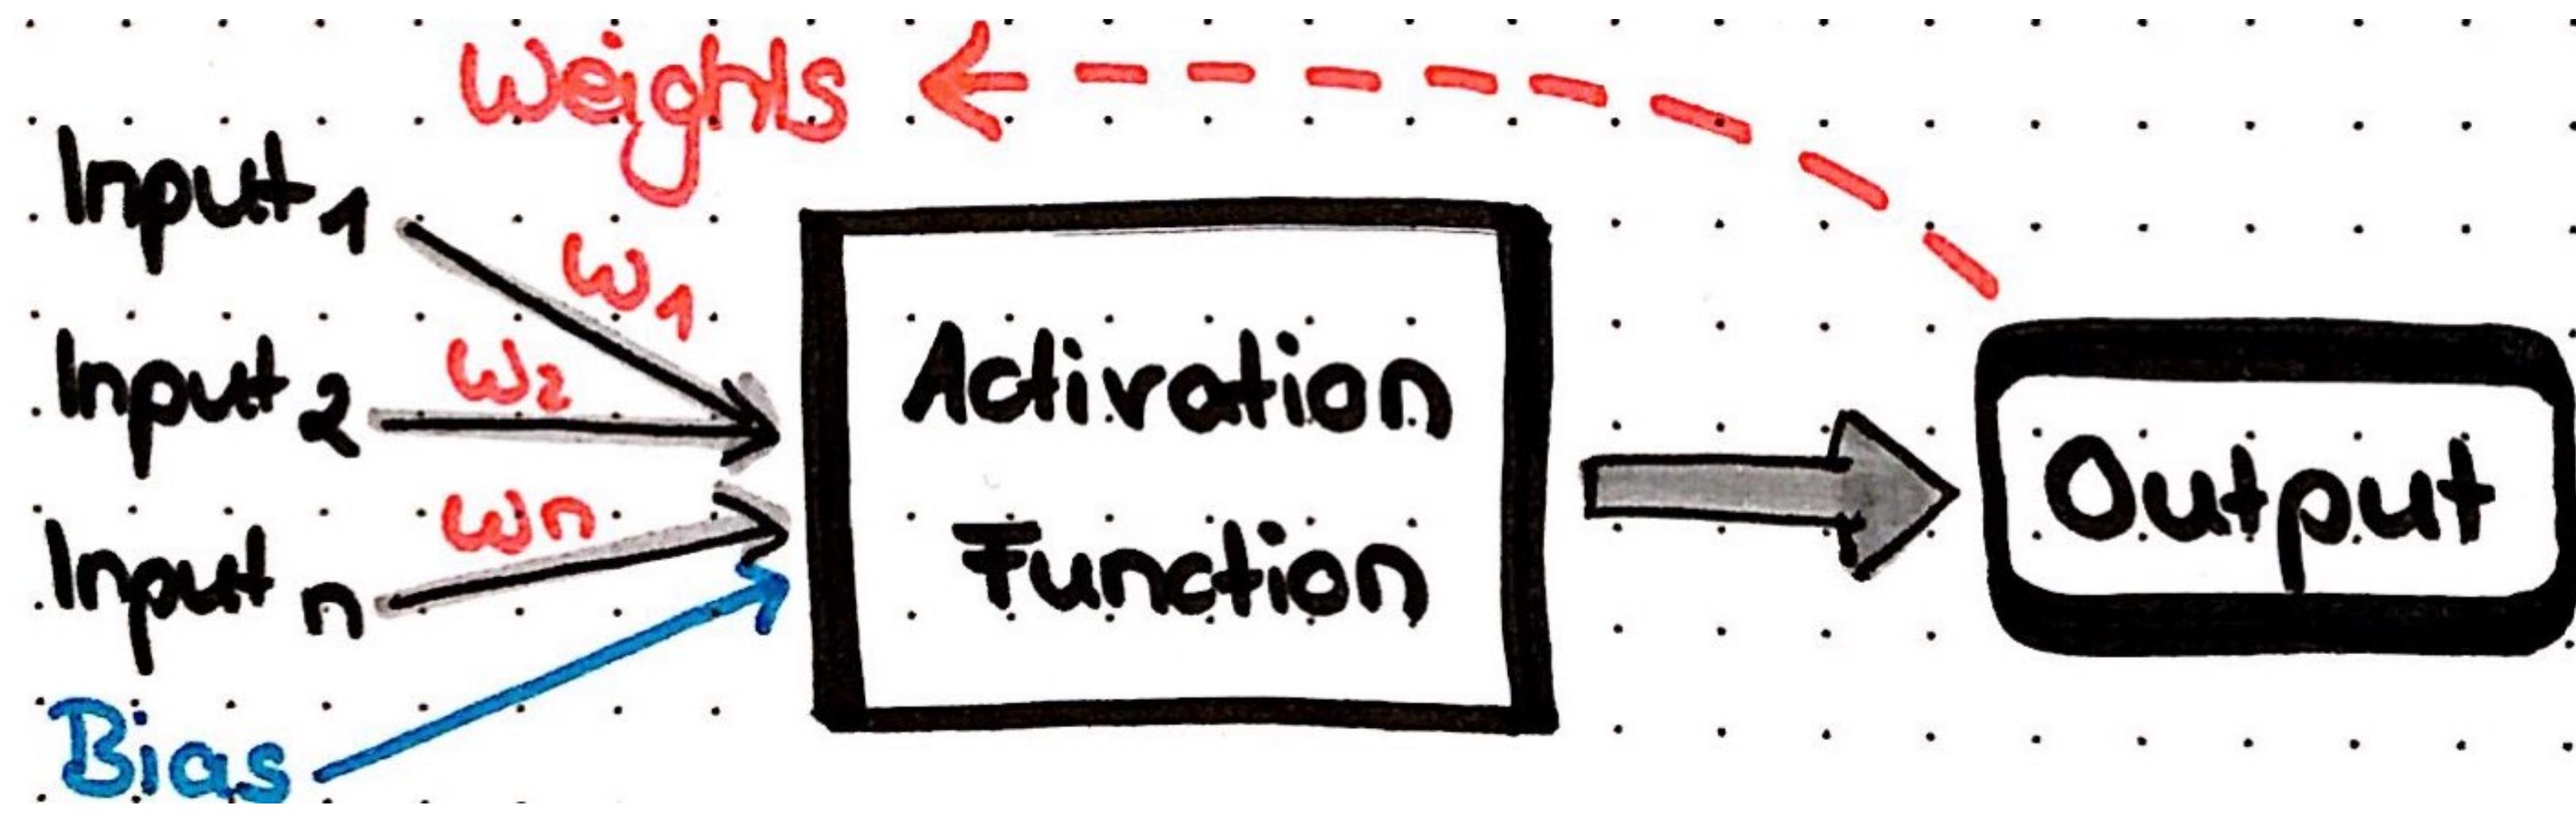
Cross-validation



What are neural nets?

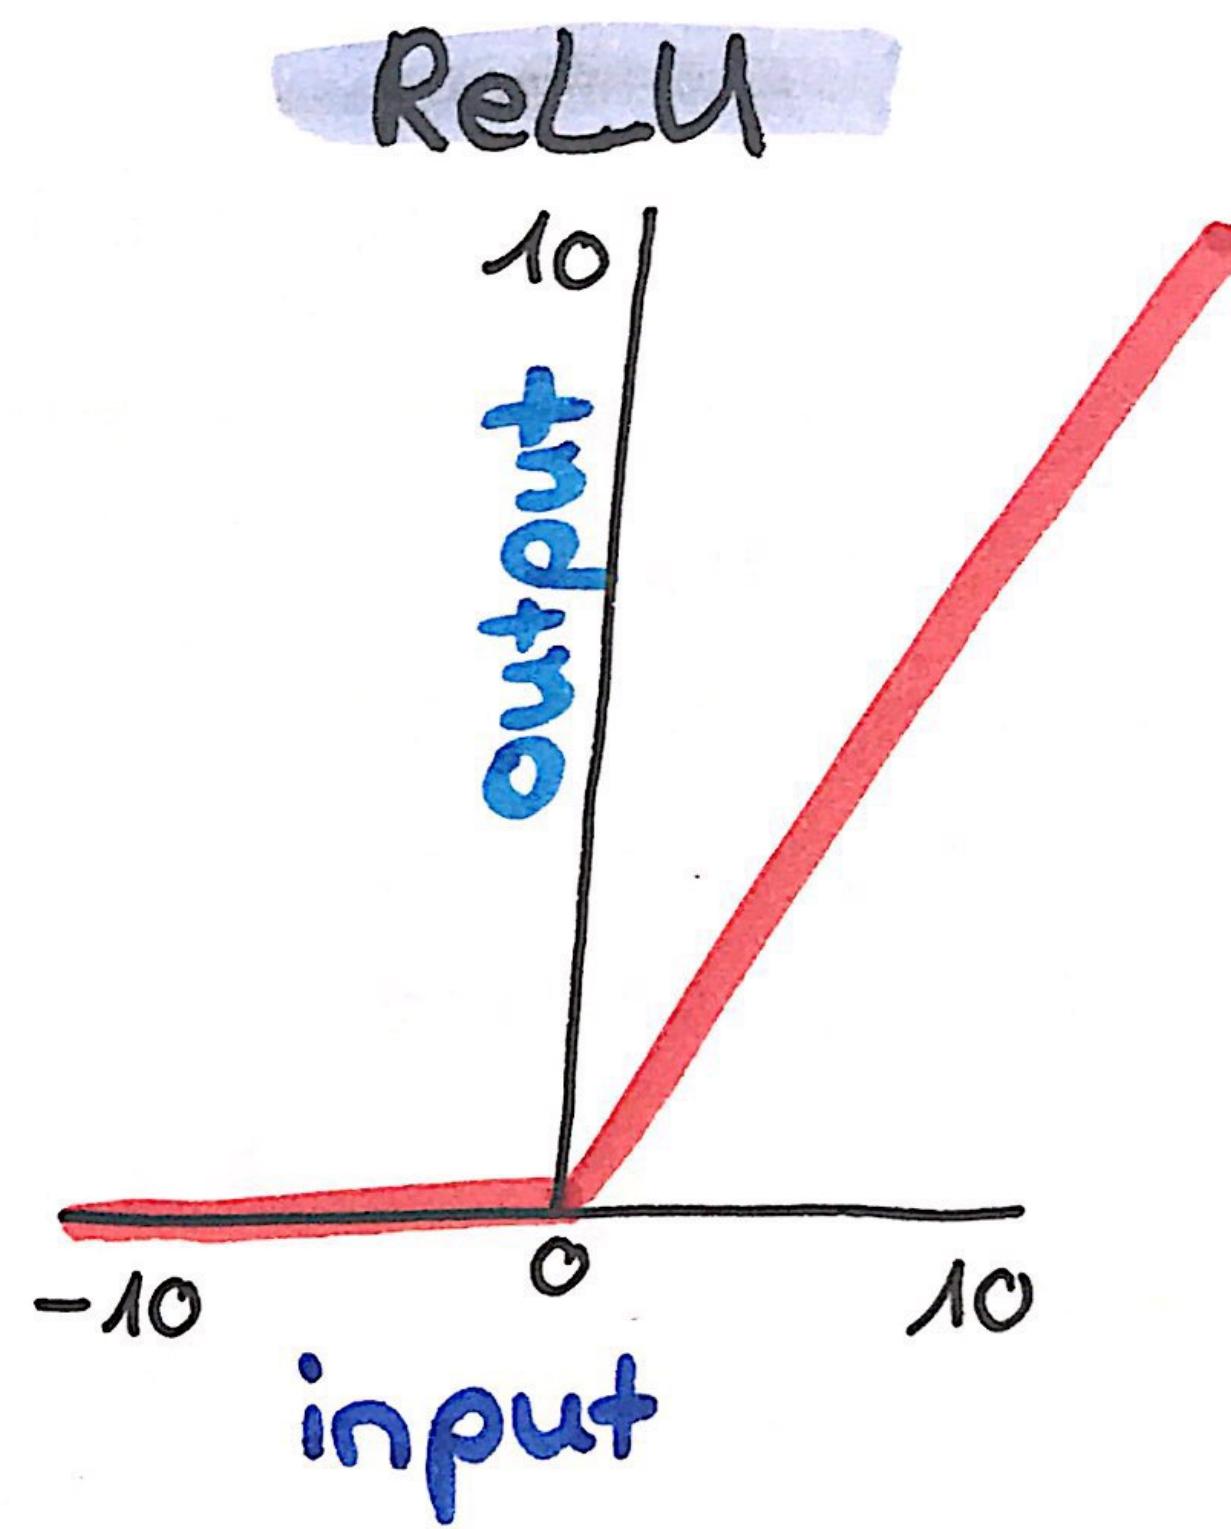
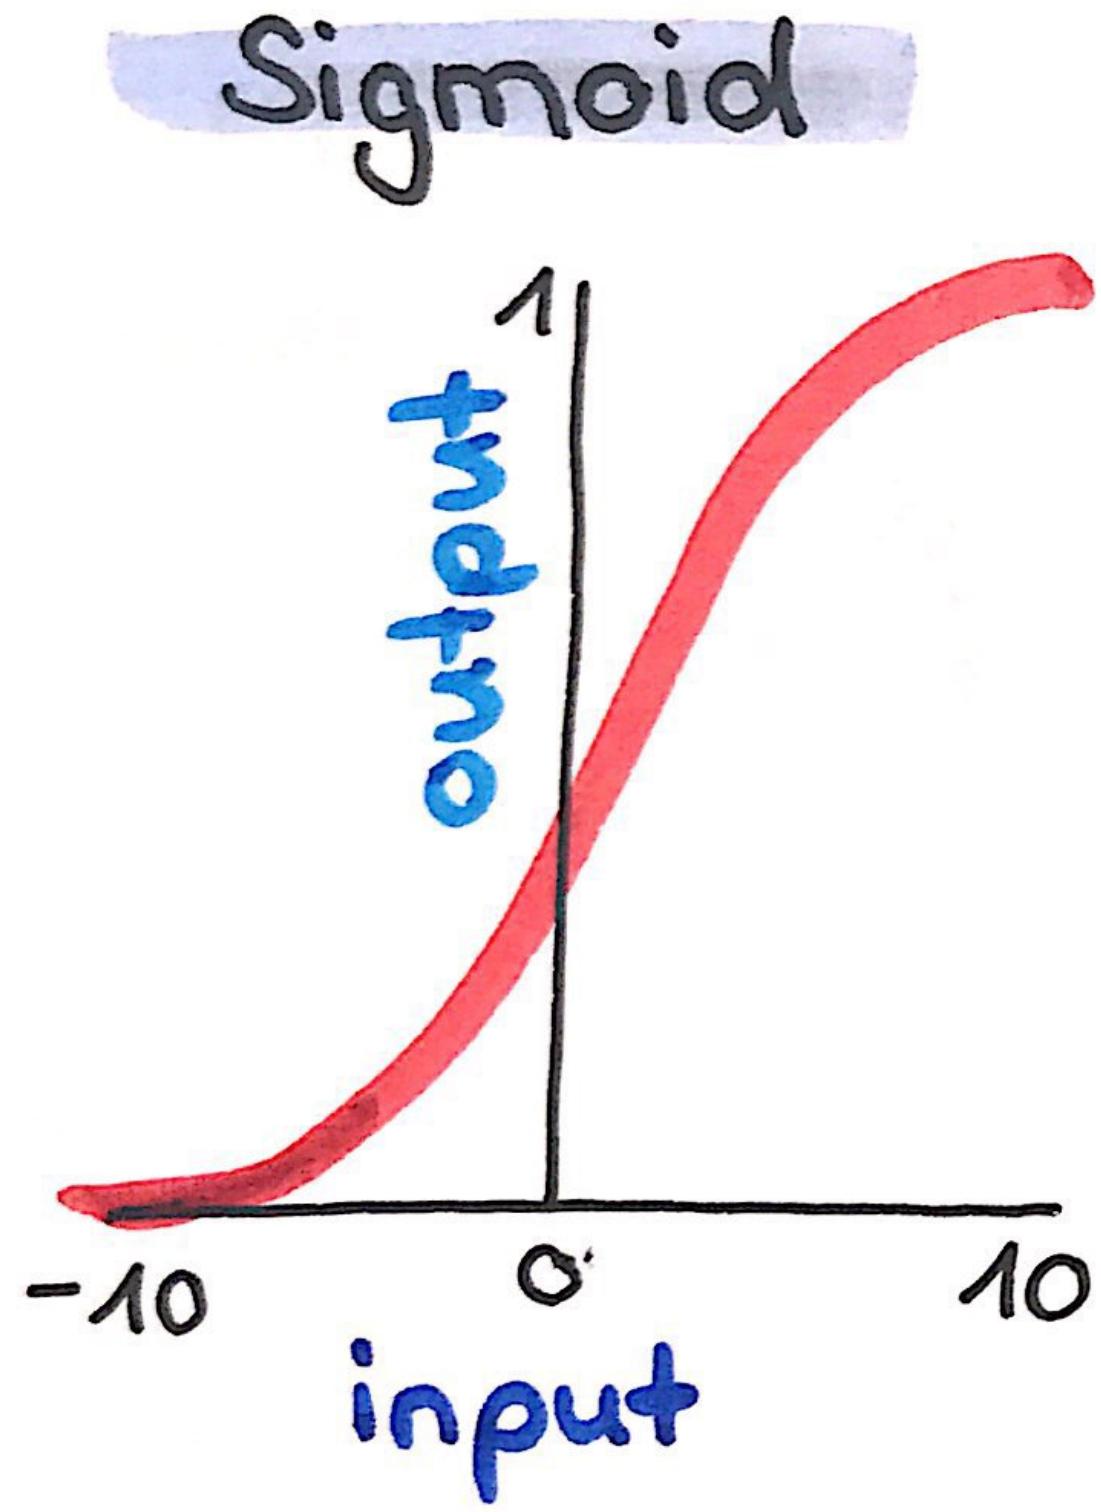


Perceptrons



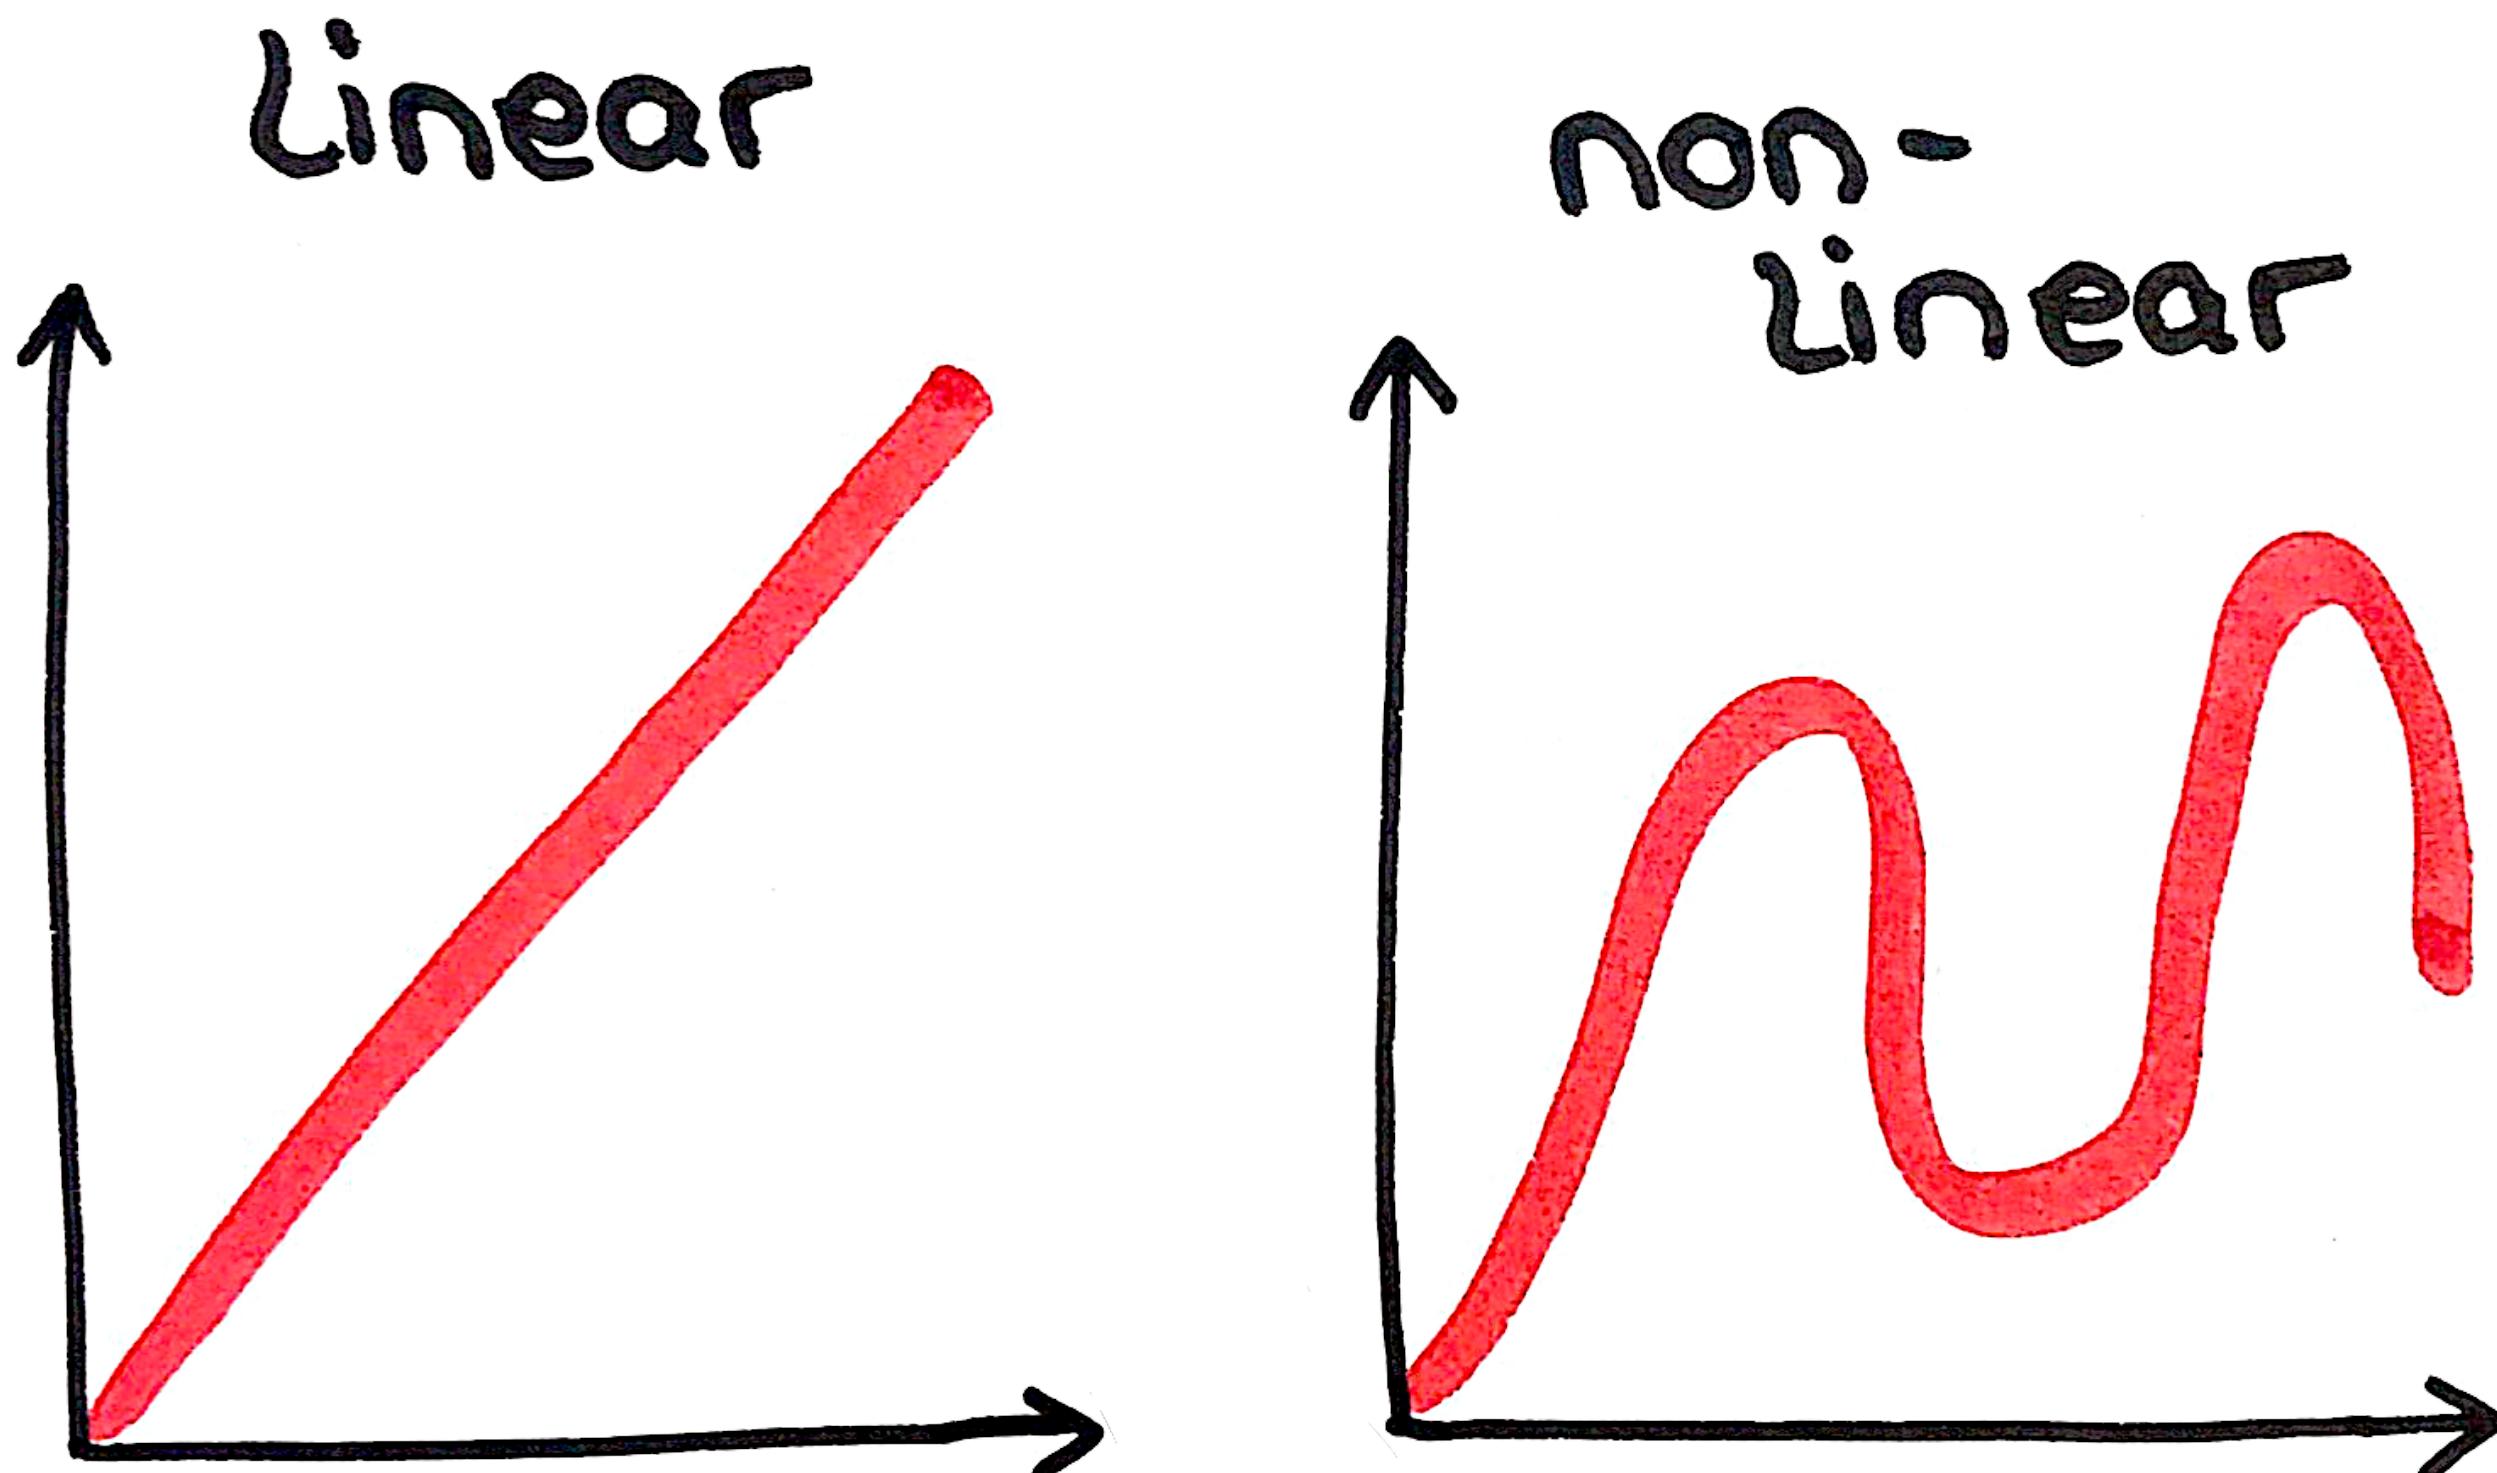
Activation functions normalise input

- every problem can be described with a mathematical function

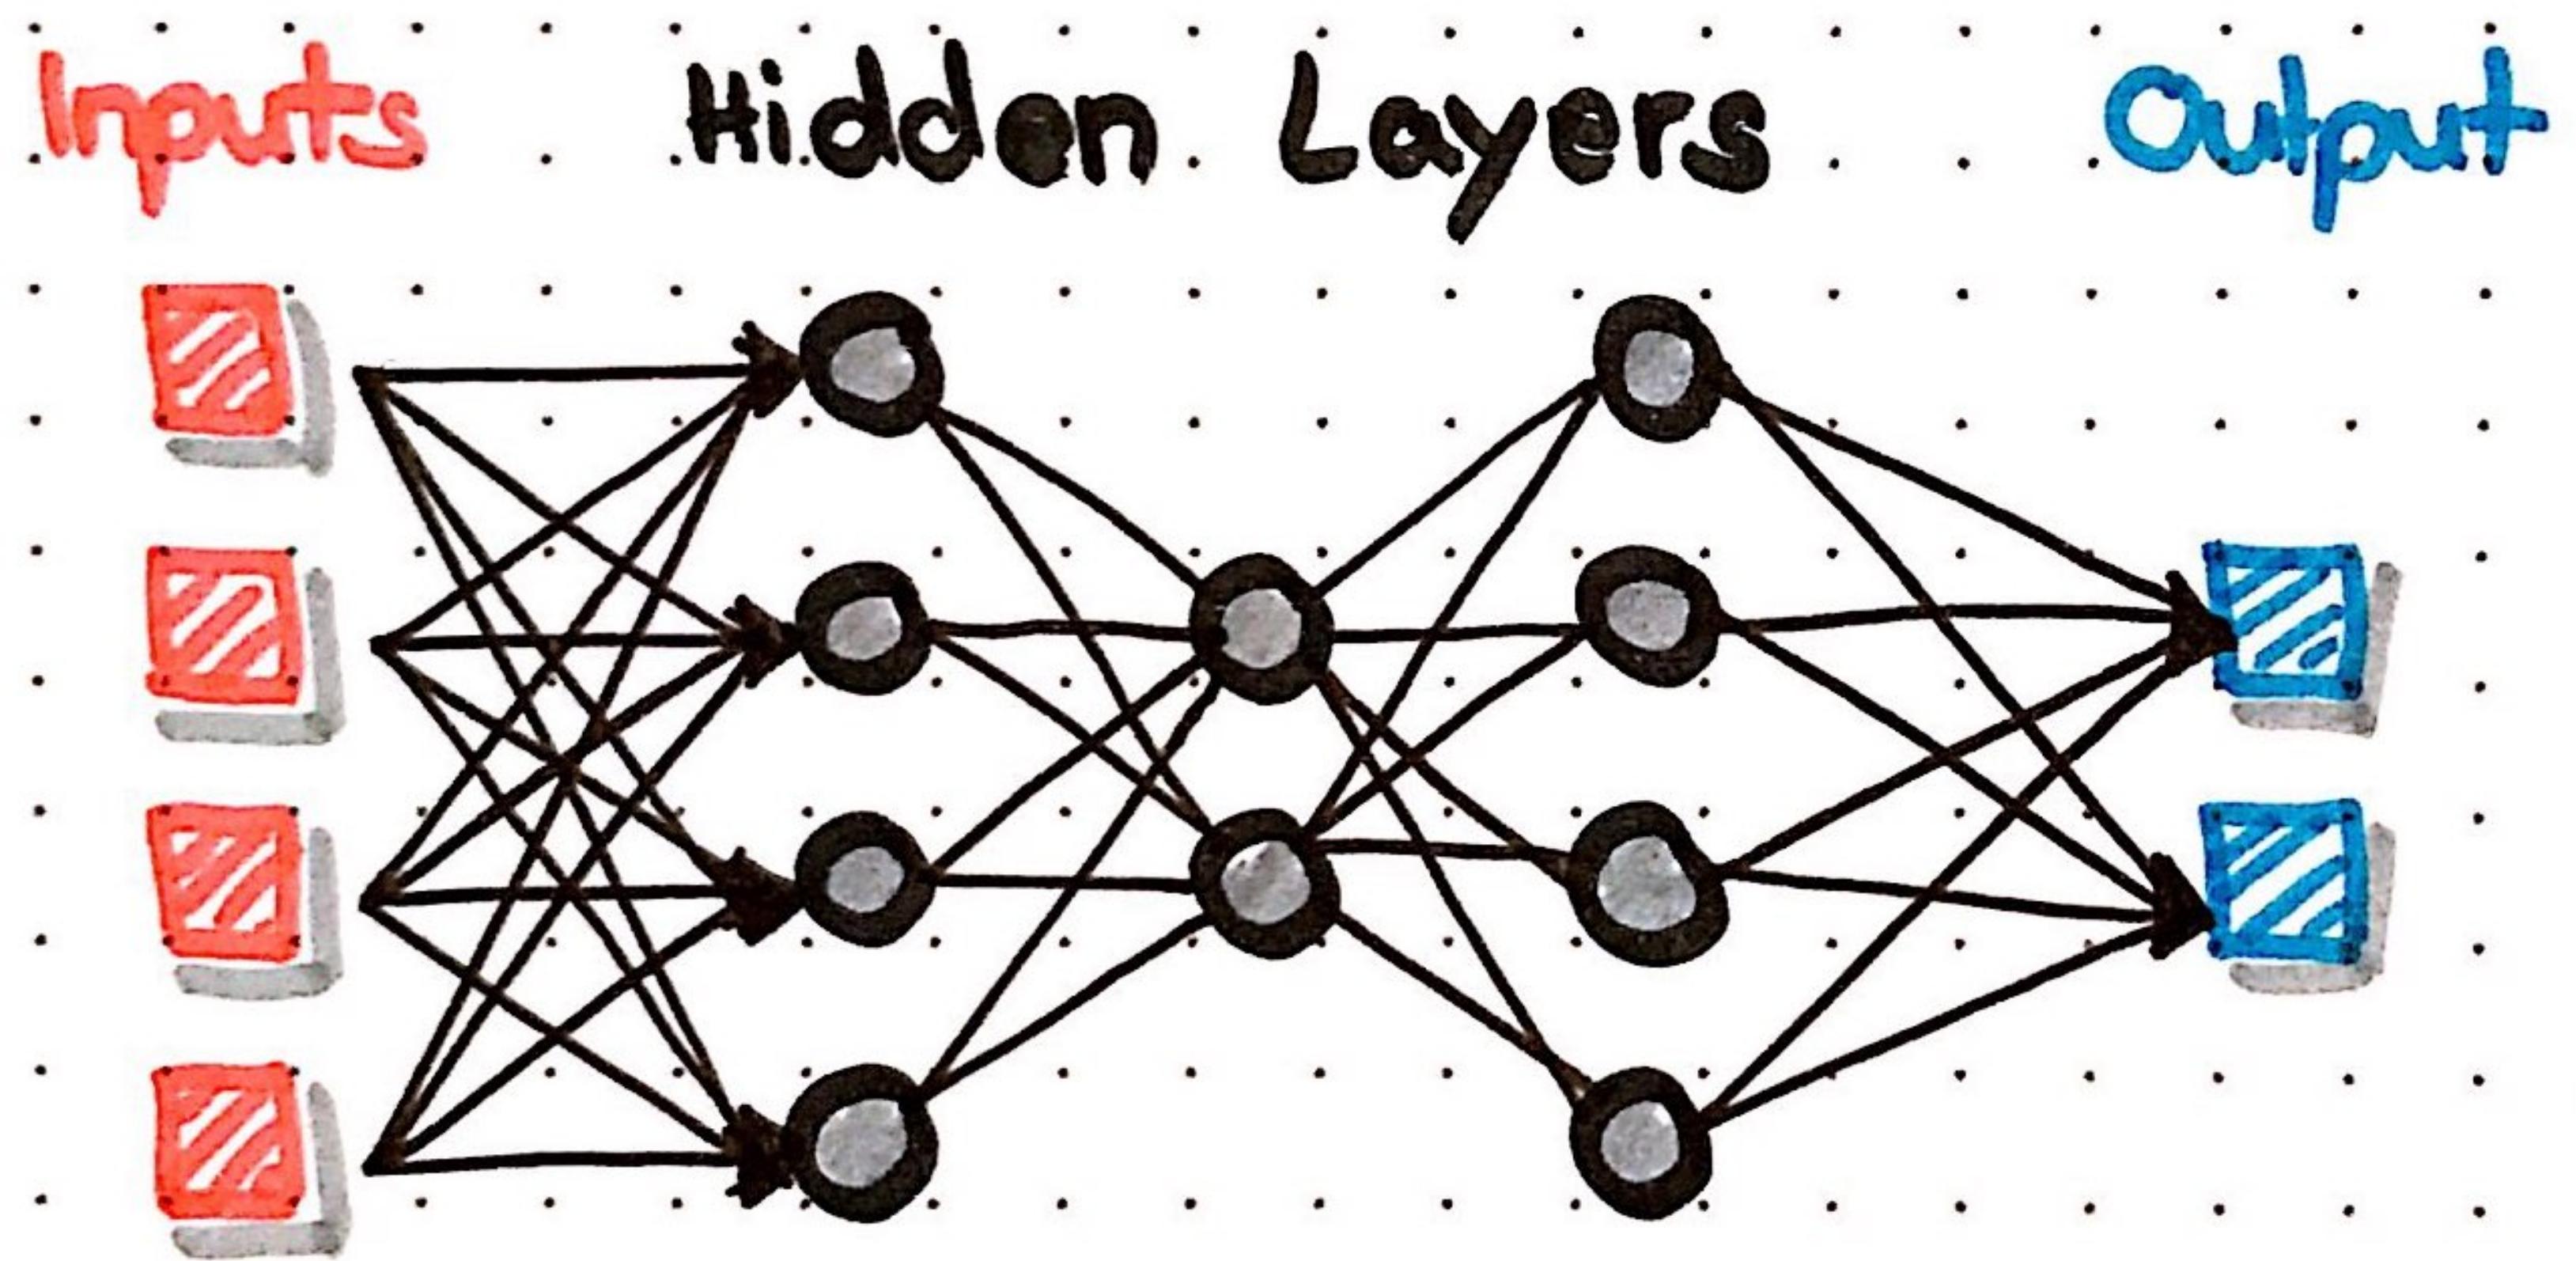


Activation functions make non-linearity possible

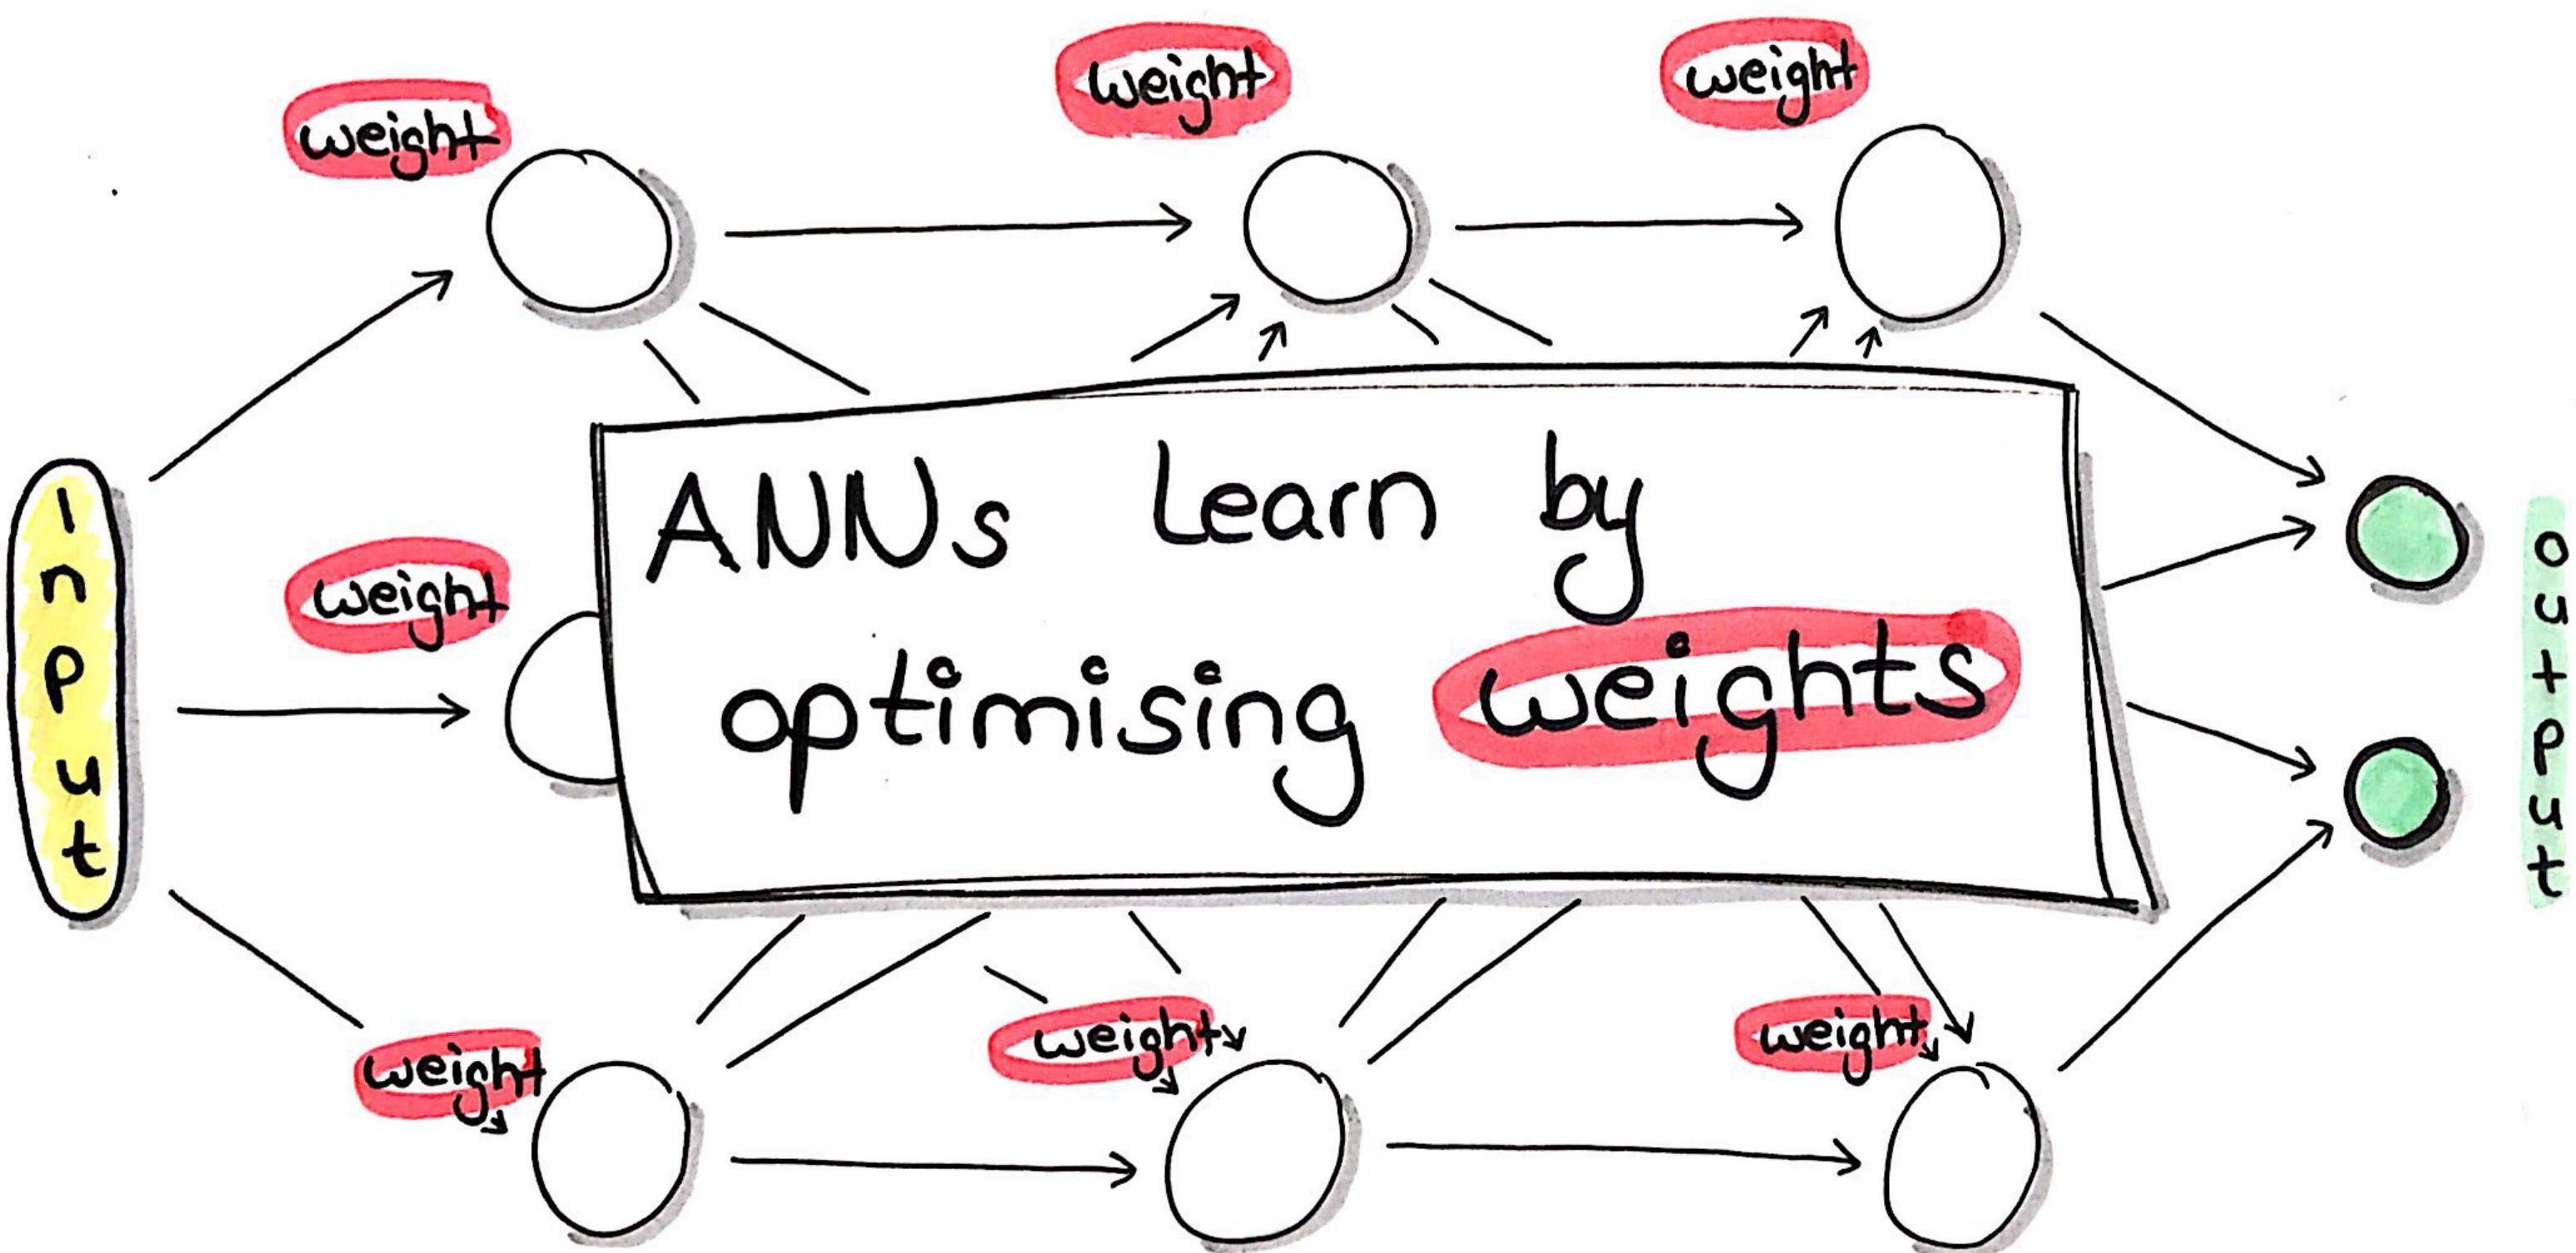
- non-linear activation functions allow us to approximate ANY mathematical formula with neural nets



Multi-Layer Perceptrons



How does a neural net learn?



How does a neural net learn?

The Softmax function

$$x * \omega + b = y$$

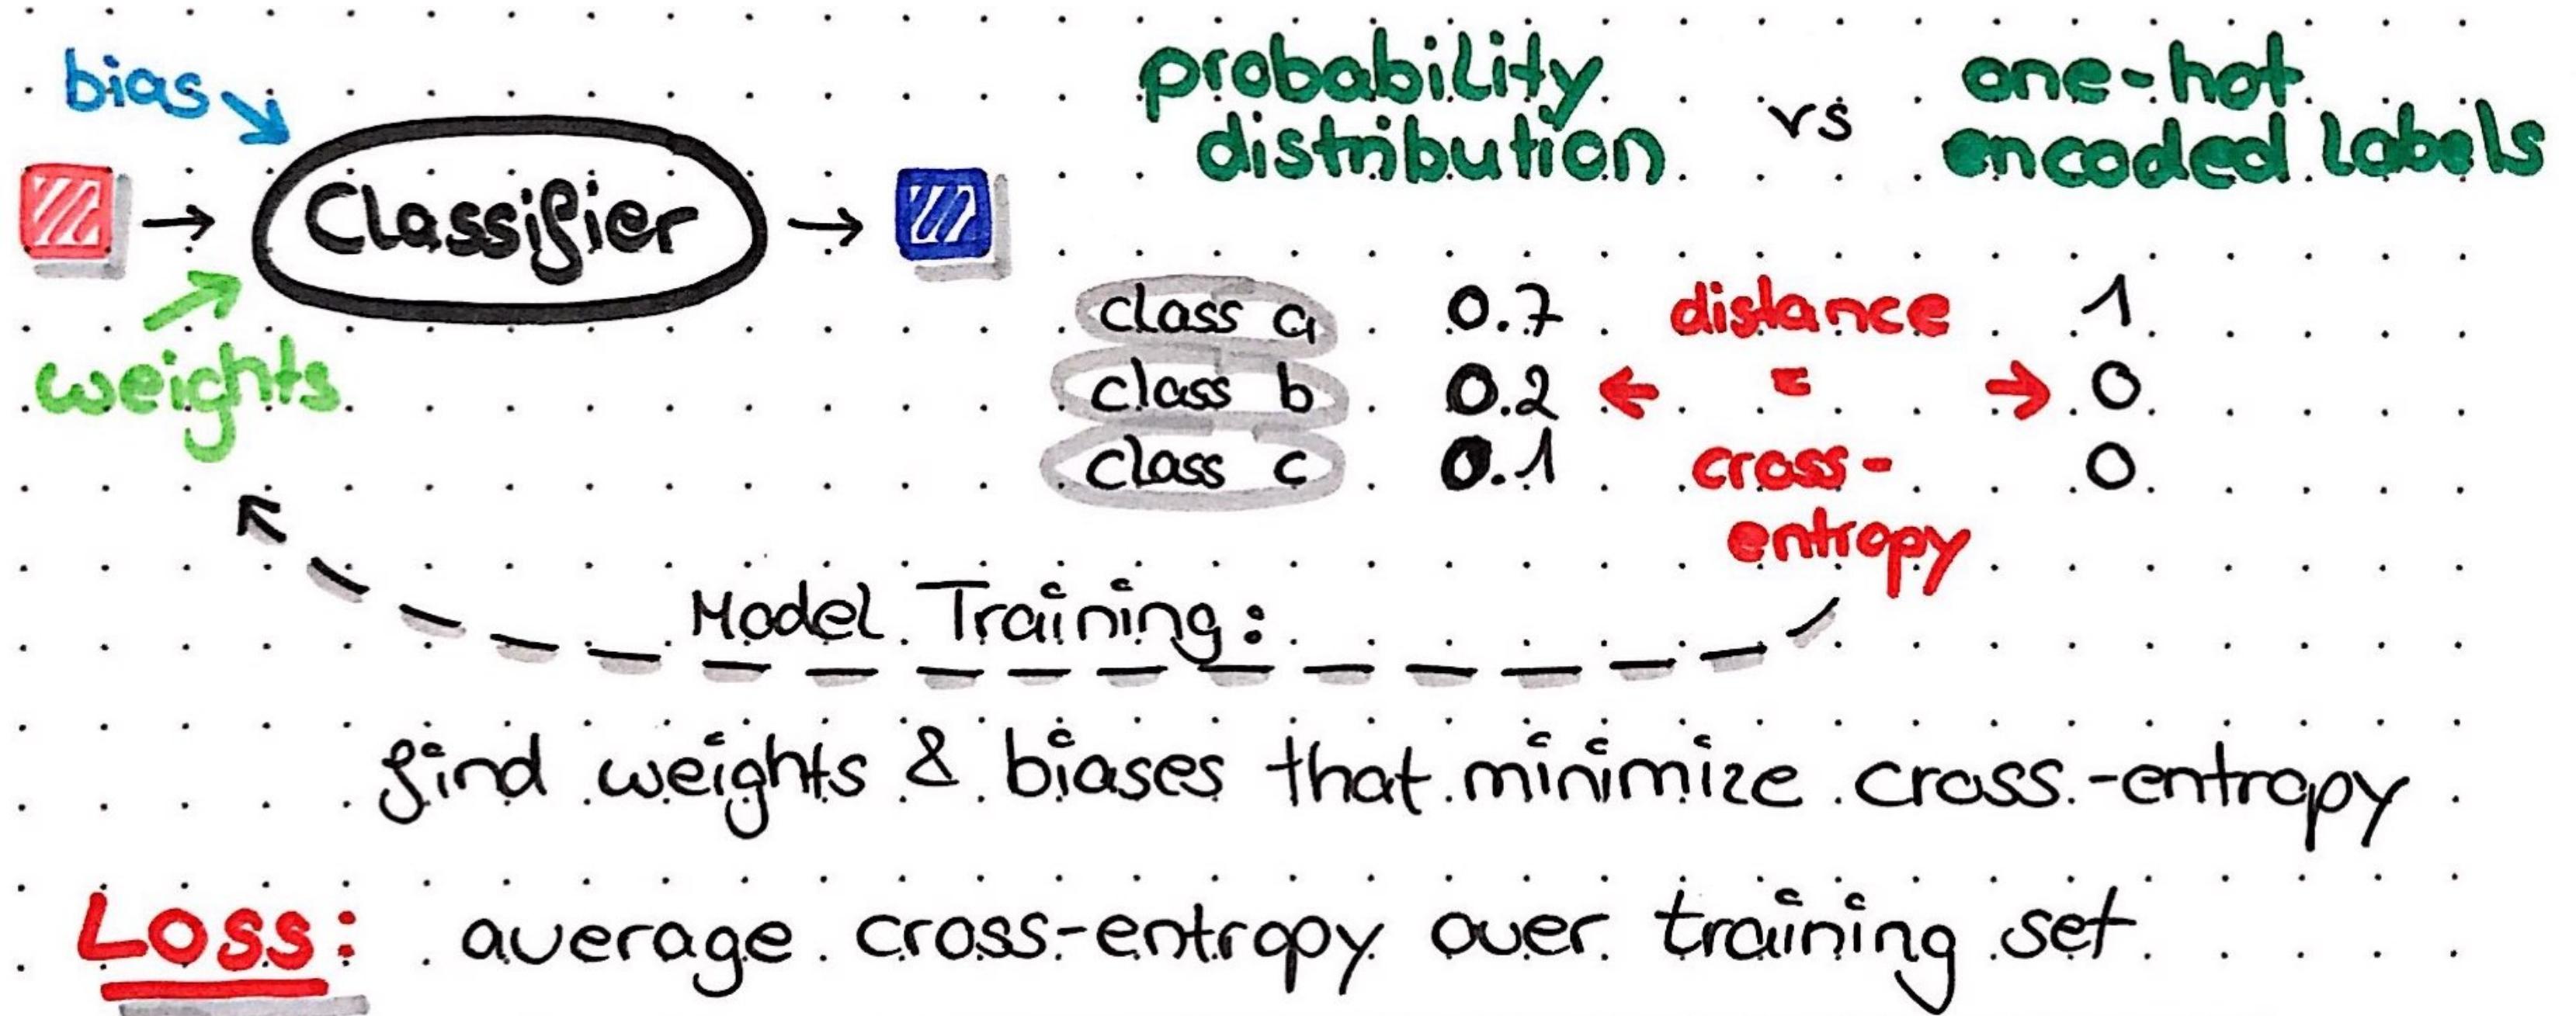
input weight bias Output

Model Training ~ finding good weights & biases

	score	probability	
class a	2	→ 0.7	correct
class b	1	→ 0.2	
class c	0.1	→ 0.1	
			soft-max

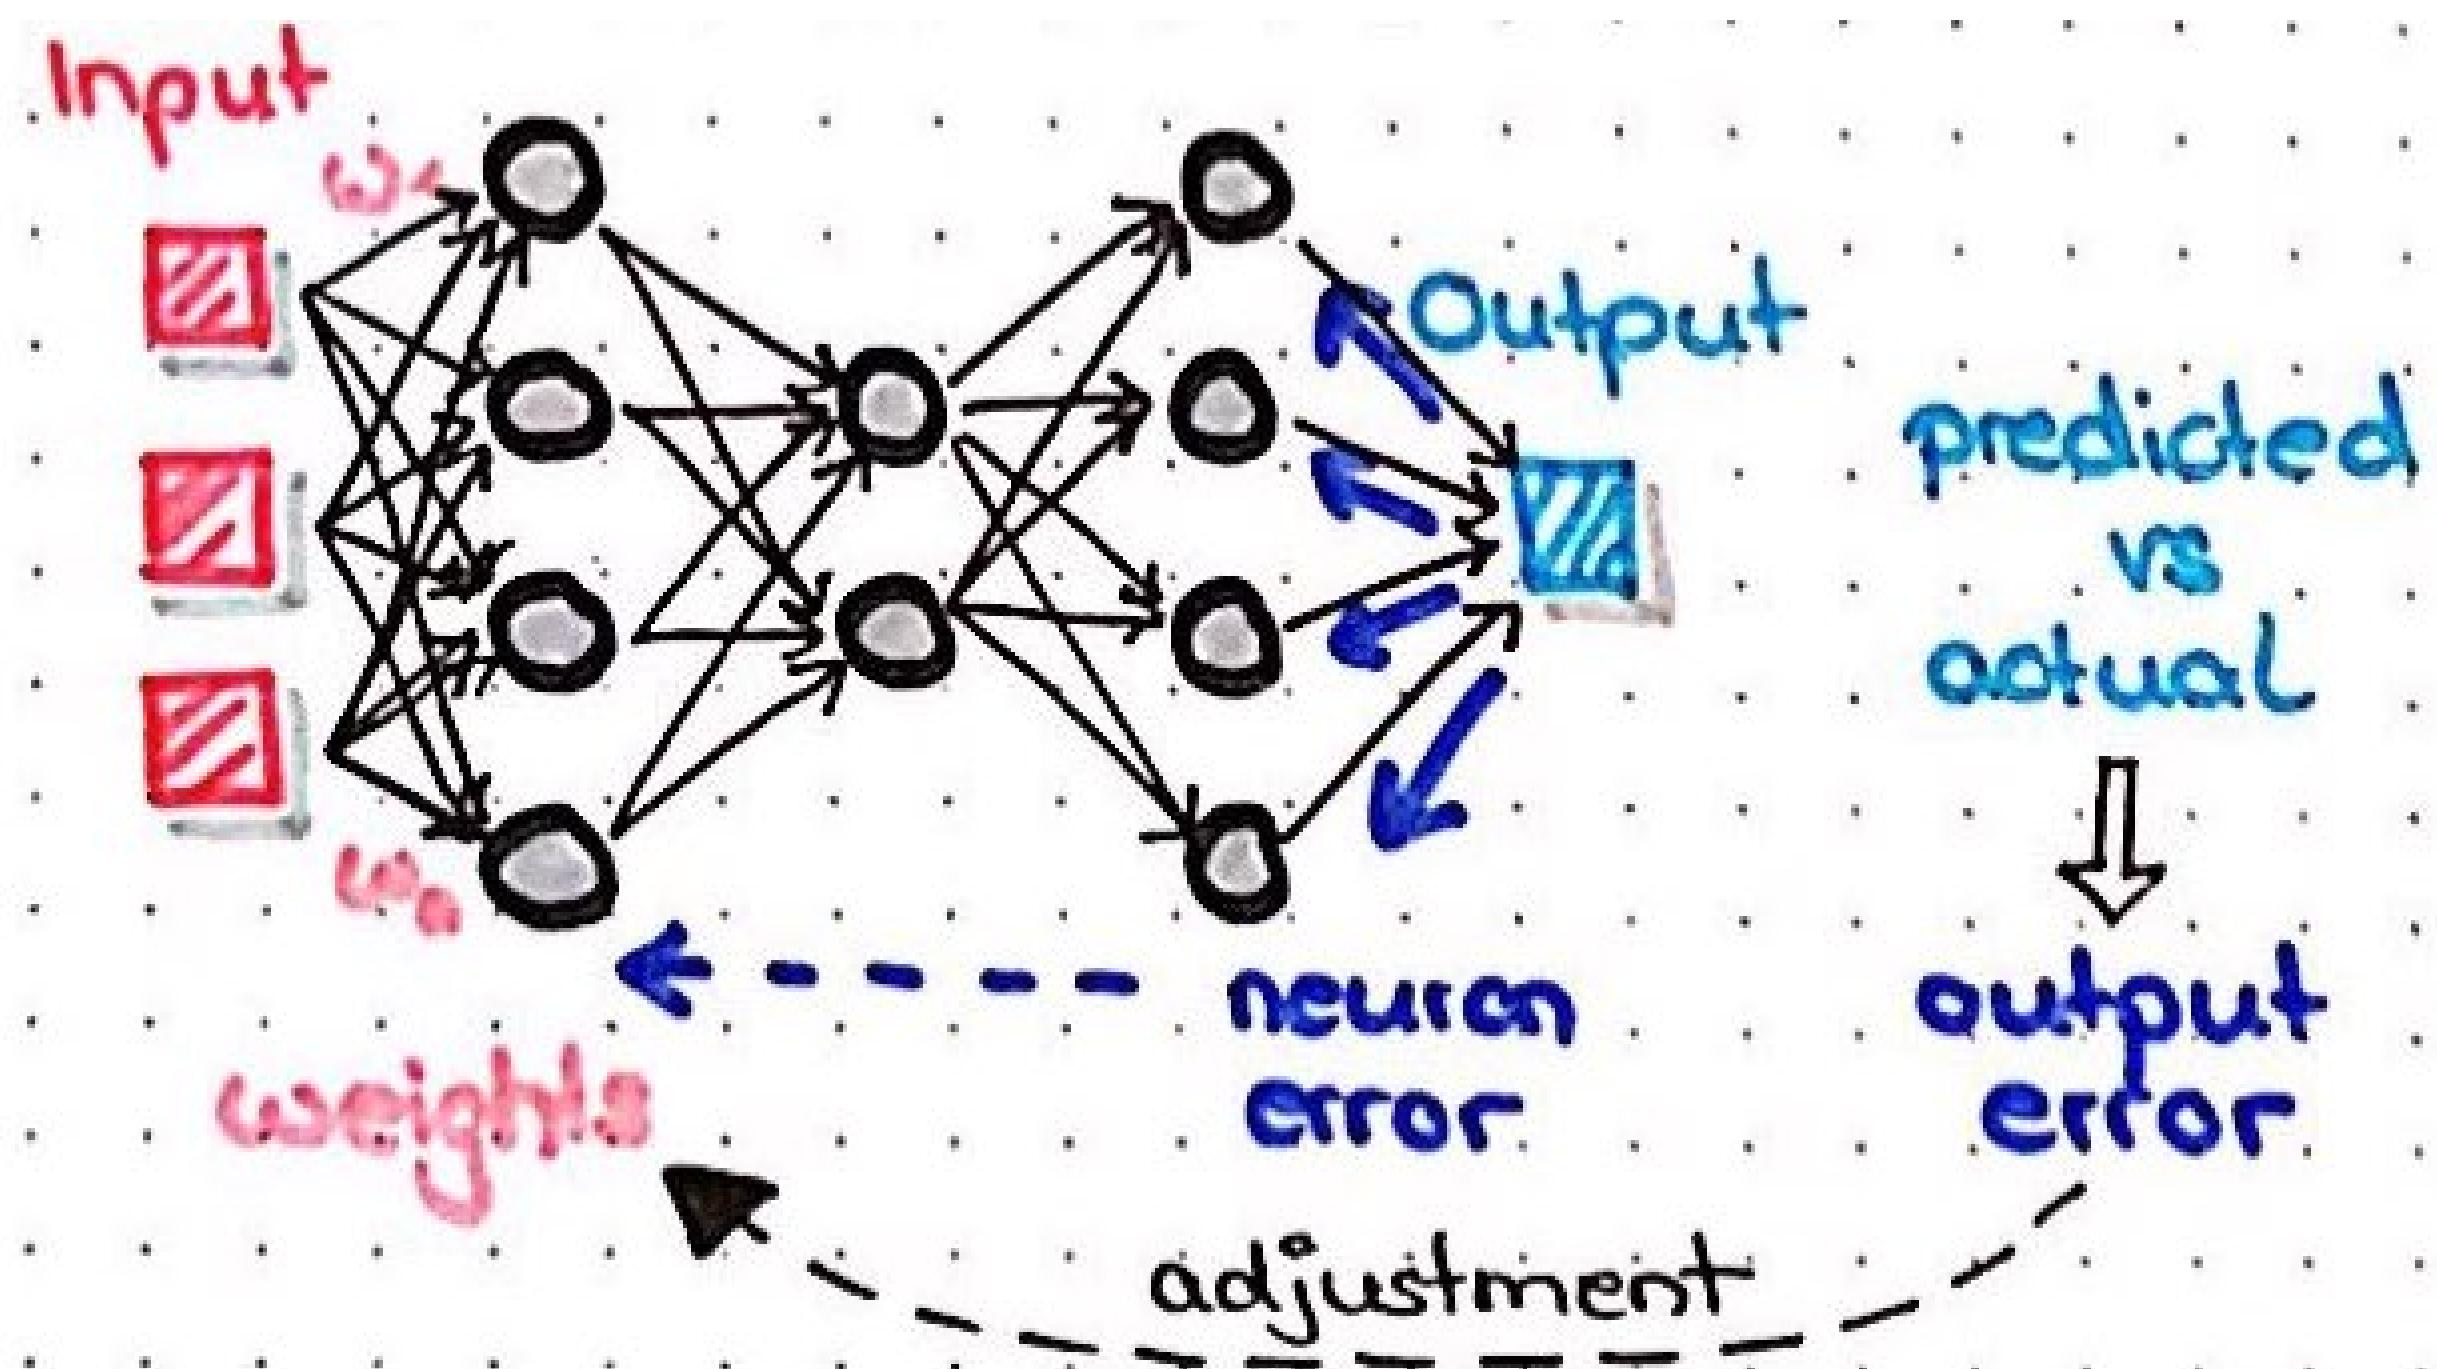
How does a neural net learn?

Cross-entropy



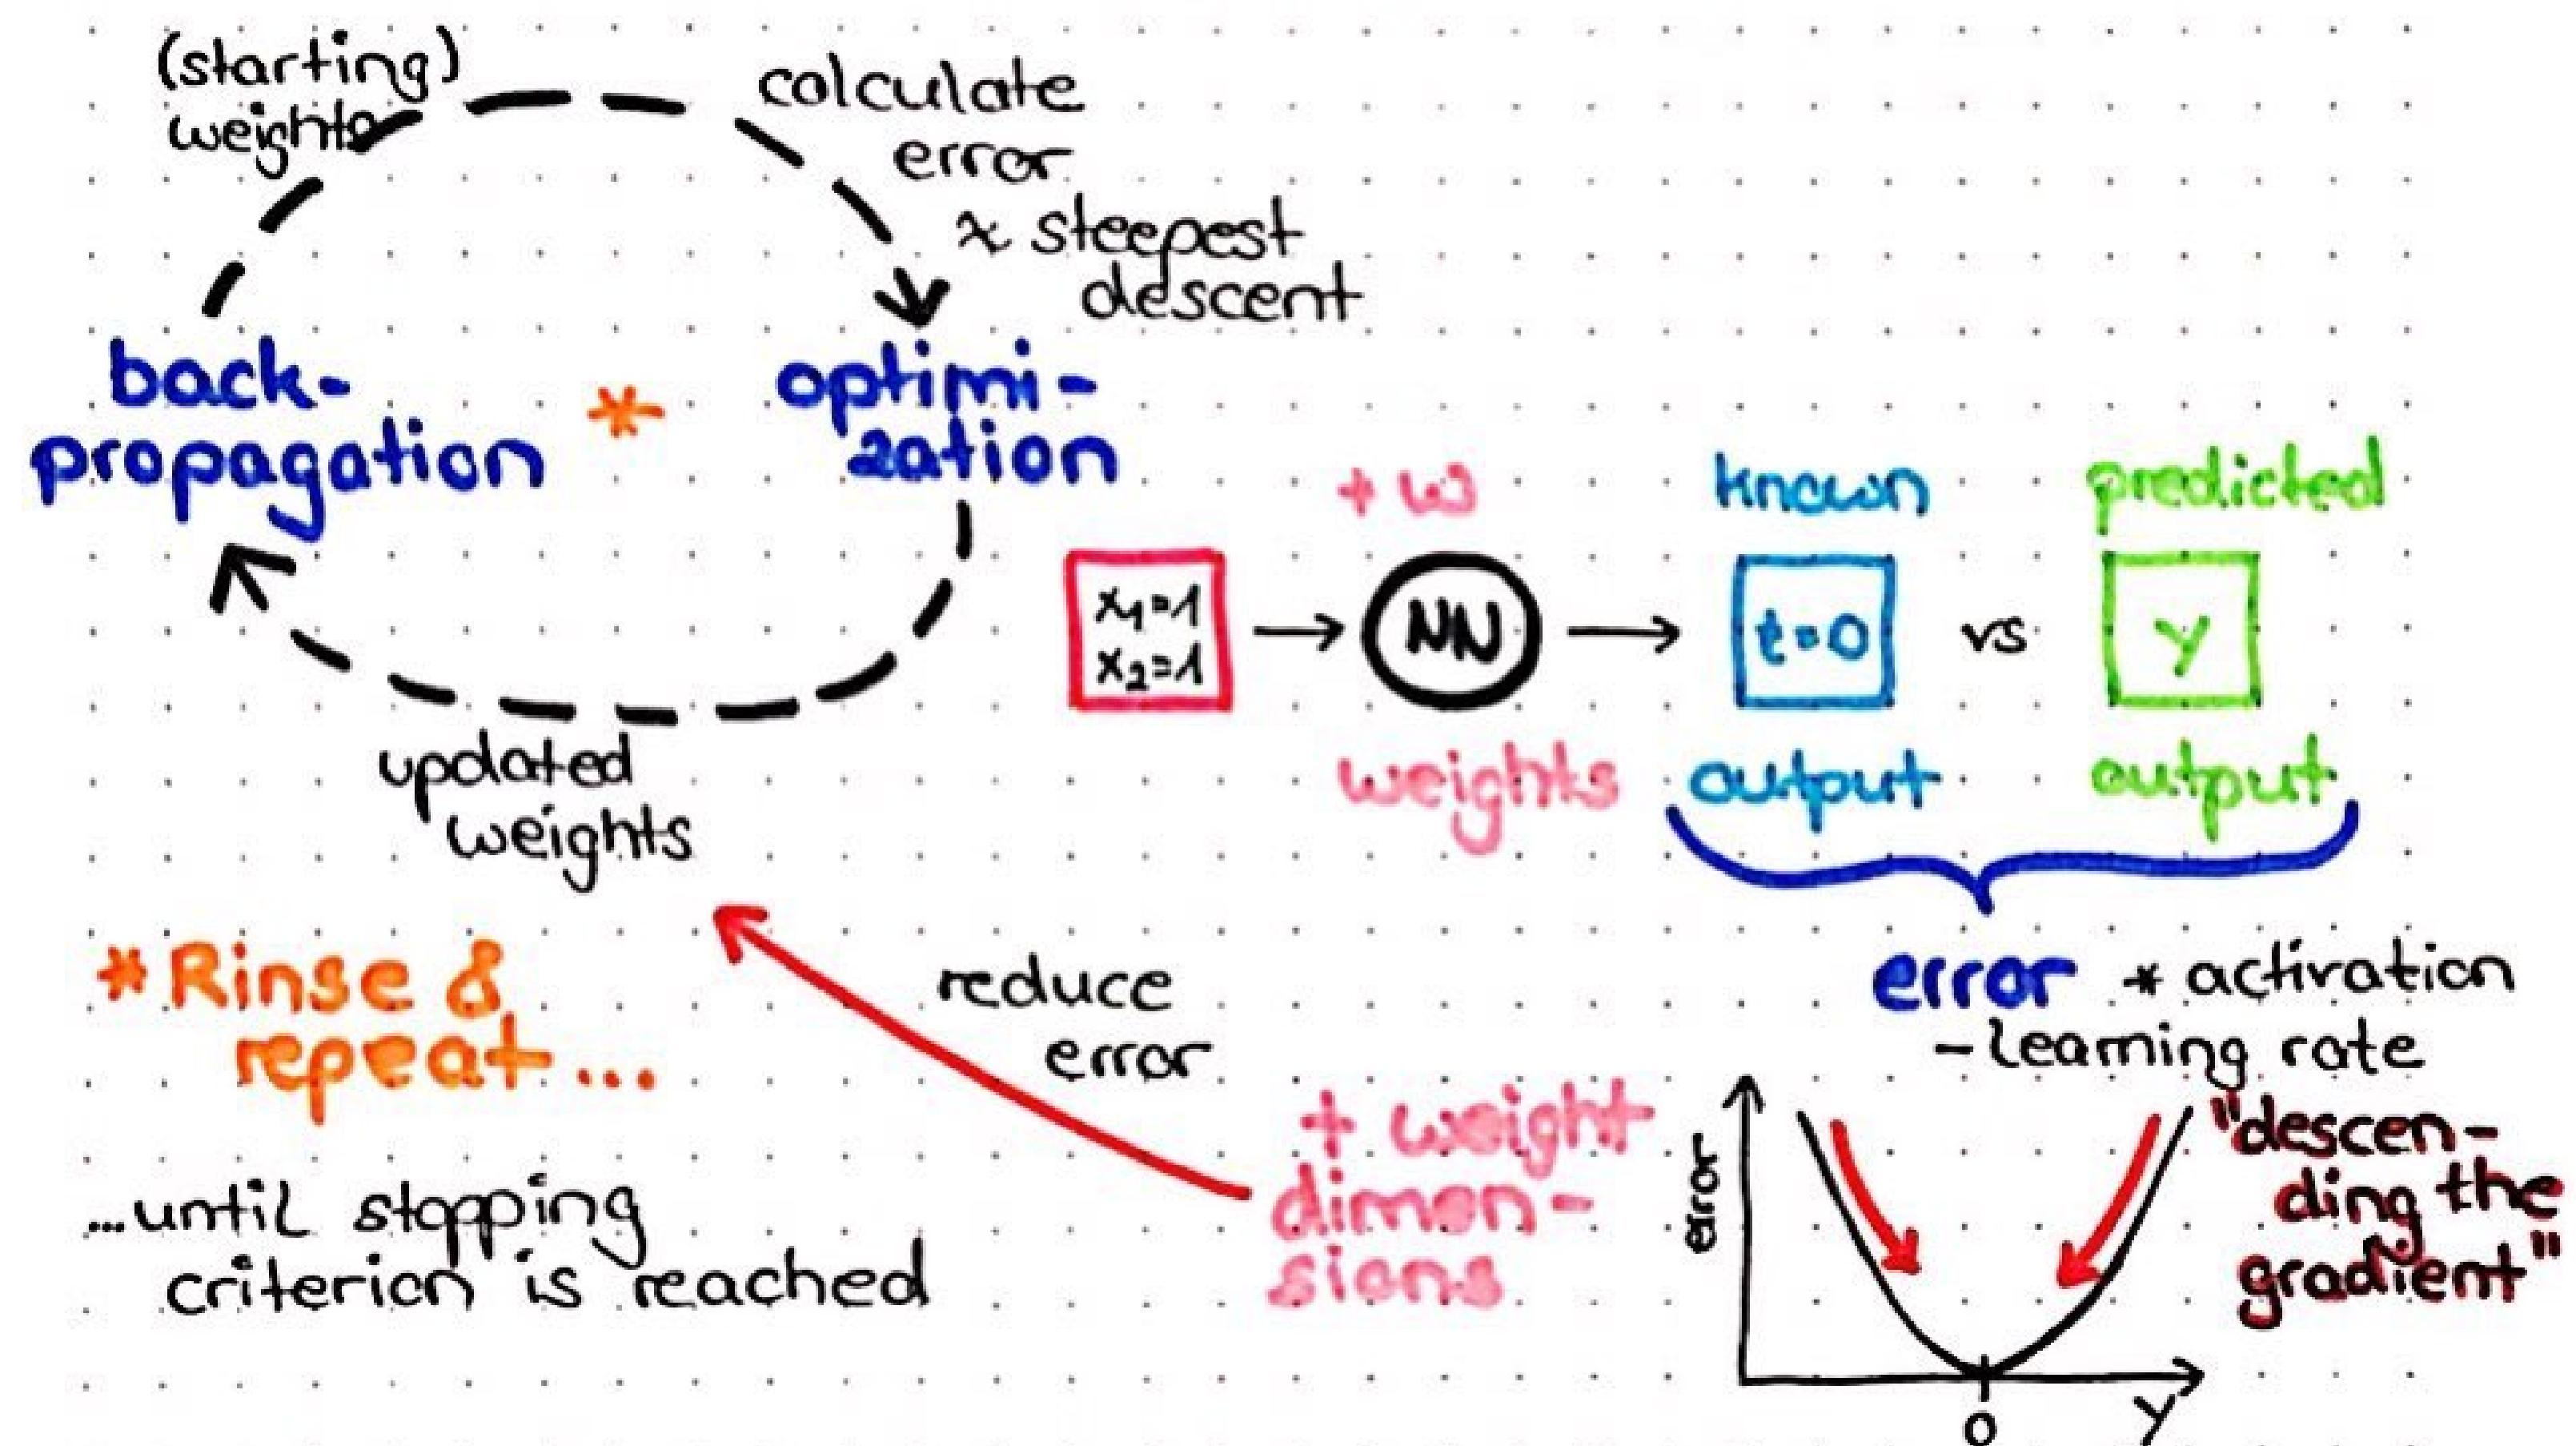
How does a neural net learn?

Backpropagation

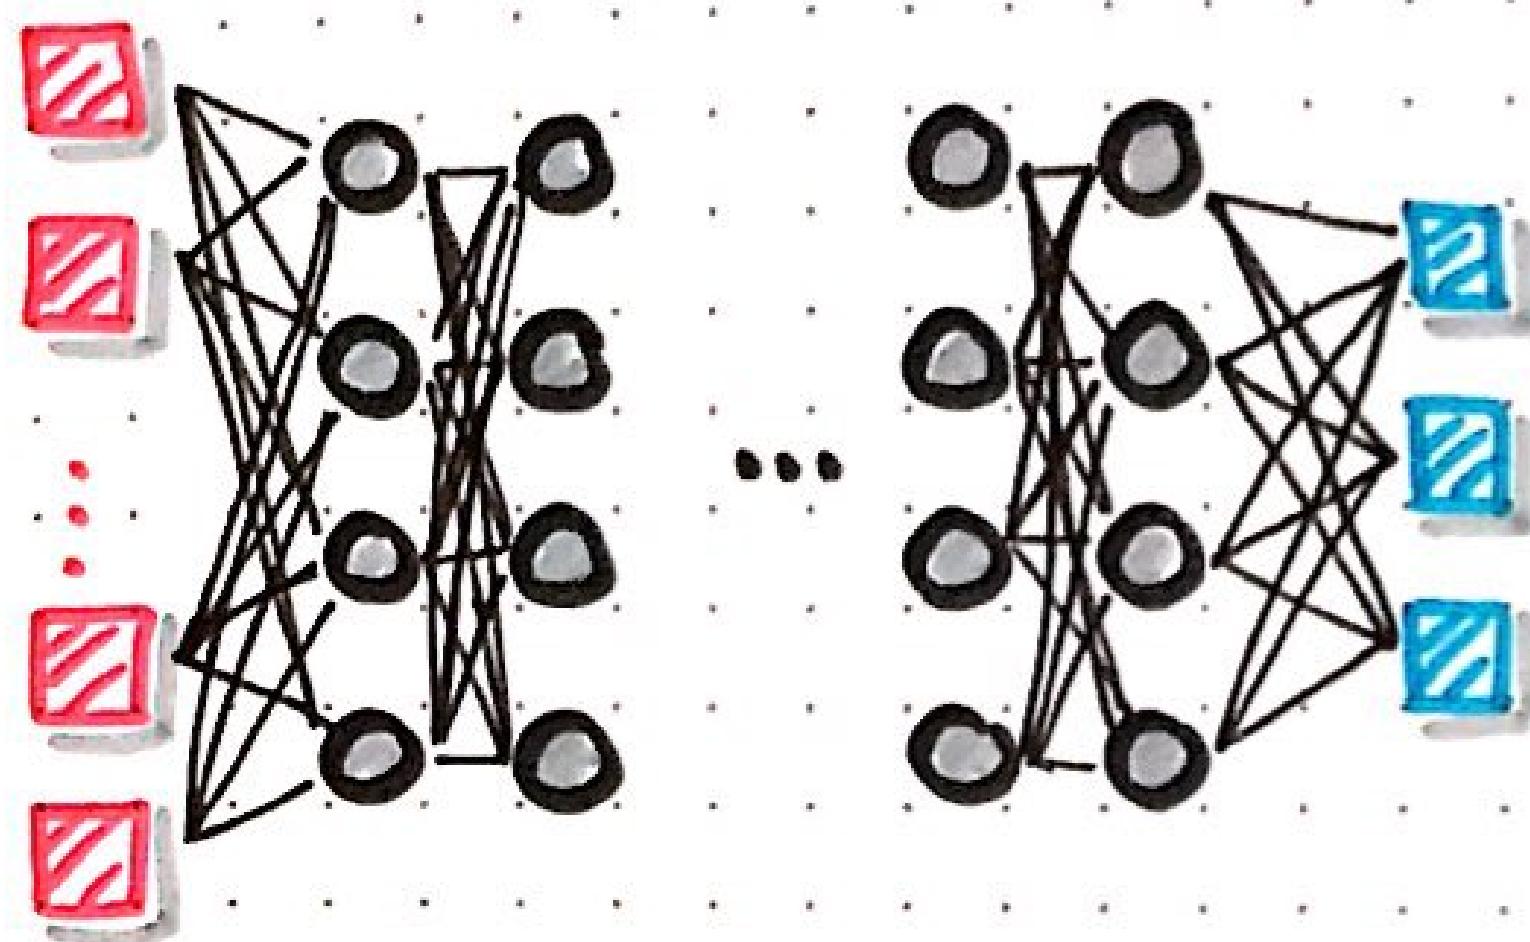


How does a neural net learn?

Gradient descent optimization



Deep Learning



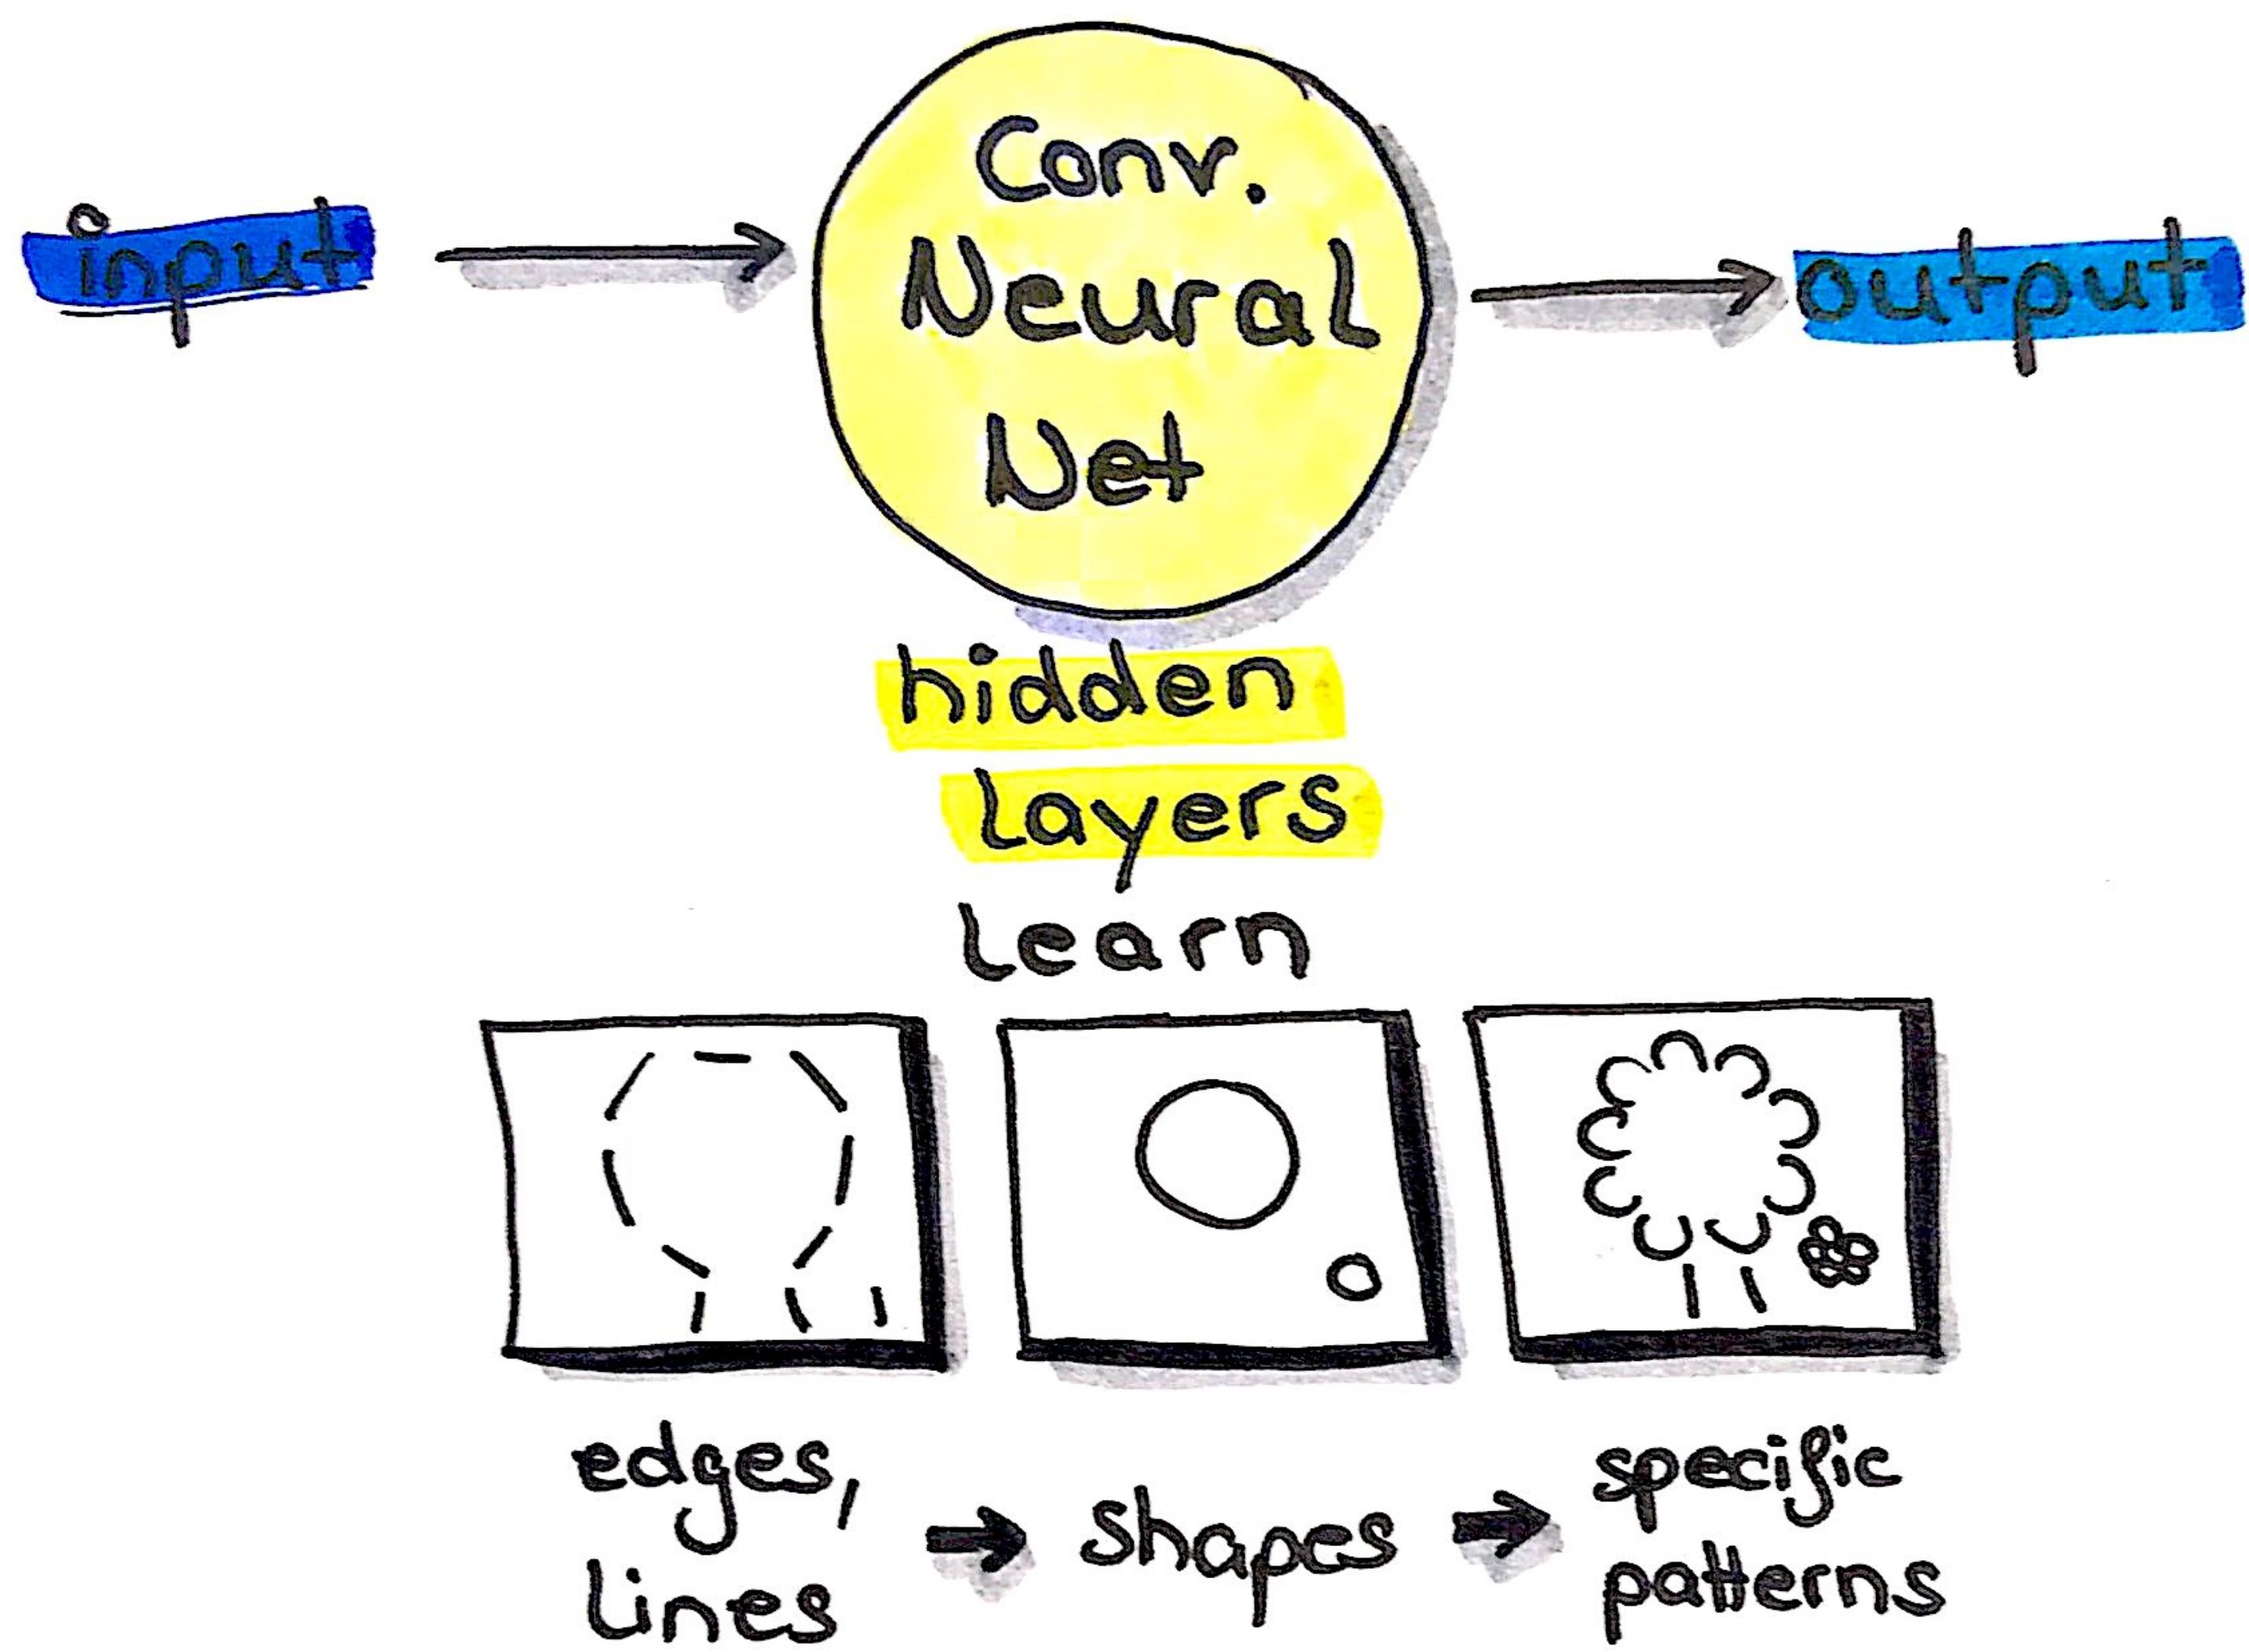
- supervised , semi - supervised or unsupervised
- NLP , speech recognition
- image recognition
- object classification
- recommender systems
- etc.

Convolutional Neural Nets

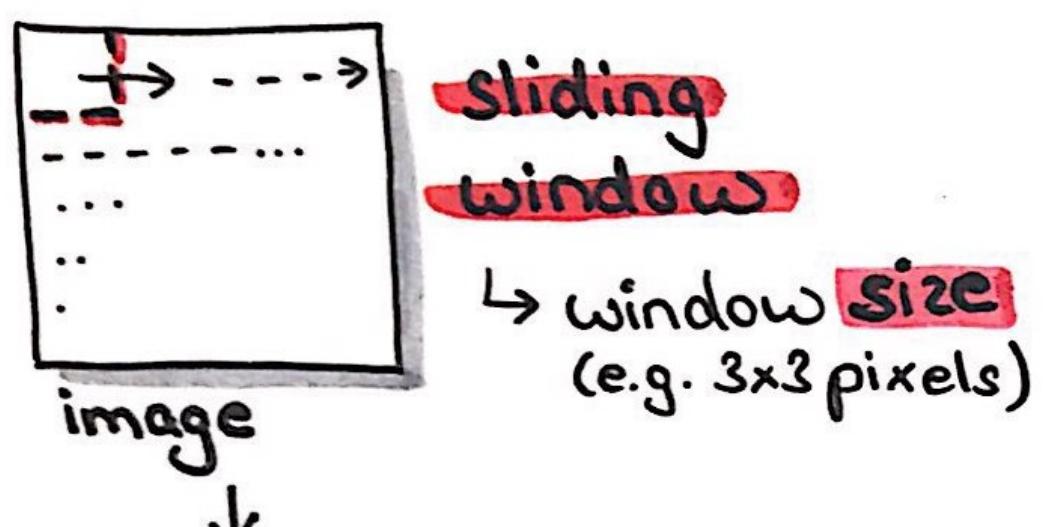
MLPs vs CNNs

- pixels are considered independent
- computationally faster
- Learned : weights
- pixels are considered as groups of connected information (context)
- analysed as chunks (= windows)
- Learned : filters

ConvNets



ConvNets

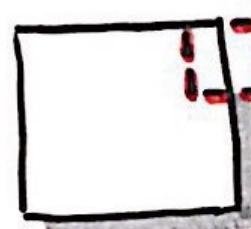


vertical Lines

horizontal Lines

- **Filters**
 - detect shapes & patterns in window chunks
 - multiple filters are combined
 - filters are learned

padding



- fake pixels are created at the edge of images to incorporate border pixels

Convolutional Layers

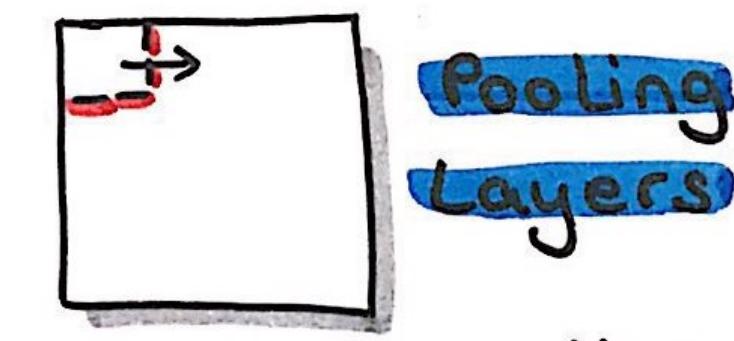
apply filters

↓

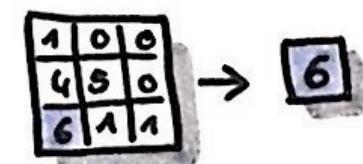
output

=
feature maps

→
stacks of feature maps



e.g. max pooling

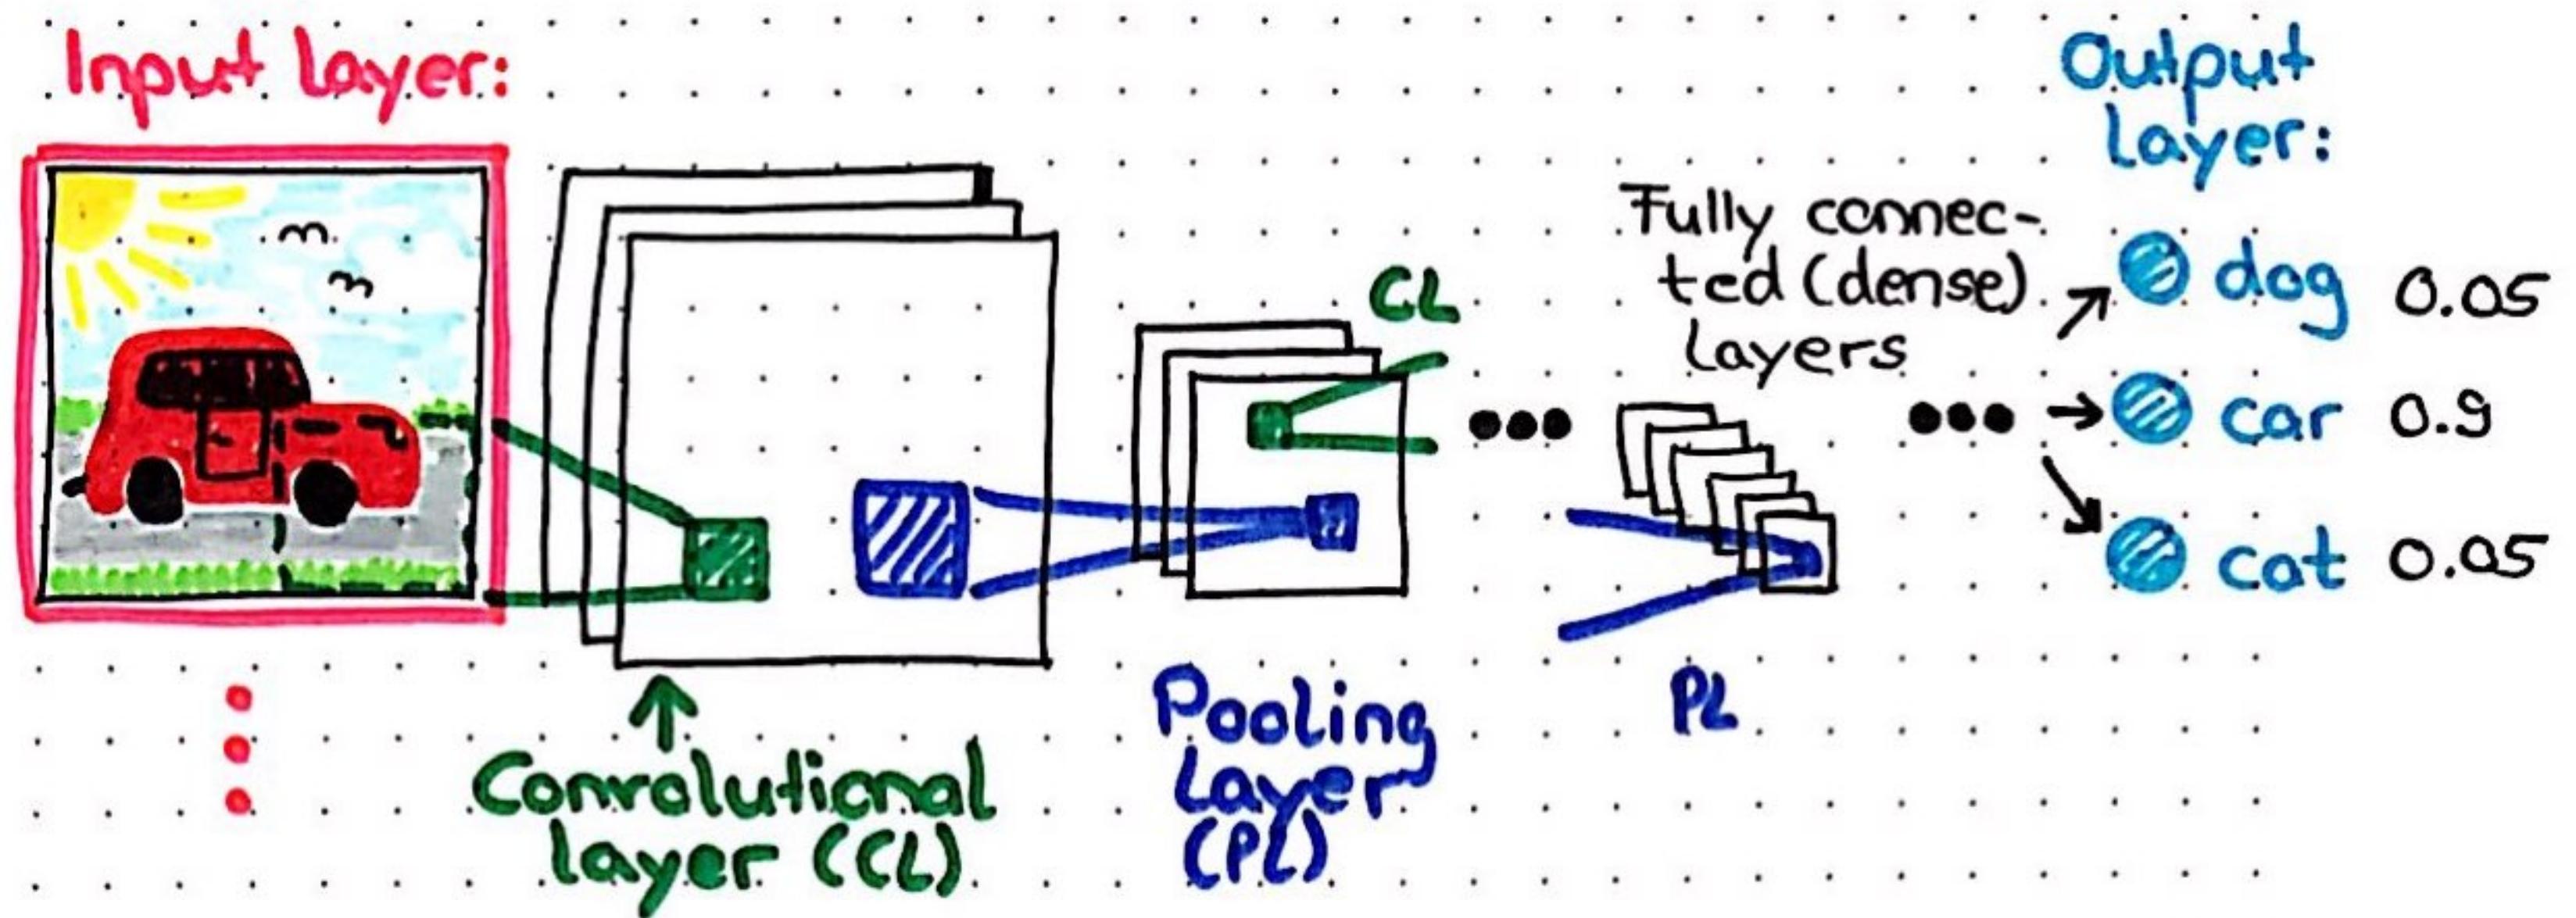


- reduces compute time
- boils information down

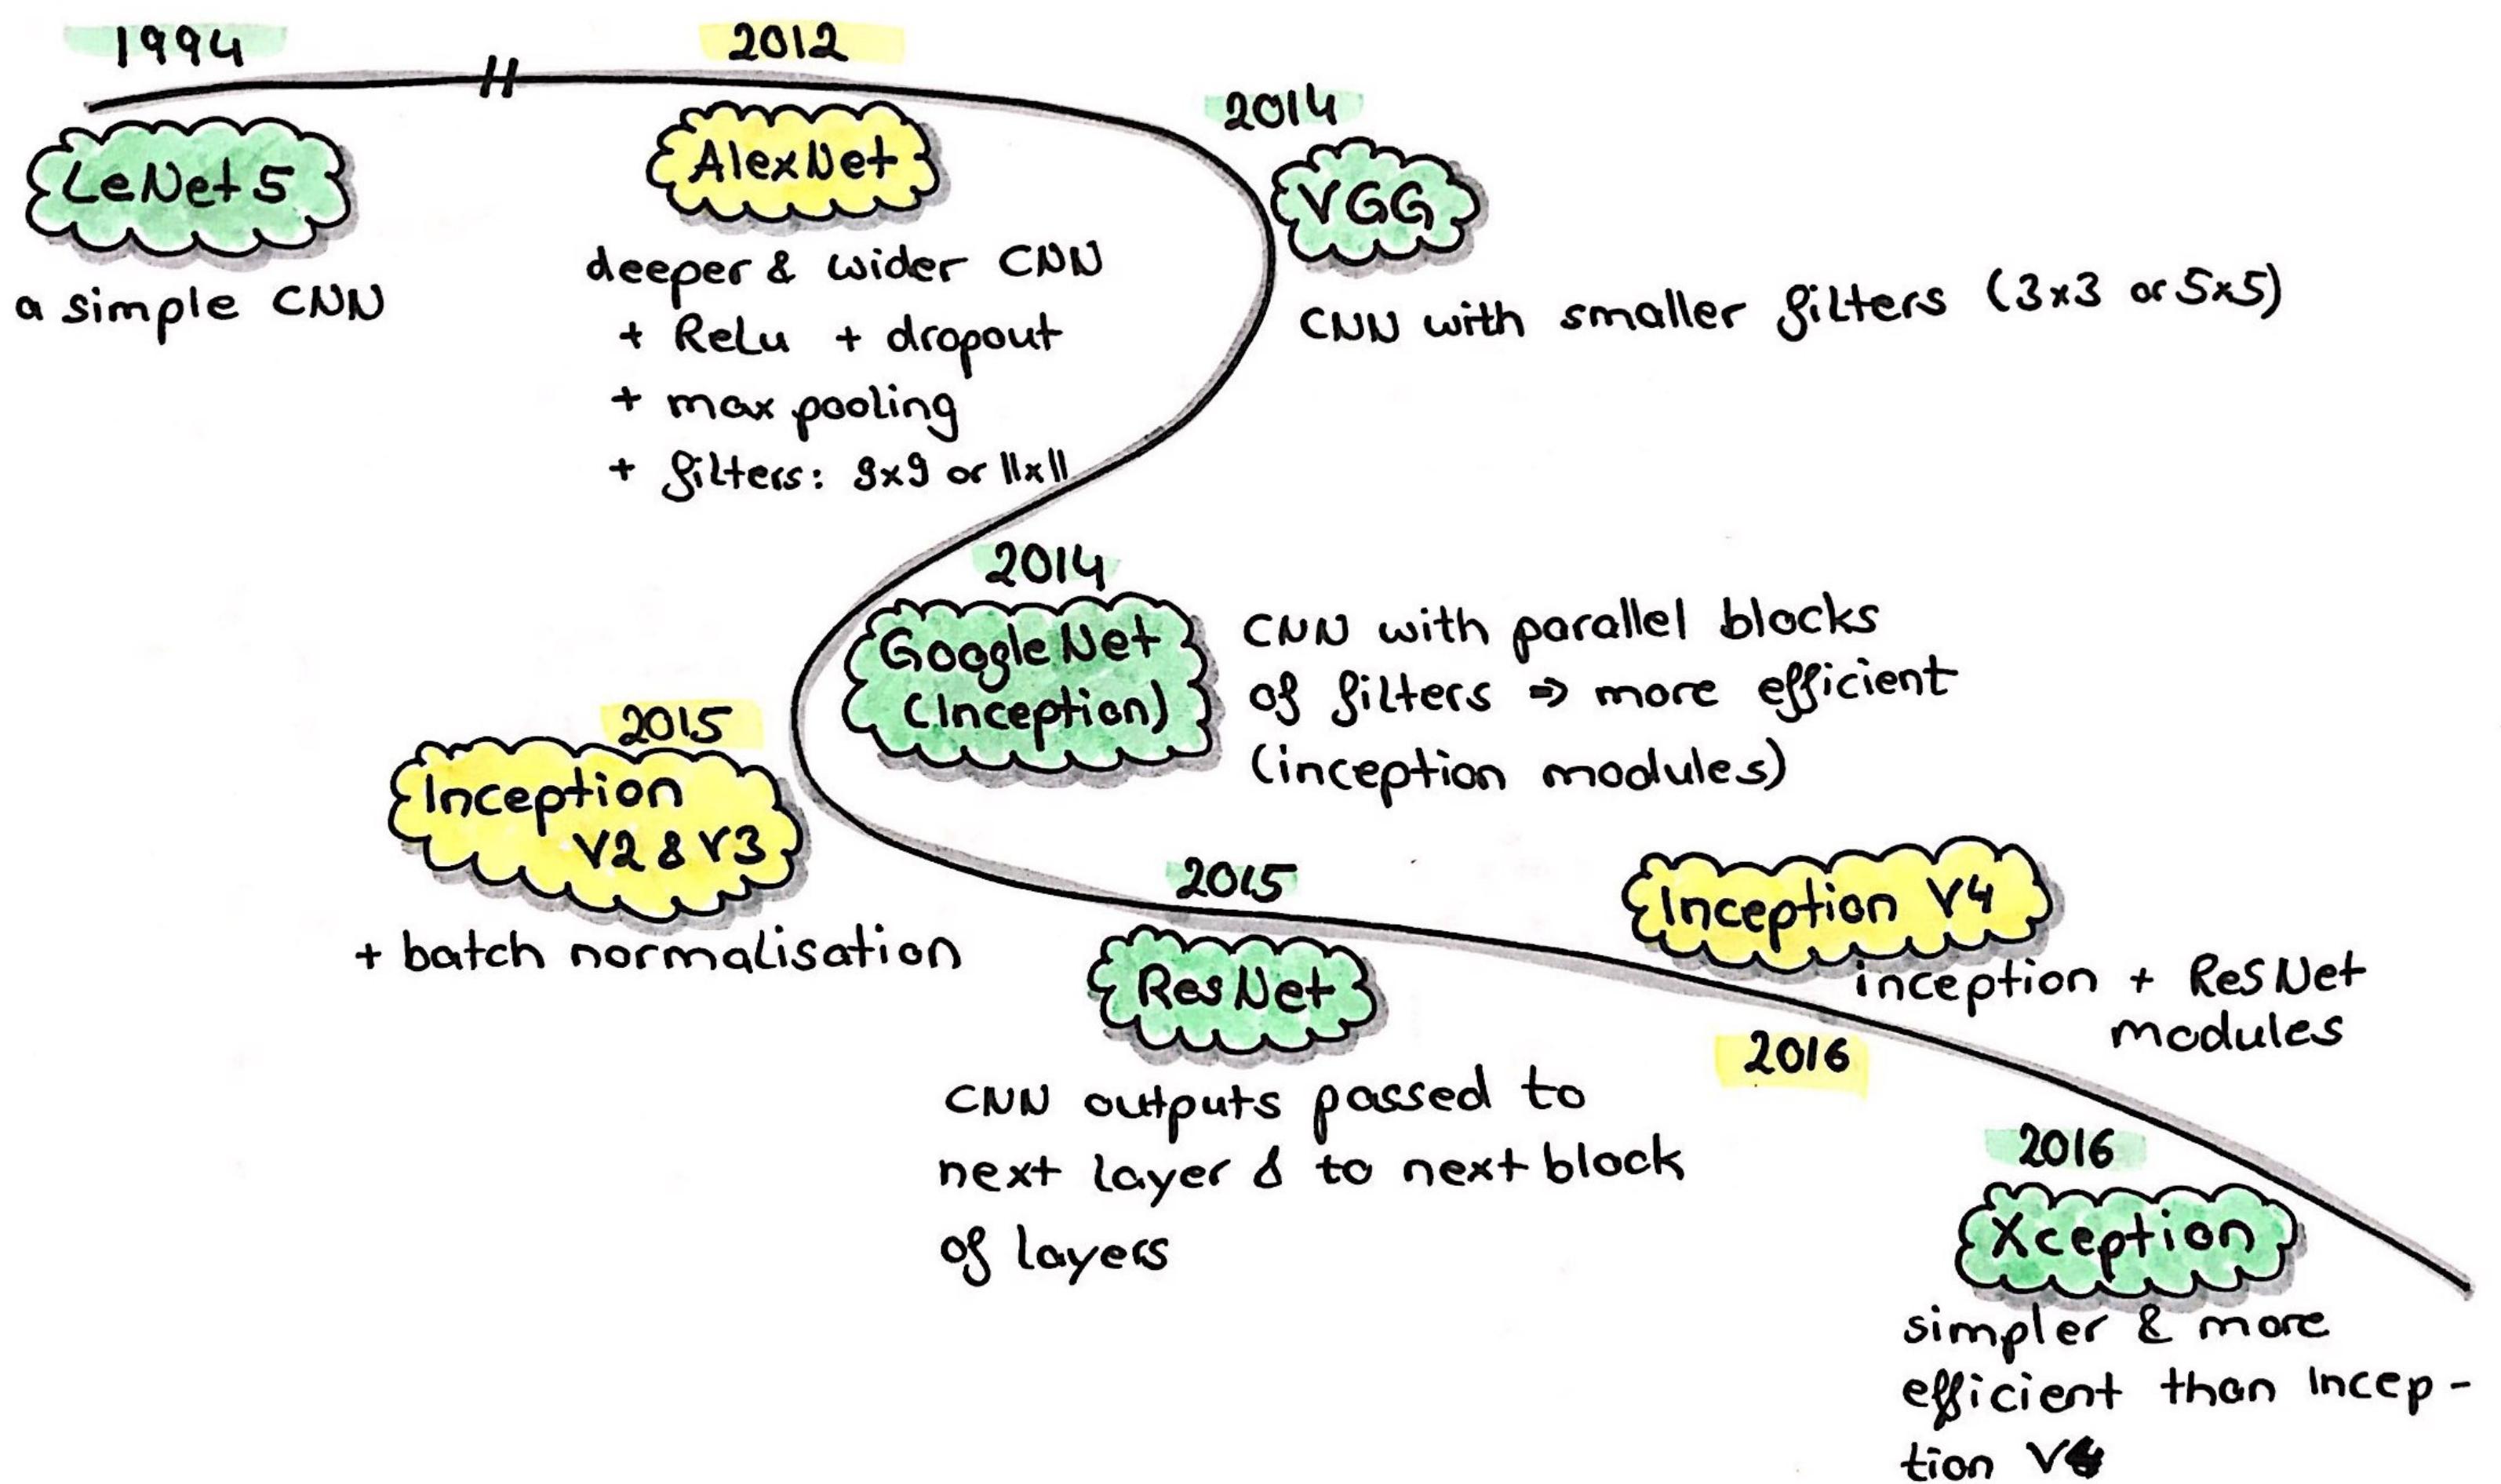
stride

- how much overlap to have in sliding window

ConvNets

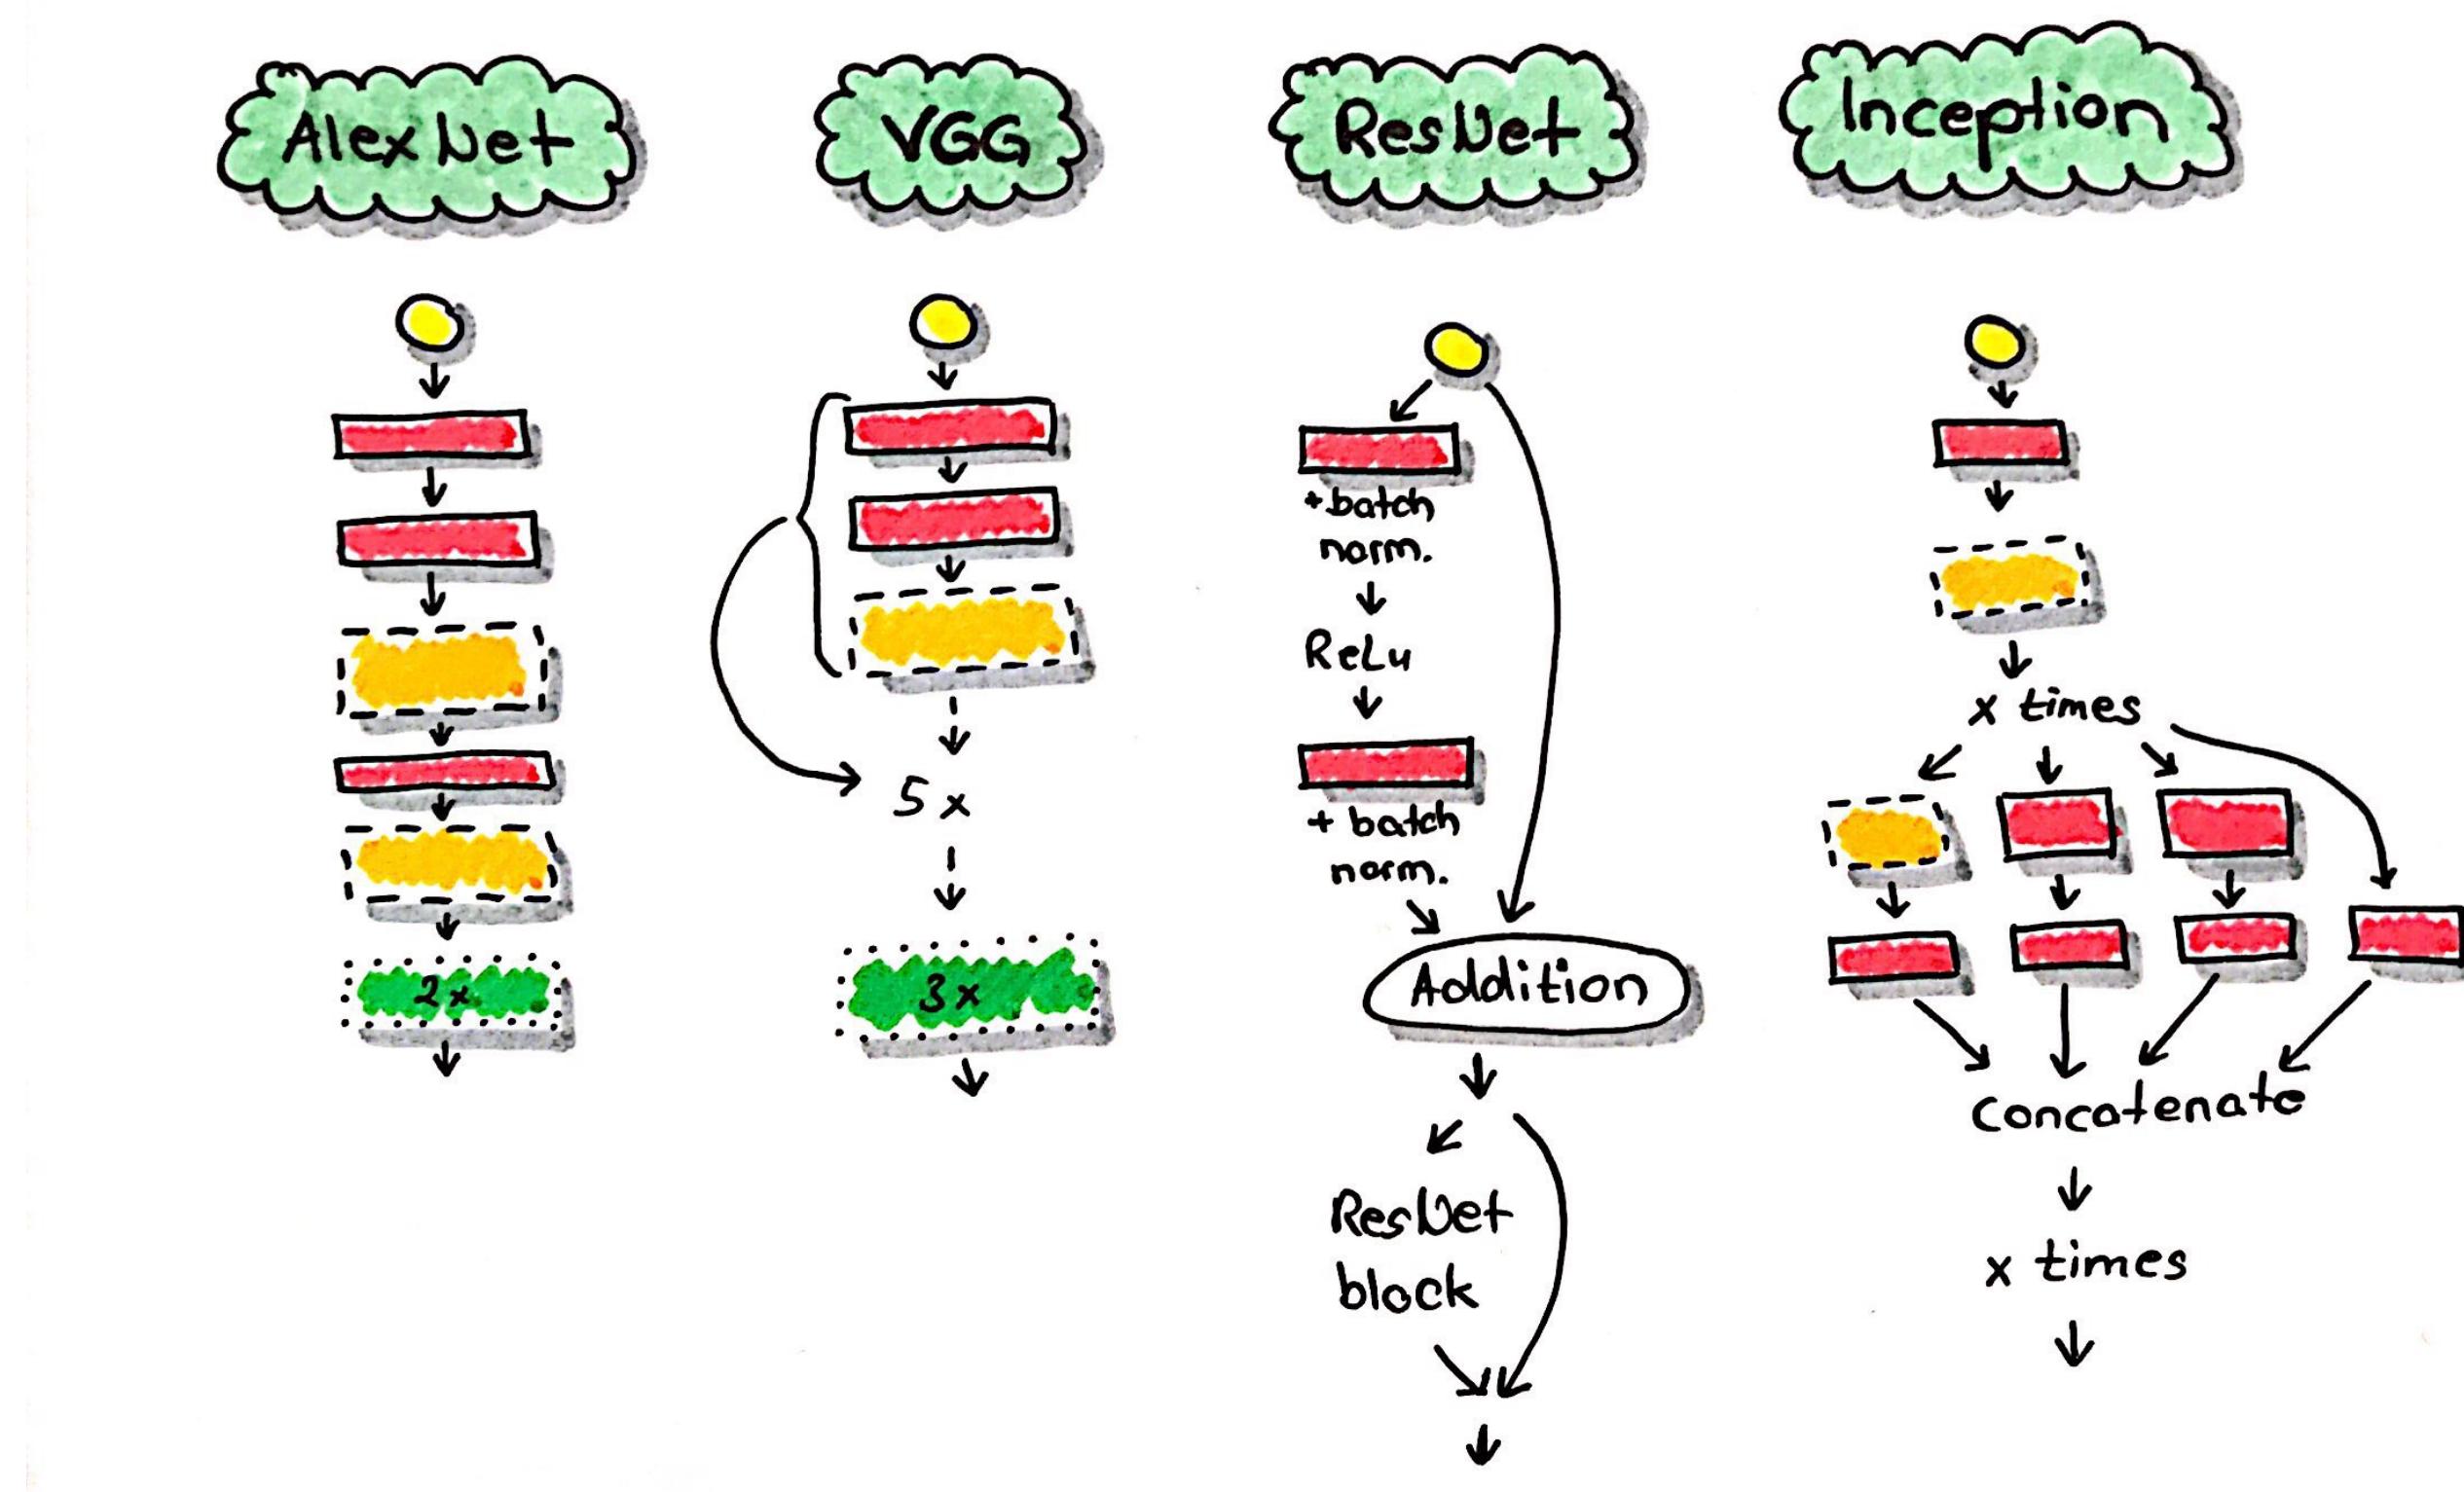


Evolution of neural nets for image recognition

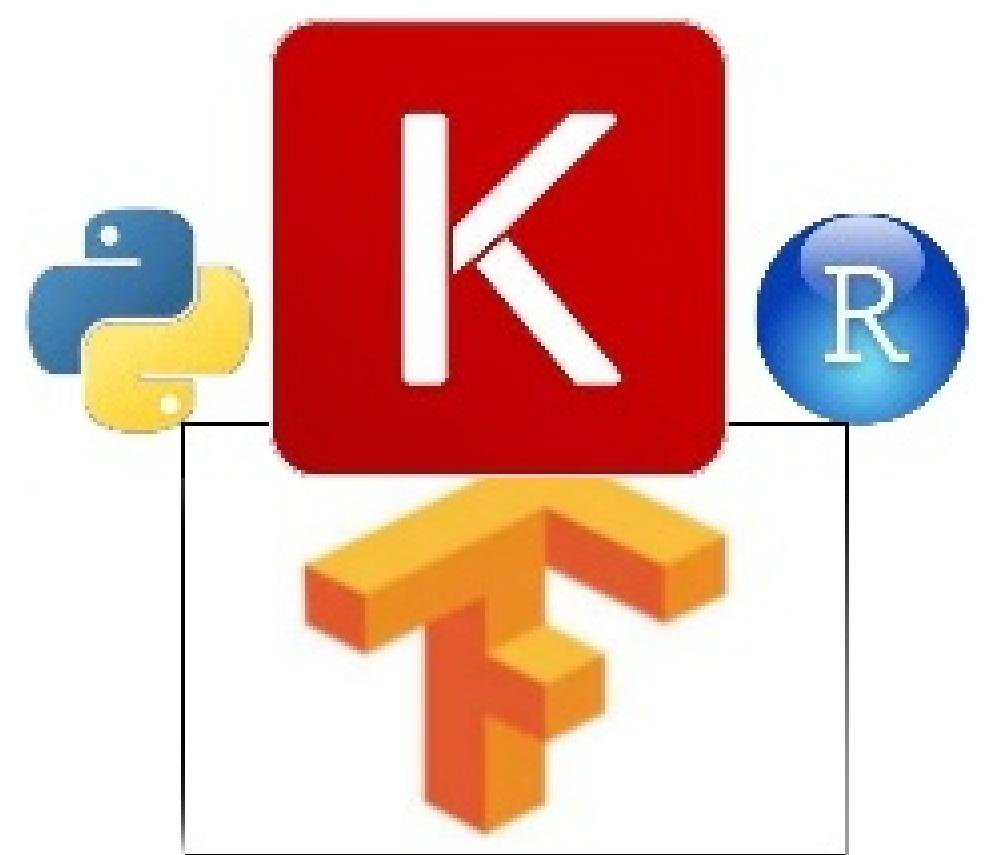


CNN architectures

- input image
- convolutional layer
- pooling layer
- dense layer



Introduction to TensorFlow



What are tensors?



Tensors = multidimensional arrays

Dimension

1

1	2	3
---	---	---

vector

2

1	2	3
4	5	6
7	8	9

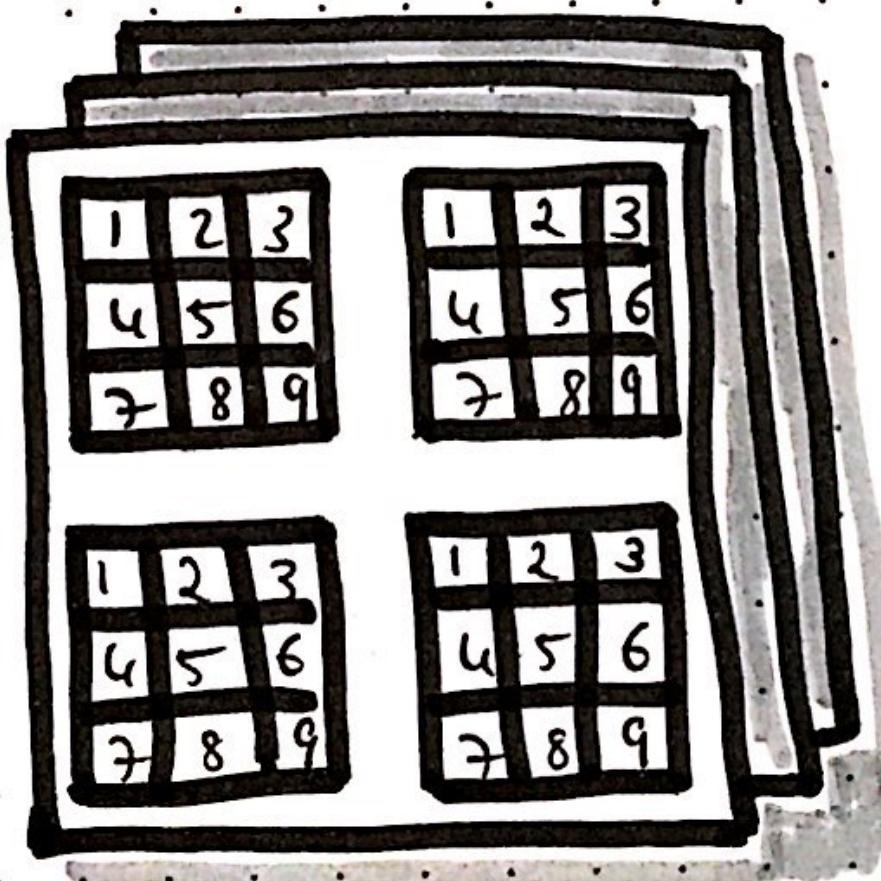
matrix

3

1	2	3
4	5	6
7	8	9

3-dim.
array

n



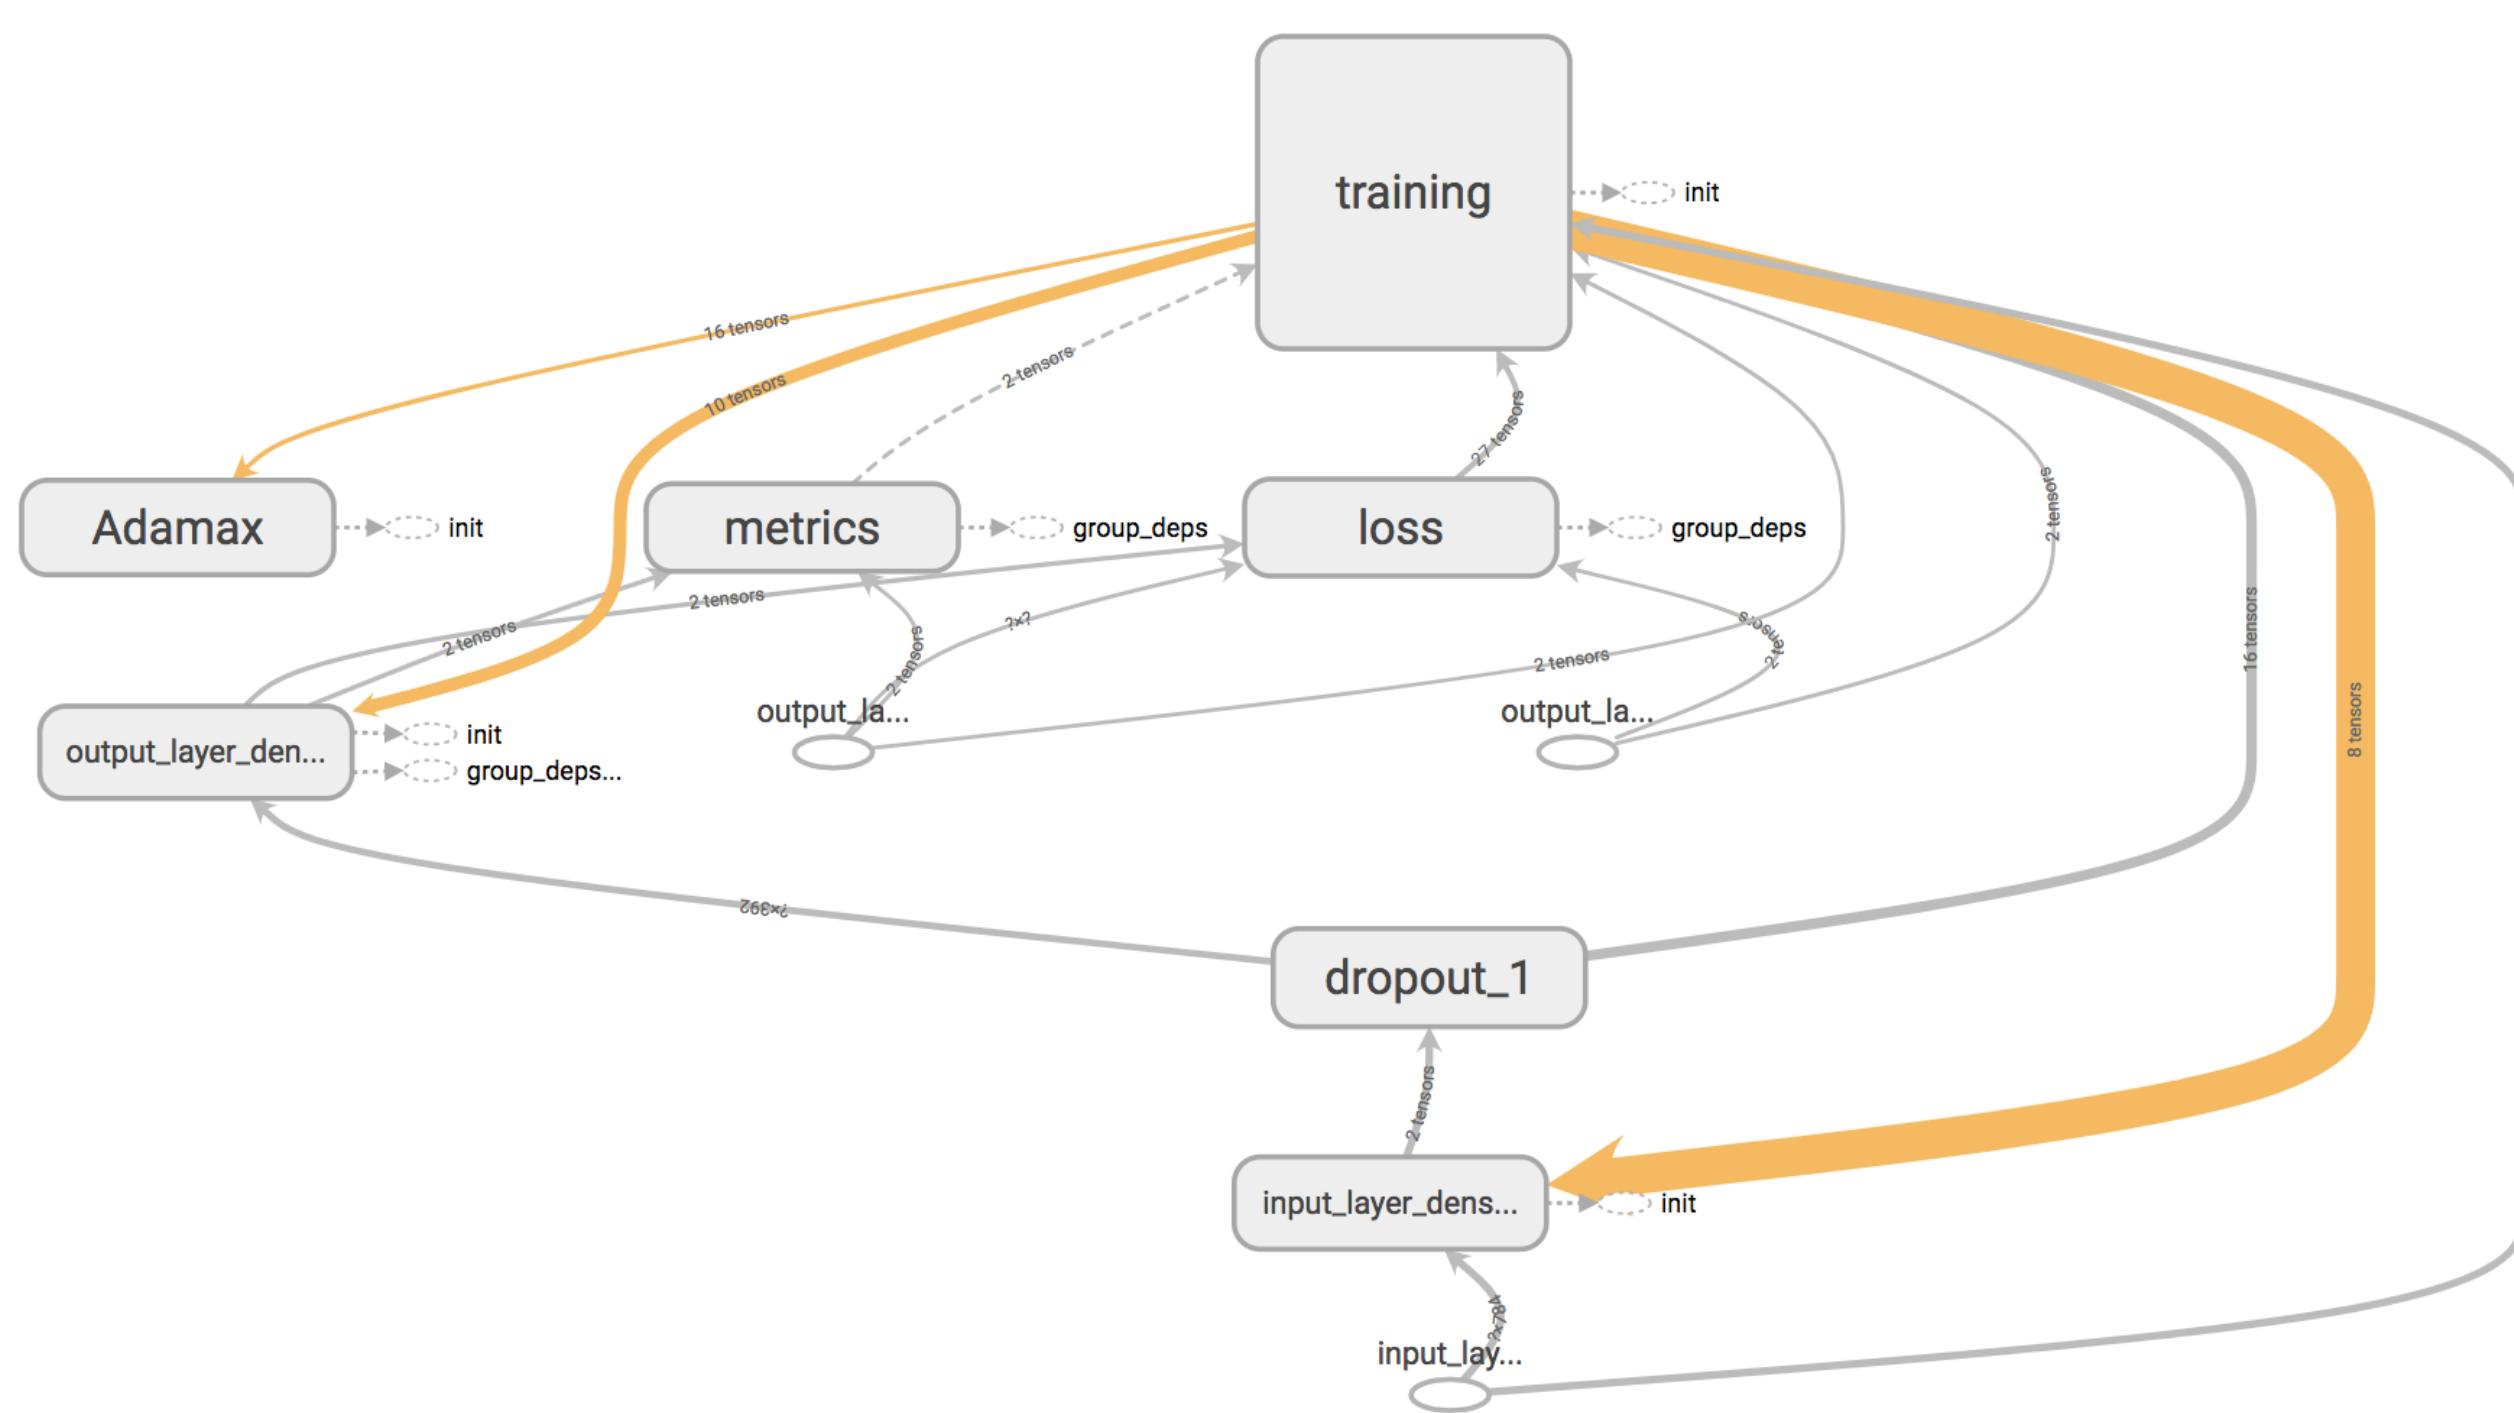
n-dim.
array

Tensor "Flow"

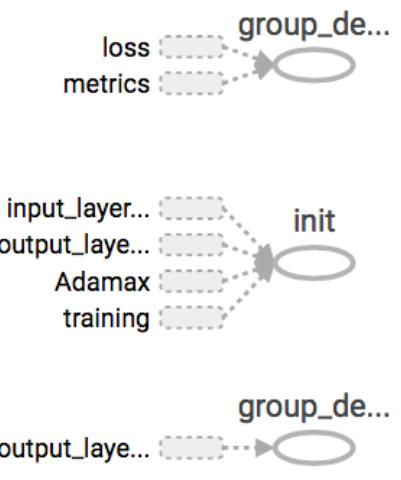


Graphs in TensorBoard

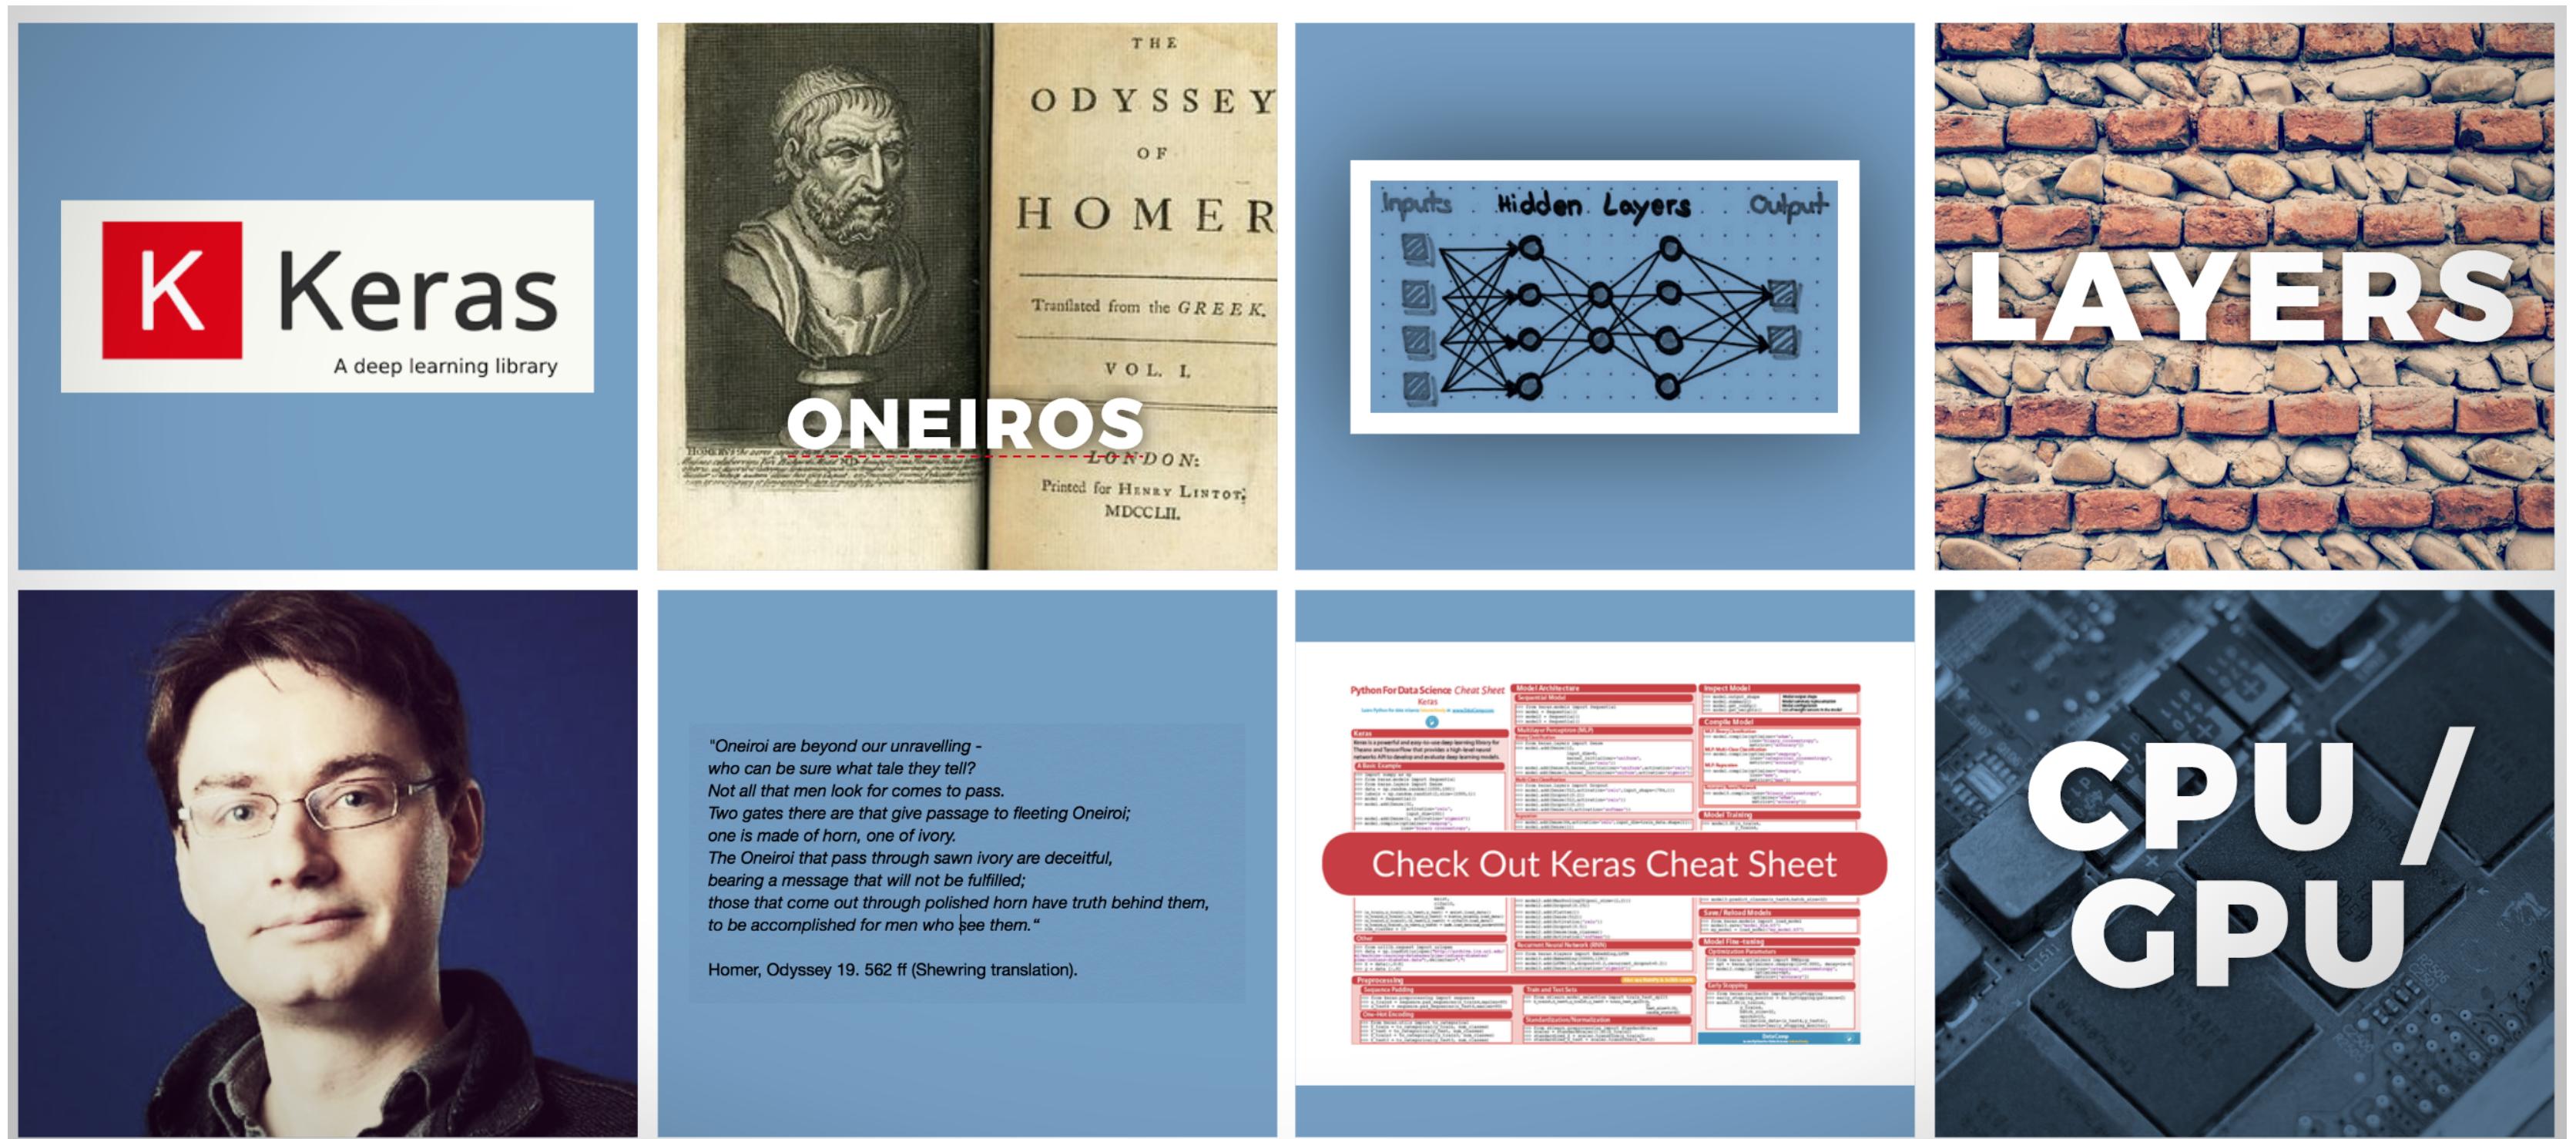
Main Graph



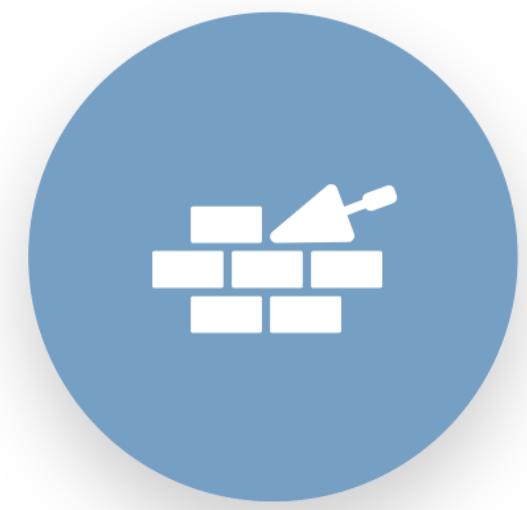
Auxiliary Nodes



Keras High-Level API for TensorFlow



Keras APIs



SEQUENTIAL MODELS

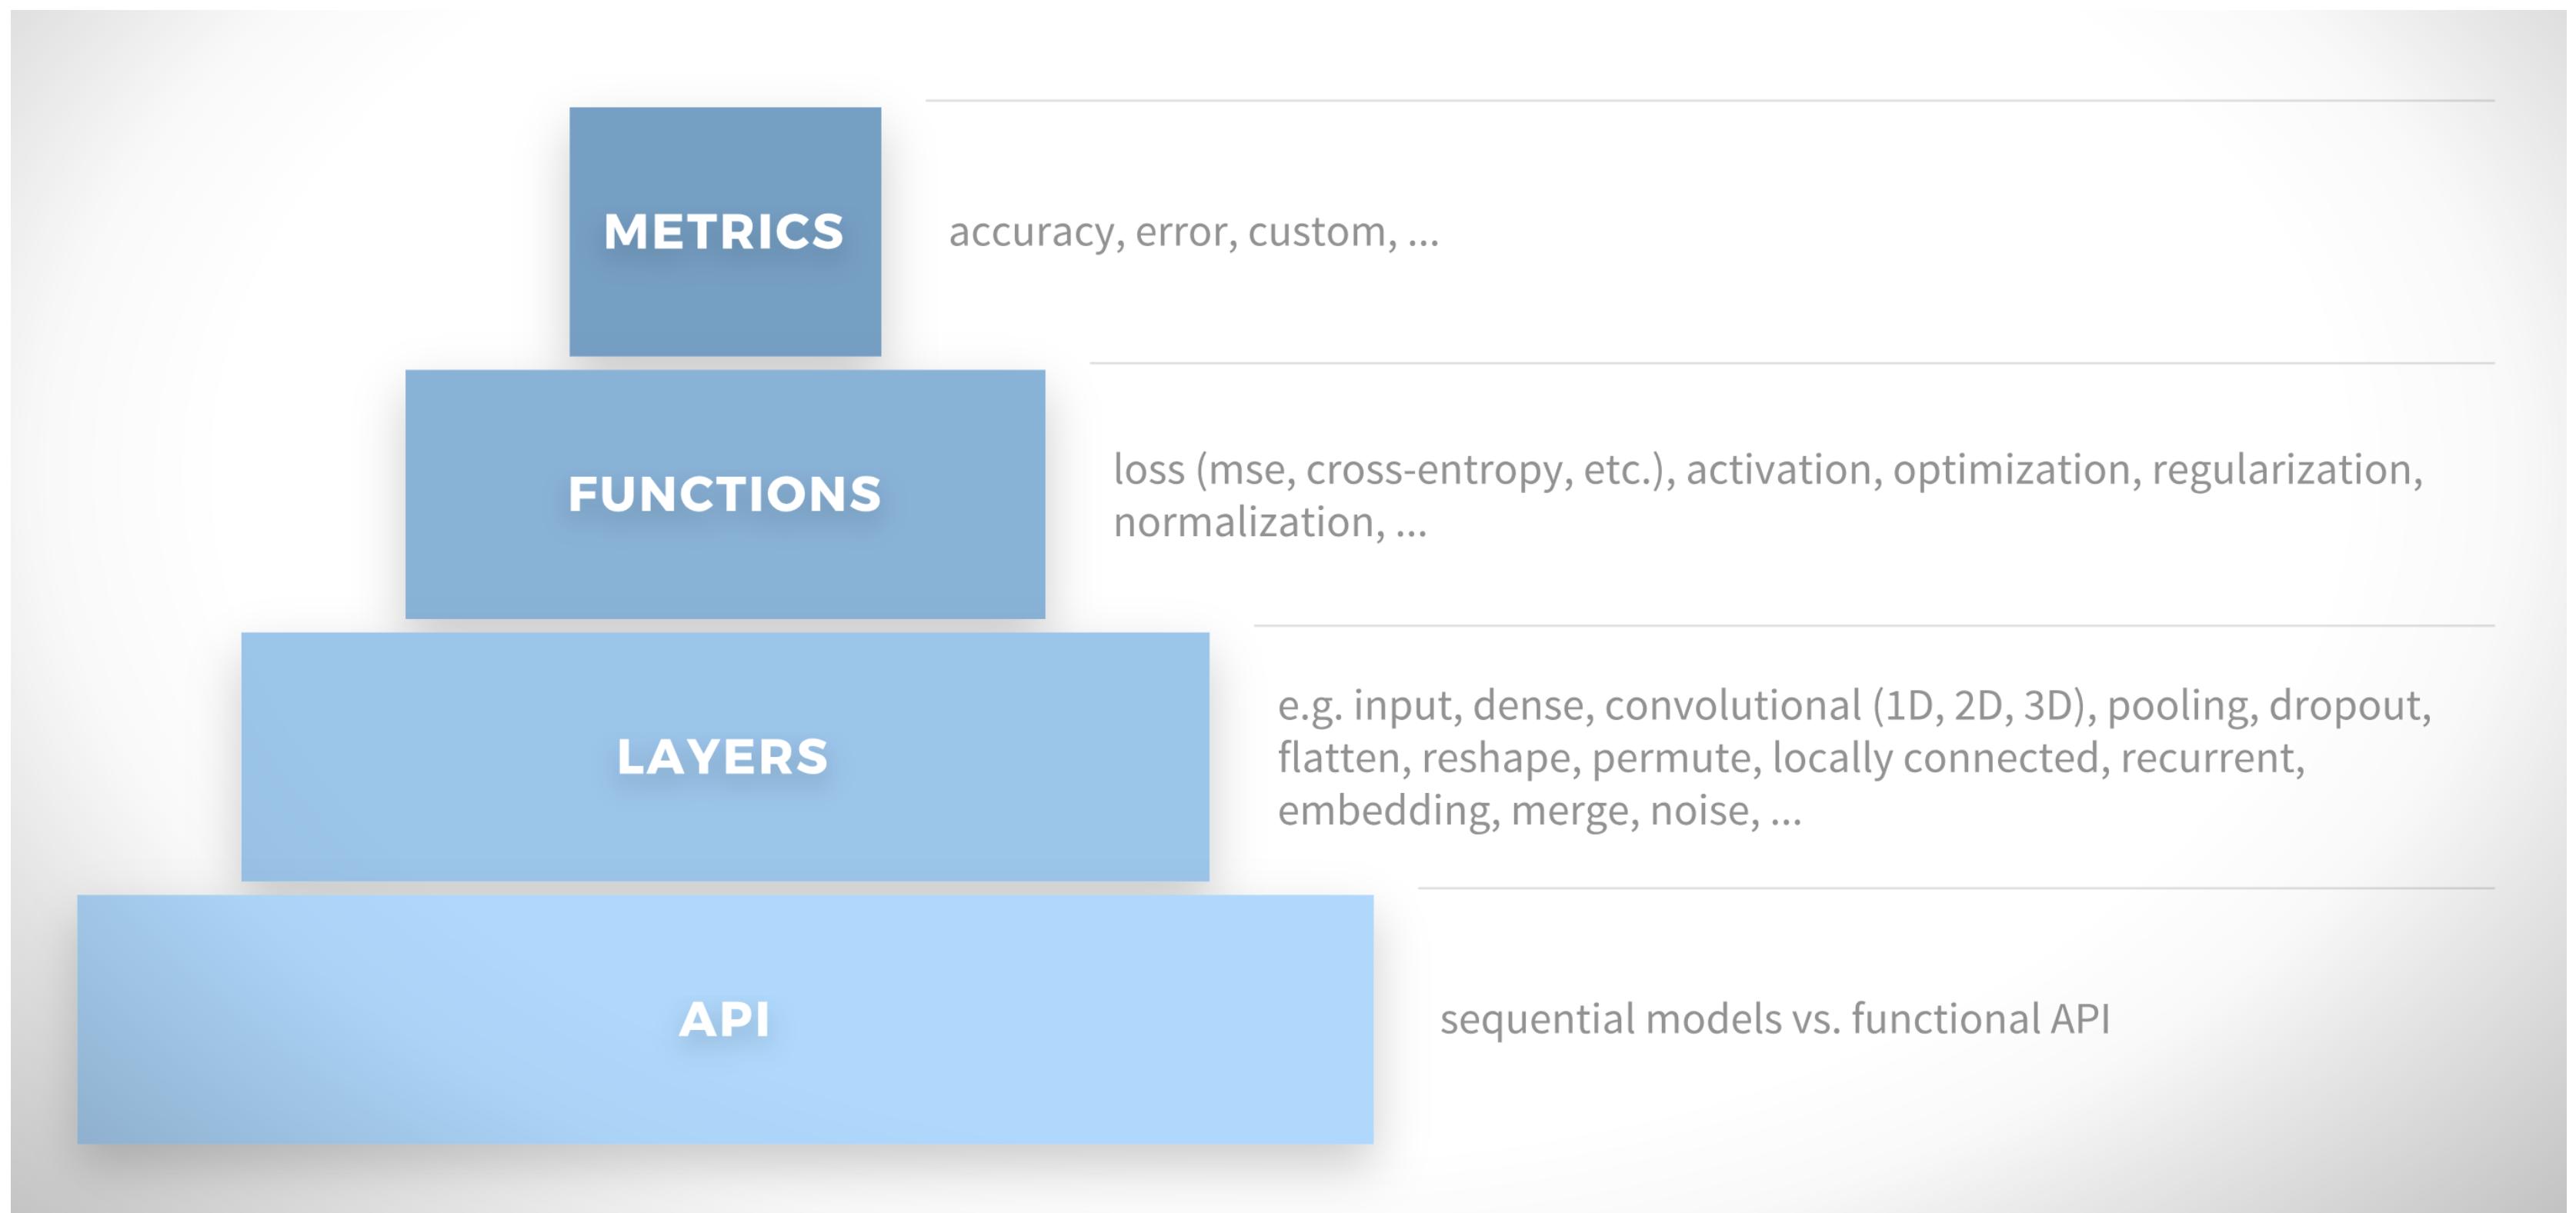
simple
suitable for most cases
linear order of layers
only one direction from input to output



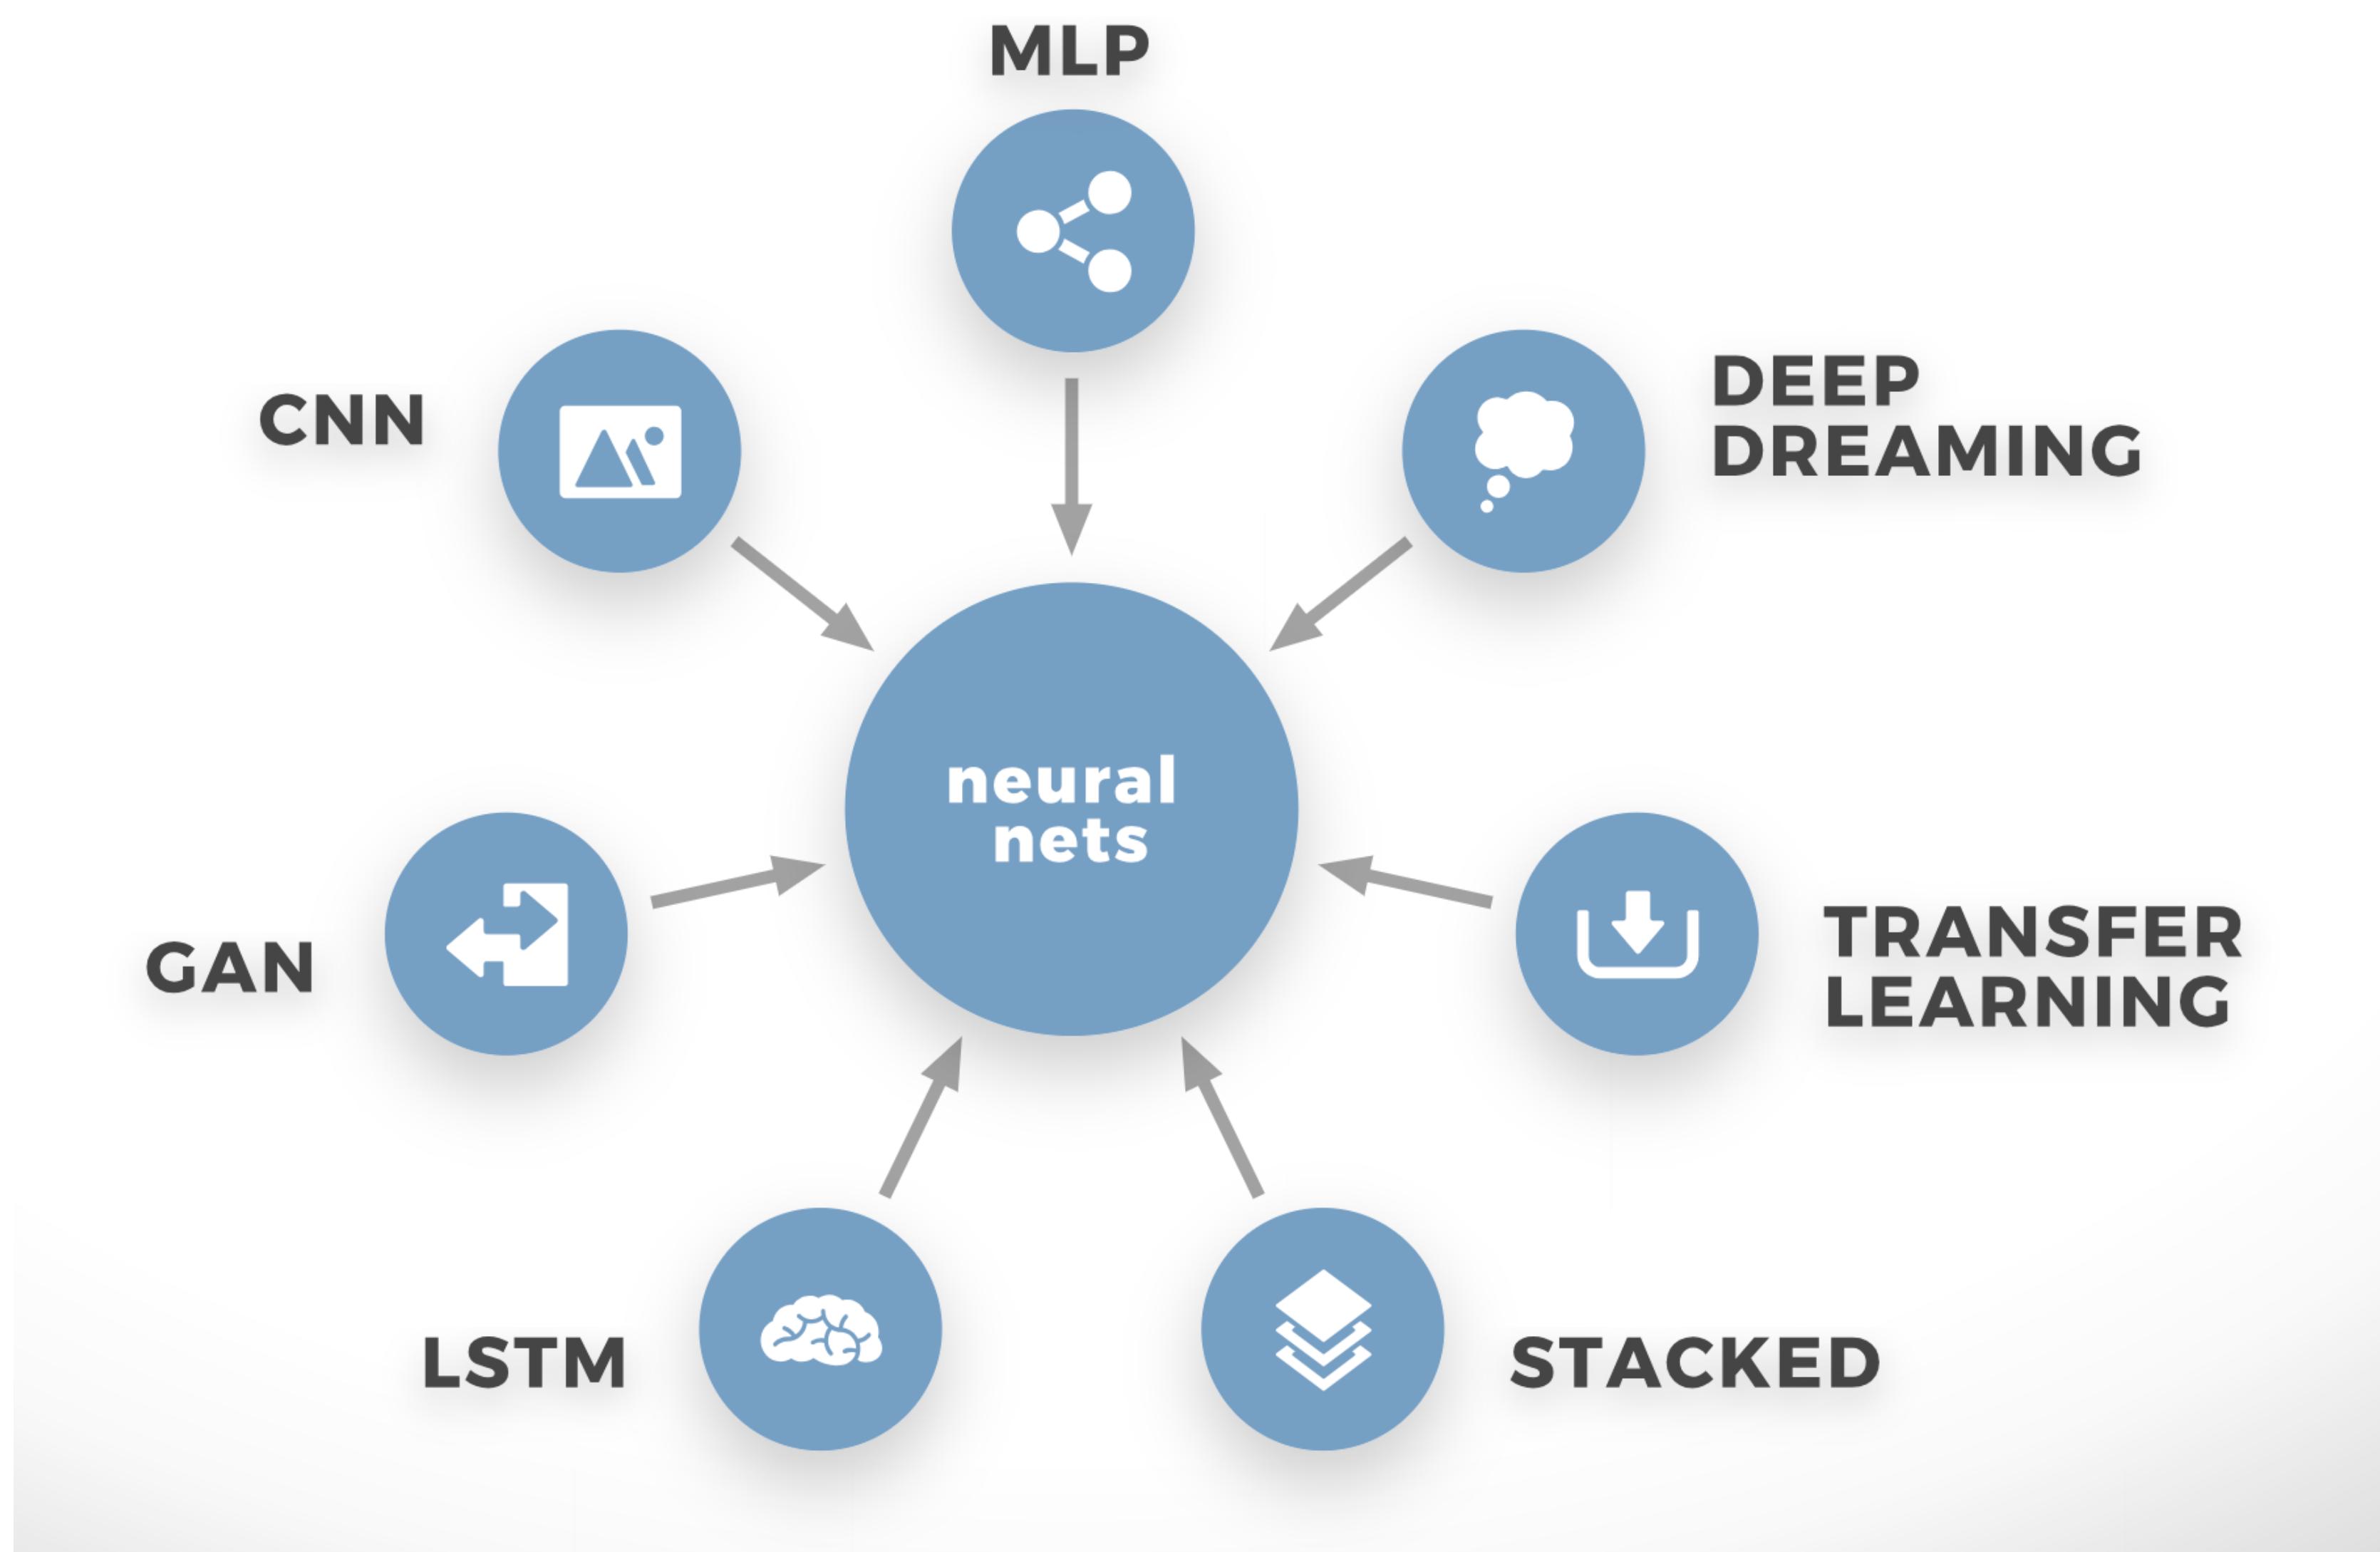
FUNCTIONAL API

more complex
suitable for complex models
can have multiple in- or outputs
layers can be non-sequential, e.g. LSTM

Keras layers



Endless possibilities



Let's get our hands dirty!

EDA

https://github.com/ShirinG/image_classification_keras_tf/blob/master/notebooks/01-eda.ipynb

```
In [10]: plot_images(imgs=train_images[:10], rows=2, columns=5)
```



Apply pretrained nets

- VGG16

https://github.com/ShirinG/image_classification_keras_tf/blob/master/notebooks/02-pretrained-1.ipynb

- Xception

https://github.com/ShirinG/image_classification_keras_tf/blob/master/notebooks/02-pretrained-2.ipynb

- Try out: ResNet50

Modify pretrained nets

https://github.com/ShirinG/image_classification_keras_tf/blob/master/notebooks/03-modify-pretrained.ipynb

Modify VGG16

- transfer learning (freeze all but the penultimate layer and re-train the last Dense layer) and
- fine tuning (un-freeze the lower convolutional layers and retrain more layers)

Validation set: fit_generator has no option validation_split

<https://keras.io/applications/#usage-examples-for-image-classification-models>

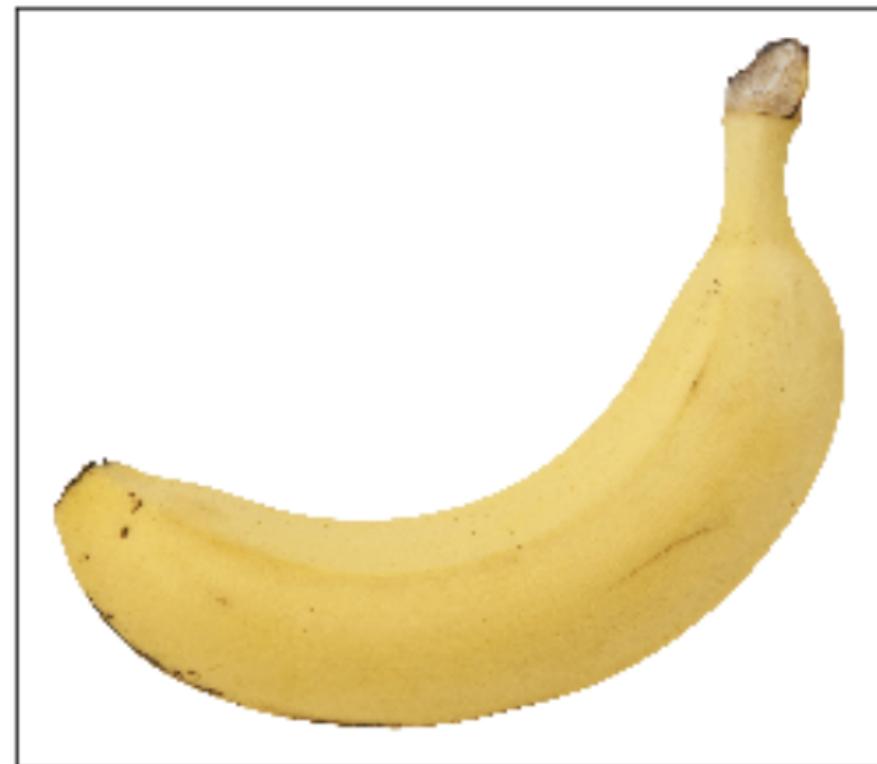
```
In [6]: # important: exclude top layers
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(75, 75, 3))
#base_model.summary()
```

```
In [7]: # Freeze the layers except the last 4 layers
for layer in base_model.layers[:-4]:
    layer.trainable = False
```

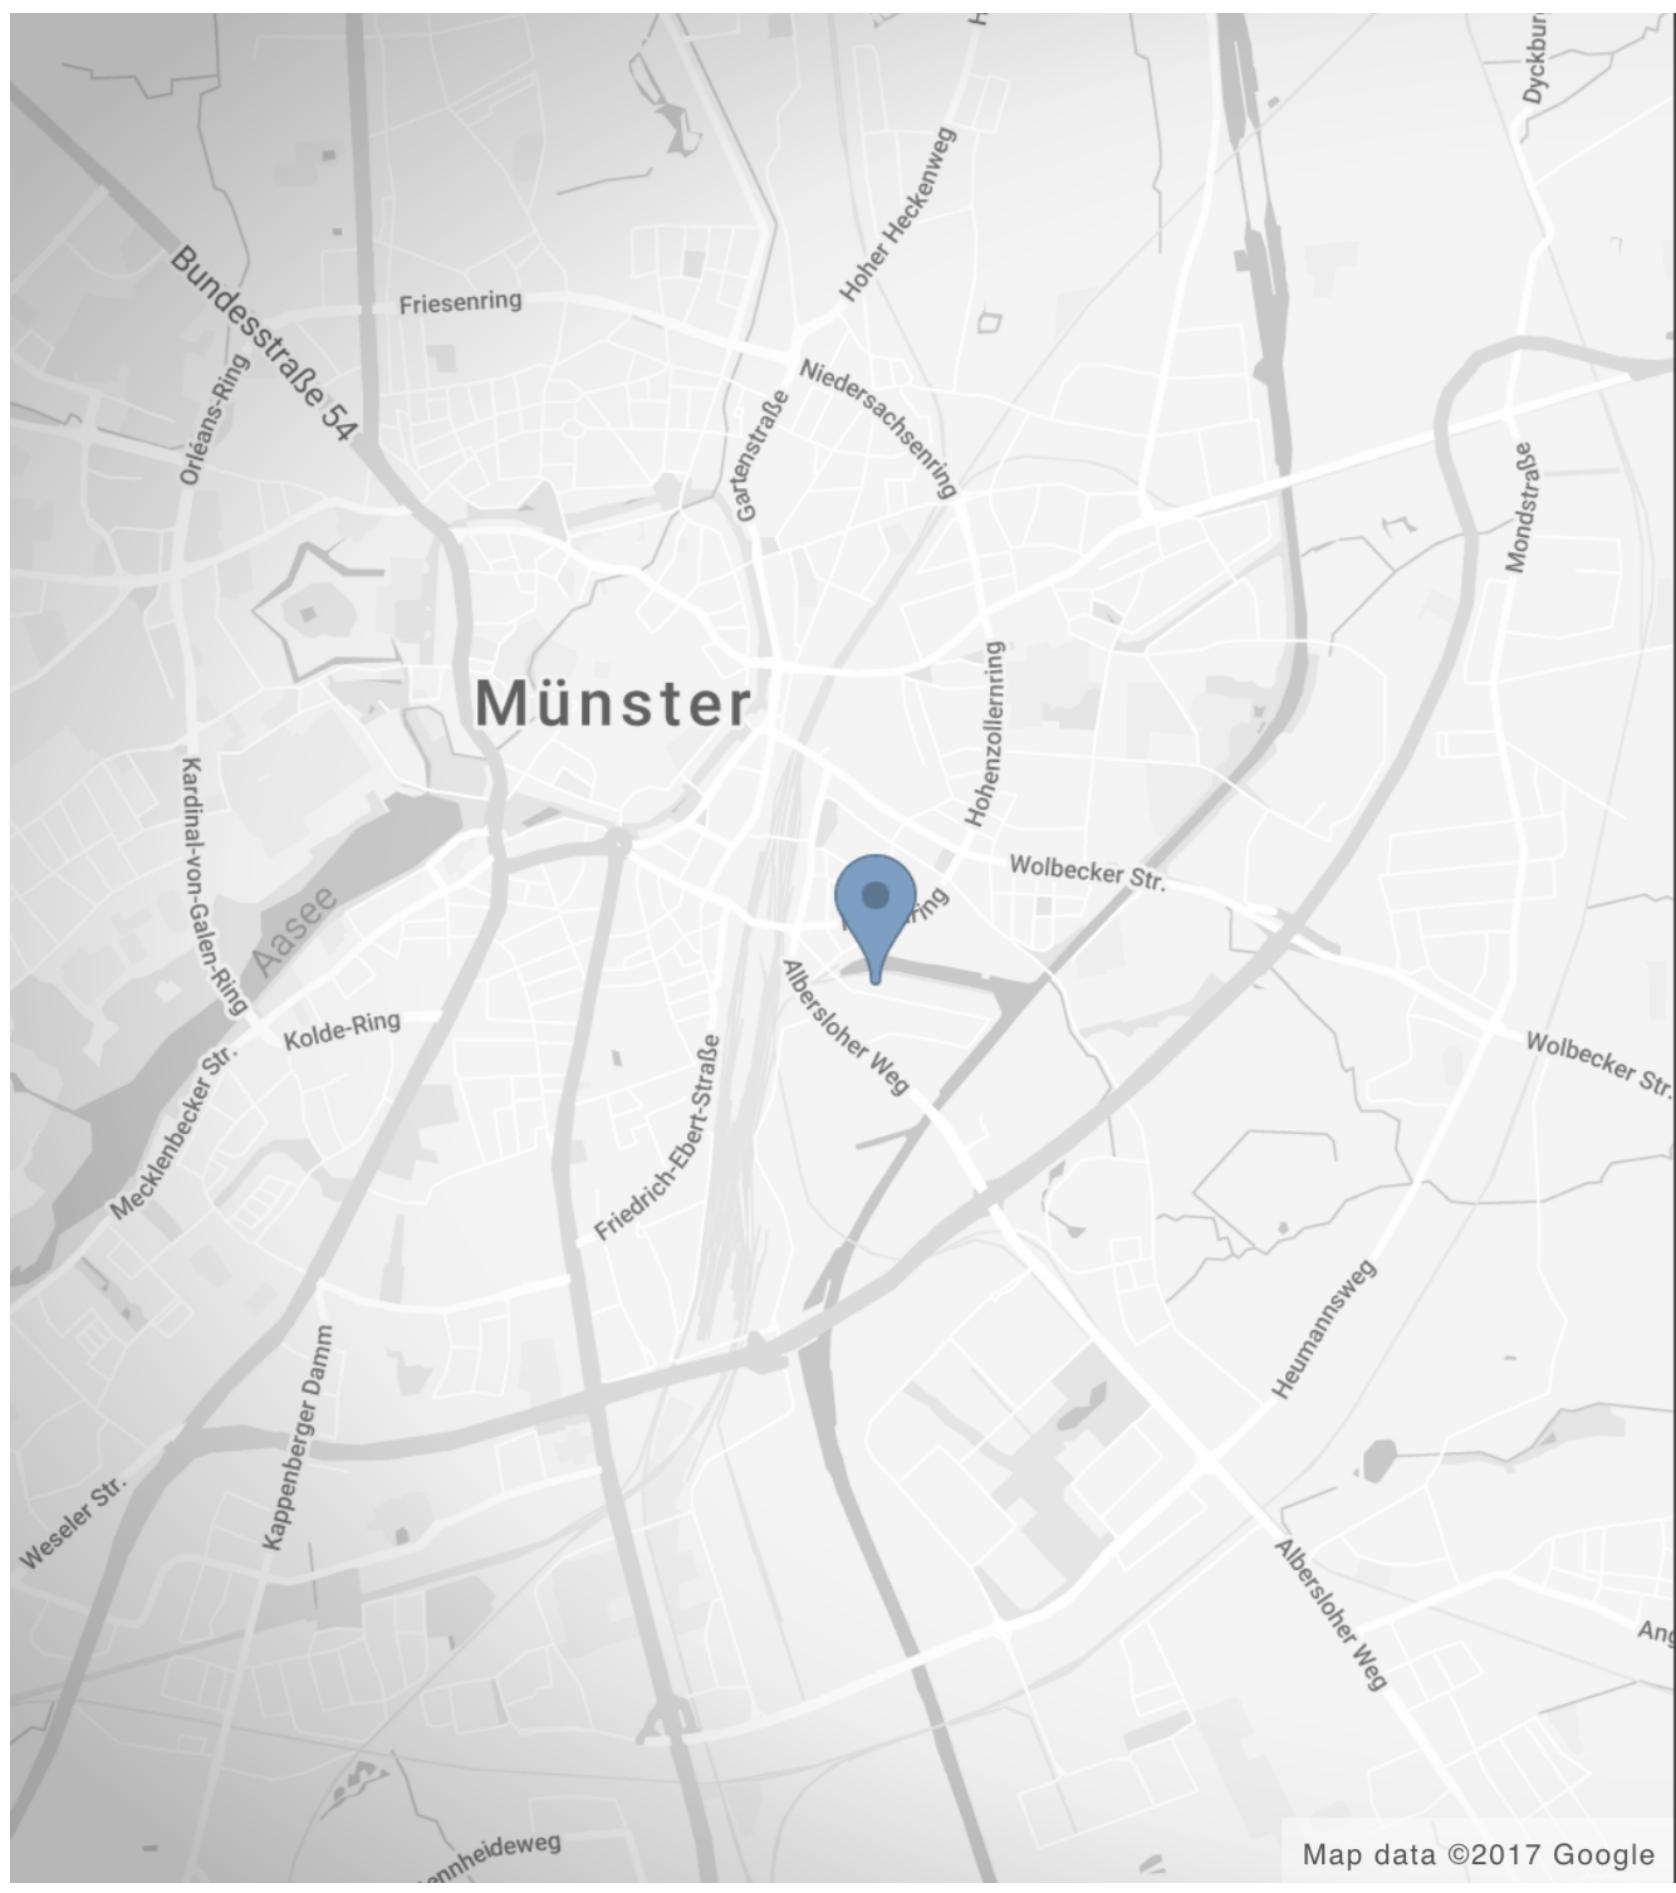
Build your own CNN

https://github.com/ShirinG/image_classification_keras_tf/blob/master/notebooks/04-fruits-cnn.ipynb

```
In [16]: classify_image_model(test_images[0])
```



Predicted class: Banana with probability 99.99985694885254%



**Thank
you!
And stay
connected**

Am Mittelhafen 14, Münster

 ShirinGlander

 shirin.glander@codecentric.de

www.codecentric.ai

<https://www.youtube.com/codecentricAI>

www.shirin-glander.de